

Genetic Diversity in Hybrid Breeding

Measuring it, constraining it, and optimising against it, with R

Eder D. B. S.

Contents

Preface	13
The problem, stated once	13
What is unusual here	13
How to read this	13
Running the code	14

I

Foundations

1	Coancestry, and the two matrices	19
1.1	From line to hybrid	19
1.2	The VanRaden matrix, and why it cannot measure diversity	19
1.3	Molecular coancestry	20
1.3.1	Computable form	20
1.3.2	Reading the diagonal	21
1.4	f is not an ad hoc substitute	21
1.5	The division of labour	22
2	Diversity in a subdivided population	23
2.1	Notation	23
2.2	The three coancestry blocks	23
2.2.1	f_{ij} , between subpopulations	23
2.2.2	f_{ii} , within a subpopulation	23
2.2.3	s_i and F_i , self-coancestry and molecular inbreeding	24
2.3	Two per-subpopulation metrics	24
2.4	Distances, and what they actually are	24
2.5	The equation that matters most in practice	25
2.6	F statistics and the partition	25
2.6.1	Why the partition changes the strategy by species	25
2.7	The worked example, reproduced	26
2.7.1	Removing an accession can <i>increase</i> diversity	26
2.7.2	Optimal contributions	27
2.8	Minimising coancestry is maximising N_e	27
2.9	The honest caveats	27
2.10	Two errata, found while verifying	28
2.11	What manages this quantity, and where	28
3	Response to selection: the other half of the objective	29
3.1	The breeder's equation, and how well it does here	29
3.2	Genic variance, and what it has to do with gene diversity	30

3.3	The decomposition this book already had	31
3.4	What weights = NULL is assuming	31
3.5	The α constraint is a restricted index	32
3.6	Where this connects	33

II

Simulation

4	Simulating lines, pools and crosses	37
4.1	What has to be produced	37
4.2	Simulating the lines	37
4.2.1	The divergence knob	38
4.3	From lines to hybrids	39
4.4	Predicted traits	40
4.5	The two matrices, both built	40
4.6	The stage-1 contract in one call	41
4.7	What the candidate pool looks like	42
5	Sizing a segregating population, F1 to F4	43
5.1	The simulated space is ancestry, not alleles	43
5.2	The assumptions, one at a time	43
5.3	Verifying the identities live	44
5.4	The results	45
5.4.1	The diversity curve saturates at 0.988, not at 1	46
5.4.2	The bulk is the worst scheme across the whole range	47
5.4.3	Block size does not depend on N	47
5.4.4	The metrics saturate at different N	48
5.5	Sizing the F1	49
5.6	Working numbers	50
5.7	Where this connects	50
6	Introgressing several genes by backcrossing	53
6.1	The operation, in one diagram	53
6.2	The model	54
6.3	Three criteria, and only one of them binds	54
6.4	Which selection index	55
6.5	How long to keep selecting recombinants	57
6.6	Allocating a fixed budget across BC1, BC2 and BC3	57
6.7	Fixing the stack: selfing against doubled haploids	58
6.8	Running the simulation	59
6.9	Assumptions, and what is not modelled	60
6.10	Where this connects	61

7	Simulating a breeding programme: AlphaSimR from the ground up	63
7.1	Step 1 – Founder genomes, and which engine to use	63
7.2	Step 2 – <code>SimParam</code> , and defining a trait	65
7.3	Step 3 – The base population, and where the variance components live . . .	66
7.4	Step 4 – The DNA lottery, and why an F1 is not a family	66
7.5	Step 5 – Selfing, doubled haploids, and what the shortcut buys	67
7.6	Step 6 – Phenotyping, selection, and checking the breeder’s equation	68
7.7	Step 7 – Two traits, and what equal weighting assumes	69
7.8	Step 8 – The parts assembled: reciprocal recurrent selection	70
7.8.1	Founding two pools with a real coalescent history	70
7.8.2	One cycle	71
7.8.3	The trajectory across cycles	71
7.8.4	The final hybrid pool through <code>stage2_select()</code> , unchanged	72
7.9	Assumptions, and what this adds over Chapter 4.	73
7.10	Where this connects	74
8	A complete plant breeding programme, and what it costs in diversity	75
8.1	The scheme	75
8.2	Why a breeding programme needs a burn-in	76
8.3	One programme, fifteen years	76
8.4	Genomic selection: two levers, not one	78
8.5	The comparison, replicated	79
8.6	The diversity ledger	80
8.7	Assumptions	80
8.8	Where this connects	81

III

Metrics

9	Diversity metrics in the genomic era	85
9.1	Three lenses, not three synonyms	85
9.2	The metrics, one at a time	85
9.2.1	Rate of inbreeding and effective size	85
9.2.2	<code>F_hom</code> , homozygosity-based	86
9.2.3	<code>F_drift</code> , drift-based	86
9.3	The classical equivalence, and the covariance that destroys it	86
9.3.1	Verifying the identity	87
9.3.2	Why each matrix generates a sign	87
9.3.3	The signature, on live data	87
9.4	The other matrices, briefly	88
9.4.1	Side by side	88
9.5	Reference results	89
9.5.1	The diagnostic nobody runs	89
9.6	What to carry into plant breeding	90

9.7	Limits of the animal-to-plant transposition	90
9.8	Operational checklist	90
10	Heterosis, dominance, and what divergence buys	93
10.1	The dominance model, minimally	93
10.1.1	The one-locus scale	93
10.1.2	What d_k is, and what it is not	94
10.1.3	The mid-parent difference	94
10.2	Expected heterosis over a pool pair	95
10.3	The identity	95
10.3.1	A third form: one minus the between-pool coancestry	96
10.4	A second identity, at the hybrid level	96
10.5	Inbreeding depression, and why it does not price α	97
10.6	When genetic distance predicts heterosis exactly	98
10.7	Divergence is not a bonus	98
10.7.1	The two terms, term by term	99
10.8	What actually buys heterosis	101
10.8.1	Alignment is a covariance	101
10.8.2	The scenarios	101
10.8.3	Isolating the divergence channel	102
10.9	Reading the S1 panel again	103
10.10	General and specific combining ability	103
10.10.1	GCA is an average effect in the other pool's background	105
10.10.2	SCA is minus twice a weighted parental inner product	105
10.10.3	The variance components	105
10.11	Divergence destroys SCA faster than it builds GCA	106
10.11.1	What this says about ranking crosses	109
10.12	What this does not change	109
10.12.1	What it would take to use d	110
10.13	Where this connects	110
11	A catalogue of metrics on $\mathbf{f}$	111
11.1	Group coancestry, and the off-diagonal problem	111
11.2	Gene diversity: θ and H_e are the same quantity	112
11.3	Status number N_s	112
11.4	Effective number of parents	113
11.5	Allelic richness	113
11.6	Distances: E-NE and A-NE	114
11.7	Effective dimensionality	114
11.8	Decomposition by heterotic group	115
11.8.1	Where this could go next	115
11.9	The correct baseline	116
11.9.1	Why not compare against the whole population	116
11.9.2	Operational definition of loss	116
11.9.3	How to set the α limit	117
11.10	Every metric against the null	117

12	Choosing metrics: collapse, cost, redundancy	119
12.1	Finding 1: the hybrid coancestry matrix is never needed	119
12.1.1	Derivation	119
12.1.2	Four consequences	120
12.1.3	Verification	120
12.2	Finding 2: what each metric costs	121
12.3	Finding 3: the metric space is nearly one-dimensional	122
12.3.1	Reproducing the argument at book scale	123
12.4	Finding 4: the encoding, not the constraint handler, decides whether DE works	125
12.5	Finding 5: the index is not the truth, and its error is not white	126
12.5.1	The correlation structure matters more than the amount	127
12.5.2	What this does to the recommendation	127
12.6	The recommendation, stated	128
12.6.1	Inside the optimiser	128
12.6.2	At validation, on the returned plan only	128
12.6.3	Not recommended	128
12.6.4	Caveats	128
12.7	Drop-in code	128
13	Metrics at each pipeline stage, and five traps	131
13.1	The kernel: one matrix, all four stages	131
13.1.1	Trap 1: modified Rogers distance is the same statistic	131
13.2	S1: conserving the heterotic groups	132
13.2.1	Trap 2: two F_{ST} definitions differ by up to 2x	132
13.2.2	Trap 3: F_{ST} is the wrong target to maximise	132
13.2.3	The S1 monitoring panel	133
13.3	S2: choosing crosses for DH induction	133
13.3.1	Three verified identities	133
13.3.2	Progeny variance, and the usefulness criterion	134
13.3.3	Trap 4: genetic distance is not a proxy for progeny variance	134
13.4	S3: selecting lines per se	135
13.4.1	Trap 5: the incremental greedy constraint silently returns nothing	136
13.5	S4: hybrid mining	137
13.6	Where the diversity actually goes	137
13.6.1	What the diversity term buys, and what it does not	138
13.7	Cost summary	139
13.8	Checklist	139
14	Predicting breeding values: GBLUP, BayesA, BayesB	141
14.1	A training population with real noise	141
14.2	The truth this chapter is predicting	141
14.2.1	The generative model is GBLUP's own assumption	142
14.2.2	A hybrid's value is the mean of two parental values	142
14.3	GBLUP: prediction from the relationship matrix	142
14.3.1	Where the mixed model equations come from	142
14.3.2	Why the ridge blend is not optional	143
14.3.3	Reliability: the accuracy you can compute without the truth	144
14.4	What the random split was measuring	144

14.5	BayesA and BayesB: prediction from marker effects	146
14.5.1	The priors, explicitly	146
14.5.2	What the sampler actually selects	146
14.5.3	What BayesB does that BayesA cannot	148
14.6	Verification: two routes agree	148
14.7	Which model wins, and when	148
14.7.1	<code>pi</code> was fixed, not tuned	150
14.8	From GEBVs to <code>stage1_build()</code>	150
14.9	Back to Finding 5: a measured accuracy, not an assumed one	151

IV

Optimisation

15	Optimal contributions and the α constraint	155
15.1	Constraint, not a weighted sum	155
15.1.1	This argument is older than it looks	155
15.2	Relation to the standard OCS formulation	156
15.2.1	The continuous relaxation as an upper bound	156
15.3	Setting $\alpha_{\max}$	158
15.3.1	How far $\alpha_{\max}$ moves when <code>f</code> moves	160
15.4	A caveat carried over from the OCS literature	161
15.5	Quadratic programming, when contributions are continuous	162
15.6	What belongs in the inner loop	162
16	Combinatorial selection	165
16.1	Four baselines before any optimiser	165
16.2	Greedy, with incremental updates	166
16.3	The encoding problem	168
16.3.1	The better encoding	168
16.4	Differential evolution	169
16.4.1	The warm start is not optional	170
16.4.2	Checking convergence against the bound	171
16.5	The full frontier	172
16.6	Which strategy to use	173
17	Advanced selection: structure, bounds and scale	175
17.1	Three facts about this problem	175
17.1.1	<code>f</code> is a Gram matrix	175
17.1.2	Maximising diversity is minimising the submatrix sum	175
17.1.3	The per-line cap is a partition matroid	175
17.2	What lazy greedy would not buy	176
17.3	The move that was already paid for	176
17.3.1	Against the differential evolution, at equal wall-clock	177
17.3.2	The encoding, wired up at last	179
17.4	The bound was not a bound	179
17.5	Rounding the relaxation into a plan	181
17.6	Exact, and how far the heuristics really are	181

17.7	Scaling to 200 of 10,000	182
17.8	Multi-objective, and why not NSGA-II	184
17.9	What not to reach for	184
17.10	Which method to use	184

V

Evidence and practice

18	What the plant literature establishes	189
18.1	The headline qualification	189
18.2	Why plant practice differs from animal breeding	190
18.2.1	Distance measures, and the absence of a standard	190
18.3	Heterotic groups	190
18.3.1	Assignment is not a solved problem	190
18.3.2	The empirical trend that justifies the constraint	191
18.4	Selecting within the pool	191
18.4.1	The constraint alone is not sufficient	191
18.4.2	Cycle length, and what goes unreported	191
18.4.3	What is hybrid-specific, and what is not	191
18.5	Genebanks, core collections and introgression	192
18.6	The animal-breeding corpus: which matrix to manage	192
18.7	Optimisation machinery reported in plants	192
18.8	What remains unestablished	193
18.9	Consequences for the pipeline	193
19	Diversity across cycles	195
19.1	The cycle	195
19.2	Freezing p_0 , and what happens if you do not	196
19.3	The trajectory	196
19.3.1	It is not an artefact of the cheap selector	198
19.4	The rate of inbreeding, and what N_s is not	199
19.5	Choosing a budget for a horizon	199
19.6	The per-line cap, at last measurable	200
19.7	Injection, and what it costs	200
19.8	What the corpus supports	201
19.9	Where this connects	201
20	The limit, and what the metrics cannot see	203
20.1	Robertson's limit	203
20.2	The R^2 column decides whether the theory applies	203
20.3	What the metrics cannot see	204
20.3.1	The experiment	204
20.3.2	With real linkage	205
20.4	The two costs of selection	205
20.5	Where this connects	206

21	End to end, and the operational checklist	207
21.1	The whole pipeline in one script	207
21.1.1	The reference values	207
21.1.2	Metrics per scenario	208
21.1.3	Both lenses, and the pool partition	208
21.1.4	Distance from the random null	209
21.1.5	The selection table	209
21.2	Reading the frontier	210
21.3	What this simulation cannot tell you	211
21.4	The operational checklist	212
21.4.1	Setting up	212
21.4.2	Choosing the budget	212
21.4.3	Running the optimiser	212
21.4.4	Validating the plan	212
21.4.5	Across cycles	212
21.5	Where to go next	212
	Appendices	215
A	Notation and glossary	217
A.1	Symbols	217
A.1.1	Matrices	217
A.1.2	Coancestry and diversity	217
A.1.3	The partition	218
A.1.4	Heterosis and cycles	218
A.1.5	Inbreeding, the three lenses	219
A.1.6	Selection and optimisation	219
A.1.7	Genomic prediction	219
A.1.8	Population sizing	220
A.1.9	Response to selection	220
A.2	Conventions	221
A.3	Terms	221
B	The hybdiv package	223
B.1	Installation	223
B.2	File map	223
B.3	The two-stage contract	224
B.4	Function index	224
B.4.1	Simulation and context	224
B.4.2	The two matrices	224
B.4.3	Diversity metrics – all <code>f(idx, ctx) -> scalar</code>	224
B.4.4	The fast line route	225
B.4.5	Selection strategies	225
B.4.6	Benchmarking	226
B.5	Runnable scripts	226
B.5.1	Pipeline stages S1-S4	226
B.5.2	Heterosis and cycles	227
B.5.3	Caballero & Toro	227
B.5.4	Population sizing	227

B.5.5	Reference oracles	228
B.6	Full export list	228
C	Verification	231
C.1	The method: slow oracles	231
C.2	The identity sweep	232
C.3	The test suite	232
C.4	Provenance of every figure and table	235
C.4.1	Computed live at render	235
C.4.2	Read from cache	236
C.5	Regenerating the cached results	237
C.6	Reproducing the book	237
C.7	Known limitations of the verification	238
D	The bibliographic corpora	239
D.1	What they are	239
D.2	Structure	239
D.3	Relevance and coverage	239
D.4	Themes in the animal corpus	241
D.5	Querying them	241
D.6	Limitations	242
D.7	Classical references	242
	References	245
	Bibliography	247

Preface

A hybrid breeding programme has two objectives that pull against each other. The first is short-term: pick the crosses with the highest predicted performance. The second is long-term: do not spend the germplasm that will be needed to make the next round of gain possible. Every programme knows this. Very few can say, in a number, how much of the second they traded away to buy the first.

This book is about producing that number, defending it, and putting it inside an optimiser.

The problem, stated once

Given a pool of N candidate F1 hybrids – each a cross between one line from heterotic group A and one from group B, each with a predicted multi-trait index – choose n_{sel} of them so that the mean index is as high as possible, subject to losing no more than a fraction α of the molecular gene diversity that a *random* set of the same size would have carried.

Three of those words carry the entire argument.

Molecular. Diversity is measured on the Caballero & Toro molecular coancestry matrix f , never on a VanRaden genomic relationship matrix G . This is not a stylistic preference. G is centered on its own population by construction, so the sum of all its elements is exactly zero; every diversity quantity computed from it is meaningless. Chapter 1. shows this happening.

Random. The reference for “no loss” is the distribution of random subsets of size n_{sel} , not the full candidate pool. Comparing a selected set of 100 against a pool of 2500 confounds a selection effect with a sampling effect, and the sampling effect alone can look like a loss of half a percent. Section 11.9 separates them.

Subject to. α is a *constraint*, not a term in a weighted sum. When you write $\alpha_{max} = 0.05$ you are saying “I accept losing five percent”, which is a statement a breeding committee can approve. A weighted sum says “diversity is worth λ index units”, which is a statement nobody can approve because nobody knows what λ means. Chapter 15. makes the distinction precise, and shows the standard optimal-contribution formulation it specialises.

What is unusual here

Most treatments of diversity management come from animal breeding, where the population is randomly mating and the unit of selection is an individual. Plant hybrid breeding breaks both assumptions:

- The parents are **inbred lines**, so their within-individual diversity is zero by construction, and the interesting heterozygosity lives in the F1 product, not in the parents.
- The population is **permanently subdivided** into heterotic groups, and the divergence between them is not a problem to be minimised – it is the engine of heterosis, to be *maintained*. A single pool-wide statistic cannot express “conserve within, hold between”. Chapter 10. derives that claim as an exact identity, and then bounds it: divergence pays only where it coincides with dominance, and drifting the pools apart buys nothing at all.
- The candidates are **crosses of fixed lines**, so the progeny variance term that drives cross-selection criteria elsewhere collapses to zero, and the hybrid-by-hybrid coancestry matrix turns out never to be needed at all (Section 12.1).

Each of these changes what should be measured. None of them changes the underlying theory, which is why the book starts with the theory.

How to read this

The book is meant to be run, not only read. Every derivation that could be checked numerically is checked numerically, in a code block you can execute. Where a result is an exact algebraic identity, the deviation from a simulated oracle is printed next to it; where it is a Monte Carlo agreement, the residual is printed.

The argument the parts assemble is a chain, and it is worth stating before the table that lists them:

1. Diversity between the pools is what makes the hybrid worth selling, and diversity within them is what the next cycle has to select on (Chapter 2., Chapter 10.).
2. Those two are in tension: every plan that raises one lowers the other, and the partition of Chapter 11. is what makes the trade visible.
3. Within-pool diversity is worth protecting because it is the raw material of *response*: genic variance is the same sum over the same loci as gene diversity, weighted by squared effect instead of equally (Chapter 3.). That is what makes a diversity constraint a variance constraint in disguise, and what makes “restrict rather than price” a sixty-year-old theorem rather than a preference.
4. So the loss of within-pool diversity is a **budget**, set against what random selection of the same size would have cost (Section 11.9), not a term traded off inside an objective (Chapter 15.).
5. How much diversity there is to spend depends on how the candidates were generated: a panel cannot be selected back up to variation it never sampled (Chapter 5.), and a pool of near-copies cannot be rescued by any optimiser (Chapter 6.). Chapter 4. and Chapter 7. build that panel two different ways, and everything downstream is blind to which – Chapter 7. from the primitives up, so the generative model is something a reader can rebuild rather than accept, and Chapter 8. pointing the same primitives at an inbred-line programme, where the diversity cost of running faster is measured directly.
6. Spending the budget optimally is a combinatorial problem with structure worth naming (Chapter 16., Chapter 17.).
7. The budget is worth paying because of what it leaves for later cycles, which is a claim about trajectories and is checked as one (Chapter 19.) – and then predicted, which turns out to work only where the constraint holds (Chapter 20.).

The parts are ordered so that each depends only on what came before:

Part	What it establishes
Foundations	The coancestry kernel, the partition of diversity in a subdivided population, and the equation that says what selection buys
Simulation	How the candidate pool is generated, and how large a segregating population has to be
Metrics	What can be measured, what it costs, and which measurements are redundant
Optimisation	How the constraint enters an optimiser, and which optimiser to use
Evidence and practice	What the literature does and does not establish, and a worked end-to-end run

Readers who only want the operational answer can read Chapter 12. (which metrics), Chapter 13. (at which stage) and Chapter 21. (the whole pipeline in one script), then come back for the derivations.

Running the code

All code is in the `hybdiv` package that ships with this book.

```
# from the repository root
install.packages(c("DEoptim", "quadprog"))
devtools::install(".")

library(hybdiv)
st1 <- stage1_simulate() # lines -> X, f, hybrids
res <- stage2_select(st1$X, st1$f, st1$hybrids, n_sel = 100, fL = st1$fL)
res$metrics
```

Every live example in the book runs at a deliberately small scale – 25 lines per pool and 2000 markers, giving 625 candidate hybrids – so that a full render takes minutes. The scale is set once, in `_common.R`:

```
str(book_cfg)
#> List of 6
```

```
#> $ n_pool_A: num 25
#> $ n_pool_B: num 25
#> $ m       : num 2000
#> $ n_traits: num 3
#> $ fst     : num 0.15
#> $ seed    : num 1
```

Results that are expensive at any scale – the differential-evolution sweeps, the cost benchmark, the population-sizing grid – are precomputed and shipped in [book/data](#) and [book/figs](#). [book/scripts/regenerate.R](#) regenerates all of them, and Appendix C. says exactly which figures and tables come from where.

A note on numbers: everything reported here is reproducible output from a simulated dataset, not a fixed constant. Change the simulation configuration or a random seed and the numbers move. The *relationships* between them – which metric dominates which, which identity holds exactly – do not.

I

Foundations

1	Coancestry, and the two matrices . . .	19
1.1	From line to hybrid	19
1.2	The VanRaden matrix, and why it cannot measure diversity	19
1.3	Molecular coancestry	20
1.4	f is not an ad hoc substitute	21
1.5	The division of labour	22
2	Diversity in a subdivided population	23
2.1	Notation	23
2.2	The three coancestry blocks	23
2.3	Two per-subpopulation metrics	24
2.4	Distances, and what they actually are	24
2.5	The equation that matters most in practice	25
2.6	F statistics and the partition	25
2.7	The worked example, reproduced	26
2.8	Minimising coancestry is maximising N_e	27
2.9	The honest caveats	27
2.10	Two errata, found while verifying	28
2.11	What manages this quantity, and where	28
3	Response to selection: the other half of the objective	29
3.1	The breeder's equation, and how well it does here	29
3.2	Genic variance, and what it has to do with gene diversity	30
3.3	The decomposition this book already had	31
3.4	What <code>weights = NULL</code> is assuming	31
3.5	The α constraint is a restricted index	32
3.6	Where this connects	33

1. Coancestry, and the two matrices

Two matrices are built from the same marker data and they are not interchangeable. One predicts traits. The other measures diversity. Using the first for the second is the most consequential mistake available in this subject, and it fails silently – it returns a number, and the number is meaningless.

This chapter builds both, shows the failure, and derives the matrix used for the rest of the book.

1.1 From line to hybrid

Let M^L be the $L \times m$ matrix of line genotypes, with $M_{ik}^L \in \{0, 2\}$ the count of the reference allele at marker k . Inbred lines are homozygous, so the value 1 never occurs – residual heterozygotes must be imputed or discarded beforehand.

The F1 hybrid between lines a and b receives one gamete from each parent. Since both parents are homozygous, the gamete from a carries $M_{ak}^L/2$ copies of the reference allele, so

$$M_{(ab)k}^H = \frac{M_{ak}^L + M_{bk}^L}{2} \in \{0, 1, 2\}. \quad (1.1)$$

The value 1 appears exactly at markers where the two lines differ. Those are the hybrid's heterozygous loci, and they are the ones carrying heterosis.

Predicting the hybrid's *genotype* is deterministic, not statistical. It is the one piece of information in the whole pipeline that requires no assumption at all – unlike its phenotype, which requires a great many.

From here on we work with the **within-individual reference-allele frequency**,

$$x_{ik} = \frac{M_{ik}^H}{2} \in \{0, 0.5, 1\}, \quad (1.2)$$

which is the natural quantity for coancestry: it is the probability that an allele drawn at random from individual i at marker k is the reference allele.

```
ctx <- simulate_data(book_cfg)
ctx$X[1:4, 1:8]
#>      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
#> [1,] 0.5 0.5 0.5 1 0 1 0.5 0.5
#> [2,] 0.5 0.0 0.0 1 0 1 1.0 0.0
#> [3,] 0.0 0.0 0.5 1 0 1 1.0 0.5
#> [4,] 0.0 0.0 0.0 1 0 1 1.0 0.5
```

Rows are hybrids, columns markers. A 0.5 is a heterozygous locus.

1.2 The VanRaden matrix, and why it cannot measure diversity

With p_k the reference-allele frequency in the hybrid population and $Z = M^H - 12p'$ the centered matrix,

$$G = \frac{ZZ'}{2 \sum_k p_k(1 - p_k)}. \quad (1.3)$$

This is VanRaden's first parametrisation. It estimates realised additive relationship and it is the right matrix for GBLUP. It is implemented as `vanraden_G()`:

```
args(vanraden_G)
#> function (M)
#> NULL
```

Now the failure. Centering on $2\mathbf{p}$ uses frequencies from the population itself, so every column of $\mathbf{Z}$ sums to zero:

$$\mathbf{Z}'\mathbf{1} = \mathbf{0} \Rightarrow \mathbf{1}'\mathbf{G}\mathbf{1} = \frac{\|\mathbf{Z}'\mathbf{1}\|^2}{2 \sum_k p_k q_k} = 0. \quad (1.4)$$

The sum of **all** elements of $\mathbf{G}$ is identically zero. Group coancestry is proportional to that sum, so it is zero for the whole population, and the status number $N_s = 1/(2\theta)$ diverges:

```
c(sum_G      = sum(ctx$G),
  theta_on_G  = mean(ctx$G) / 2,
  Ns_on_G     = 1 / (2 * mean(ctx$G) / 2))
#>      sum_G      theta_on_G      Ns_on_G
#> -4.178879e-11 -5.348298e-17 -9.348769e+15
```

That is not a rounding artefact to be worked around. It is exact, and it is what centering means. In a *subset* the sum is not zero, but the quantity is still centered: it oscillates around zero, takes negative values, is not a probability, and admits no identity-by-descent interpretation.

! The root cause is the implicit base population

$\mathbf{G}$ measures relatedness *relative to the current population*. Diversity requires an *absolute* scale. Any matrix whose centering constant is estimated from the same data whose diversity you are measuring will have this problem, whatever it is called.

1.3 Molecular coancestry

Malécot's coancestry between i and j is the probability that two randomly sampled alleles, one from each individual at the same locus, are identical. Replacing identity *by descent* with identity *by state* – which is what markers actually observe – yields molecular coancestry [1], [2]:

$$f_{ij} = \frac{1}{m} \sum_{k=1}^m [x_{ik} x_{jk} + (1 - x_{ik})(1 - x_{jk})] . \quad (1.5)$$

The first term is the probability that both sampled alleles are the reference allele; the second, that both are the alternative. Three properties follow immediately:

- $f_{ij} \in [0, 1]$ always. **A negative value is a symptom of the wrong matrix.**
- Self-coancestry is $f_{ii} = (1 + F_i)/2$, with F_i the individual's molecular homozygosity. For a fully heterozygous F1, $f_{ii} = 0.5$.
- No centering, no scaling, no arbitrary base population.

1.3.1 Computable form

The double sum is $O(n^2m)$ if computed naively. Expanding,

$$\sum_k [x_{ik} x_{jk} + (1 - x_{ik})(1 - x_{jk})] = 2 \sum_k x_{ik} x_{jk} - \sum_k x_{ik} - \sum_k x_{jk} + m, \quad (1.6)$$

so with $\mathbf{s} = \mathbf{X}\mathbf{1}$ the row sums,

$$\mathbf{f} = \frac{1}{m} \left(2 \mathbf{X}\mathbf{X}' - \mathbf{s}\mathbf{1}' - \mathbf{1}\mathbf{s}' + m \mathbf{J} \right), \quad (1.7)$$

a single matrix product – the same cost as building $\mathbf{G}$.

```
molecular_coancestry
#> function (X)
#> {
#>   m <- ncol(X)
```

```
#> s <- rowSums(X)
#> f <- (2 * tcrossprod(X) - outer(s, rep(1, nrow(X))) - outer(rep(1,
#> nrow(X)), s) + m)/m
#> dimnames(f) <- NULL
#> f
#> }
#> <bytecode: 0x94da9fa48>
#> <environment: namespace:hybdiv>
```

```
fmt(c(minimum      = min(ctx$f),
      maximum      = max(ctx$f),
      diagonal_mean = mean(diag(ctx$f)),
      offdiag_mean  = mean(ctx$f[upper.tri(ctx$f)])))
#>      minimum      maximum diagonal_mean offdiag_mean
#>      0.6448      0.8318      0.8192      0.6717
```

Everything is inside $[0, 1]$, as it must be.

1.3.2 Reading the diagonal

The value 0.5 is the floor of the diagonal, reached by a hybrid heterozygous at *every* marker. The excess above 0.5 measures exactly the fraction of markers at which the two parent lines carry the same allele – loci where the F1 is homozygous and therefore contributes nothing to heterosis.

```
fd <- mean(diag(ctx$f))
fmt(c(diagonal_mean      = fd,
      molecular_F        = 2 * fd - 1,      # f_ii = (1 + F_i) / 2
      heterozygous_fraction = 2 * (1 - fd)))
#>      diagonal_mean      molecular_F heterozygous_fraction
#>      0.8192      0.6383      0.3617
```

This fraction depends directly on the divergence between the heterotic pools (`fst` in the configuration): more divergent pools give more heterozygous F1 hybrids and a diagonal closer to 0.5. That relationship is what Chapter 4. puts under experimental control.

1.4 **f** is not an ad hoc substitute

Two independent lines of work converge on this matrix, which is worth stating because it is otherwise easy to read **f** as “the thing we used because **G** broke”.

Nejati-Javaremi et al. (1997) proposed the *allelic-similarity* method, $f_{ij} = \frac{1}{4} \sum_{k=1}^2 \sum_{l=1}^2 \delta_{kl}$ with $\delta_{kl} = 1$ when the k -th allele of i equals the l -th of j . That is exactly the identity-by-state average above, generalised from single loci to a genome-wide mean. F. Gómez-Romano, B. Villanueva, J. Fernández, J. A. Woolliams, and R. Pong-Wong [4] contrast it directly against crossproduct-based matrices: the allelic-similarity matrix is guaranteed to stay in $[0, 1]$ by construction, while crossproduct matrices normalised by *current* allele frequencies are not – precisely the failure demonstrated above.

The same matrix arrives from a second direction. M. A. Toro, B. Villanueva, and J. Fernández [5] and T. H. E. Meuwissen, A. K. Sonesson, G. Gebregiorgis, and J. A. Woolliams [6] show that fixing the VanRaden centering frequency at $p_k = 0.5$ for every marker, instead of the sample frequency, produces a matrix – written $\mathbf{G}_{0.5}$ – that is proportional to homozygosity-based inbreeding and molecular coancestry. Structurally that is **f**: replacing $2p_k$ with 1 removes the population-dependent term that forced $\mathbf{1}'\mathbf{G}\mathbf{1} = 0$.

So **f** is the member of the VanRaden family that has an absolute scale, reached independently from two unrelated derivations. That it also happens to be the one that does not break is not a coincidence.

1.5 The division of labour

Matrix	Centering			Range	Sum of all elements	Use	
G (VanRaden)	sample	allele	fre-	centered,	exactly 0	trait	prediction
	quencies			signed		(GBLUP)	
f (molecular coancestry)	none			[0, 1]	positive, informative	every	diversity quantity

This split is enforced throughout `hybdiv`: every metric takes `ctx$f`, and `ctx$G` appears only where traits are predicted. The test suite asserts `sum(G) == 0` precisely so that a future change which silently starts feeding `G` to a diversity metric fails loudly.

The next chapter takes `f` and asks what happens when the population it describes is not one population, but several.

2. Diversity in a subdivided population

A breeding programme is never one population. It is heterotic groups, or accessions, or landraces, or cycles – and the diversity of the whole is not the average of the diversity of the parts. [A. Caballero and M. A. Toro \[1\]](#) give the partition that makes this precise, and it is the analytical backbone of everything that follows: the constraint in Chapter 15. is a special case of it, and the within/between asymmetry that makes plant hybrid breeding different from animal breeding is a statement about which term of it you are allowed to minimise.

The paper's central claim is uncomfortable and worth stating first: **a method that ranks accessions by distance alone can invert the correct conclusion**. An accession can be the most distant *because* it is internally fixed, and removing it can *raise* the diversity of the pool. We reproduce that result numerically below.

2.1 Notation

- A metapopulation of n subpopulations; subpopulation i has N_i individuals, $N_T = \sum_i N_i$.
- A tilde, $\tilde{x}$, is a weighted mean **within** subpopulations: $\tilde{f} = \sum_i f_{ii} N_i / N_T$.
- A bar, $\bar{x}$, is a mean over **all pairs** in the metapopulation, between-subpopulation pairs included: $\bar{f} = \sum_{i,j} f_{ij} N_i N_j / N_T^2$.
- Always $\bar{f} \leq \tilde{f} \leq \tilde{s}$. The differences between those three numbers *are* the partition of diversity.
- Everything is computed per locus and then averaged over loci. Never average frequencies first.

⚠ The subtle point, and the one most often got wrong

Every coancestry here includes the pairs of an individual with **itself** – group coancestry in the sense of Cockerham. That is what makes $1 - f_{ii}$ exactly Nei's expected heterozygosity, and it is what makes the algebra close. Drop the diagonal and none of the identities below hold.

2.2 The three coancestry blocks

2.2.1 f_{ij} , between subpopulations

The probability that two alleles, one sampled at random in subpopulation i and one in j , are identical. From markers, $f_{ij} = \sum_k \tilde{p}_{k,i} \tilde{p}_{k,j}$.

Read it as **redundancy** between two materials. If you cross plants of i with plants of j , the progeny has expected inbreeding $F = f_{ij}$. It answers: what do I gain by recombining these two accessions?

In maize, f_{ij} between a Stiff Stalk pool and a Non-Stiff Stalk pool is typically low – and that low between-group coancestry is exactly what sustains the heterotic pattern. Between two elite lines of the same family it can exceed 0.6, and the cross generates nothing new.

2.2.2 f_{ii} , within a subpopulation

The case $i = j$. From markers $f_{ii} = \sum_k \tilde{p}_{k,i}^2$, the expected homozygosity, so

$$1 - f_{ii} = H_{e,i} = \text{Nei's gene diversity of accession } i. \quad (2.1)$$

A landrace accession regenerated twenty times with forty ears per cycle accumulates coancestry and watches f_{ii} rise silently. It is the cheapest alarm for genetic erosion *inside* a genebank – and it is precisely what a distance tree cannot see.

2.2.3 s_i and F_i , self-coancestry and molecular inbreeding

$$s_i = \sum_k \sum_l \frac{p_{k,l,i}^2}{N_i}, \quad F_i = 2s_i - 1, \quad 1 - s_i = H_{o,i} \quad (2.2)$$

with $p_{k,l,i}$ the dosage of allele k in individual l .

! F_i here is molecular, and is not F_{IS}

It is measured against a hypothetical base in which all alleles are distinct. It is **not** comparable across studies using different marker panels. For the question “is this population mating at random?”, use α_i below.

In soybean, rice and wheat, pure lines have $s_i \rightarrow 1$, so $F_i \rightarrow 1$ and $H_o \rightarrow 0$. That is expected, not a symptom. It changes only *where* the diversity resides.

2.3 Two per-subpopulation metrics

$$\alpha_i = \frac{F_i - f_{ii}}{1 - f_{ii}}, \quad G_i = \frac{s_i - f_{ii}}{1 - f_{ii}} = \frac{1 + \alpha_i}{2} \quad (2.3)$$

α_i is the subpopulation’s analogue of F_{IS} : positive means an excess of homozygotes (selfing, mating among relatives, a Wahlund effect from internal substructure, null alleles); negative means an excess of heterozygotes (dioecy, self-incompatibility, selection, or a freshly formed F1 or synthetic).

G_i is the proportion of the accession’s diversity that sits **between** individuals rather than inside them, and it is probably the most underused quantity in the paper. It is a mating-system meter:

G_i	Situation	Typical material
~ 0.5	perfect panmixia	maize, sunflower, open populations
< 0.5	excess of heterozygotes	F1 synthetics, heterozygous clones, self-incompatibles
-> 1.0	all homozygous	inbred lines, doubled haploids, selfing species

Operationally, $1 - G_i$ is the fraction of an accession’s diversity stored *inside* each plant, and it decides the sampling design. With $G_i \rightarrow 1$ (selfers), one individual captures almost nothing and you need many plants per accession. With $G_i \approx 0.5$ (maize), a few plants already capture half.

2.4 Distances, and what they actually are

Nei’s minimum distance between subpopulations is

$$D_{ij} = \frac{f_{ii} + f_{jj}}{2} - f_{ij} = \sum_k \frac{1}{2} (\tilde{p}_{k,i} - \tilde{p}_{k,j})^2. \quad (2.4)$$

Note what that is: **half the sum of squared allele-frequency differences** – a variance component, not an evolutionary distance calibrated in time. It is the piece that enters F_{ST} , and it is legitimate. The error in a distance-only method is not using D_{ij} ; it is using *only* D_{ij} .

💡 A free $O(N^2)$ saving

The paper’s Eq. 13 computes D_{ij} by a double sum over all pairs of individuals, at cost $O(N_i N_j m)$. That is unnecessary: the identity $D_{ij} = (s_i + s_j)/2 - f_{ij}$ is **exact** and reduces the cost to $O(Nm)$. For a 50k-SNP panel and 5000 accessions this is the difference between hours and seconds.

2.5 The equation that matters most in practice

$$\bar{f} = \sum_i \frac{N_i}{N_T} \left[\underbrace{f_{ii}}_{\text{internal coancestry}} - \underbrace{\sum_j D_{ij} N_j / N_T}_{\text{mean distance to the others}} \right] \quad (2.5)$$

Each accession's contribution to global coancestry – that is, to the *loss* of diversity in the pool – has **two terms of opposite sign**:

- **Term 1 (+)**. The more internally fixed an accession is, the more it *raises* global coancestry: it dilutes the pool with copies of the same allele.
- **Term 2 (-)**. The more distant it is from the others, the more it *reduces* global coancestry.

A distance-only method sees term 2 alone. That is the whole mechanism of the inverted conclusions.

2.6 F statistics and the partition

$$F_{IS} = \frac{\tilde{F} - \tilde{f}}{1 - \tilde{f}}, \quad F_{ST} = \frac{\tilde{f} - \bar{f}}{1 - \bar{f}} = \frac{\bar{D}}{1 - \bar{f}}, \quad F_{IT} = \frac{\tilde{F} - \bar{f}}{1 - \bar{f}} \quad (2.6)$$

with $(1 - F_{IT}) = (1 - F_{IS})(1 - F_{ST})$.

The partition itself is the heart of the paper:

$$(1 - \bar{f}) = \underbrace{(1 - \tilde{s})}_{GD_{WI}} + \underbrace{(\tilde{s} - \tilde{f})}_{GD_{BI}} + \underbrace{(\tilde{f} - \bar{f})}_{GD_{BS}} \quad (2.7)$$

Component	Definition	Biologically
GD_T = 1 - f_bar	total diversity	He of the whole pool if everything were recombined
GD_WI = 1 - s_til	within individuals	heterozygosity carried inside each plant
GD_BI = s_til - f_til	between individuals, within accession	segregating variation available to within-population selection
GD_WS = 1 - f_til	within subpopulations	GD_WI + GD_BI
GD_BS = f_til - f_bar	between subpopulations	divergence; the only component a distance method sees

with the clean identities

$$\frac{GD_{WI}}{GD_T} = (1 - G)(1 - F_{ST}), \quad \frac{GD_{BI}}{GD_T} = G(1 - F_{ST}), \quad \frac{GD_{WS}}{GD_T} = 1 - F_{ST}, \quad \frac{GD_{BS}}{GD_T} = F_{ST} \quad (2.8)$$

2.6.1 Why the partition changes the strategy by species

- **Outcrossers** (maize, sunflower, forages, forest trees): GD_{WI} is large. The diversity is literally *inside* the plants. Conserving means maintaining heterozygosity – regenerate with many individuals and balanced contributions.
- **Selfers** (soybean, wheat, rice, beans) **and DH or line collections**: $GD_{WI} \approx 0$. All diversity is $GD_{BI} + GD_{BS}$. Conserving means maintaining the *number of distinct lines*; heterozygosity is not a target. Regenerating with few plants per line is acceptable; losing lines is not.
- **Clonal crops** (potato, sugarcane, fruit trees): GD_{WI} is high and **frozen** – it does not decay by drift while the clone exists. The N_e algebra does not apply without adaptation.

This is the reason a hybrid maize programme cannot manage diversity with a single pool-wide number: it needs GD_{WS} defended and GD_{BS} held, and those two pull in opposite directions under selection.

2.7 The worked example, reproduced

The paper's metapopulation: four subpopulations of $N = 12, 4, 8, 6$, three loci with 5, 3 and 3 alleles. Its published coancestry matrix:

```
f <- matrix(c(0.3391, 0.3021, 0.2535, 0.3229,
             0.3021, 0.8750, 0.3073, 0.3889,
             0.2535, 0.3073, 0.4766, 0.2483,
             0.3229, 0.3889, 0.2483, 0.3704), 4, 4, byrow = TRUE,
           dimnames = list(paste0("Sub", 1:4), paste0("Sub", 1:4)))
N <- c(12, 4, 8, 6)
w <- N / sum(N)
```

Note subpopulation 2: $f_{22} = 0.875$, so $H_e = 0.125$. It is nearly fixed internally. Now the global quantities, with the paper's published values beside them:

```
f_bar <- drop(w %*% f %*% w)
f_til <- sum(w * diag(f))
s_til <- 0.7056 # published; needs individual genotypes

data.frame(
  quantity = c("f_bar", "f_til", "GD_T", "GD_WI", "GD_BI", "GD_BS", "F_ST"),
  computed = round(c(f_bar, f_til, 1 - f_bar, 1 - s_til, s_til - f_til,
                    f_til - f_bar, (f_til - f_bar) / (1 - f_bar)), 4),
  published = c(0.3256, 0.4535, 0.6744, 0.2944, 0.2521, 0.1279, 0.1897)
)
```

quantity	computed	published
f_bar	0.3256	0.3256
f_til	0.4535	0.4535
GD_T	0.6744	0.6744
GD_WI	0.2944	0.2944
GD_BI	0.2521	0.2521
GD_BS	0.1279	0.1279
F_ST	0.1897	0.1897

Every value reproduces exactly. Note that GD_{BS} is only 19% of the total: the between-subpopulation divergence – the only thing a distance method measures – accounts for less than a fifth of the diversity in this metapopulation.

2.7.1 Removing an accession can *increase* diversity

```
res <- list(f = f, by_subpop = data.frame(N = N))
cbind(ct_loss_gain(res),
      delta_pct_equal = round(ct_loss_gain(res, "equal")$delta_pct, 2))
```

pop	GD_without	delta_pct	delta_pct_equal
Sub1	0.6296049	-6.6468853	-6.92
Sub2	0.6868734	1.8444502	6.47
Sub3	0.6480760	-3.9081294	-6.94
Sub4	0.6689694	-0.8102105	-3.56

The number to look at is subpopulation 2. A Weitzman-style analysis calls it the most valuable in the set – losing it would cost 65.5% of the diversity, because it is the most divergent. The coancestry analysis says removing it **raises** pool diversity by 1.8% (census weights) or 6.5% (equal weights).

It is distant *because* it is fixed. Alleles fixed in a sample of four individuals are not a genetic resource; they are the result of drift. The same pattern appears in the paper’s reanalysis of eleven European pig breeds: the breed a distance method nominates as most valuable turns out to raise diversity by 0.67% when removed.

What this does and does not license

For a genebank: a rare accession with few plants and high homozygosity – the archetypal “unique” accession on a distance tree – may be contributing *negatively* to the pool. That does **not** mean discard it. It means the correct priority may be to **regenerate it at larger N** , not to replicate it into a core collection. Storing seed is cheap; extinction is irreversible.

2.7.2 Optimal contributions

Maximising $GD_T = 1 - \sum_{i,j} f_{ij}c_i c_j$ subject to $c_i \geq 0$ and $\sum_i c_i = 1$ is a quadratic programme:

```
oc <- optimal_contributions(f)
oc$c_i
#>      Sub1      Sub2      Sub3      Sub4
#> 0.522398 0.000000 0.269890 0.207711
c(GD_before = 1 - f_bar, GD_after = oc$GD_max,
  gain_pct = 100 * (oc$GD_max - (1 - f_bar)) / (1 - f_bar))
#> GD_before GD_after gain_pct
#> 0.6744338 0.6873675 1.9177220
```

Subpopulation 2 receives weight **zero**, and total diversity rises from 0.6744 to 0.6874. This reproduces the paper’s published solution $\mathbf{c} = (0.522, 0, 0.270, 0.208)$ exactly.

Three direct uses in plant breeding:

- **Base population for recurrent selection.** c_i gives the seed proportion from each source that maximises long-term additive variance.
- **Core collection.** Instead of clustering heuristics plus proportional sampling, solve the quadratic programme directly with operational constraints ($c_i \leq$ multiplication capacity, $c_i \geq$ a policy minimum for accessions of cultural importance).
- **Optimal contribution selection.** Replace the objective with $\max \mathbf{c}'\mathbf{g} - \lambda \mathbf{c}'\mathbf{F}\mathbf{c}$ and you have Meuwissen’s framework, now standard in genomic programmes. This chapter is the case $\lambda \rightarrow \infty$, pure diversity. A breeder operates in the middle of the curve – which is Chapter 15..

2.8 Minimising coancestry is maximising N_e

Since $N_s = 1/(2\bar{f})$, minimising group coancestry and maximising effective size are the same operation. The paper’s Eq. 14 makes the operational consequence explicit. With **equal** contributions ($S_W^2 = S_B^2 = 0$), $N_e \approx 2Nn/(1 - F_{IT})$ – roughly **double** the N_e obtained under random (Poisson) contributions.

Harvesting the same number of seeds from every plant, and the same number of plants from every family, doubles the effective size. No genotyping, no technology, no cost.

It is probably the best cost-benefit intervention available in germplasm management. The opposite – bulk harvest from a plot, which in practice samples the most productive plants disproportionately – inflates S_W^2 and collapses N_e .

2.9 The honest caveats

The paper makes these, and a breeder should take them seriously.

Criteria based on within-population variation systematically **favour larger populations**. A $\bar{f}$ criterion can push a genebank toward abundant elite material – the opposite of the conservation objective. Beyond that:

- Neutral markers do not measure adaptive value, disease resistance, stress adaptation or cultural value. A fixed accession may carry the only resistance allele to an emerging rust.
- **Allelic richness** is more sensitive to bottlenecks than heterozygosity and is often the more relevant criterion for pre-breeding. The argument that minimising $\bar{f}$ maximises allelic richness in the long run (via the effective number of alleles being the inverse of mean coancestry) is valid asymptotically but does **not** protect rare alleles in the short term. This is why allelic richness survives the redundancy analysis in Chapter 12. as the one genuinely independent axis.
- Weighting between-group variation more heavily is arbitrary. A more defensible fix is to optimise c with per-accession minimum constraints.

Rule of thumb: use $\bar{f}$ to decide *proportions and regeneration effort*; never use $\bar{f}$ alone to decide *discard*.

2.10 Two errata, found while verifying

Worth recording, since the numbers above were checked against the published tables.

- Table 2b lists $F_2 = 0.8383$, but $F_i = 2s_i - 1 = 2(0.9167) - 1 = 0.8333$. A typographical error. Every other value in Tables 2 and 3 reproduces exactly.
- Conceptually, the Discussion states that equal contributions ($S_B^2 = S_W^2 = 0$) are required “to minimise effective size” and “maximise mean coancestry”. That is inverted: Eq. 15 and the entire argument of the paper show that equal contributions **maximise** N_e and **minimise** $\bar{f}$.

2.11 What manages this quantity, and where

This chapter is the book’s keystone, and it is easy to read it as a self-contained treatment of Caballero & Toro. The partition it derives is not background: it is the object every later chapter operates on. Stated once, so the rest of the book can refer back rather than re-introduce it:

Term	What it is	Where it is managed
GD_{WI}	diversity within individuals	exactly zero with inbred parents (Chapter 11.); the hybrid-level GD_{WI-hyb} is what carries signal
GD_{BI}	between individuals, within pool	the quantity α budgets (Chapter 15.), and the one recurrent selection erodes (Chapter 19.)
GD_{BS}	between pools	the heterosis engine, to be pre-served not minimised – derived as an exact identity in Chapter 10.

Three consequences follow, and each is a chapter.

A single pool-wide statistic cannot express “conserve within, hold between”. That is why Chapter 11. reports the partition rather than only GD_T , and why Chapter 13. makes S1 a separate monitoring stage instead of folding pool conservation into the hybrid-selection constraint.

The between term is not a bonus to be maximised. Section 10.7 shows that under neutral drift every unit of GD_{BS} gained is paid for out of the shared term, so pushing the pools apart buys nothing. The partition tells you where diversity is, not which direction to move it.

The within term is the one with a future. GD_{BI} is the raw material of next cycle’s response, which Chapter 3. makes precise: it is the same sum over the same loci as genic variance, weighted equally instead of by squared effect. That is the entire reason a diversity constraint is worth writing down, and Chapter 20. is where the equivalence is pushed until it breaks.

3. Response to selection: the other half of the objective

Every chapter so far measures what selection *spends*. None of them writes down what it *buys*.

That asymmetry is not cosmetic. Chapter 19. will show a constrained programme overtaking an unconstrained one after eight cycles, and the book can report that number because it simulated it – not because anything in the preceding chapters predicts it. A reader who wants to know whether the crossover arrives at cycle 8 in *their* programme has no way to find out except to re-run the simulation with their own parameters.

The quantitative genetics that answers such questions is older than any of the genomic machinery here, and this chapter installs it: the variance selection acts on, the equation that turns that variance into gain, and the classical index theory that turns out to be the direct ancestor of this book's α constraint.

```
sim <- simulate_pools(n_A = book_cfg$n_pool_A, n_B = book_cfg$n_pool_B,
                     m = book_cfg$m, fst = 0.15, n_qtl = 300, h2 = 0.5, seed = 1)
X <- sim$X # N x m hybrids, 0/0.5/1
gv <- drop(X %*% sim$beta) # true genetic value, known here by construction
I <- sim$index # what the breeder actually sees
c(hybrids = nrow(X), markers = ncol(X), qtl = sum(sim$beta != 0))
#> hybrids markers qtl
#> 625 2000 300
```

3.1 The breeder's equation, and how well it does here

Response to one round of truncation selection is

$$R = i r \sigma_A \quad (3.1)$$

with i the selection intensity (a property of the selected fraction alone), r the accuracy of the criterion as a predictor of genetic value, and σ_A the additive standard deviation. `sel_intensity()` has been in this package since Chapter 13.; it has never once been used to compute a response.

```
n_sel <- 60
N <- nrow(X)
i_norm <- sel_intensity(n_sel / N)
r <- cor(I, gv)
sel <- order(I, decreasing = TRUE)[seq_len(n_sel)]

fmt(data.frame(i = i_norm, r = r, sigma_A = sd(gv),
               predicted = i_norm * r * sd(gv),
               realised = mean(gv[sel]) - mean(gv)), 3)
```

i	r	sigma_A	predicted	realised
1.774	0.702	4.742	5.908	5.116

The equation overpredicts, by about 13 percent. That is worth taking apart rather than rounding away, because both halves of the discrepancy are instructive.

```
i_real <- (mean(I[sel]) - mean(I)) / sd(I)
b <- cov(gv, I) / var(I) # regression of value on criterion
fmt(data.frame(
  i_normal_theory = i_norm,
  i_realised = i_real,
```

```

R_with_realised_i = i_real * r * sd(gv),
R_from_regression = b * (mean(I[sel]) - mean(I)),
R_actual          = mean(gv[sel]) - mean(gv),
index_skew        = mean(scale(I)^3), 3)

```

i_normal_theory	i_realised	R_with_realised_i	R_from_regression	R_actual	index_skew
1.774	1.701	5.662	5.662	5.116	-0.17

Two separate failures, and they are different in kind.

The realised intensity is below normal theory. $i = \varphi(z)/q$ is exact for a normal criterion. This index is slightly left-skewed, so its upper tail is thinner than a normal's and the top 60 individuals sit closer to the mean than the formula assumes. This is a property of *this* population, not an error.

The regression is exact, and still wrong. `R_from_regression` and `R_with_realised_i` agree to machine precision – they are the same quantity written twice, since $b S_I = i r \sigma_A$ identically. Both still exceed the realised response. What fails is not the algebra but the assumption underneath it: $R = b S_I$ requires $E[g | I]$ to be linear in I in the selected tail, which is a bivariate-normality claim, and bivariate normality is exactly what a two-pool hybrid population does not have.

! Read the breeder's equation as an upper bound here

Every response number in this book's cycle simulations is a *realised* response, measured from the simulation. Where a predicted response is quoted, it will run above the realised one for the reason above. The gap is not noise and does not shrink with more markers; it shrinks when the criterion becomes more normal and the population less structured.

3.2 Genic variance, and what it has to do with gene diversity

The additive variance in the equation above is not a fixed property of the population. Split it in two.

Genic variance is what the loci would contribute if they were independent:

$$V_{\text{genic}} = \sum_k \beta_k^2 \text{Var}(x_k) \quad (3.2)$$

Realised variance is what the population actually has, $\text{Var}(\mathbf{X}\beta)$, which is the sum of the *whole* covariance matrix rather than only its diagonal. The difference is exactly the gametic-phase disequilibrium:

$$\text{Var}(\mathbf{X}\beta) - V_{\text{genic}} = 2 \sum_{k < l} \beta_k \beta_l \text{Cov}(x_k, x_l) \quad (3.3)$$

```

c(realised = var(gv),
  genic     = genic_var(X, sim$beta),
  gap       = var(gv) - genic_var(X, sim$beta),
  oracle_deviation = abs(genic_var(X, sim$beta) - ref_genic_var(X, sim$beta)))
#>      realised      genic      gap oracle_deviation
#> 2.248375e+01 2.520510e+01 -2.721357e+00 1.065814e-14

```

Now the sentence that ties this chapter to every one that follows.

V_{genic} is $\sum_k \beta_k^2 p_k q_k$ up to a coding constant, and gene diversity is $\frac{1}{m} \sum_k 2p_k q_k$. **They are the same sum over the same loci, weighted differently** – one by squared effect, one by nothing at all. That is the precise sense in which managing gene diversity manages the raw material of response, and it is why a diversity constraint is a defensible proxy for a variance constraint in the first place.

It is also the precise sense in which the proxy is incomplete. Gene diversity weights every locus equally, so it cannot distinguish a marker carrying half the genetic variance from one carrying none. And – more seriously –

⚠ No metric in this book can see the disequilibrium term

θ , GD , N_s , allelic richness, F_{hom} and F_{drift} are all functions of allele frequencies **one locus at a time**. The covariance term above does not move a single allele frequency. A flat gene diversity trajectory therefore does not imply flat additive variance, and Chapter 20. measures how far apart the two can travel.

3.3 The decomposition this book already had

Additive variance is the variance of *breeding values*, not of genotypic values, and in a two-pool hybrid population the average effect of an allele substitution depends on which pool it is substituted into. That sounds like it needs new machinery. It does not: Section 10.10 already built the exact orthogonal split, and general combining ability is the additive part.

```
set.seed(7)
d <- rnorm(ncol(X), 0.3, 0.1) # dominance deviations
U <- sim$GL[sim$pool == "A", ]
V <- sim$GL[sim$pool == "B", ]
gs <- heterosis_gca_sca(U, V, d)
H <- heterosis_value(X, d)

nA <- nrow(U); nB <- nrow(V)
gA <- rep(gs$gca_A, times = nB) # ped varies pool A fastest
gB <- rep(gs$gca_B, each = nA)
sc <- as.vector(gs$sca)

c(reconstruction_dev = max(abs(H - (gs$mu + gA + gB + sc))),
  var_H = var(H),
  sum_of_components = var(gA) + var(gB) + var(sc),
  identity_dev = abs(var(H) - (var(gA) + var(gB) + var(sc))))
#> reconstruction_dev      var_H      sum_of_components      identity_dev
#>      5.115908e-13      3.811311e+01      3.811311e+01      9.237056e-14
```

The three components are orthogonal, so their variances add exactly. Read the shares:

```
p <- c(gca_A = var(gA), gca_B = var(gB), sca = var(sc))
fmt(data.frame(component = names(p), variance = p,
  share = p / sum(p)), 3)
```

	component	variance	share
gca_A	gca_A	9.064	0.238
gca_B	gca_B	8.673	0.228
sca	sca	20.376	0.535

gca_A and gca_B are the part of heterosis an additive model can express and that therefore behaves like additive variance under selection. sca is the part it cannot, and it does not accumulate across cycles the way the general components do. A programme selecting on an additive index is selecting on the first two and getting the third by luck.

3.4 What weights = NULL is assuming

`stage2_select()` defaults to equal weight on every trait, and every chapter of this book including Chapter 21. takes that default. It is a choice, not a neutral one.

The classical answer is the Smith-Hazel index (Smith 1936; Hazel 1943): the weights maximising the correlation between the index and the aggregate genotype $H = \mathbf{a}'\mathbf{g}$ are $\mathbf{b} = \mathbf{P}^{-1}\mathbf{Ga}$.

```

P <- matrix(c(4.0, 1.2, 1.2, 1.0), 2)      # phenotypic covariance
G <- matrix(c(2.0, 0.9, 0.9, 0.3), 2)      # genetic covariance
a <- c(1, 1)                               # equal economic weight

fmt(data.frame(trait = c("yield", "moisture"),
               economic_weight = a,
               smith_hazel     = index_smith_hazel(P, G, a)), 3)

```

trait	economic_weight	smith_hazel
yield	1	0.570
moisture	1	0.516

Equal economic weights do **not** give equal index weights. They coincide only when $\mathbf{P}^{-1}\mathbf{G}$ is proportional to the identity, which requires the traits to be equally variable, equally heritable and mutually uncorrelated. Nothing guarantees that, and where it fails the index being maximised is not the index anyone intended.

3.5 The α constraint is a restricted index

Chapter 15. will argue that a diversity *budget* is approvable where a weighted sum is not: “I accept losing five percent” is a sentence a breeding committee can sign, and “diversity is worth λ index units” is not.

That argument is usually presented as a point about human decision-making. It is older and harder than that. Kempthorne & Nordskog (1959) posed exactly the same problem – maximise aggregate gain subject to **zero expected change** in a specified direction – and solved it:

$$\mathbf{b} = \mathbf{P}^{-1}[\mathbf{G}\mathbf{a} - \mathbf{C}(\mathbf{C}'\mathbf{P}^{-1}\mathbf{C})^{-1}\mathbf{C}'\mathbf{P}^{-1}\mathbf{G}\mathbf{a}], \quad \mathbf{C} = \mathbf{G}\mathbf{R} \quad (3.4)$$

```

R <- cbind(c(0, 1))                               # hold moisture at zero gain
b_r <- index_restricted(P, G, a, R)
b_u <- index_smith_hazel(P, G, a)

fmt(data.frame(
  index = c("Smith-Hazel", "restricted"),
  w_yield = c(b_u[1], b_r[1]),
  w_moisture = c(b_u[2], b_r[2]),
  gain_in_restricted_trait = c(drop(crossprod(R, G %*% b_u)),
                                drop(crossprod(R, G %*% b_r))), 4)

```

index	w_yield	w_moisture	gain_in_restricted_trait
Smith-Hazel	0.5703	0.5156	0.668
restricted	-0.1207	0.3621	0.000

The restricted index moves the restricted trait by exactly zero – not approximately, not on average. That is the property distinguishing a restriction from a penalty, and it is the property α inherits.

i What is genuinely new in this book, then

Not the idea of constraining rather than pricing; that is 1959. What is new is the restricted direction. Kempthorne and Nordskog restricted a *trait*. This book restricts the loss of molecular gene diversity, which is not a trait, has no economic weight, and is measured on a matrix rather than a phenotype. The algebra of restriction carries over; the object being restricted does not.

Stating the lineage makes Chapter 15. stronger, not weaker. A sixty-year-old theorem stands behind the move.

3.6 Where this connects

Chapter 9. and Chapter 11. measure the frequencies whose weighted sum is the genic variance defined here, which is why a diversity constraint is a variance constraint in disguise. Chapter 14. measures the r this chapter took as given, and Finding 5 of Chapter 12. shows what happens when it is assumed rather than measured. Chapter 15. is where the restricted index becomes the α constraint, and Chapter 20. is where the disequilibrium term this chapter isolated turns out to matter more than the equation suggests.

II

4.4	Predicted traits	40
4.5	The two matrices, both built	40
4.6	The stage-1 contract in one call	41
4.7	What the candidate pool looks like	42
5	Sizing a segregating population, F1 to F4	43
5.1	The simulated space is ancestry, no alleles	43
5.2	The assumptions, one at a time	43
5.3	Verifying the identities live	44
5.4	The results	45
5.5	Sizing the F1	49
5.6	Working numbers	50
5.7	Where this connects	50
6	Introgressing several genes by backcrossing	53
6.1	The operation, in one diagram	53
6.2	The model	54
6.3	Three criteria, and only one of them binds	54
6.4	Which selection index	55
6.5	How long to keep selecting recombinants	57
6.6	Allocating a fixed budget across BC1, BC2 and BC3	57
6.7	Fixing the stack: selfing against doubled haploids	58
6.8	Running the simulation	59
6.9	Assumptions, and what is not modelled	60
6.10	Where this connects	61
7	Simulating a breeding programme: AlphaSimR from the ground up	63
7.1	Step 1 – Founder genomes, and which engine to use	63
7.2	Step 2 – SimParam, and defining a trait	65
7.3	Step 3 – The base population, and where the variance components live	66
7.4	Step 4 – The DNA lottery, and why an F1 is not a family	66
7.5	Step 5 – Selfing, doubled haploids, and what the shortcut buys	67
7.6	Step 6 – Phenotyping, selection, and checking the breeder’s equation	68
7.7	Step 7 – Two traits, and what equal weighting assumes	69
7.8	Step 8 – The parts assembled: reciprocal recurrent selection	70
7.9	Assumptions, and what this adds over Chapter 4.	73
7.10	Where this connects	74
8	A complete plant breeding programme, and what it costs in diversity	75
8.1	The scheme	75
8.2	Why a breeding programme needs a burn-in	76
8.3	One programme, fifteen years	76
8.4	Genomic selection: two levers, not one	78
8.5	The comparison, replicated	79
8.6	The diversity ledger	80
8.7	Assumptions	80
8.8	Where this connects	81

4. Simulating lines, pools and crosses

Everything downstream of this chapter consumes three objects and knows nothing about where they came from. That is deliberate: it is the seam at which a simulation is swapped for real data, and it is the only place that has to change when it is.

This chapter builds those objects. It also builds the candidate pool the rest of the book selects from, so the choices made here – how divergent the pools are, how the index is constructed – determine what the later chapters have to work with.

4.1 What has to be produced

Object	Shape	Content	Needed for
<code>X</code>	$N \times m$	hybrid marker matrix, values in 0/0.5/1	every frequency-based metric
<code>f</code>	$N \times N$	molecular coancestry of the hybrids	every coancestry-based metric
<code>hybrids</code>	$N \times (3 + n_traits)$	hybrid id, the two parent lines, predicted traits	the index, and the line-usage vector
<code>fL</code>	$L \times L$, optional	coancestry BETWEEN LINES	only the heterotic-group decomposition

4.2 Simulating the lines

Two heterotic pools are generated with a Balding-Nichols model. Ancestral allele frequencies are drawn uniform on $[0.05, 0.95]$, then each pool's frequency at marker k is drawn as

$$p_{A,k} \sim \text{Beta}\left(p_k \frac{1 - F_{ST}}{F_{ST}}, (1 - p_k) \frac{1 - F_{ST}}{F_{ST}}\right), \quad (4.1)$$

and independently for pool B. The Beta has mean p_k and variance $F_{ST} p_k (1 - p_k)$, so the parameter `fst` in the configuration is the expected between-pool divergence – it is not a proxy for it.

Genotypes are then drawn as $2 \times \text{Bernoulli}(p)$, which makes every line fully homozygous, as an inbred line should be.

```
simulate_lines
#> function (cfg)
#> {
#>   set.seed(cfg$seed)
#>   p_anc <- stats::runif(cfg$m, 0.05, 0.95)
#>   bn <- function(p, fst) {
#>     stats::rbeta(length(p), p * (1 - fst)/fst, (1 - p) *
#>       (1 - fst)/fst)
#>   }
#>   pA <- bn(p_anc, cfg$fst)
#>   pB <- bn(p_anc, cfg$fst)
#>   draw <- function(n, p) {
#>     matrix(2L * stats::rbinom(n * length(p), 1, rep(p, each = n)),
#>       nrow = n)
#>   }
#>   L <- rbind(draw(cfg$n_pool_A, pA), draw(cfg$n_pool_B, pB))
#>   storage.mode(L) <- "double"
```

```
#>      L
#> }
#> <bytecode: 0xbcc919508>
#> <environment: namespace:hybdiv>
```

```
L <- simulate_lines(book_cfg)
dim(L)
#> [1]   50 2000
table(as.vector(L[1:5, 1:200]))
#>
#>    0    2
#> 480 520
```

Only 0 and 2 appear: no residual heterozygotes. With real data those must be imputed or dropped before this point, because the hybrid-genotype identity in Chapter 1. assumes homozygous parents.

4.2.1 The divergence knob

fst is the single most consequential parameter in the simulation, because it sets how much of the pools' diversity sits *between* them rather than within. What that buys is less obvious than it looks:

```
fst_grid <- c(0.02, 0.05, 0.10, 0.15, 0.25, 0.40)
div <- vapply(fst_grid, function(fst) {
  cfg <- modifyList(book_cfg, list(fst = fst, m = 1500))
  Lx <- simulate_lines(cfg)
  fL <- molecular_coancestry(Lx / 2)
  A <- seq_len(cfg$n_pool_A); B <- cfg$n_pool_A + seq_len(cfg$n_pool_B)
  # a hybrid A x B is heterozygous wherever its parents differ
  c(between_pool_f = mean(fL[A, B]),
    within_pool_f = mean(c(fL[A, A], fL[B, B])),
    het_fraction = 1 - mean(fL[A, B]))
}, numeric(3))

d <- data.frame(fst = fst_grid, t(div))
plot(d$fst, d$het_fraction, type = "b", pch = 19, ylim = c(0, 1),
     xlab = expression(F[ST]~"between pools"),
     ylab = "expected heterozygous fraction of the F1")
grid()
```




Figure 4.1: Driving the pools apart leaves the expected heterozygous fraction of the F1 flat. What moves is the within-pool coancestry, which climbs throughout. Section 10.7 derives why the two terms must trade one for one.

```
fmt(d)
```

fst	between_pool_f	within_pool_f	het_fraction
0.02	0.6349	0.6564	0.3651
0.05	0.6362	0.6684	0.3638
0.10	0.6364	0.6877	0.3636
0.15	0.6326	0.7052	0.3674
0.25	0.6370	0.7387	0.3630
0.40	0.6274	0.7898	0.3726

The heterozygous fraction does not move. It is flat across a twentyfold range of `fst`, while `within_pool_f` climbs steadily: the divergence is being paid for out of within-pool diversity, unit for unit.

That is not an artefact of this grid. The expected heterozygous fraction of an A x B hybrid is exactly $1 - \theta_{AB}$, and Section 10.7 shows that under neutral drift every unit gained by the between-pool term is subtracted from the shared one, leaving the total fixed. Neutral divergence relocates diversity; it does not create heterosis. Chapter 2. is why `GD_BS` is nonetheless held rather than minimised, and Section 10.8 is what has to be true of the divergence for it to be worth anything.

4.3 From lines to hybrids

Every A x B combination is formed. With 25 lines per pool that is 625 candidates; at the production scale of 71 per pool it is 5041.

```
cfg <- book_cfg
a_ids <- seq_len(cfg$n_pool_A)
b_ids <- cfg$n_pool_A + seq_len(cfg$n_pool_B)
```

```
ped <- expand.grid(a = a_ids, b = b_ids)
M <- (L[ped$a, , drop = FALSE] + L[ped$b, , drop = FALSE]) / 2 # 0/1/2
dim(M)
#> [1] 625 2000
head(ped, 3)
```

a	b
1	26
2	26
3	26

The factorial structure matters more than it looks. Because every line appears in many candidate hybrids, selecting hybrids is really selecting *line usage* – and that observation, made precise in Section 12.1, is what makes the whole problem tractable at production scale.

4.4 Predicted traits

Marker effects are drawn correlated across traits, through a Cholesky factor of a fixed correlation matrix, and the predicted values are the centered marker matrix times those effects:

```
set.seed(cfg$seed + 99)
R <- matrix(c(1, 0.3, -0.3,
              0.3, 1, 0.1,
              -0.3, 0.1, 1), 3, 3)
vr <- vanraden_G(M)
B <- matrix(rnorm(cfg$m * cfg$n_traits), cfg$m) %%% chol(R)
traits <- scale(vr$Z %%% B)
round(cor(traits), 3)
#>      [,1] [,2] [,3]
#> [1,] 1.000 0.259 -0.440
#> [2,] 0.259 1.000 0.049
#> [3,] -0.440 0.049 1.000
```

The realised trait correlations recover the target R , which is the point of the Cholesky step: without it, a multi-trait index would be a weighted sum of independent noise, and the trade-off it induces against diversity would be unrealistically easy.

i What this simulation deliberately does not model

These predicted values are constructed from *known* marker effects, so they carry no prediction error. Real GEBVs do. The book therefore treats `ctx$index` as given throughout, and the reader should read every gain figure as an upper bound on what a real programme would realise. The additive model also carries no dominance, so it cannot express specific combining ability – see the caveats in Chapter 21..

4.5 The two matrices, both built

```
X <- M / 2
ctx <- build_ctx(X = X, f = molecular_coancestry(X), G = vr$G, traits = traits,
                ped = ped, n_lines = nrow(L),
                pool = rep(c("A", "B"), c(cfg$n_pool_A, cfg$n_pool_B)),
                fL = molecular_coancestry(L / 2), cfg = cfg)

c(N = ctx$N, m = ctx$m, n_lines = ctx$n_lines,
  sum_G = sum(ctx$G), min_f = min(ctx$f), max_f = max(ctx$f))
#>      N      m      n_lines      sum_G      min_f
#> 6.250000e+02 2.000000e+03 5.000000e+01 -4.178879e-11 6.447500e-01
```

```
#>      max_f
#> 8.317500e-01
```

`sum(G)` is zero to machine precision and `f` is inside $[0, 1]$: the two properties from Chapter 1., on this dataset.

The context also freezes two things that must never be recomputed per scenario:

```
str(ctx[c("p0", "maf_pop", "index")], max.level = 1)
#> List of 3
#> $ p0      : num [1:2000] 0.44 0.3 0.56 0.9 0.12 0.98 0.98 0.7 0.64 0.22 ...
#> $ maf_pop: num [1:2000] 0.44 0.3 0.44 0.1 0.12 0.02 0.02 0.3 0.36 0.22 ...
#> $ index   : num [1:625] -0.291 0.498 -0.361 -0.186 -0.348 ...
```

`p0` is the candidate-pool allele frequency vector. It is the **base** against which `F_hom` and `F_drift` are measured in Chapter 9., and recomputing it per scenario would silently zero out the historical memory of erosion – the single most common error in applying these metrics.

4.6 The stage-1 contract in one call

All of the above is `stage1_simulate()`:

```
stage1_build
#> function (X, ped, traits, fL = NULL)
#> {
#>   if (max(X) > 1)
#>     X <- X/2
#>   stopifnot(nrow(X) == nrow(ped), nrow(X) == nrow(traits))
#>   hybrids <- data.frame(hybrid = seq_len(nrow(X)), line_A = as.integer(ped$a),
#>     line_B = as.integer(ped$b), traits)
#>   list(X = X, f = molecular_coancestry(X), hybrids = hybrids,
#>     fL = fL)
#> }
#> <bytecode: 0xbca44ab68>
#> <environment: namespace:hybdiv>
```

```
st1 <- stage1_simulate(book_cfg)
str(st1, max.level = 1)
#> List of 4
#> $ X      : num [1:625, 1:2000] 0.5 0.5 0 0 0 0.5 0.5 0 0 0 ...
#> $ f      : num [1:625, 1:625] 0.818 0.74 0.743 0.738 0.745 ...
#> $ hybrids:'data.frame': 625 obs. of  6 variables:
#> $ fL      : num [1:50, 1:50] 1 0.677 0.692 0.674 0.684 ...
head(st1$hybrids, 3)
```

hybrid	line_A	line_B	trait1	trait2	trait3
1	1	26	-0.6880637	0.1907069	0.0153955
2	2	26	-0.6102738	-0.0320787	1.4661012
3	3	26	-0.3094099	-0.5013539	0.2137806

For real data the same function is called with your own matrices:

```
st1 <- stage1_build(
  X      = hybrid_markers, # N x m, coded 0/1/2 or 0/0.5/1 (auto-detected)
  ped    = pedigree,      # data.frame with columns a, b
  traits = predicted_traits, # N x n_traits
  fL     = line_coancestry # optional
)
```

`stage1_build()` detects the marker coding and rescales, so both conventions give an identical result:

```
tr <- stl$hybrids[, grep("^trait", names(stl$hybrids))]
pd <- data.frame(a = stl$hybrids$line_A, b = stl$hybrids$line_B)
identical(stage1_build(stl$X, pd, tr)$f,
          stage1_build(stl$X * 2, pd, tr)$f)
#> [1] TRUE
```

4.7 What the candidate pool looks like

```
op <- par(mfrow = c(1, 2), mar = c(4.2, 4.2, 2.5, 1))
hist(ctx$index, breaks = 40, col = "grey85", border = "white",
     main = "multi-trait index", xlab = "index (s.d. units)")
gd_random <- replicate(400, gene_diversity(sample.int(ctx$N, 60), ctx))
hist(gd_random, breaks = 30, col = "grey85", border = "white",
     main = "GD of random subsets of 60", xlab = "gene diversity")
abline(v = gene_diversity(seq_len(ctx$N), ctx), col = "#b2182b", lwd = 2)
```




Figure 4.2: The candidate pool. Left: the index the selection maximises. Right: gene diversity of random subsets of 60, the baseline every later chapter measures against.

```
par(op)
```

The red line is the gene diversity of the *whole* pool. It does not sit at the centre of the distribution of random subsets of 60 – it sits above it. That gap is pure sampling, it has nothing to do with selection, and mistaking it for a loss is the subject of Section 11.9.

```
c(GD_population = gene_diversity(seq_len(ctx$N), ctx),
  GD_random_60 = mean(gd_random),
  apparent_loss_pct = 100 * (gene_diversity(seq_len(ctx$N), ctx) - mean(gd_random)) /
                           gene_diversity(seq_len(ctx$N), ctx))
#>      GD_population      GD_random_60 apparent_loss_pct
#>      0.3280928      0.3259023      0.6676555
```

5. Sizing a segregating population, F1 to F4

The previous chapters took the candidate pool as given. This one asks where it comes from: a wide cross between two divergent parents, carried F1 → F2 → F3 → F4, and the question every programme has to answer before it plants anything – **how many plants?**

The answer turns out to depend entirely on what you are sizing *for*. There is no single optimal N , and the most common failure is to size for the objective that saturates earliest and then be surprised when a different objective was never met.

5.1 The simulated space is ancestry, not alleles

Each genome position is coded 0 (from parent P1) or 1 (from parent P2). For a wide biparental cross between two divergent heterotic groups this is the correct abstraction: essentially every locus polymorphic between the parents is biallelic, so ancestry is a sufficient statistic. Throughout this chapter, “allele frequency” means “frequency of P1 ancestry”.

5.2 The assumptions, one at a time

These are the assumptions the numbers rest on. Read them before trusting any of the numbers.

P1. Ancestry, not alleles. As above.

P2. Parents fully homozygous and fixed for alternative alleles. Consequently every F1 plant is genetically identical and the genetic N_e of the F1 is 1. Sizing the F1 is a **logistical** decision, not a genetic one – unless the parents carry residual heterozygosity, which is handled separately below.

P3. The genetic map. Ten chromosomes, lengths from a classical composite maize map, about 1670 cM (~17 Morgans). Do **not** use IBM-type map lengths from intermated RILs: they are inflated 5-10x by construction and will make every answer wrong.

```
map <- build_map(maize_chr_len, spacing_cM = 2)
c(chromosomes = length(map$chr_len), markers = map$m,
  total_cM = map$total_cM, morgans = map$total_cM / 100)
#> chromosomes    markers    total_cM    morgans
#>         10.0         846.0       1670.0         16.7
```

P4. Haldane recombination, no interference, identical in male and female. The probability of a haplotype switch between adjacent markers d cM apart is $r = 0.5(1 - e^{-2d/100})$; the first marker of each chromosome gets $r = 0.5$, which produces independent segregation between chromosomes automatically.

Limitation: without interference the variance in crossover number is inflated (Poisson). Maize has strong positive interference, so the simulated block-length distribution has a slightly longer tail than reality. Mean block lengths and allele frequencies are unaffected.

P5. Segregation distortion by rejection sampling, therefore exact rather than approximate. Two mechanisms: *gametophytic* (ga1/Ga1-s type, pollen incompatibility) acting only on the male gamete, and *zygotic* (hybrid lethality or subvitality) as genotype fitnesses. Because rejection acts on the **whole gamete**, the region linked to the distorter is dragged along – which is exactly the phenomenon of interest.

f2_distorters

chr	pos_cM	type	k	w00	w01	w11	label
4	75	gametic	0.80	NA	NA	NA	ga1-like (4.05)
1	160	zygotic	NA	1	1	0.45	partial lethal (1.10)
5	60	gametic	0.68	NA	NA	NA	weak gametic (5.03)

⚠ The distorters are plausible, not measured

This is the parameter that moves the results most. Replace `f2_distorters` with your own F2 genotyping data before using any number here to decide field area.

P6. Five schemes, all ending with $H = 0.125$ so they are directly comparable: `ssd` (one seed per plant per generation – the neutral reference, no selection, no drift beyond the meiotic), `ssd_attr` (SSD with 15% random loss per generation: vigour, partial sterility, flowering asynchrony – common in wide crosses with exotic material), `bulk` (differential seed contribution, a Wright-Fisher model with weights $w_i = e^{\beta(z_i - \bar{z})}$ where z_i is the share of genome from the adapted parent – a *phenomenological* model of polygenic natural selection, not a model of adaptation loci), and `syn1/syn2` (one or two random intercrossing cycles in the F2 before the two selfings).

P7. The F4 panel is not a RIL panel. Residual heterozygosity is $H = 0.125$, and that changes every theoretical expectation:

Quantity	Formula	Value in F4
Var(dosage per line)	$(1 - H) / 4$	0.21875
SE(p_hat)	$0.5 \sqrt{(1 - H) / N}$	$0.468 / \sqrt{N}$
Diversity retained D	$1 - (1 - H) / N$	$1 - 0.875 / N$
P(P2 homozygote per locus)	$(1 - H) / 2$	0.4375
Junctions per haplotype	$A_2 = L; A_{t+1} = A_t + L H_t$	$1.75 L$

P8. The standard error is estimated between replicates, not between markers. This is essential: linked markers are strongly autocorrelated, so the spread across markers within one replicate is not an estimator of sampling error.

P9. What is not modelled, and therefore cannot be answered here: deliberate artificial selection, QTL architecture, phenotypic inbreeding depression, mutation (irrelevant on this timescale), and structural variation. If you know there is a large inversion between your parents, zero the recombination rate over that interval in `build_map()`.

Known bias: at 2 cM spacing, double crossovers within an interval are not counted. The bias is -2.0% in junction number, always downward. Use `spacing_cM = 1` if you need a faithful count.

5.3 Verifying the identities live

The theory above is not decoration – it is checkable, and the simulator is only trustworthy if it reproduces it.

```
map10 <- build_map(maize_chr_len, spacing_cM = 10)
neutral <- make_distortion(map10, NULL, active = FALSE)

set.seed(4)
F1 <- make_F1(map10, 400)
F2 <- cross_indices(F1, 1:400, 1:400, map10, neutral)
F3 <- cross_indices(F2, 1:400, 1:400, map10, neutral)
F4 <- cross_indices(F3, 1:400, 1:400, map10, neutral)

data.frame(
  generation = c("F1", "F2", "F3", "F4"),
  H_observed = round(c(panel_het(F1), panel_het(F2), panel_het(F3), panel_het(F4)), 4),
  H_expected = c(1, 0.5, 0.25, 0.125)
)
```

generation	H_observed	H_expected
F1	1.0000	1.000
F2	0.4968	0.500

generation	H_observed	H_expected
F3	0.2541	0.250
F4	0.1268	0.125

Heterozygosity halves at every selfing, as it must. And the diversity identity:

```
set.seed(5)
check <- do.call(rbind, lapply(c(20, 50, 100, 200), function(N) {
  pop <- run_scheme("ssd", N, map10, neutral, f2_opt)
  H <- panel_het(pop)
  data.frame(N = N,
             D_observed = diversity_retained(panel_freq(pop)),
             D_expected = 1 - (1 - H) / N)
}))
fmt(check, 4)
```

N	D_observed	D_expected
20	0.9582	0.9569
50	0.9845	0.9825
100	0.9904	0.9912
200	0.9954	0.9956

Map expansion is an algebraic prediction, and the payoff is a result many programmes get wrong:

```
Lm <- map$total_cM / 100
rbind(SSD = expected_junctions("ssd", map),
      Syn1 = expected_junctions("syn1", map),
      Syn2 = expected_junctions("syn2", map))
#>      junctions mean_block_cM
#> SSD      29.225      42.57489
#> Syn1     37.575      35.10247
#> Syn2     45.925      29.86142
```

Each Syn cycle adds only $0.5L$ junctions, not $1.0L$. The extra meiosis creates a junction only where the parent is heterozygous, and $H = 0.5$ at the F2/Syn stage. Budgeting an intercrossing cycle as “one more map length of recombination” overstates what it buys by a factor of two.

5.4 The results

The full grid – five schemes, 20 population sizes, adaptive replication – takes several minutes and is shipped precomputed.

N	D	SE(p_hat)	skewed genome	min(p, 1-p)	N equivalent
10	0.903	0.148	42.8%	0.069	9
30	0.961	0.084	14.4%	0.187	23
50	0.972	0.066	7.9%	0.209	32
70	0.978	0.056	5.6%	0.221	40
90	0.981	0.047	4.3%	0.233	47
110	0.982	0.044	4.2%	0.232	48
150	0.984	0.039	3.8%	0.236	56
190	0.985	0.032	3.7%	0.227	60
290	0.988	0.027	3.3%	0.240	70

N	D	SE(p_hat)	skewed genome	min(p, 1-p)	N equivalent
390	0.989	0.023	3.1%	0.242	77

5.4.1 The diversity curve saturates at 0.988, not at 1

```
schemes <- unique(f2$scheme)
cols <- c(ssd = "#1b6ca8", ssd_attr = "#7fb3d5", bulk = "#c0392b",
         syn1 = "#27ae60", syn2 = "#16a085")
plot(NA, xlim = range(f2$N), ylim = c(0.88, 1.0),
     xlab = expression(N[F2]~"plants"), ylab = "diversity retained, D")
grid()
for (s in schemes) {
  d <- f2[f2$scheme == s, ]
  lines(d$N, d$ddiv, type = "b", pch = 19, cex = 0.6, col = cols[s], lwd = 1.8)
}
curve(1 - 0.875 / x, add = TRUE, lty = 2, col = "grey30", lwd = 1.5)
legend("bottomright", c(schemes, "theory, no distortion"),
      col = c(cols[schemes], "grey30"), lwd = c(rep(1.8, length(schemes)), 1.5),
      lty = c(rep(1, length(schemes)), 2), pch = c(rep(19, length(schemes)), NA),
      bty = "n", cex = 0.8)
```




Figure 5.1: Diversity retained against population size, five schemes. The theoretical curve is $1 - (1 - H)/N$ under no distortion.

The gap between the observed curve and the theoretical one is **segregation distortion** – a deterministic bias that does not shrink with more plants. Adding plants buys you drift protection; it does not buy you distortion protection.

5.4.2 The bulk is the worst scheme across the whole range

comparison	N	D
bulk, largest N	390	0.954
SSD, N = 30	30	0.961

Bulk at the largest population size retains **less** diversity than SSD at 30 plants. With a less adapted donor parent, natural selection in the bulk removes precisely the germplasm the wide cross was made to introduce. If the objective is introgression, the bulk is not a cheaper SSD – it is a different experiment.

5.4.3 Block size does not depend on N

```
map4 <- build_map(maize_chr_len, spacing_cM = 4)
dist <- make_distortion(map4, f2_distorters)
set.seed(6)
small <- panel_freq(run_scheme("ssd", 20, map4, dist, f2_opt))
large <- panel_freq(run_scheme("ssd", 300, map4, dist, f2_opt))

plot(small, type = "l", col = "#e08214", ylim = c(0, 1),
     xlab = "genome (chromosomes 1-10)", ylab = "P1 ancestry frequency", xaxt = "n")
lines(large, col = "#1b6ca8", lwd = 1.6)
abline(h = 0.5, lty = 2, col = "grey50")
# chromosome boundaries and distorter positions
bnd <- which(map4$first)
axis(1, at = (c(bnd, map4$m + 1)[-1] + bnd) / 2, labels = seq_along(map4$chr_len))
abline(v = bnd[-1] - 0.5, col = "grey85")
for (i in seq_len(nrow(f2_distorters)))
  abline(v = marker_index(map4, f2_distorters$chr[i], f2_distorters$pos_cM[i]),
        col = "#b2182b", lty = 3)
legend("topleft", c("N = 20", "N = 300", "distorter"), col = c("#e08214", "#1b6ca8",
"#b2182b"),
      lty = c(1, 1, 3), lwd = c(1, 1.6, 1), bty = "n", cex = 0.8, horiz = TRUE)
```




Figure 5.2: Ancestry frequency along the ten chromosomes, at $N = 20$ and $N = 300$. At $N = 20$ the whole curve is noisy: that is drift, and more plants remove it. At $N = 300$ the noise is gone and three peaks remain, exactly at the distorters. Those are the deterministic bias, and no number of plants removes them.

Block length depends only on the number of meioses. Against linkage drag the instrument is an **inter-crossing cycle, not field area** – which is what the Syn curves in the theory table above quantify.

5.4.4 The metrics saturate at different N

This is the central practical result, and the reason the question “what is the optimal N ?” has no single answer.

Objective	Saturates near
Diversity retained (D)	100-120
Worst region of the genome, $\min(p, 1-p)$	60-80
Precision of the panel's p_{hat}	never – falls as $1/\sqrt{N}$
Recovery of 5 loci	200-260
Recovery of 8-10 loci	not saturated at the largest N tested

```

kk <- sort(unique(ml$k))
plot(NA, xlim = range(ml$N), ylim = c(0, 1), xlab = expression(N[F2]),
     ylab = "P(at least one line)")
grid()
for (i in seq_along(kk)) {
  d <- ml[ml$k == kk[i], ]
  lines(d$N, d$P, type = "b", pch = 19, cex = 0.6, col = i + 1, lwd = 1.8)
}
legend("topright", paste("k =", kk), col = seq_along(kk) + 1, lwd = 1.8,
      pch = 19, bty = "n")

```




Figure 5.3: Multilocus recovery: $P(\text{at least one line homozygous for P2 at } k \text{ loci})$. This is the objective that keeps paying for larger N .

The marginal gain in D tells the same story from the other side:

step	gain_in_D
10 -> 30	0.0586
30 -> 50	0.0110
50 -> 70	0.0057
70 -> 90	0.0034
90 -> 110	0.0005
110 -> 130	0.0018

The knee is early. Past it, twenty more plants buy thousandths.

5.5 Sizing the F1

With fully homozygous parents the F1 is genetically uniform and the decision is logistical. With residual heterozygosity – $F \approx 0.95$ is common in programme lines – each parent still segregates at a fraction of loci, and each F1 plant receives one random gamete from each parent. At a segregating locus the chance of losing one of the two parental alleles with n F1 plants is $2(1/2)^n$.

```
s <- fl_sizing(n_F1 = 1:20, h_res = 0.05, n_loci_eff = 5000)
s[s$n_F1 %in% c(4, 8, 10, 12, 15, 20), ]
```

	n_F1	p_loss_per_locus	expected_loci_lost
	4	0.1250000	31.2500000
	8	0.0078125	1.9531250
	10	0.0019531	0.4882812

	n_F1	p_loss_per_locus	expected_loci_lost
	12	0.0004883	0.1220703
	15	0.0000610	0.0152588
	20	0.0000019	0.0004768

Twelve selfed F1 plants already put the expected number of lost loci below one. That is the whole F1 sizing argument.

5.6 Working numbers

Gen- eration	Plants	Rationale
F1	8-15 selfed	$N_e = 1$, a logistical call. With 5% residual heterozygosity, 12 plants put expected lost loci below 1
F2	120 to retain diversity 250-400 for multilocus selection or panel frequencies	Above this, +280 plants buy +0.005 of D These are the objectives that keep gaining with N
F3	same N (SSD, one seed per plant)	Does not create diversity, only avoids losing it
F4	lines x 3-8 plants	Seed multiplication. Plants of the same line are not independent replicates: they share 87.5% of the genome

🔥 On the “equivalent N”

$N_{eq} = (1 - H)/(1 - D)$ answers “how many F4 lines without distortion would give the diversity actually observed”. It is a **communication device, not a Wright effective size**: it converts a deterministic bias into a drift equivalent. Use it to compare scenarios; never substitute it into an N_e formula.

5.7 Where this connects

The panel this chapter sizes is the raw material for everything else, and the claim that its diversity is a *ceiling* should be measured rather than asserted. The rest of the book works on one kernel – molecular coancestry f and $GD = 1 - \theta$ (Chapter 1.). The F4 panel can be pushed straight into it:

```
set.seed(3)
map4 <- build_map(spacing_cM = 10)
dist4 <- make_distortion(map4, active = TRUE)

bridge <- do.call(rbind, lapply(c(30, 120, 480), function(n) {
  pop <- run_scheme("SSD", n, map4, dist4, opt = f2_opt)
  X <- t((pop$H1 + pop$H2) / 2) # lines x markers, 0/0.5/1
  th <- mean(molecular_coancestry(X))
  data.frame(n_F2 = n, D = diversity_retained(panel_freq(pop)),
             theta = th, GD = 1 - th)
}))
fmt(bridge, 4)
```

n_F2	D	theta	GD
30	0.9643	0.5179	0.4821

n_F2	D	theta	GD
120	0.9795	0.5103	0.4897
480	0.9899	0.5050	0.4950

D , the chapter's own diversity-retained statistic, and GD , the quantity every later chapter constrains, move together: they are the same loss in two coordinate systems. The sizing decision made here therefore sets the GD that Chapter 12., Chapter 15. and Chapter 16. spend the rest of the book budgeting, and it sets it *once*. An optimiser cannot select back up to variation the panel never sampled; the best it can do is spend what is left more slowly.

i N_s is not informative inside a biparental cross

Computing the status number on these panels returns approximately 1 whatever the sizing, and that is correct rather than broken. Every line descends from two parents, so the panel has the genetic base of a single diploid individual, which is exactly what $N_s = 1/(2\theta)$ reports. Section 11.3 warns that N_s describes a base rather than projecting drift; here the base is genuinely one cross.

Use D and GD to compare sizing scenarios within a cross, and N_s only once material from several crosses has been pooled.

Chapter 13. puts a number on how the loss is distributed across the whole pipeline, and this stage is where the cheapest insurance is available – more plants, once, at the start. Chapter 20. is where the same argument is made about the *variance* rather than the frequencies, and reaches a different conclusion about what is recoverable.

6. Introgressing several genes by backcrossing

Chapter 5. sized a segregating population for *diversity*: the objective was to sample as much of the cross as possible. This chapter is the mirror image. A donor line carries k genes worth having and an elite line carries everything else worth having, and the job is to move the k genes across while **throwing the rest of the donor away**. Here donor genome is contamination, and the only diversity that is wanted is the k target loci themselves.

The operation is marker-assisted backcrossing (MABC). The numbers below come from a simulator of it: a marker grid at 5 cM on ten maize chromosomes (1670 cM), one selected plant advanced per generation, $k = 1$ to 4 recessive targets from a single donor, carried to BC3. It ships with the package – the engine and the selection index are `R/08_mabc.R`, the experiment driver is `inst/scripts/run_mabc.R` – and Section 6.8 shows how to run it.

6.1 The operation, in one diagram

Before any sizing question can be asked, the crossing operations have to be unambiguous – which plant is pollinated by which, and what is genotyped at each step.




Figure 6.1: The MABC operation. Left column: the cross made in each generation. Right column: what is genotyped and selected on, and the population size that criterion demands. The recurrent parent E re-enters at every backcross; the fork at the bottom is the fixation step, where doubled haploids and selfing differ by more than an order of magnitude in plants required.

Two features of that diagram do all the work in the rest of the chapter. The recurrent parent re-enters at every backcross, so donor content halves every generation *whether or not anyone selects for it*; and the fork at the bottom is where the plant numbers explode, because fixation asks for a homozygote and backcrossing never does.

6.2 The model

Two states, one haplotype. The recurrent parent is a fully homozygous inbred, so the gamete it contributes is invariant and only the donor-derived haplotype has to be tracked: a binary vector over 334 markers. Meiosis is a per-interval Bernoulli switch with Haldane's

$$r = \frac{1}{2} (1 - e^{-2d/100}), \quad (6.1)$$

which is exact for this chain precisely because Haldane is memoryless (see the limitation in Section 6.9).

Donor dosage and IBD are different quantities. Donor and elite already share a fraction s of the genome by descent – they are both maize, and often both from the same heterotic group. Donor genome inside a shared tract costs nothing. So with D the donor dosage,

$$\text{IBD}_E = 1 - (1 - s)D, \quad \text{IBD}_E + \text{donor-specific} = 1. \quad (6.2)$$

Both identities hold exactly in the simulator and are asserted in its test battery. Everything below uses $s = 0.6$.

Coupling phase. All k targets come from one donor, on the same haplotype. Two consequences, and the second is easy to get wrong. Linked targets travel *together*, so linkage drag is the **union** of the donor segments, not their sum. And conditional on both flanking targets being donor-derived, an *even* number of crossovers is forced in the interval between them, so the probability that they sit in one continuous donor segment is not e^{-L} but

$$P_{\text{same}} = \frac{e^{-L}}{(1 + e^{-2L})/2}, \quad L = d/100 \text{ Morgans}. \quad (6.3)$$

```
data.frame(d_cM = c(20, 60, 400),
           naive = exp(-c(20, 60, 400) / 100),
           conditional = p_same_segment(c(20, 60, 400)))
```

d_cM	naive	conditional
20	0.8187308	0.9803280
60	0.5488116	0.8435507
400	0.0183156	0.0366190

At 20 cM that is 0.98 against a naive 0.82. Published stacking simulations get this wrong often enough that it is worth stating separately.

! If the alleles come from different donors, none of this applies

Coupling phase is an assumption, not a convenience. Two linked targets in repulsion, from two donors, cannot both be recovered without a recombinant in the interval, and the cost explodes. That programme needs an intercross between the donors before any backcrossing starts, and this model does not describe it.

6.3 Three criteria, and only one of them binds

A conversion programme is usually written with three success conditions at BC3: do not lose the genes, recover the recurrent parent, and do not drag a chunk of donor chromosome along with each target. They are not remotely equal in cost.

Not losing the genes is exact and cheap. A plant is foreground-positive with probability $q = (1/2)^k$, so n plants fail with probability $(1 - q)^n$:

```
data.frame(genes = mb_sz$genes, q = mb_sz$q_foreground,
           n_BC_per_gen = vapply(mb_sz$q_foreground, n_for_count, integer(1)))
```

genes	q	n_BC_per_gen
1	0.5000	5
2	0.2500	11
3	0.1250	23
4	0.0625	47

Fewer than 50 plants per generation at four genes, and that is the whole foreground problem.

Recovering the elite is free by BC3. With no background selection at all, expected donor dosage after 3 backcrosses is $(1/2)^4 = 0.0625$, so

$$IBD_E = 1 - (1 - s)(1/2)^4 = 0.975 \quad (6.4)$$

at $s = 0.6$ – comfortably above a 90% target before a single background marker is scored.

genes	N = 20	N = 80	N = 200	N = 400
1	0.984	0.995	0.997	0.998
2	0.970	0.985	0.990	0.992
3	0.956	0.973	0.980	0.984
4	0.944	0.954	0.969	0.975

That leaves **linkage drag** – the donor chromosome still attached to each target – as the only criterion that a large population can buy. It is the criterion the rest of the chapter is about.

The reframing

An index that spends its decisive stage on background recovery is spending it on the criterion that was already satisfied. This is not a hypothetical: it was the defect the technical review found in the first version of this simulator, and correcting it reverses the headline conclusion. The next section is that correction.

6.4 Which selection index

Four policies, all applied to foreground-positive plants only:

Key	Rule
bg	minimise background donor content
rec12	recombinant selection in BC1-BC2 only (classical practice), then background
rec_bg	recombinant selection throughout, background decides among survivors
rec_drag	recombinant selection throughout, linkage drag decides, background breaks ties

rec_bg is the natural-looking two-stage rule: truncate on drag, then pick the cleanest background among the survivors. Stage 2 discards stage 1. The plant with the least drag survives the truncation and then loses to a sibling that happens to carry less donor genome somewhere irrelevant.




Figure 6.2: Paired comparison (common random numbers, 300 replicates, $s = 0.60$) of the two-stage background-priority index against the drag-priority index. Bars are Clopper-Pearson 95% intervals. The criterion is drag ≤ 30 cM per gene AND IBD ≥ 0.90 at BC3.

genes	N/gen	drag cM	IBD	P(both)	drag cM	IBD	P(both)
3	80	78.7	0.972	0.723	64.6	0.962	0.930
3	150	68.8	0.979	0.813	49.4	0.966	0.997
3	250	65.7	0.982	0.860	41.5	0.968	1.000
3	400	59.0	0.984	0.930	36.0	0.969	1.000
4	80	138.9	0.956	0.340	130.2	0.952	0.417
4	150	120.9	0.965	0.530	105.7	0.958	0.737
4	250	107.9	0.971	0.667	87.6	0.960	0.923
4	400	103.9	0.973	0.767	75.3	0.961	0.987

The trade is 1.1 percentage points of IBD for 19.2 cM of drag, and since the IBD constraint was met in every replicate with margin, it is free. Four genes reach 0.987 at 400 plants per generation. Under the background-priority index the same cell reaches 0.767, which is what produced the original – and wrong – conclusion that the ceiling is unreachable at four genes and the programme must be extended to BC4-BC5 or split in two. **The wall was a property of the index, not of maize.**

6.5 How long to keep selecting recombinants




Figure 6.3: Linkage drag at BC3 (left) and probability of meeting the ceiling (right) under the four policies, $s = 0.60$. Solid = 1 gene, dashed = 3 genes. Extending recombinant selection into BC3 costs no extra plants, only the flanking markers on plants that are being genotyped anyway.

Classical practice stops recombinant selection after BC2, on the reasoning that BC3 is for background clean-up. By BC3 the background needs no clean-up, and the expected distance to the nearest crossover from a target shrinks as $1/g$ with generation g – so BC3 is the generation where a plant buys the *most* drag reduction. Both of the free changes in Figure 6.3 point the same way: genotype the flanking markers in BC3 as well, and let drag decide.

6.6 Allocating a fixed budget across BC1, BC2 and BC3

Given a total number of plants to grow, how should they be split between the three backcrosses?




Figure 6.4: $P(\text{criterion met})$ against the share of a fixed budget spent early, for 1 and 3 genes, binned by total plants. Front-loading BC1 is the worst use of a fixed budget.

A paired re-run at a fixed budget of 360 plants, three genes, makes the ordering explicit:

BC1/BC2/BC3	drag cM	P(both)	lo	hi	drag cM	P(both)	lo	hi
300/40/20	93.0	0.448	0.385	0.512	75.6	0.656	0.594	0.715
120/120/120	72.6	0.772	0.715	0.823	53.5	0.988	0.965	0.998
20/150/190	78.7	0.640	0.577	0.700	58.9	0.876	0.829	0.914
40/120/200	74.3	0.748	0.689	0.801	58.1	0.940	0.903	0.966
200/120/40	79.7	0.680	0.618	0.737	59.3	0.932	0.893	0.960

The balanced schedule wins under both indices, and 300/40/20 – the intuitive “cast a wide net early” schedule – is the worst. The mechanism is the $1/g$ shrinkage again: a plant grown in BC3 is offered a much shorter donor segment to trim than a plant grown in BC1, so it converts into drag reduction far more efficiently.

6.7 Fixing the stack: selfing against doubled haploids

Backcrossing only ever asks for a heterozygote. Fixation asks for a homozygote at all k targets, and that is where the plant numbers move by an order of magnitude. Selfing a BC3F1 gives $q = (1/4)^k$; a doubled haploid derived from a BC3F1 gamete is homozygous everywhere, so it carries all k targets with $q = (1/2)^k$.

And the requirement is not one plant but 3, which is a *cumulative binomial* $P(\text{Bin}(n, q) \geq 3) \geq 0.95$, not $1 - (1 - q)^n$:

genes	q (selfing)	BC3F2 plants	q (DH)	DH lines
1	0.25000	23	0.5000	11
2	0.06250	99	0.2500	23
3	0.01562	401	0.1250	49
4	0.00391	1610	0.0625	99




Figure 6.5: Points: simulation. Dashed: the cumulative binomial. Log scale. Each extra gene multiplies the BC3F2 requirement by four.




Figure 6.6: Analytical sizing by stage, log scale. The two blue bars are backcross generations; red is fixation by selfing; green is fixation by doubled haploids.

At four genes that is 1610 BC3F2 plants against 99 DH lines, a factor of 16.3. For maize, where DH is routine, the alarming number is not the real bottleneck.

Without DH there is a staggered alternative. Do not ask one generation to fix everything. Select BC3F2 plants homozygous for two targets and heterozygous for the other two (roughly 330 plants), then close out the remaining two at BC3F3 (roughly 100). Two stages of hundreds instead of one of thousands, at the cost of one generation.

6.8 Running the simulation

Everything above is reproducible from the package. The engine and the decision layer are in `R/08_mabc.R`; the experiment driver that produced the cached tables is `inst/scripts/run_mabc.R`.

One programme, interactively. Build a map, place the targets, draw the donor-elite sharing, and run the three backcrosses:

```
map <- build_map(maize_chr_len, spacing_cM = 5)
w <- marker_weights(map)
tg <- target_indices(map, k = 3) # the first three of maize_targets
sh <- make_shared_ibd(map, s = 0.60, mw = w, target_idx = tg)

pr <- run_bc_program(n_sched = c(120, 120, 120), s = 0.60,
  strategy = "rec_drag", map = map, opt = mabc_opt,
  tg = tg, w = w, shared = sh)

fmt(pr$traj, 3)
```

gen	ibd	dose	drag	n_fg	ibd_pop_mean	fail
1	0.878	0.251	162.5	13	0.891	FALSE
2	0.934	0.127	97.5	12	0.925	FALSE
3	0.951	0.094	70.0	15	0.952	FALSE

`n_sched` is the schedule, one entry per backcross, so `c(300, 40, 20)` and `c(120, 120, 120)` are the two allocations compared above. Passing the same `shared` mask to two calls is what makes a comparison between selection indices paired rather than confounded with the donor-elite pair.

Comparing indices on common random numbers. Reset the seed before each strategy so the two programmes see the same meioses:

```
drag_of <- function(strategy) {
  set.seed(7)
  run_bc_program(c(120, 120, 120), 0.60, strategy, map, mabc_opt, tg, w,
    shared = sh)$traj$drag[mabc_opt$n_bc]
}
c(rec_bg = drag_of("rec_bg"), rec_drag = drag_of("rec_drag"))
#>   rec_bg rec_drag
#>      60      45
```

A single replicate is noise; the tables above are 300 of these per cell, which is the driver's job.

The analytical sizing needs no simulation at all. It is a cumulative binomial, so it is exact and instant:

```
sizing_table(k_max = 4, flank_cM = 5, P = 0.95, n_keep = 3)
```

genes	q_foreground	n_BC_1plant	n_BC_1rec	q_F2	n_F2_1plant	n_F2_3plants	q_DH	n_DH_3lines
1	0.5000	5	125	0.2500000	11	23	0.5000	11
2	0.2500	11	251	0.0625000	47	99	0.2500	23
3	0.1250	23	503	0.0156250	191	401	0.1250	49
4	0.0625	47	1006	0.0039063	766	1610	0.0625	99

The full scans, from the shell. The driver runs the whole grid – four policies by four gene numbers by eleven population sizes by three values of s , plus the allocation grid and the BCnF2 stage – and writes the CSVs, the nine figures and a self-contained HTML report:

```
Rscript inst/scripts/run_mabc.R --test          # 35 checks, ~10 s
Rscript inst/scripts/run_mabc.R --quick --out=DIR # reduced grid, ~30 s
Rscript inst/scripts/run_mabc.R --out=report/mabc # the full run, ~3 min
```

The cached tables this chapter reads are the full run's `results_uniform.csv`, `results_allocation.csv`, `results_bcnf2.csv` and `analytical_sizing.csv`, copied into `book/data` as `mabc_*.csv`; the header of the driver lists the copy commands.

Pointing it at your own programme. Three things change and nothing else:

```
# 1. your map
map <- build_map(my_chr_len_cM, spacing_cM = 5)

# 2. your targets -- same columns as maize_targets
my_targets <- data.frame(gene = c("A", "B"), chr = c(1, 5),
  pos_cM = c(40, 110), label = c("A", "B"))
tg <- target_indices(map, k = 2, targets = my_targets)

# 3. your criteria
opt <- utils::modifyList(mabc_opt,
  list(n_bc = 4, drag_per_gene = 20, target_ibd = 0.95))
```

s – how much of the genome your donor and your elite already share – is the parameter worth measuring rather than assuming. It scales the whole cost: at $s = 0.7$ most of the donor genome that survives is free, at $s = 0$ none of it is.

6.9 Assumptions, and what is not modelled

These matter more than usual here, because unlike the diversity chapters this one is used to decide field area.

No crossover interference (material, and it cuts against the results above). Meiosis is Haldane, so crossovers are independent between intervals. Maize has strong positive interference. For the event this simulation depends on most – a double recombinant in the two 5 cM intervals flanking a target, which is what turns a long donor segment into a short one – Haldane gives 2.26×10^{-3} against 4.90×10^{-4} under a coincidence model. **The model over-produces tight recombinants by about 4.6x**, so every population size quoted for close-in trimming is a lower bound. Note the direction: the index correction makes the problem easier, interference makes it harder, and they do not cancel in any principled way. The correct fix is a crossover-position sampler – chiasmata as a stationary gamma renewal process with shape ν – which reduces to Haldane exactly at $\nu = 1$ and so is backward compatible; maize estimates put ν near 3.

One plant advanced per generation. Real programmes advance 3-10 and cross each to the recurrent parent. Every probability in this chapter is therefore a **single-lineage** probability: it understates achievable progress and it removes lineage-level variance from every distribution shown.

Uniform recombination rate, no distortion, no genotyping error. Maize has 5-10x rate variation and strongly suppressed pericentromeric regions; a target in a cold region carries drag no feasible population will trim. Gametophyte factors shift the foreground frequency away from 1/2. Map positions here are illustrative and not pinned to a named consensus map – and IBM-type maps must be rescaled before use, exactly as in Chapter 5..

Triploid endosperm genetics is not modelled. The illustrative targets are endosperm genes. The endosperm is 2 maternal : 1 paternal, so a recessive kernel phenotype does not appear on the ear of a heterozygous plant: phenotypic confirmation lags the genotype by one generation, and xenia means it requires a self-pollination, not open pollination.

The gene set is illustrative, not a breeding recommendation. *sh2* and *bt2* encode the two subunits of the same enzyme, so stacking both is largely redundant – dropping one turns the hard case ($k = 4$) into the comfortable one ($k = 3$). *o2* alone gives opaque endosperm, not QPM: the hardness modifiers are polygenic, unlinked and unmodelled here, which changes the problem class from “stack k Mendelian genes” to “stack k genes plus a QTL background”.

⚠ Two statistical habits this chapter had to be repaired for

Select on the interval, not the point estimate. The first version of the sizing tables took the first grid point whose *point estimate* crossed 0.95. With 30-150 replicates the standard error is 0.02-0.09, so that is a winner’s curse: the tables were non-monotone in N (0.97, 0.97, then 0.93 at a larger N), which is a monotone quantity reading noise. Selection is now on the Clopper-Pearson lower bound, and intervals are reported everywhere.

Never type a number into the prose. Every figure in the original narrative was a string constant written against an earlier configuration; re-runs put the true values outside the reported confidence intervals, and the same document contradicted itself in three places. Every number in this chapter is interpolated from the cached tables at render time. Appendix C. says which table each one comes from.

6.10 Where this connects

Chapter 5. sized a population to *keep* diversity; this chapter sized one to discard it everywhere except at k loci. The two share their machinery – the same map, the same Haldane meiosis, the same distinction between what is segregating and what is fixed – and they share the lesson that dominates both: **size for the objective that binds**. In Chapter 5. that meant not sizing for gene diversity, which saturates early. Here it means not sizing for background recovery, which is already satisfied at BC3, and putting the plants into the one criterion that a population can actually buy.

The converted lines this chapter produces re-enter the pool the rest of the book optimises over. A conversion programme run for its own sake reduces diversity by construction: it makes an elite line into a near-copy of itself. If several conversions are run in parallel off the same recurrent parent, the resulting lines are close to each other.

That has been asserted for a chapter and a half. It is cheap to measure, and the size of it is the argument. Eight BC3 conversions off one recurrent parent, against eight unrelated elite lines, put through the kernel of Chapter 1.:

```

set.seed(11)
m_c <- 2000
elite <- matrix(rbinom(40 * m_c, 1, 0.5), 40, m_c) # unrelated elite lines
rec <- elite[1, ] # one recurrent parent
donor <- elite[40, ]
ibd <- expected_ibd(3, 0.5) # recurrent genome at BC3

conv <- t(apply(1:8, function(i)
  ifelse(rbinom(m_c, 1, ibd) == 1, rec, donor)))

cost <- data.frame(
  panel = c("8 BC3 conversions, one recurrent parent", "8 unrelated elite lines"),
  theta = c(mean(molecular_coancestry(conv)),
    mean(molecular_coancestry(elite[1:8, ])))
)
cost$GD <- 1 - cost$theta
fmt(cost, 4)

```

panel	theta	GD
8 BC3 conversions, one recurrent parent	0.9714	0.0286
8 unrelated elite lines	0.5617	0.4383

The converted panel carries about 0.07 times the gene diversity of the unrelated one – not a marginal erosion but a different order of magnitude, and it follows directly from `expected_ibd(3, 0.5) = 0.969`. Conversion is *designed* to make lines identical everywhere except the target, and it succeeds.

Two consequences for the rest of the book. A conversion programme must be counted against the diversity budget of Chapter 15. rather than run alongside it, because θ cannot distinguish eight conversions from one line entered eight times. And Chapter 13. ‘per-line usage cap is the mechanism that already handles this: capping how often a line is used caps how often its near-copies are, provided the conversions are recognised as near-copies rather than entered as eight independent parents.

⚠ Do not enter a converted panel as independent lines

The failure mode is administrative, not statistical. Conversions arrive with distinct line codes and distinct pedigree entries, so a pipeline that trusts the codes will treat them as eight parents. The coancestry matrix will not, and Chapter 12. ‘s recommendation to compute θ on markers rather than on pedigree is what protects against exactly this.

7. Simulating a breeding programme: AlphaSimR from the ground up

Chapter 4. built the `{X, f, hybrids, fL}` contract with a deliberately fast simulator: Balding-Nichols allele frequencies, independent markers, no recombination, no dominance. That is enough to exercise the diversity and selection machinery, but it is not how a breeding programme is actually run, and it is not how anyone would build one from scratch.

This chapter does both. Steps 1 to 7 are a **step-by-step construction of a simulation** with AlphaSimR (R. C. Gaynor, G. Gorjanc, and J. M. Hickey [7]), written out in the primitive calls – `runMacs()`, `SimParam`, `addTraitA()`, `makeCross()`, `makeDH()`, `setPheno()`, `selectInd()` – with a plant breeder’s reading of each argument. Step 8 then assembles those primitives into **reciprocal recurrent selection** (RRS; R. E. Comstock, H. F. Robinson, and P. H. Harvey [8]), the standard commercial scheme for two heterotic pools, and shows that the same stage-1 contract comes out the other end, unchanged, so `stage2_select()` runs on it exactly as it does on the book’s own simulated data.

The order matters. Everything in Step 8 is a wrapper in `R/13_alphasimr_bridge.R` – `rrs_config`, `alphasimr_founder_pools()`, `alphasimr_rrs_cycle()`, `alphasimr_make_hybrids()`, `alphasimr_to_stage1()`, `alphasimr_pipeline()` – and a wrapper is only readable once you know what it is wrapping. AlphaSimR is a `Suggests` dependency, guarded exactly like `quadprog` and `highs` elsewhere in the package; this chapter, and Chapter 8., are the only ones that need it.

Chapter 8. takes the same primitives in the other direction: a complete *inbred-line* programme with staged yield trials, rather than a hybrid one.

7.1 Step 1 – Founder genomes, and which engine to use

Every AlphaSimR simulation starts with founder haplotypes. There are two engines, and the choice is not cosmetic.

`runMacs()` runs a backward-in-time coalescent (G. K. Chen, P. Marjoram, and J. D. Wall [9]) under a species-specific demography – `species = "MAIZE", "WHEAT", "CATTLE"` – so the founders carry a realistic allele-frequency spectrum and realistic linkage disequilibrium. `quickHaplo()` samples 0s and 1s directly: instant, and with no history behind it.

`inbred = TRUE` is the argument a plant breeder almost always wants. It makes every founder fully homozygous, which is what a panel of inbred lines or released varieties actually is. Without it the founders are outbred individuals, which is the right model for cattle and the wrong one for wheat.

```
tm <- function(expr) round(as.numeric(system.time(expr)["elapsed"]), 2)

set.seed(7); t_q <- tm(qa <- quickHaplo(nInd = 20, nChr = 10, segSites = 50,
                                       inbred = TRUE))
set.seed(7); qb <- quickHaplo(nInd = 20, nChr = 10, segSites = 50,
                              inbred = TRUE)

set.seed(7); t_m <- tm(ma <- runMacs(nInd = 20, nChr = 10, segSites = 50,
                                    inbred = TRUE, species = "MAIZE"))
set.seed(7); mb <- runMacs(nInd = 20, nChr = 10, segSites = 50,
                          inbred = TRUE, species = "MAIZE")

fmt(data.frame(engine = c("quickHaplo", "runMacs"),
               seconds = c(t_q, t_m),
               same_twice = c(identical(qa@geno, qb@geno),
                             identical(ma@geno, mb@geno))))
```

engine	seconds	same_twice
quickHaplo	0.00	TRUE
runMacs	2.84	FALSE

Read the third column again. **runMacs()** is not reproducible from **set.seed()**. It hands the coalescent off to the MaCS binary, which seeds itself outside R's random number stream, so two calls with an identical seed return different founders. **quickHaplo()** is pure R sampling and is exactly reproducible.

(The **seconds** column is wall clock and will differ on your machine and between renders; the ratio between the two engines is the part that transfers.)

⚠ The founders are not seeded, and everything downstream is

This is the single most expensive trap in AlphaSimR, because it fails silently: the code looks reproducible, runs without complaint, and returns different numbers. It is already documented in **alphasimr_founder_pools()**'s help page – the note there says **runMacs2()**, and the table above shows it is not **runMacs2**-specific but true of MaCS generally.

Two consequences:

1. **Found once, then branch.** Any experiment comparing two schemes must draw the founders a single time and start both arms from that one object. Draw them twice and the comparison measures the founders as much as the schemes. Chapter 20's linked arm is built this way for exactly this reason.
2. **Save what you cannot reseed.** **saveRDS(founderGenomes, ...)** is the only way to get the same coalescent history back in a later session.

Everything after the founders – the crosses, the doubled haploids, the phenotyping error, the recombination – is controlled by **set.seed()**.

So what does **quickHaplo()** give up? Linkage disequilibrium among the founders. The honest way to measure that is against the sampling floor: with n individuals, two completely independent loci still show r^2 of about $1 / (n - 1)$ by chance alone.

```
ldstat <- function(fg) {
  SPx <- SimParam$new(fg)
  g <- pullSegSiteGeno(newPop(fg, simParam = SPx), simParam = SPx)
  g <- g[, apply(g, 2, stats::sd) > 0, drop = FALSE]
  r <- stats::cor(g)
  c(adjacent = mean(diag(r[-nrow(r), -1])^2), # neighbouring sites
    overall = mean(r[upper.tri(r)]^2),        # all pairs, mostly distant
    noise = 1 / (nrow(g) - 1))                # what pure chance gives
}
set.seed(3); lq <- ldstat(quickHaplo(nInd = 100, nChr = 5, segSites = 200,
                                     inbred = TRUE))
set.seed(3); lm_ <- ldstat(runMacs(nInd = 100, nChr = 5, segSites = 200,
                                   inbred = TRUE, species = "MAIZE"))
fmt(rbind(quickHaplo = lq, runMacs = lm_))
```

	adjacent	overall	noise
quickHaplo	0.0103	0.0101	0.0101
runMacs	0.1230	0.0119	0.0101

quickHaplo()'s adjacent-site r^2 sits exactly on the noise floor: there is no founder LD at all, as designed. **runMacs()**'s is an order of magnitude above it, and its *overall* r^2 is barely above the floor – so the LD it creates is local, decaying with map distance, which is what real LD does.

The **runMacs** row of that table is itself not reproducible, for the reason the row above it establishes; the **quickHaplo** row is, to the last digit.

That contamination spreads further than it looks. **runMacs()** consumes an unpredictable amount of R's own random stream on its way to and from the MaCS binary, so *everything executed after it in the same session* inherits an unpredictable seed – an earlier draft of this chapter had reproducible-looking Steps 2

to 7 that quietly moved whenever Step 1 was re-executed. The remedy is the rule stated above, applied to the chapter itself: each of the steps below calls `set.seed()` immediately before drawing its own founders, rather than trusting a seed set once at the top. With that in place, two independent executions agree on every number in Steps 1 to 7 except the wall-clock seconds and the `runMacs` row itself.

That measurement decides the policy this chapter follows:

- **quickHaplo(inbred = TRUE) for Steps 2–7 and all of Chapter 8.** They teach polygenic response arithmetic, where founder LD is not the point, and in exchange every number printed is reproducible from a clean clone. AlphaSimR still applies a real genetic map and real meiosis to every descendant, so linkage and the Bulmer effect of Chapter 20. still operate – only the founders' *history* is stylised.
- **runMacs2(split =) in Step 8.** There the divergence between two heterotic pools *is* the object of study, and no amount of speed compensates for not having it.

Founder LD does matter where marker-QTL association is the mechanism, which is genomic prediction – see Chapter 14., and note that the accuracies there are computed on a different simulator for this reason.

7.2 Step 2 – SimParam, and defining a trait

`SimParam` holds everything global: the genetic map, which sites are QTL, the trait effects, the error variances. It is an R6 object, so its methods are called with `$`, and a variable named `SP` is picked up automatically by every AlphaSimR function that needs it.

```
# Seeded here, not once at the top of the chapter: the runMacs() calls in Step 1
# consume an unpredictable amount of R's random stream, so anything downstream of
# them is only reproducible by accident unless the seed is reset. This is the
# "found once, then branch" rule of Step 1 applied to the chapter itself.
set.seed(101)
founders <- quickHaplo(nInd = 60, nChr = 10, segSites = 100, inbred = TRUE)
SP <- SimParam$new(founders)
SP$addTraitA(nQtlPerChr = 100, mean = 10, var = 1)
SP$setVarE(h2 = 0.4)
```

There are four trait constructors, and choosing between them is choosing a model of gene action:

Method	Adds	When a plant programme needs it
<code>SP\$addTraitA()</code>	additive only	line per se performance; the default teaching case
<code>SP\$addTraitAD()</code>	1. dominance	testcross and hybrid value. Without dominance a line's cross performance is just its own breeding value and there is nothing for RRS to exploit (Chapter 10.). This is what Step 8 uses
<code>SP\$addTraitAG()</code>	1. genotype-by-environment	multi-location yield trials, where the same genotype is scored in several places. Chapter 8. uses it
<code>SP\$addTraitADE()</code>	1. epistasis	available; used nowhere in this book

`nQtlPerChr` sets how polygenic the trait is. Setting it equal to `segSites` makes *every* segregating site a QTL, which is what Chapter 7.'s own configuration does – convenient, because then the marker matrix and the causal loci are the same set and there is no marker-versus-QTL gap to reason about. Real programmes have that gap; Chapter 9. is where the book takes it seriously.

`mean` and `var` fix the trait's mean and *genetic* variance in the founder population by rescaling the sampled effects. `corA` (Step 7) sets the genetic correlation between traits.

The environmental variance has two routes, and plant breeders use the second:

- `SP$setVarE(h2 = 0.4)` fixes an error variance once, from the founders' genetic variance, and every later `setPheno()` reuses it.
- `setPheno(varE = , reps = )` sets it per call. **reps is the plant idiom for locations**: a genotype scored in `reps` environments has its error variance divided by `reps`. That is how a staged programme raises selection accuracy from one stage to the next without ever changing the trait's heritability – the mechanism Chapter 8. is built on.

7.3 Step 3 – The base population, and where the variance components live

`newPop()` turns founder haplotypes into individuals and, because a trait is now defined, computes their genetic values immediately.

```
pop <- setPheno(newPop(founders, simParam = SP), simParam = SP)
c(n = pop@nInd, meanG = as.numeric(meanG(pop)), varG = as.numeric(varG(pop)),
  varP = as.numeric(varP(pop)),
  h2_realised = as.numeric(varG(pop) / varP(pop))) |> fmt(3)
#>           n      meanG      varG      varP h2_realised
#>      60.000      10.000       1.000      3.376         0.296
```

The realised heritability is near but not equal to the requested 0.4: `setVarE()` fixes a variance, and the ratio it produces in any finite sample is itself a random variable. This is a general and useful habit – **check what the simulation realised, do not assume what you requested**.

`genParam()` returns the full decomposition, and two of its fields are the ones this book cares about most:

```
mixed <- pop
for (k in 1:3) mixed <- randCross(mixed, nCrosses = 1000, simParam = SP)

vrow <- function(p, lab) {
  g <- genParam(p, simParam = SP)
  data.frame(population = lab,
             varA       = as.numeric(g$varA),
             genicVarA  = as.numeric(g$genicVarA),
             ratio      = as.numeric(g$varA / g$genicVarA))
}
fmt(rbind(vrow(pop, "founders (fully inbred)"),
          vrow(mixed, "after 3 random matings")))
```

population	varA	genicVarA	ratio
founders (fully inbred)	1.000	0.4247	2.3544
after 3 random matings	0.462	0.4240	1.0895

`varA` is the realised additive variance; `genicVarA` is what it would be if the loci were independent and in Hardy-Weinberg proportions. Their ratio is the whole subject of Chapter 20., and the package's `genic_var()` computes the same quantity from a marker matrix directly.

The founders' ratio above 1 is not disequilibrium – it is **inbreeding**. A fully homozygous population has per-locus genotypic variance $4pq$ where a Hardy-Weinberg one has $2pq$, so `varA` runs up to twice the genic value by construction. Three rounds of random mating restore Hardy-Weinberg and pull the ratio back towards 1; what remains is the disequilibrium among only 60 founders' worth of gametes. A ratio away from 1 in an *outbred* population is the gametic-phase disequilibrium Chapter 3. defines and Chapter 20. measures; a ratio away from 1 in an inbred one is mostly just the homozygosity.

7.4 Step 4 – The DNA lottery, and why an F1 is not a family

The course this chapter's structure follows demonstrates Mendelian sampling by crossing outbred founders and showing that full sibs vary. With **inbred** parents the result is different, and it is the fact plant breeding is organised around.

```
lottery <- do.call(rbind, lapply(list(c(1, 2), c(1, 3), c(2, 3)), function(ij) {
  plan <- matrix(ij, ncol = 2)
  f1 <- makeCross(pop, plan, nProgeny = 200, simParam = SP)
  dh <- makeDH(makeCross(pop, plan, nProgeny = 1, simParam = SP),
               nDH = 200, simParam = SP)
  better <- max(gv(pop)[ij])
  data.frame(cross      = paste(ij, collapse = " x "),
             parent_mean = mean(gv(pop)[ij]),
             better_parent = better,
             f1_sd       = as.numeric(sqrt(varG(f1))),
             dh_sd       = as.numeric(sqrt(varG(dh))),
             dh_best     = max(gv(dh)),
             transgressive = mean(gv(dh) > better))
}))
fmt(lottery, 3)
```

cross	parent_mean	better_parent	f1_sd	dh_sd	dh_best	transgressive
1 x 2	10.655	12.193	0	0.765	12.697	0.015
1 x 3	10.912	12.193	0	0.548	12.384	0.010
2 x 3	9.374	9.632	0	0.684	11.789	0.340

f1_sd is **exactly zero**. Two fully inbred parents are homozygous everywhere, so every gamete each one makes is identical, and every F1 is the same heterozygote. Two hundred F1 plants from a cross of two inbred lines are two hundred copies of one genotype. There is no lottery yet.

The lottery starts one meiosis later. **dh_sd** is the standard deviation among doubled haploids drawn from that F1, and it is substantial – this is the segregation that F2, DH and RIL populations exist to sample. It is also why Chapter 5. is a chapter: the number of DH lines you pull from a cross is the number of tickets you buy, and no downstream selection stage can recover a combination that was never sampled.

transgressive is the fraction of DH lines exceeding the *better* parent. It is not small. Crossing two mediocre parents whose favourable alleles sit at different loci produces progeny beyond either of them – transgressive segregation, which [I. J. Mackay, J. Cockram, and P. Howell \[10\]](#) argues is the same phenomenon as heterosis and inbreeding depression viewed from a third angle. Chapter 10. is where this book takes up that argument.

7.5 Step 5 – Selfing, doubled haploids, and what the shortcut buys

A plant breeder has to get from a heterozygous F1 to a homozygous line. There are two routes, and AlphaSimR implements both.

```
homozygosity <- function(p)
  mean(pullSegSiteGeno(p, simParam = SP) %in% c(0, 2))

f1 <- makeCross(pop, matrix(c(1, 2), ncol = 2), nProgeny = 200, simParam = SP)

gen <- f1
traj <- data.frame(generation = "F1", homozygosity = homozygosity(f1))
for (k in 2:7) {
  gen <- self(gen, nProgeny = 1, simParam = SP)
  traj <- rbind(traj, data.frame(generation = paste0("F", k),
                                homozygosity = homozygosity(gen)))
}
traj <- rbind(traj, data.frame(generation = "DH",
                              homozygosity = homozygosity(makeDH(f1, nDH = 1,
                                                                    simParam = SP))))

fmt(traj)
```

generation	homozygosity
F1	0.5010
F2	0.7493
F3	0.8755
F4	0.9374
F5	0.9692
F6	0.9849
F7	0.9927
DH	1.0000

`self()` halves the remaining heterozygosity every generation – the familiar $1 - (1/2)^{(t-1)}$ approach to fixation, which never quite arrives. `makeDH()` gets there in one step, at 1.0 exactly, by doubling a haploid gamete.

The saving is seasons, and seasons are the currency of a breeding programme: six or seven selfing generations against one DH cycle. Chapter 8’s scheme uses `makeDH()` for precisely this reason, and the cost that never shows up in the simulation is that DH production is a laboratory operation with its own failure rate and its own genotype-dependent response.

The other operations a scheme needs are all one call each: `randCross()` for random mating within a pool, `makeCross()` for a specified crossing plan, `hybridCross()` for a full factorial between two groups (Step 8), and `mergePops()` or `c()` to combine populations.

7.6 Step 6 – Phenotyping, selection, and checking the breeder’s equation

Selection needs a criterion. `setPheno()` supplies the noisy one; `selectInd()` acts on it.

```
mixed <- setPheno(mixed, simParam = SP)

sweep_q <- do.call(rbind, lapply(c(0.05, 0.10, 0.20), function(q) {
  n_sel <- round(q * mixed@nInd)
  sel <- selectInd(mixed, nInd = n_sel, use = "pheno", simParam = SP)
  prog <- randCross(sel, nCrosses = 1000, simParam = SP)

  S <- as.numeric(meanP(sel) - meanP(mixed))
  vA <- as.numeric(varA(mixed, simParam = SP))
  h2r <- vA / as.numeric(varP(mixed))
  i_emp <- S / sqrt(as.numeric(varP(mixed)))
  r_emp <- stats::cor(gv(mixed)[, 1], pheno(mixed)[, 1])

  data.frame(q = q,
             i_empirical = i_emp, i_normal = sel_intensity(q),
             R_S_h2 = S * h2r,
             R_i_r_sdA = i_emp * r_emp * sqrt(vA),
             R_observed = as.numeric(meanG(prog) - meanG(mixed)))
}))
fmt(sweep_q)
```

q	i_empirical	i_normal	R_S_h2	R_i_r_sdA	R_observed
0.05	2.0325	2.0627	0.6671	0.6358	0.7256
0.10	1.7512	1.7550	0.5748	0.5478	0.5674
0.20	1.4209	1.3998	0.4664	0.4445	0.4488

Both forms of the breeder’s equation are here. $R = S * h^2$ uses the selection differential directly; $R = i * r * \sigma_A$ decomposes it into intensity, accuracy and available variation – the form Chapter 3. builds the book’s objective side on. R_{observed} is the realised change in *genetic* mean between the selected parents’ progeny and the starting population, which is what the prediction is a prediction of.

They agree to a few percent, and the agreement degrades as q rises, which is the expected direction: the normal-theory intensity assumes an infinite population and a normal criterion, and a 20% selected fraction of a finite population departs from both.

The book exports its own selection intensity, and AlphaSimR exports one too. They are independent implementations of $\phi(z) / q$:

```
q <- 0.05
c(hybdiv = sel_intensity(q), AlphaSimR = selInt(q),
  difference = abs(sel_intensity(q) - selInt(q))) |> fmt(10)
#>      hybdiv  AlphaSimR difference
#>    2.062713    2.062713    0.000000
```

`selectInd()` takes `use = "pheno", "gv" or "ebv"`. Selecting on "gv" is selecting with perfect accuracy – impossible in a real programme, and the correct way to compute the ceiling a scheme is being measured against. `use = "ebv"` selects on whatever has been written into the population’s `ebv` slot, which is how genomic selection enters a simulation (Chapter 8.). `selectTop = FALSE` selects downwards, for traits where less is better.

`selectWithinFam()` is the plant-specific one. A programme that truncates globally on a set of families will take every entry from the two best crosses and nothing from the rest; capping the number retained per family holds the programme’s genetic base open at some cost in immediate gain. It is the same trade-off as the `max_use` parental cap of Chapter 16., applied one stage earlier.

7.7 Step 7 – Two traits, and what equal weighting assumes

`corA` makes the traits genetically correlated, which is where selection stops being one-dimensional.

```
set.seed(202)
f2 <- quickHaplo(nInd = 80, nChr = 10, segSites = 100, inbred = TRUE)
SP2 <- SimParam$new(f2)
SP2$addTraitA(nQtlPerChr = 100, mean = c(10, 10), var = c(1, 4),
  corA = matrix(c(1, -0.4, -0.4, 1), 2, 2))
SP2$setVarE(h2 = c(0.6, 0.1))

mt <- newPop(f2, simParam = SP2)
for (k in 1:3) mt <- randCross(mt, nCrosses = 1000, simParam = SP2)
mt <- setPheno(mt, simParam = SP2)

G <- varG(mt); P <- varP(mt); a <- c(1, 1)
c(cor_genetic = stats::cov2cor(G)[2, 1],
  cor_pheno = stats::cov2cor(P)[2, 1],
  h2_trait1 = as.numeric(diag(G)[1] / diag(P)[1]),
  h2_trait2 = as.numeric(diag(G)[2] / diag(P)[2])) |> fmt(3)
#> cor_genetic cor_pheno h2_trait1 h2_trait2
#>    -0.416    -0.084    0.470    0.072
```

The genetic correlation lands near the requested -0.4; the phenotypic one is much weaker, because the environmental components were left uncorrelated (`SP$setCorE()` is where that would change). The two traits differ in variance and, sharply, in heritability – the ordinary situation.

The book’s pipeline defaults to `weights = NULL`, meaning equal weight on every trait. Step 7 is where that default gets examined, using the index functions Chapter 3. introduced:

```
cor_with_H <- function(w)
  drop(crossprod(w, G %*% a)) /
  sqrt(drop(crossprod(w, P %*% w)) * drop(crossprod(a, G %*% a)))

B <- rbind(equal = a,
  smith_hazel = index_smith_hazel(P, G, a),
```

```
restricted = index_restricted(P, G, a, R = cbind(c(0, 1)))
dG <- t(apply(B, 1, function(w) drop(G %*% w)))
colnames(dG) <- c("gain_trait1", "gain_trait2")
fmt(cbind(B, dG, cor_with_H = apply(B, 1, cor_with_H)))
```

	Trait1	Trait2	gain_trait1	gain_trait2	cor_with_H
equal	1.0000	1.0000	0.0677	2.2681	0.2434
smith_hazel	0.0789	0.0589	0.0163	0.1226	0.2438
restricted	0.1986	0.0386	0.1011	0.0000	0.2080

Three readings, in the order they matter:

1. **Equal weights are not the Smith-Hazel solution.** Smith-Hazel puts 0.46 units of weight on trait 2 for every unit on trait 1, not 1.00. Equal weighting is optimal only when $P^{-1}G$ is proportional to the identity, which needs the traits equally variable, equally heritable and uncorrelated. None of that holds here, and it rarely holds anywhere.
2. **The cost of getting it wrong is modest here, and the direction of the error is not.** Equal weighting loses about 2% of the attainable correlation with the aggregate genotype, but it drives a large unintended gain in trait 2 simply because trait 2 has four times the variance. An index that is nearly as good on the objective can still be badly wrong about where the gain lands.
3. **The restricted index holds trait 2 at exactly zero.** That is Kempthorne and Nordskog's construction (O. Kempthorne and A. W. Nordskog [11]) and, as Chapter 3. argues, the classical ancestor of this book's alpha constraint – restrict a direction rather than price it.

Smith-Hazel is defined as the weight vector maximising that correlation, so the claim is checkable against a numerical optimum rather than taken on faith:

```
opt <- stats::optim(c(1, 1), function(w) -cor_with_H(w), method = "BFGS")
c(numerical = -opt$value, smith_hazel = cor_with_H(B["smith_hazel", ]),
  difference = abs(-opt$value - cor_with_H(B["smith_hazel", ]))) |> fmt(10)
#> numerical smith_hazel difference
#> 0.2438417 0.2438417 0.0000000
```

Multi-trait genomic selection in wheat, balancing yield against protein content and quality, is the standing plant example – S. Michel, F. Löschnerberger, C. Ametz, B. Pachler, E. Sperry, and H. Bürstmayr [12].

7.8 Step 8 – The parts assembled: reciprocal recurrent selection

Everything above was one operation at a time. A breeding scheme is those operations in a loop, and the loop for two heterotic pools is reciprocal recurrent selection: each pool is improved using the *other* pool as tester. Doubled-haploid lines are drawn from pool A, testcrossed to pool B (and vice versa), the testcross hybrids are phenotyped, and the lines with the best mean testcross performance – a proxy for general combining ability, GCA – are recombined to found the next cycle's pool. Elite lines from both pools, at the end, are crossed directly to produce the commercial hybrids. It is exactly the asymmetry Chapter 19. argues for: selection happens on the hybrid, recycling happens within pool.

7.8.1 Founding two pools with a real coalescent history

This is the one place the chapter pays for `runMacs2()` rather than `quickHaplo()`. `AlphaSimR::runMacs2(..., split = )` draws founder haplotypes from a population-split coalescent history: one ancestral population, then two demes diverging `split` generations ago. That is the realistic analogue of `simulate_lines()`'s Balding-Nichols draw in Chapter 4., and between-pool divergence is the entire subject of Chapter 2., so a stylised substitute would not do.

```
fp <- alphasimr_founder_pools(rrs_cfg)
c(pool_A = fp$pop_A@nInd, pool_B = fp$pop_B@nInd,
  markers = rrs_cfg$n_chr * rrs_cfg$seg_sites)
#> pool_A pool_B markers
#> 12 12 750
```

```

alphasimr_gd(fp$pop_A, fp$SP)
#> [1] 0.2149259
alphasimr_gd(fp$pop_B, fp$SP)
#> [1] 0.207463

```

`alphasimr_gd()` is $1 - \text{mean}(f)$, the exact definition `gene_diversity()` uses, computed directly from AlphaSimR's genotype pull – so the two pools are already comparable to every diversity number in the rest of the book on the same scale.

Reading `alphasimr_founder_pools()` against Steps 1 and 2, it is `runMacs2(nInd = , nChr = , segSites = , split = )` followed by `SP$addTraitAD(nQtlPerChr = seg_sites, meanDD = 0.4, varDD = 0.1)` and `SP$setVarE(h2 = )`. The requested `nInd` individuals are sampled across both demes – the first `n_pool_A` become pool A, the rest pool B. There is no separate marker panel: every one of the 150 segregating sites per chromosome is a QTL, so the sites that carry the trait are the ones `pullSegSiteGeno()` reads for `X`, `f` and `fL` throughout. And per Step 2's table, the trait is `addTraitAD` because without dominance there is nothing for testcross selection to exploit that within-pool selection could not already get.

Because these founders come from MaCS, the numbers in the rest of this section change whenever the chapter is re-executed – see Step 1's warning. The numbers in Steps 1 to 7, and in Chapter 8., do not.

7.8.2 One cycle

```

cyl <- alphasimr_rrs_cycle(fp$pop_A, fp$pop_B, fp$SP, rrs_cfg)
fmt(cyl$summary, 4)

```

mean_tc_A	mean_tc_B	gd_A	gd_B
2.2786	2.447	0.1787	0.182

`alphasimr_rrs_cycle()` runs the same six steps for each pool, symmetrically (pool A tested on B, pool B tested on A), and every one is a primitive from above:

1. **Double haploids** (Step 5). `makeDH()` draws 3 fully homozygous lines from *each* individual in the current pool – at book scale, 12 parents x 3 DH each = 36 DH lines per pool, per cycle.
2. **Pick reciprocal testers.** The first 2 individuals of the *opposite* pool are used as testers – a fixed slice, not a random or merit-based draw. This is the RRS-defining move: a line's value is judged by its cross onto the other heterotic pool, not by itself.
3. **Testcross and phenotype** (Step 5, Step 6). `hybridCross()` forms the full factorial between every DH line and every tester – 36 DH x 2 testers = 72 testcross hybrids per pool – and `setPheno()` adds phenotyping error sized to hit the heritability set by `SP$setVarE(h2 = )`.
4. **GCA proxy.** Testcross phenotypes are grouped by DH-line parent (`tapply(pheno[, 1], mother, mean)`) and averaged across the 2 testers each line was crossed to – a general combining ability estimate, following [R. E. Comstock, H. F. Robinson, and P. H. Harvey \[8\]](#). Only trait 1 is used for this ranking even though this chapter's config simulates 2 traits per line; the second trait rides along in `X/hybrids` but plays no role in selection here. Step 7 is what it would take to use both.
5. **Keep the elites** (Step 6). The 6 DH lines per pool with the highest mean testcross phenotype are kept as `elite_A/elite_B`.
6. **Recombine** (Step 5). `randCross()` random-mates the elites back up to 12 individuals, founding the next cycle's pool A (symmetrically for B).

`elite_A/elite_B` are what step 6 mates – and, in the *last* cycle, what the final commercial cross uses instead of being recombined again. The cycle's `summary` mixes two different populations on purpose: `mean_tc_A/mean_tc_B` average over *all* DH lines tested that cycle (the population's testcross performance before selection), while `gd_A/gd_B` are gene diversity of only the 6 surviving elites (the population after selection). That asymmetry is exactly what the next section plots – response measured pre-selection, diversity measured post-selection.

7.8.3 The trajectory across cycles

`alphasimr_pipeline()` just runs `alphasimr_rrs_cycle()` in a loop for 4 cycles, carrying `pop_A/pop_B` forward from one cycle's output to the next's input – there is no reseeding inside the loop, only the founder draw at the start is seeded, so the whole trajectory is one continuous random process. Only the

last cycle's `elite_A/elite_B` – not an intermediate cycle's – go on to found the commercial hybrid pool below; every earlier cycle's elites are spent entirely on `randCross()`ing the next cycle's pool.

```
res <- alphasimr_pipeline(rrs_cfg)
fmt(res$cycles, 4)
```

cycle	mean_tc_A	mean_tc_B	gd_A	gd_B
1	2.6772	2.3858	0.1667	0.1876
2	3.3216	3.8968	0.1273	0.1501
3	3.9435	4.5341	0.0907	0.1296
4	5.2887	5.4234	0.0767	0.1245




Figure 7.1: Reciprocal recurrent selection over four cycles, book scale. Left: mean testcross phenotype (selection response) per pool. Right: gene diversity of the elite lines retained each cycle.

Both pools respond – mean testcross phenotype rises – and both pools lose gene diversity in their retained elite lines, cycle over cycle. Nobody pays this cost for its own sake: Chapter 19. makes the same point with the package's own faster simulator, over a longer trajectory and a grid of selection intensities, and Chapter 8. makes it a third time in an inbred-line programme. This section shows it holds with actual meiosis and dominance behind it too.

7.8.4 The final hybrid pool through `stage2_select()`, unchanged

`alphasimr_make_hybrids()` testcrosses the final cycle's elites directly against each other – a full factorial diallel between the 6 pool-A elites and 6 pool-B elites, so 36 commercial hybrids – and phenotypes them the same way every testcross in this chapter was phenotyped. `alphasimr_to_stage1()` then does the bookkeeping that makes this look like any other stage-1 output: it remaps AlphaSimR's internal line ids into the `1..n_A, n_A+1..n_A+n_B` convention `hybrids$line_A/line_B` and `fL` use elsewhere in the book, and builds `fL` by stacking the elite lines from *both* pools' genotypes into one matrix and computing a single joint coancestry matrix over all of them – not two separate per-pool matrices – which is what lets `theta_pools()` and the rest of the pipeline index into pool A and pool B lines with one consistent numbering.

```
st1 <- res$stage1
dim(st1$X)
#> [1] 36 750
sel <- stage2_select(st1$X, st1$f, st1$hybrids, n_sel = 8, fL = st1$fL)
#> [1/4] null distribution by resampling (B = 2000 cheap, 200 full)...
```

```
#>      GD reference = 0.23794 (sd 0.00302)
#>      Monte-Carlo s.e. 6.7e-05 -> alpha differences below 0.028% are noise
#>      bias if using the population as reference: 2.13%
#> [2/4] maximum attainable alpha (pure truncation) = -1.14%
#> [3/4] warm start (greedy) and DE per scenario...
#>      DE_alpha_0.00      index 1.323 | alpha achieved -1.141% | Ns 0.6585 | Ne_par 7.1
#>      DE_alpha_-0.29     index 1.323 | alpha achieved -1.141% | Ns 0.6585 | Ne_par 7.1
#>      DE_alpha_-0.57     index 1.323 | alpha achieved -1.141% | Ns 0.6585 | Ne_par 7.1
#>      DE_alpha_-0.86     index 1.323 | alpha achieved -1.141% | Ns 0.6585 | Ne_par 7.1
#>      DE_alpha_-1.14     index 1.323 | alpha achieved -1.141% | Ns 0.6585 | Ne_par 7.1
#> [4/4] metrics per scenario...
#>      F_hom is what the constraint pins; F_drift is what it does not:
#>      DE_alpha_0.00      F_hom +0.0104 | F_drift 0.0241 | gap -0.0138 | GD_BS 0.1504
#>      DE_alpha_-0.29     F_hom +0.0104 | F_drift 0.0241 | gap -0.0138 | GD_BS 0.1504
#>      DE_alpha_-0.57     F_hom +0.0104 | F_drift 0.0241 | gap -0.0138 | GD_BS 0.1504
#>      DE_alpha_-0.86     F_hom +0.0104 | F_drift 0.0241 | gap -0.0138 | GD_BS 0.1504
#>      DE_alpha_-1.14     F_hom +0.0104 | F_drift 0.0241 | gap -0.0138 | GD_BS 0.1504
#>      ref_truncation     F_hom +0.0104 | F_drift 0.0241 | gap -0.0138 | GD_BS 0.1504
#>      ref_trunc_cap      F_hom +0.0104 | F_drift 0.0241 | gap -0.0138 | GD_BS 0.1504
#>      ref_max_div        F_hom -0.0129 | F_drift 0.0079 | gap -0.0208 | GD_BS 0.1414
#>      ref_random         F_hom +0.0337 | F_drift 0.0195 | gap +0.0142 | GD_BS 0.1466
fmt(sel$metrics[, c("scenario", "alpha", "GD", "GD_BS", "F_ST", "index")], 4)
```

scenario	alpha	GD	GD_BS	F_ST	index
DE_alpha_0.00	-0.0114	0.2407	0.1504	0.6251	1.3233
DE_alpha_-0.29	-0.0114	0.2407	0.1504	0.6251	1.3233
DE_alpha_-0.57	-0.0114	0.2407	0.1504	0.6251	1.3233
DE_alpha_-0.86	-0.0114	0.2407	0.1504	0.6251	1.3233
DE_alpha_-1.14	-0.0114	0.2407	0.1504	0.6251	1.3233
ref_truncation	-0.0114	0.2407	0.1504	0.6251	1.3233
ref_trunc_cap	-0.0114	0.2407	0.1504	0.6251	1.3233
ref_max_div	-0.0229	0.2434	0.1414	0.5812	-0.3250
ref_random	0.0022	0.2374	0.1466	0.6173	-0.1526

Nothing here is AlphaSimR-specific: `stage2_select()` never sees where `X`, `f`, `hybrids` and `fL` came from, only that they satisfy the contract stage 1 promises. `GD_BS`, the between-pool divergence that carries the heterosis signal (Chapter 12.), is computed from `fL` exactly as it is everywhere else in the book.

7.9 Assumptions, and what this adds over Chapter 4.

What is more realistic here. Actual recombination and linkage (a genetic map per chromosome, not independent markers), a dominance-driven trait so testcross selection has something real to select on, and multi-generation recurrent selection rather than one static pool.

What is still simplified. Testers are a small fixed subset of the opposite pool (2 individuals) taken by index, not chosen at random or by merit, and not the single unrelated inbred tester a real programme would typically use – a fixed pool slice makes the same testers reused every cycle, which a production programme would treat as a confound; there is one phenotyping environment (`setPheno()`'s error variance, no genotype-by-environment term – Chapter 8. adds one); the population-split coalescent history (`runMacs2(split = )`) is a stylised divergence process, not calibrated to a named heterotic-pool pair; doubled-haploid production is instantaneous, skipping the field seasons Step 5 counted; and selection acts on trait 1 only – the 2 traits this chapter's config simulates are all carried through `X/hybrids`, but only the first ever enters the GCA ranking, so the rest are along for the ride, not selected on.

Scale. This chapter runs live at a book scale (12 + 12 founders, 750 markers, 4 cycles) – AlphaSimR is fast enough at this size that no caching is needed, unlike the MABC scans in Chapter 6.. A heavier, production-scale run is `inst/scripts/run_alphasimr.R`:

```
Rscript inst/scripts/run_alphasimr.R --test          # smoke run, ~10 s
Rscript inst/scripts/run_alphasimr.R --out=DIR       # full run, ~2 min
```

It runs `alphasimr_pipeline()` at a larger configuration, feeds the resulting hybrid pool through `stage2_select()`, and writes the cycle trajectory and selection metrics as CSVs.

7.10 Where this connects

Chapter 4. and Step 8 build the same three-object contract from two different generative models – the fast one and the realistic one – which is the point of drawing that seam where the book does: everything past it, `build_ctx()` through `stage2_select()`, cannot tell which one it is looking at. Chapter 10. explains why the dominance term here is what makes testcross selection informative at all. Chapter 3. supplies the breeder’s equation Step 6 checks and the index theory Step 7 uses; Chapter 20. is where Step 3’s `varA`-against-`genicVarA` gap becomes the subject rather than a diagnostic. Chapter 19. runs the same diversity-erosion argument at a longer horizon with the package’s own faster simulator; this chapter is its cross-check against an engine with real meiosis behind it.

Chapter 8. is the other half of the method: the same primitives assembled into an inbred-line programme with staged yield trials, where the question is not which hybrids to sell but how fast a variety pipeline can run and what it spends to run faster.

8. A complete plant breeding programme, and what it costs in diversity

Chapter 7. assembled AlphaSimR's primitives into a *hybrid* scheme: two heterotic pools, reciprocal testers, a commercial cross at the end. This chapter assembles the same primitives into the other kind of programme – an **inbred-line** pipeline, where the product is a variety rather than a hybrid, and where the organising problem is not which cross to sell but how many genotypes to carry through how many stages of field testing.

The scheme is the wheat programme of R. C. Gaynor *et al.* [13], reduced in scale. It is worth building because it is the shape most public and many commercial line-breeding programmes actually have, and because it exposes something the hybrid chapter cannot: **the trade between the number of genotypes tested and the precision with which each is measured**. A breeder has a fixed number of plots. Spending them on more entries at one location or fewer entries at eight locations is the same budget spent two ways, and the staged trial is the standard compromise.

Then the chapter does what the course this scheme comes from does not: it puts the book's diversity machinery over the result, and asks what the faster programme spends to be faster.

8.1 The scheme




Figure 8.1: The staged inbred-line programme. Entries fall by roughly an order of magnitude at each stage while replication – the number of locations each entry is scored in – rises. Parents for the next cycle are recycled from the late trials (solid) or, under genomic selection, from the preliminary trial as well (dashed).

Seven stages, and the logic connecting them is a single idea: **at every step, entries fall and replication rises**. Four hundred doubled-haploid lines are scored once, cheaply, in unreplicated headrows; sixty survive into a preliminary yield trial at one location; twelve into an advanced trial at four; four into an elite trial at eight. Selection accuracy at the last stage is high because the error variance has been divided by eight locations – not because the trait became more heritable.

That is what `reps` does in `setPheno()`, and it is the argument Chapter 7.'s Step 2 flagged as the plant idiom. The configuration:

```
sch <- list(
  n_founders = 40, n_chr = 10, n_qtl = 60, # founder genomes
  n_crosses = 20, n_dh = 20, fam_max = 5, # crossing and DH production
  n_pyt = 60, n_ayt = 12, n_eyt = 4, # entries per yield trial
  rep_hdrw = 1/4, rep_pyt = 1, rep_ayt = 4, rep_eyt = 8, # locations per stage
```

```
var_e = 0.4, gs_boost = 3, n_burn = 7, n_years = 15
)
```

`rep_hdrw = 1/4` is the course's device for a stage *less* precise than a single plot: headrows are scored visually on a short row, which carries less information than one replicated plot, so the error variance is multiplied by four. `fam_max` caps how many DH lines any one cross may contribute to the preliminary trial – `selectWithinFam()` from Chapter 7's Step 6, holding the genetic base open against a truncation that would otherwise take the whole trial from the two best crosses.

8.2 Why a breeding programme needs a burn-in

The loop below advances every stage by one year, simultaneously. That only means something if all seven stages are already *occupied*, by cohorts of different ages: this year's elite trial entered the programme six years ago, this year's headrows entered last year. Starting the loop from an empty programme would compare a year-1 elite trial that does not exist against a year-1 headrow that does.

The fix is the staggered fill in the course's vignette, and it is the most confusing seven lines in it:

```
for (year in 1:7) {
  F1 <- randCross(Parents, n_crosses)
  if (year < 7) DH <- makeDH(F1, n_dh)
  if (year < 6) HDRW <- setPheno(DH, reps = rep_hdrw)
  if (year < 5) PYT <- ...
  if (year < 4) AYT <- ...
  if (year < 3) EYT <- ...
}
```

Read it as a triangle rather than a loop. On the first pass every stage runs, so every stage holds a cohort. On the second pass the elite trial is skipped, so the cohort sitting in it is a year older than the one behind it. By the seventh pass only the crossing block still runs. What comes out is a programme in which each stage holds a cohort of the right age relative to its neighbours – a *steady state*, which is the only state in which a per-year genetic gain is a meaningful number.

8.3 One programme, fifteen years

The whole scheme is one function. It is written out rather than packaged because the loop is the thing being taught, and because a reader adapting it to their own programme will edit the loop, not call it.

```
gd_of <- function(pop, SP)
  1 - mean(molecular_coancestry(pullSegSiteGeno(pop, simParam = SP) / 2))

run_programme <- function(cfg, seed, genomic = FALSE, gs_from = 1) {
  set.seed(seed)
  fg <- quickHaplo(nInd = cfg$n_founders, nChr = cfg$n_chr,
    segSites = cfg$n_qtl, inbred = TRUE)
  SP <- SimParam$new(fg)
  SP$addTraitAG(nQtlPerChr = cfg$n_qtl, mean = 4, var = 0.1, varGxE = 0.2)
  ph <- function(p, r) setPheno(p, varE = cfg$var_e, reps = r, simParam = SP)

  # --- burn-in: fill every stage with a cohort of the right age ---
  Parents <- newPop(fg, simParam = SP)
  for (year in 1:cfg$n_burn) {
    F1 <- randCross(Parents, cfg$n_crosses, simParam = SP)
    if (year < 7) DH <- makeDH(F1, cfg$n_dh, simParam = SP)
    if (year < 6) HDRW <- ph(DH, cfg$rep_hdrw)
    if (year < 5) {
      PYT <- selectWithinFam(HDRW, nInd = cfg$fam_max, use = "pheno", simParam = SP)
      PYT <- ph(selectInd(PYT, nInd = cfg$n_pyt, use = "pheno", simParam = SP),
        cfg$rep_pyt)
    }
    if (year < 4)
```

```

    AYT <- ph(selectInd(PYT, nInd = cfg$n_ayt, use = "pheno", simParam = SP),
             cfg$rep_ayt)
  if (year < 3)
    EYT <- ph(selectInd(AYT, nInd = cfg$n_eyt, use = "pheno", simParam = SP),
             cfg$rep_eyt)
}

# --- advancing years, last stage first so each reads the previous cohort ---
out <- data.frame(year = 1:cfg$n_years, meanG = NA_real_, varG = NA_real_,
                 acc_pyt = NA_real_, gd_pyt = NA_real_, gd_parents = NA_real_)
for (year in 1:cfg$n_years) {
  gs <- genomic && year >= gs_from

  Variety <- selectInd(EYT, nInd = 1, use = "pheno", simParam = SP)
  EYT <- ph(selectInd(AYT, nInd = cfg$n_eyt, use = "pheno", simParam = SP),
           cfg$rep_eyt)

  if (gs) {
    # lever 1: a more accurate criterion at the PYT stage
    PYT@ebv <- setPheno(PYT, varE = cfg$var_e,
                      reps = cfg$rep_pyt * cfg$gs_boost,
                      onlyPheno = TRUE, simParam = SP)
    AYT <- ph(selectInd(PYT, nInd = cfg$n_ayt, use = "ebv", simParam = SP),
             cfg$rep_ayt)
  } else {
    AYT <- ph(selectInd(PYT, nInd = cfg$n_ayt, use = "pheno", simParam = SP),
             cfg$rep_ayt)
  }

  PYT <- selectWithinFam(HDRW, nInd = cfg$fam_max, use = "pheno", simParam = SP)
  PYT <- ph(selectInd(PYT, nInd = cfg$n_pyt, use = "pheno", simParam = SP),
           cfg$rep_pyt)
  HDRW <- ph(DH, cfg$rep_hdrw)
  DH <- makeDH(F1, cfg$n_dh, simParam = SP)

  if (gs) {
    # lever 2: recycle parents a year earlier, from the PYT
    PYT@ebv <- setPheno(PYT, varE = cfg$var_e,
                      reps = cfg$rep_pyt * cfg$gs_boost,
                      onlyPheno = TRUE, simParam = SP)
    EYT@ebv <- pheno(EYT); AYT@ebv <- pheno(AYT)
    Parents <- selectInd(c(EYT, AYT, PYT), nInd = cfg$n_eyt + cfg$n_ayt,
                       use = "ebv", simParam = SP)
  } else {
    Parents <- c(EYT, AYT)
  }
  F1 <- randCross(Parents, cfg$n_crosses, simParam = SP)

  out$meanG[year] <- meanG(PYT)
  out$varG[year] <- varG(PYT)
  out$acc_pyt[year] <- stats::cor(PYT@gv[, 1], PYT@pheno[, 1])
  out$gd_pyt[year] <- gd_of(PYT, SP)
  out$gd_parents[year] <- gd_of(Parents, SP)
}
out
}

```

Two details are worth pausing on.

The advancing loop runs backwards. The elite trial is filled *before* the advanced trial is refilled, which is before the preliminary trial is refilled. Written in the other order, EYT would be selected from an AYT that had already been overwritten with this year's cohort, and the programme would advance two stages per year without anyone noticing. Ordering is the whole correctness argument of a staged simulation.

Progress is reported at the PYT. It is the first stage with a yield phenotype and the last with enough entries for a variance to mean anything – `meanG` on four elite lines is noise. This is the course’s choice and it is the right one.

```
ps <- run_programme(sch, seed = 1, genomic = FALSE)
fmt(ps[c(1, 5, 10, 15), ])
```

	year	meanG	varG	acc_py	gd_py	gd_parents
1	1	4.1209	0.1227	0.4863	0.4647	0.4353
5	5	4.6397	0.0868	0.1795	0.3985	0.4018
10	10	4.9367	0.0772	0.4830	0.3899	0.3644
15	15	5.0549	0.1296	0.4062	0.3778	0.3524

Genetic mean rises, genetic variance falls, and gene diversity of the preliminary trial falls with it. Nothing here is surprising; what matters is that all three are now being measured on the same run.

i These numbers are reproducible

The founders come from `quickHaplo()`, which is seeded by `set.seed()`, so `run_programme(sch, seed = 1)` returns the same data frame every time – unlike Chapter 7’s Step 8, whose `runMacs2()` founders are not reproducible. That is what makes the paired comparison below legitimate: the phenotypic and genomic arms are run from the *same* seed, so they share founders, and every difference between them is attributable to the scheme.

8.4 Genomic selection: two levers, not one

Genomic selection is usually described as buying accuracy. It buys two things, and the second is normally the larger (R. C. Gaynor *et al.* [13]; T. H. E. Meuwissen, B. J. Hayes, and M. E. Goddard [14]):

1. **Accuracy earlier in the cycle.** A prediction is available on a line that has been genotyped but barely phenotyped, so the advance from preliminary to advanced trial can be made on better information than one location’s yield.
2. **A shorter generation interval.** Parents can be recycled from the preliminary trial instead of waiting for the elite trial – and the breeder’s equation of Chapter 3. is a rate *per unit time*, so cutting years off the cycle raises gain per year even at unchanged accuracy per selection.

This chapter mimics the prediction rather than fitting one: `ebv` is filled with a phenotype simulated at `gs_boost` times the replication, which is a stand-in for a genomic prediction of the corresponding accuracy. Chapter 14. is where the book actually fits prediction models and measures what accuracy is attainable; the honest statement here is that this chapter *assumes* an accuracy and studies its consequences, rather than earning it.

```
gs <- run_programme(sch, seed = 1, genomic = TRUE, gs_from = 6)
fmt(gs[c(1, 5, 10, 15), ])
```

	year	meanG	varG	acc_py	gd_py	gd_parents
1	1	4.1209	0.1227	0.4863	0.4647	0.4353
5	5	4.6397	0.0868	0.1795	0.3985	0.4018
10	10	4.9378	0.0663	0.3026	0.3969	0.3517
15	15	5.5676	0.0686	0.4497	0.3148	0.3038

Years 1 to 5 are identical to the phenotypic arm, digit for digit, because genomic selection only switches on in year 6. That is a matched control in the sense Chapter 20. uses the term: the two arms share a history, so the comparison is paired rather than between two independent runs.

8.5 The comparison, replicated

One run of one seed is an anecdote. Both arms are cheap enough to replicate.

```
seeds <- 1:24
reps <- do.call(rbind, lapply(seeds, function(s) {
  a <- run_programme(sch, s, genomic = FALSE)
  b <- run_programme(sch, s, genomic = TRUE, gs_from = 6)
  rate <- function(d) stats::coef(stats::lm(meanG ~ year,
                                             data = d[d$year >= 6, ])) [2]

  data.frame(seed = s,
             gain_ps = rate(a), gain_gs = rate(b),
             gd_ps = a$gd_pyt[nrow(a)], gd_gs = b$gd_pyt[nrow(b)],
             end_ps = a$meanG[nrow(a)], end_gs = b$meanG[nrow(b)])
}))

paired <- function(x, y) {
  d <- y - x
  c(mean_difference = mean(d), t = mean(d) / (stats::sd(d) / sqrt(length(d))))
}

fmt(rbind(`gain per year` = paired(reps$gain_ps, reps$gain_gs),
         `gene diversity, yr 15` = paired(reps$gd_ps, reps$gd_gs),
         `genetic mean, yr 15` = paired(reps$end_ps, reps$end_gs)))
```

	mean_difference	t
gain per year	0.0285	8.0851
gene diversity, yr 15	-0.0280	-2.8310
genetic mean, yr 15	0.2771	6.0496




Figure 8.2: Twenty-four paired runs, genomic selection switched on in year 6. Left: genetic mean of the preliminary yield trial. Right: gene diversity of the same population. The faster programme is also the one spending diversity faster.

Genomic selection raises the rate of gain by roughly half, and the paired *t*-statistic leaves no doubt about the sign. That is the result [R. C. Gaynor *et al.* \[13\]](#) reports and the reason two-part strategies were adopted.

The second row is the one this book exists for. **Gene diversity in the preliminary trial is significantly lower under genomic selection at year 15**, and it is lower for a reason that is structural rather than incidental: recycling parents from the preliminary trial means recycling them a year before the trial data could have distinguished full sibs, so the parents chosen are more closely related than the

ones the phenotypic arm would have chosen. A shorter generation interval is a higher rate of coancestry accumulation per unit time, not merely per cycle.

8.6 The diversity ledger

The measurement above used gene diversity of the preliminary trial. The parent set is the population that actually transmits, and it is the one an optimal contribution scheme would act on:

```
ledger <- data.frame(
  year      = yr,
  gd_pyt_ps = avg("ps", "gd_pyt"), gd_pyt_gs = avg("gs", "gd_pyt"),
  gd_par_ps = avg("ps", "gd_parents"), gd_par_gs = avg("gs", "gd_parents")
)
fmt(ledger[c(1, 5, 10, 15), ])
```

	year	gd_pyt_ps	gd_pyt_gs	gd_par_ps	gd_par_gs
1	1	0.4681	0.4681	0.4418	0.4418
5	5	0.4243	0.4243	0.4158	0.4158
10	10	0.3901	0.3891	0.3812	0.3758
15	15	0.3684	0.3404	0.3628	0.3327

```
loss <- function(v) (ledger[[v]][1] - ledger[[v]][nrow(ledger)]) / ledger[[v]][1]
fmt(c(`PYT, phenotypic` = loss("gd_pyt_ps"), `PYT, genomic` = loss("gd_pyt_gs"),
      `parents, phenotypic` = loss("gd_par_ps"), `parents, genomic` = loss("gd_par_gs")))
#>      PYT, phenotypic      PYT, genomic parents, phenotypic parents, genomic
#>      0.2131          0.2729          0.1788          0.2470
```

Those four numbers are exactly the quantity this book calls **alpha**: the proportional loss of gene diversity, here measured against the programme's own starting point rather than against a random-selection baseline of the same size (Section 11.9 is where that distinction is made, and it matters – a same-size random baseline would attribute part of this loss to drift rather than to selection).

i Status number, and why it is not the headline here

`status_number()` on a set of fully inbred lines returns a value near or below 1, because a homozygous line has self-coancestry $f_{ii} = 1$ and the diagonal dominates the mean. That is the same arithmetic Chapter 5. explains for a biparental cross, and it is correct rather than broken – but it makes **N_s** an unhelpful headline for a line-breeding programme. Gene diversity is the metric that stays interpretable across both inbred and outbred populations, which is why the ledger is written in it.

What would an **alpha** budget mean here? Not a constraint on which hybrids to sell, as in Chapter 16., but a constraint on **which parents to recycle**: the same optimal-contribution problem of Chapter 15., applied to the **c(EYT, AYT, PYT)** candidate set that the genomic arm selects from. The candidate set is small, the coancestry matrix is cheap, and the choice being constrained is annual. Chapter 19. is the argument for why that constraint pays back, and Chapter 20. is where the payback is predicted rather than observed.

8.7 Assumptions

What this models. Staged testing with replication rising across stages, genotype-by-environment interaction as a variance component (`SP$addTraitAG(varGxE = )`), doubled-haploid line development, within-family caps, a steady-state programme reached through a burn-in, and two distinct genomic-selection levers separated so their effects can be told apart.

What it does not. Genomic prediction is *mimicked* by adding replication, not fitted – so the accuracy is assumed rather than earned, and the training population, its size, and its relationship to the selection candidates play no role (Chapter 14. is where they do). There is one trait, so the index questions of Chapter 7.'s Step 7 do not arise. Locations differ only through a variance component: there is no correlated environment structure, no year effect, and no location-by-year term, so a "location" here is an independent

draw rather than a place. Doubled-haploid production is instantaneous and never fails. Founders come from `quickHaplo()`, so they carry no linkage disequilibrium – Chapter 7.’s Step 1 measures what that costs and argues it is the right trade for a chapter about response rather than about marker-QTL association.

Scale. 40 founders, 600 loci, 15 years, 24 paired seeds, all live at render. [R. C. Gaynor *et al.* \[13\]](#)’s programme is an order of magnitude larger in every dimension; the qualitative results here are the ones that survive the reduction, and the absolute rates are not comparable to a real wheat programme.

8.8 Where this connects

Chapter 7. supplies every primitive this chapter uses and builds them into the hybrid scheme instead; the two chapters are the same toolkit pointed at the two kinds of programme. Chapter 3. is where the breeder’s equation says why a shorter generation interval raises gain per year. Chapter 14. is where the accuracy this chapter assumes is actually measured, and Chapter 9. is where marker-versus-QTL, the reason that accuracy is finite, is taken seriously.

The diversity result belongs to the book’s spine. Chapter 19. shows the same tension – faster response, faster erosion – with the package’s own simulator over a longer horizon; Chapter 20. explains why the erosion is what bounds the programme and when it can be predicted; Chapter 15. is the machinery for spending the budget deliberately rather than discovering afterwards what was spent.

III

	exactly	98
10.7	Divergence is not a bonus	98
10.8	What actually buys heterosis	101
10.9	Reading the S1 panel again	103
10.10	General and specific combining ability	103
10.11	Divergence destroys SCA faster than it builds GCA	106
10.12	What this does not change	109
10.13	Where this connects	110
11	A catalogue of metrics on f	111
11.1	Group coancestry, and the off-diagonal problem	111
11.2	Gene diversity: θ and H_e are the same quantity	112
11.3	Status number N_s	112
11.4	Effective number of parents	113
11.5	Allelic richness	113
11.6	Distances: E-NE and A-NE	114
11.7	Effective dimensionality	114
11.8	Decomposition by heterotic group	115
11.9	The correct baseline	116
11.10	Every metric against the null	117
12	Choosing metrics: collapse, cost, redundancy	119
12.1	Finding 1: the hybrid coancestry matrix is never needed	119
12.2	Finding 2: what each metric costs	121
12.3	Finding 3: the metric space is nearly one- dimensional	122
12.4	Finding 4: the encoding, not the constraint handler, decides whether DE works	125
12.5	Finding 5: the index is not the truth, and its error is not white	126
12.6	The recommendation, stated	128
12.7	Drop-in code	128
13	Metrics at each pipeline stage, and five traps	131
13.1	The kernel: one matrix, all four stages	131
13.2	S1: conserving the heterotic groups	132
13.3	S2: choosing crosses for DH induction	133
13.4	S3: selecting lines per se	135
13.5	S4: hybrid mining	137
13.6	Where the diversity actually goes	137
13.7	Cost summary	139
13.8	Checklist	139
14	Predicting breeding values: GBLUP, BayesA, BayesB	141
14.1	A training population with real noise	141
14.2	The truth this chapter is predicting	141
14.3	GBLUP: prediction from the relationship matrix	142
14.4	What the random split was measuring	144
14.5	BayesA and BayesB: prediction from marker effects	146
14.6	Verification: two routes agree	148
14.7	Which model wins, and when	148
14.8	From GEBVs to <code>stage1_build()</code>	150
14.9	Back to Finding 5: a measured accuracy, not an assumed one	151

9. Diversity metrics in the genomic era

In classical theory, inbreeding measured as **loss of heterozygosity** and inbreeding measured as **drift of allele frequencies** are the same thing. T. H. E. Meuwissen, A. K. Sonesson, G. Gebregiorgis, and J. A. Woolliams [6] show that inside genomic optimal-contribution schemes they stop being the same thing, and can diverge badly, whenever diversity is controlled through **identity by state** on a marker panel. Only when control is through **identity by descent** is the classical equivalence restored.

The practical consequence is uncomfortable:

You hit exactly the target you wrote into the optimiser's constraint, and you pay an unmonitored price in the metric you did not write.

For a plant breeder: if you constrain $\frac{1}{2}\mathbf{c}'\mathbf{G}\mathbf{c}$ using a VanRaden $\mathbf{G}$, you are constraining **drift**. The loss of heterozygosity in your germplasm may be running at twice the rate you believe you are controlling.

9.1 Three lenses, not three synonyms

Lens	What the metric actually sees	Inbreeding metric	Relationship matrix
Drift	squared allele-frequency change relative to a base	F_drift	G_VR2, G_VR1, G_i(p)
Ho-mozygosity	proportion of heterozygosity remaining	F_hom	G_0.5 (molecular coancestry)
IBD	segments inherited by descent from a defined base	F_IBD	A, G_LA
Hybrid	haplotype homozygosity used as a proxy for IBD	F_ROH	G_ROH

! Every inbreeding metric is a ratio to a reference population

Without explicitly declaring your base – cycle 0 of the recurrent programme? a landrace panel? the founding heterotic pool? – none of the numbers below mean anything. This is the commonest error in plant applications: using the *current* generation's frequencies as the base, which zeroes the historical memory of erosion.

9.2 The metrics, one at a time

9.2.1 Rate of inbreeding and effective size

What you manage is not F but its rate of accumulation. F is a stock; ΔF is a flow.

$$\Delta F = \frac{F_t - F_{t-1}}{1 - F_{t-1}}, \quad N_e = \frac{1}{2\Delta F} \quad (9.1)$$

Since $1 - F_t = (1 - \Delta F)^t$, a regression of $\log(1 - F_t)$ on t has slope $\approx -\Delta F$. **The linearity of that regression is itself a diagnostic:** in the paper, $\log(1 - F_{drift})$ was approximately linear under every

scheme, while $\log(1 - F_{hom})$ showed marked curvature under some, notably the ROH-based one. Curvature means a non-constant rate – your N_e target is not being met steadily.

A tropical maize recurrent programme targeting $N_e = 50$ per cycle needs $\Delta F = 0.01$. Measuring F only at cycle 8 and dividing by 8 hides the possibility that cycles 1-3 were conservative and cycles 6-8 collapsed the base. Always fit the log-linear regression and look at the residual.

9.2.2 F_{hom} , homozygosity-based

Wright defined an inbreeding coefficient to run from 0 to 1 as heterozygosity under random mating falls from its initial state to zero. So the metric is the fraction of heterozygosity already lost relative to the base:

$$F_{hom} = 1 - \frac{1}{N_{SNP}} \sum_k \frac{2p_{t,k}(1 - p_{t,k})}{2p_{0,k}(1 - p_{0,k})} = 1 - \frac{1}{N_{SNP}} \sum_k \frac{H_{t,k}}{H_{0,k}} \quad (9.2)$$

It can be negative. If management pushes frequencies toward 0.5, heterozygosity rises above the base and the “inbreeding coefficient” goes below zero. That is not a numerical error: it is the expected behaviour of any scheme that maximises heterozygosity on the monitored panel.

It protects against inbreeding depression and the expression of deleterious recessives in homozygosis. It is the right lens when the dominant risk is *vigour*.

In maize, heterozygosity is the product itself. Within a heterotic pool under reciprocal recurrent selection, F_{hom} computed **within pool** anticipates the decline in line-per-se vigour; it says nothing about between-pool divergence, which is what generates heterosis in the hybrid. In selfing species, F_{hom} measured on fixed lines is ≈ 1 by construction and is **useless for management** – it has to be computed on allele frequencies in the crossing population, not on the homozygosity of individuals.

9.2.3 F_{drift} , drift-based

$$F_{drift} = \frac{1}{N_{SNP}} \sum_k \frac{(p_{t,k} - p_{0,k})^2}{p_{0,k}(1 - p_{0,k})} \quad (9.3)$$

Never negative. It is formally analogous to Holsinger & Weir’s F_{ST} applied to *one population through time* rather than to a set of populations in space. It is exactly the quantity constrained when G_{VR2} goes into the optimiser.

It protects against random change in phenotypic value for traits **outside** the selection objective, against rare deleterious recessives rising to dangerous frequencies, and against irreversible loss of rare alleles.

There is also a selection-intensity reading: the term $\delta p^2 / [p_0(1 - p_0)]$ approximates the **squared total intensity** applied at that marker, with $i \approx \delta p / \sqrt{p_0(1 - p_0)}$. Constraining F_{drift} constrains how much effective selection the whole genome underwent – including at loci you never meant to select.

A wheat programme running six cycles of genomic selection with a VanRaden G and reporting $\Delta F \approx 1\%$ has reported ΔF_{drift} on the training panel. Quantitative disease-resistance loci not yet prioritised may have hitch-hiked to extreme frequencies alongside yield QTL, and none of that appears in the reported metric.

9.3 The classical equivalence, and the covariance that destroys it

Under purely random drift, with $E[\delta p_t | p_0] = 0$ for every p_0 ,

$$\frac{\text{var}(p_{t,k})}{p_0(1 - p_0)} = 1 - \frac{H_t}{H_0} \Rightarrow F_{drift} = F_{hom}. \quad (9.4)$$

The paper’s central result quantifies the departure:

$$F_{hom} - F_{drift} = 2 \text{cov} \left(\frac{\delta p_{t,k}}{\sqrt{p_{0,k}(1 - p_{0,k})}}, \frac{p_{0,k} - 1/2}{\sqrt{p_{0,k}(1 - p_{0,k})}} \right) \quad (9.5)$$

In words: **if the direction of frequency change is correlated with the starting frequency, the two measures diverge.** The non-obvious part is that this covariance is induced by the *diversity management*

itself, not by directional selection – the paper demonstrates this by running the scheme with random GEBVs, and the divergence persists.

9.3.1 Verifying the identity

`freq_metrics()` returns all three quantities in one marker sweep. The identity is asserted in the test suite; here it is on live data:

```
set.seed(31)
idx <- sample.int(ctx$N, 60)
fm <- freq_metrics(idx, ctx)
fmt(fm[c("F_hom", "F_drift", "cov_diag")], 6)
#>      F_hom F_drift cov_diag
#> 0.011458 0.005192 0.006267

c(F_hom_minus_F_drift = unname(fm[["F_hom"]] - fm[["F_drift"]]),
  cov_diag             = unname(fm[["cov_diag"]]),
  residual              = unname(fm[["cov_diag"]] - (fm[["F_hom"]] - fm[["F_drift"]]))
#> F_hom_minus_F_drift      cov_diag      residual
#>      6.266505e-03      6.266505e-03      8.673617e-18
```

Exact to machine precision. And when the selection is the base, both lenses must read zero:

```
fmt(freq_metrics(seq_len(ctx$N), ctx)[c("F_hom", "F_drift", "cov_diag")], 12)
#>      F_hom F_drift cov_diag
#>      0      0      0
```

9.3.2 Why each matrix generates a sign

Ma-Covari- trixance	Mechanism
G_VR2	Positive: the constraint penalises Δp^2 , so moving an allele to the <i>far</i> extreme costs more than to the near one; the optimiser pushes rare alleles to 0 and common ones to 1
G_0.5	Negative: frequencies are measured as deviations from 0.5, so the optimiser buys ‘diversity’ by pushing everything to intermediate frequencies
G_VR2 Positive, at- tenu- ated	G_VR2 reweighted by $2p_0(1-p_0)$, giving more weight to already-intermediate loci and less room at the extremes
G_i(p)	Negative: uses accumulated intensity instead of Δp^2 ; since $di/dp = [p(1-p)]^{-1/2}$, moves toward the extremes get more expensive

⚠ Your MAF spectrum decides how bad this is

The magnitude depends on the minor-allele-frequency spectrum. The effect is dramatic with sequence data dominated by rare alleles, and much milder with SNP chips designed for intermediate MAF.

Directly translated to plants: GBS or skim-seq with the full MAF spectrum puts you in the severe scenario; a 20-50k chip filtered to $MAF > 0.05$ puts you in the mild one. This is the parameter that decides whether you can live with G_VR2.

9.3.3 The signature, on live data

The pure max-diversity greedy is the G_0.5 scheme in miniature: it constrains molecular coancestry and nothing else. It should drive `F_hom` negative while paying the largest drift bill.

```

panel <- rbind(
  random      = freq_metrics(sample.int(ctx$N, 60), ctx),
  max_div     = freq_metrics(sel_greedy(ctx, 60, w = 0), ctx),
  balanced    = freq_metrics(sel_greedy(ctx, 60, w = 1), ctx),
  truncation  = freq_metrics(sel_truncation(ctx, 60), ctx)
)
fmt(as.data.frame(panel[, c("He", "F_hom", "F_drift", "cov_diag", "rare_retained")]),
4)

```

	He	F_hom	F_drift	cov_diag	rare_retained
random	0.3257	0.0048	0.0083	-0.0034	0.9286
max_div	0.3302	-0.0113	0.0060	-0.0174	0.9481
balanced	0.3124	0.0426	0.0558	-0.0131	0.7338
truncation	0.3124	0.0426	0.0558	-0.0131	0.7338

Read the **F_hom** and **F_drift** columns together. Maximising diversity on the homozygosity lens produces a *negative* **F_hom** – the selection is more heterozygous than the base – while **F_drift** is strictly positive and larger than under random sampling. The bill is real, it is being paid, and a report carrying only **F_hom** would not show it.

9.4 The other matrices, briefly

F_IBD via G_LA. Instead of assuming a 50/50 transmission probability for each parental allele (as the pedigree matrix **A** does), linkage analysis over the markers infers which parental haplotype was actually transmitted along the chromosome. It requires pedigree **and** markers. It measures IBD relative to a base explicitly defined as unrelated and non-inbred, so frequency change is determined by the base population rather than by linkage disequilibrium generated during selection – and the problematic covariance disappears. It was the **only scheme tested that met both the ΔF_{hom} and ΔF_{drift} targets**, on the managed panel and at unmonitored neutral loci, and it converted inbreeding into gain most efficiently.

For plants this is the point of greatest friction. In biparental populations and backcross families **G_LA** is natural and cheap – phase is known and the number of meioses since the base is small. In recurrent-selection populations with full pedigree (common in commercial maize) it is feasible. In germplasm panels *without* pedigree it is inapplicable, and the honest alternative is segment-based IBD.

F_ROH. A run of homozygosity is an uninterrupted stretch of homozygous markers; the logic is that many consecutive markers are unlikely to be IBS by chance, so the stretch is probably IBD. The fundamental ambiguity is that **F_ROH** is IBD when runs are **long** (recent inbreeding) and homozygosity when they are **short** (ancient inbreeding, accumulated IBS). The minimum-length threshold is therefore a knob that selects the *age* of inbreeding you are measuring, and there is no well-defined reference base.

Empirically, **G_ROH** behaved more like **G_0.5** than like an IBD measure. It also is **not guaranteed positive semi-definite** – its elements are computed pairwise – and required substantial diagonal additions for inversion, which made it diagonally dominant and artificially inflated the number of parents selected.

In a soybean inbred the entire genome is one ROH: **F_ROH** $\approx$ 1 and the metric saturates. In maize it is useful for diagnosing erosion of elite heterotic pools; for *management* inside an optimiser, check positive definiteness carefully.

9.4.1 Side by side

Matrix	Class	Behaviour in the optimiser
A	IBD (expectation)	high gain, but exceeded targets by ~40% by not accounting for within-family LD
G_VR2drift		meets the drift target <i>only on the managed panel</i> ; F_hom runs away
G_VR1drift		reweighted version; smaller bias, smaller gain

Matrix	Class	Behaviour in the optimiser
G_0.5	homozygosity	F_hom near zero, drift 4x the target; loses genetic variance early
G_i(p)	drift	inverts the sign of the bias; constraint only indirectly tied to F
G_LA	IBD	the only one meeting both targets; best gain per unit of inbreeding
G_ROH	Hybrid	high gain, uncontrolled drift, numerical problems

i An implementation note that is routinely ignored

G_VR1, G_VR2, G_0.5 and G_i(p) are only **semi**-definite positive (a zero eigenvalue from the centering). You must add α to the diagonal ($\alpha = 0.01$ in the paper) to invert them and for the optimal-contribution algorithm to work. A and G_LA are positive definite by construction.

9.5 Reference results

Target $\Delta F = 0.005$ per generation ($N_e = 100$). Values at the **unmonitored neutral loci**, generation 20.

Scheme	Class	dF_hom	dF_drift	gap	gain
G_VR2(M,M)	drift	0.0103	0.0068	0.054	7.124
G_VR2(M,D)	drift	0.0101	0.0068	0.050	7.107
G_VR1(M,M)	drift	0.0080	0.0069	0.021	6.680
G_i(p)(M,M)	drift	0.0065	0.0077	-0.031	7.111
G_0.5(M,M)	homozygosity	0.0073	0.0176	-0.170	6.734
G_ROH(M,M)	homozygosity	0.0054	0.0088	-0.070	9.099
G_LA(M,M)	IBD	0.0043	0.0049	-0.010	7.188
A(M,~)	IBD	0.0070	0.0084	-0.021	9.890

Four readings:

1. **No IBS scheme met both targets.** G_LA did, within about 2%.
2. Pedigree A gave the highest absolute gain but overshoot by ~40% – neutral loci hitch-hiking with QTL through within-family LD, which the pedigree expectation of IBD cannot capture.
3. Using a separate panel for management instead of the GEBV panel changed gain by **less than 0.3%**. Not worth the complexity.
4. At equal mean inbreeding, G_LA beat A in 65 of 100 replicates and G_ROH in 62 of 100.

9.5.1 The diagnostic nobody runs

Schemes based on IBD selected many parents; VanRaden-style schemes selected few.

Rule of thumb. If your optimal-contribution run is delivering high gain with a suspiciously small number of parents and ΔF on target, compute F_hom and F_drift separately on a panel **not used** in the management. The difference between them is your hidden bill.

The book's implementation of that diagnostic is `ne_parents()`, reported alongside the two lenses in every scenario – see Chapter 21., where a selection meeting its diversity target while collapsing from 40 effective parents to 12 is visible in one row of the output.

9.6 What to carry into plant breeding

Declare the base and never redefine it. Redefining the reference frequencies every generation is common in plant pipelines because it is the *default* in several **G**-construction functions. Since frequency changes in successive cycles are positively correlated, constraining the variance *within* each cycle lets the *accumulated* variance overshoot the target. Freeze p_0 from the founding cycle and propagate it. In this book that is `ctx$p0`, set once in `build_ctx()` and never recomputed.

Do not maximise marker heterozygosity as an objective. Raising heterozygosity on the monitored panel does not raise it at unmonitored loci – and it increases real IBD, because raising an allele’s frequency means multiplying copies of a subset of base alleles.

Heterotic pools need a two-directional constraint. Within pool you want to limit ΔF . Between pools you want to *preserve* divergence, which is directed drift and is desirable. A single global $\frac{1}{2}\mathbf{c}'\mathbf{G}\mathbf{c}$ cannot express that. Use multiple matrices with simultaneous constraints, or per-subpopulation constraints – which is what `gd_partition()` measures and Chapter 13. operationalises.

Selfing species need inbreeding by selfing separated from inbreeding by drift. In soybean, wheat or rice, individual F is ≈ 1 by design and carries no management information. The manageable quantity is allele frequency **in the crossing population** and coancestry among the chosen parents. Apply the whole $F_{\text{hom}}/F_{\text{drift}}$ framework at population level, never at line level.

Region-specific targets rarely earn their complexity. Causal variants are well distributed and what happens at a subset of loci does not predict what happens outside it. The legitimate plant exception is R-gene-type resistance loci, where allelic diversity has direct and known value.

Deleterious recessives are a load problem, not a diversity problem. If you have mapped the defect, putting it in the selection index is more efficient than controlling it indirectly through a diversity constraint.

In pure conservation schemes the problem is worse, not better. Without directional selection the inbreeding-management term dominates the optimiser and the covariance tends to be larger. The IBD recommendation is stronger there.

9.7 Limits of the animal-to-plant transposition

Assumption in the paper	Situation in plants
discrete generations, two sexes, one contribution per individual	hermaphroditism, selfing, cloning, overlapping cycles, repeated use of the same parent
reliable, deep pedigree	often truncated, or absent in genebanks
base = simulated founder population at mutation-drift equilibrium	base is an administrative choice (cycle 0) with pre-existing structure
target $N_e = 100$ over 20 generations	shorter cycles with GS, but effective elite N_e often below 30
one population, one target	multiple target environments, G x E, separate heterotic pools
heterozygosity is always desirable	in selfers transitory; in hybrids the final product; in lines undesirable

None of this invalidates the central conclusion. It is an **algebraic** property of the optimiser, not a biological property of the species. What it changes is *which matrix is practicable*.

9.8 Operational checklist

- Is the base p_0 frozen at the founding cycle, and documented?
- Are F_{hom} and F_{drift} computed **separately** and reported together?
- Is there a panel of loci **not used in management**, for audit?
- Is the regression of $\log(1 - F)$ on cycle linear (constant rate)?
- Is the covariance $2\text{cov}(\delta p/s, (p_0 - \frac{1}{2})/s)$ reported as a diagnostic?
- Is the number of selected parents biologically plausible?
- Is true genetic variance (or the best estimate of it) being tracked?

- Is $\mathbf{G}$ positive definite before inversion, and is the added α documented?
- Where possible, is an IBD alternative compared in parallel?

10. Heterosis, dominance, and what divergence buys

Since Chapter 2, this book has repeated a claim without proving it: that the divergence between the heterotic pools is the engine of heterosis, and is therefore to be maintained rather than minimised. Chapter 11. will decompose diversity on that basis and Chapter 13. will build a monitoring panel around it. The claim has carried a lot of weight for an assertion.

This chapter derives it, and then corrects it. The derivation is short and the identity is exact. The correction is that the usual reading of the identity – *push the pools apart and you will get more heterosis* – is false: under neutral divergence expected heterosis does not move at all, and the simulation shows it.

The chapter then goes one step past the mean. Expected heterosis is what a pool pair is worth on average; what a breeder selects on is how heterosis is *distributed* across the crosses of the factorial. Those are different questions with different answers, and the second one is where general and specific combining ability come from.

Three premises carry everything below. They are stated here once, because every result in the chapter uses all three and two of them are restrictive.

Parents are fully inbred. Every parent line is homozygous at every locus, so an F1's marker row is the average of its two parents' rows and a locus is heterozygous in the F1 exactly when the parents differ. This is what makes the additive effects cancel from the mid-parent difference and what makes GD_WI zero. Nothing here transfers to a programme crossing non-inbred parents without re-derivation.

The genotypic model is dominance within loci and additivity between them. There is no epistatic term anywhere below: the genotypic value is a sum over loci, H_i is a sum of per-locus indicators, and the combining-ability expansion of Section 10.10 is bilinear only because no term couples two loci. [M. R. Labroo, A. J. Studer, and J. E. Rutkoski \[15\]](#) review epistasis alongside dominance as a proposed mechanism for heterosis; this model excludes it, and no result below should be read as evidence about it. What d_k does absorb is every *within-locus* story, as the next section sets out.

The marker panel is the set of dominance loci. d_k is indexed by the same k as the markers, so every locus carrying dominance is genotyped and the frequency statistics are computed on exactly the loci that generate heterosis. This is the strongest assumption of the three, and it is the one that carries the distance results of Section 10.6. A real panel tags causal loci through linkage disequilibrium rather than containing them, which weakens every statement below about markers predicting heterosis but does not change the algebra, since the algebra is written over whatever set k indexes.

10.1 The dominance model, minimally

10.1.1 The one-locus scale

The classical parameterisation puts the origin at the midpoint of the two homozygotes: one homozygote is worth $+a$, the other $-a$, and the heterozygote d . Under random mating at frequency p , with $q = 1 - p$, the population mean is [\[16, Eq. 7.2\]](#)

$$M = a(p - q) + 2pqd \quad (10.1)$$

The first term is the additive contribution, and the mid-parent of any cross already carries it. The second is the dominance contribution. It is the only part of M that depends on how many heterozygotes there are, and therefore the only part a cross can add over its own mid-parent.

i a is a scale, not a quantity this chapter computes

The symbols in that display are the textbook ones. This book's additive marker effects are β_k (Chapter 14.), and its α is the diversity-loss statistic of Chapter 11., not the classical average effect of a substitution. The per-locus a_k reappears three times below and never as a result: in the genotypic-value model used to show that it cancels, in the definition of overdominance as $|d| > |a|$, and in the average effect of a substitution $a + d(q - p)$. No quantity this chapter reports depends on it.

10.1.2 What d_k is, and what it is not

d_k is *functional* dominance: the departure of the heterozygote from the homozygote midpoint at locus k . It is a property of the locus. It does not move when allele frequencies move.

Fisher's *statistical* dominance deviation is a different object: the residual left after projecting genotypic values onto the best additive model **in a particular population**. It is a property of the population as much as of the locus, and it changes when frequencies change.

Both appear in this chapter, and keeping them apart is the whole trick. d_k is the input. The specific combining ability of Section 10.10 is the statistical dominance deviation that a given d_k *induces* over a given pair of pools. One is a constant; the other is an answer that depends on where the pools happen to sit.

The distinction also buys the model what generality it has, and marks where that generality stops. Dominance, overdominance (a heterozygote outside both homozygotes, $|d| > |a|$) and pseudo-overdominance (two dominant favourable alleles in repulsion phase, too close for any marker to separate) are three different biological stories, and [M. R. Labroo, A. J. Studer, and J. E. Rutkoski \[15\]](#) review the evidence for each. At marker resolution they are the same arithmetic: a locus at which the F1 gains something the mid-parent does not. A single d_k absorbs all three, and nothing below can distinguish them – nor needs to. Epistasis is the fourth story on [M. R. Labroo, A. J. Studer, and J. E. Rutkoski \[15\]](#)'s list and d_k does not absorb it, because a single per-locus number cannot represent a term that depends on two loci at once. That is the second premise above, and it is an exclusion rather than an approximation.

10.1.3 The mid-parent difference

Code a genotype by $x_k \in \{0, 0.5, 1\}$, the frequency of the reference allele within the individual, and take the genotypic value to be the additive part plus the dominance deviation at every heterozygous locus:

$$G(\mathbf{x}) = \sum_k a_k (2x_k - 1) + \sum_k d_k \mathbb{1}\{x_k = 0.5\} \quad (10.2)$$

The parents are fully inbred, so each is homozygous at every locus and the F1's row is the average of its two parents' rows. The additive part is linear in $\mathbf{x}$, so the two parents' additive values average exactly to the F1's own and cancel from the mid-parent difference. What survives is the dominance deviation at the loci where the F1 is heterozygous:

$$H_i = \sum_k d_k \mathbb{1}\{x_{ik} = 0.5\} \quad (10.3)$$

The additive effects never appear. That is the result, not an omission, and it is why `heterosis_value()` takes no additive-effect argument:

```
heterosis_value
#> function (X, d)
#> {
#>   stopifnot(length(d) == ncol(X))
#>   drop((X == 0.5) %*% d)
#> }
#> <bytecode: 0xc10846d60>
#> <environment: namespace:hybdiv>
```

Against the literal oracle of Appendix C., which loops over loci and compares the two parents' genotypes one at a time:

```
max(abs(heterosis_value(ctx$X[1:25, ], d) -
      ref_heterosis(L, head(ctx$ped, 25), d)))
#> [1] 0
```

The cancellation is worth seeing directly. Draw additive effects at random, score a real F1 and its two parents under $G(\mathbf{x})$, and subtract:

```
a <- rnorm(ctx$m, 0, 0.2)
val <- function(x) sum(a * (2 * x - 1)) + sum(d * (x == 0.5))
i <- 7
mid <- (val(L[ctx$ped$a[i], ], ) / 2) + val(L[ctx$ped$b[i], ], ) / 2
c(mid_parent_difference = val(ctx$X[i, ]) - mid,
  heterosis_value = heterosis_value(ctx$X[i, ], drop = FALSE, d))
#> mid_parent_difference      heterosis_value
#>                362.3386                362.3386
```

The additive effects entered the first number and cancelled out of it; they never entered the second.

10.2 Expected heterosis over a pool pair

Write p_{Ak} and p_{Bk} for the reference-allele frequencies at locus k in the two pools, each computed over that pool's lines. A locus is heterozygous in the F1 when the two gametes carry different alleles, which for a cross drawn at random between the pools happens with probability $p_A(1 - p_B) + (1 - p_A)p_B$. The expected heterosis over all $A \times B$ crosses therefore follows from the pool frequencies alone:

$$\mathbb{E}[H] = \sum_k d_k [p_{Ak} + p_{Bk} - 2p_{Ak}p_{Bk}] \quad (10.4)$$

The candidate pool here *is* the full factorial, every $A \times B$ pairing exactly once, and the frequencies are the means of those same lines. The closed form is then the mean of the realised values as an identity rather than an approximation, and the residual is at machine precision:

```
abs(heterosis_expected(pA, pB, d) - mean(heterosis_value(ctx$X, d)))
#> [1] 1.136868e-13
```

That is a mean over $N = 625$ crosses, and a mean is not a ranking. Hold on to the question of how the individual crosses are spread around it; Section 10.10 answers it exactly.

10.3 The identity

Now split the bracket. Writing $\bar{p} = (p_A + p_B)/2$ and $y = p_A - p_B$, the per-locus F1 heterozygosity separates exactly:

$$p_A + p_B - 2p_Ap_B = \underbrace{2\bar{p}(1 - \bar{p})}_{\text{shared}} + \underbrace{y^2/2}_{\text{divergence}} \quad (10.5)$$

The first term is the heterozygosity a single merged pool at frequency $\bar{p}$ would produce under random mating. The second is what keeping the pools apart adds on top of it. Average both over loci with $d_k \equiv 1$, so that the weighting does not obscure the algebra, and they are exactly the GD_T and GD_{BS} of the Caballero & Toro partition that Chapter 2. `derived` and `gd_partition()` computes:

```
hp <- heterosis_partition(pA, pB, rep(1, ctx$m)) / ctx$m
gp <- gd_partition(seq_len(ctx$N), ctx)
fmt(data.frame(term      = c("shared", "divergence"),
               heterosis = c(hp[["shared"]], hp[["divergence"]]),
               partition = c(gp[["GD_T"]],   gp[["GD_BS"]]),
               residual  = c(abs(hp[["shared"]] - gp[["GD_T"]]),
                             abs(hp[["divergence"]] - gp[["GD_BS"]]))), 12)
```

term	heterosis	partition	residual
shared	0.3280928	0.3280928	0
divergence	0.0335968	0.0335968	0

GD_{BS} weighted by dominance deviations is the divergence half of expected heterosis. “Between-pool divergence is the heterosis engine” is an identity, not a slogan.

One thing to be clear about before reading on. GD_T is the gene diversity of the *merged* population, and by Chapter 2, it already contains GD_{BS} : with inbred lines, $GD_T = GD_{BI} + GD_{BS}$. Expected F1 heterozygosity is therefore $GD_{BI} + 2GD_{BS}$, the between-pool component counted twice – once inside the shared term, once as the divergence term. That factor of two is where the classical result sits. Falconer & Mackay’s heterosis of a cross between two *random-mating* populations is $\sum_k d_k y_k^2$ [16, Ch. 14], exactly twice the divergence term here. The two do not conflict; they use different mid-parents. There is two random-mating populations, each carrying its own within-pool heterozygosity, so their H is the total above $\sum_k d_k (p_{Ak}q_{Ak} + p_{Bk}q_{Bk})$. Ours is two inbred lines carrying none, so the whole F1 heterozygosity is heterosis, and the shared term is the mean of the two pools’ random-mating heterozygosities plus the half of $\sum_k d_k y_k^2$ that GD_T already holds.

10.3.1 A third form: one minus the between-pool coancestry

The sum of the two components has a name of its own. $\theta_{AB} = \mathbf{w}'_A \mathbf{f}^L \mathbf{w}_B$ is the between-pool coancestry of Chapter 11.: on the full factorial, where the line weights are uniform, it is the probability that an allele drawn from a pool-A line and an allele drawn from a pool-B line are identical in state. Its complement is the whole bracket:

$$GD_T + GD_{BS} = 1 - \theta_{AB} \quad (10.6)$$

```
tp <- theta_pools(seq_len(ctx$N), ctx)
c(shared_plus_divergence = gp[["GD_T"]] + gp[["GD_BS"]],
  one_minus_theta_AB     = 1 - tp[["theta_AB"]])
#> shared_plus_divergence    one_minus_theta_AB
#>          0.3616896          0.3616896
```

This one is algebra rather than simulation. `gd_partition()` defines $GD_T = 1 - \bar{f}$ and $GD_{BS} = \bar{f} - \tilde{f}$ with $\bar{f} = (\theta_A + 2\theta_{AB} + \theta_B)/4$ and $\tilde{f} = (\theta_A + \theta_B)/2$, and the sum collapses to $1 - \theta_{AB}$. Which is the shortest statement the chapter has to make: **expected F1 heterozygosity is one minus between-pool coancestry**. Everything else here is a weighting of that quantity, or a way of splitting it.

10.4 A second identity, at the hybrid level

Chapter 11. reports `GD_WI_hyb`, the within-individual gene diversity actually carried by the F1s, separately from the line-level partition, because with inbred parents `GD_WI` is zero by construction and carries no signal. A hybrid’s self-coancestry is $1 - h_i/2$, with h_i its heterozygous fraction, so `GD_WI_hyb` is half the mean F1 heterozygosity and the two levels are linked:

$$2 \cdot GD_{WI-hyb} = GD_T + GD_{BS} \quad (10.7)$$

```
c(lhs = 2 * gp[["GD_WI_hyb"]], rhs = gp[["GD_T"]] + gp[["GD_BS"]],
  GD_WI = gp[["GD_WI"]])
#>      lhs      rhs      GD_WI
#> 3.616896e-01 3.616896e-01 -4.440892e-16
```

Exact on the full factorial, because every $A \times B$ pairing is present exactly once. On a *selected* subset it holds only approximately. There GD_T and GD_{BS} are computed from the line-usage weights, which is what a random re-pairing of the selected parents would give, while `GD_WI_hyb` scores the pairs actually chosen. The gap is the pairing structure of the subset. `tests/testthat/test-dominance.R` asserts the

identity to $1e-12$ on the full factorial and bounds the gap at $5e-3$ on a *random* 40-hybrid subset, at one scale (12 lines per pool, 800 markers). A subset chosen by an optimiser, which concentrates line usage deliberately, is not covered by that bound and has not been measured here.

10.5 Inbreeding depression, and why it does not price α

The same model run backwards gives inbreeding depression. A locus at frequency p in a population with inbreeding coefficient F is heterozygous with probability $2pq(1 - F)$, so the mean declines linearly in F [16, Ch. 14]:

$$ID = 2F \sum_k d_k p_k q_k \quad (10.8)$$

At $F = 1$ that is exactly the *shared* term of `heterosis_partition()` evaluated at $p_A = p_B = p$. So the depression of a pool driven to complete homozygosity and the expected heterosis of a cross made *within* that same pool are one quantity seen from two sides. I. J. Mackay, J. Cockram, and P. Howell [10] treat heterosis and inbreeding depression together on that basis, alongside transgressive segregation, which this chapter does not take up. The identity needs no new machinery:

```
c(inbreeding_depression = inbreeding_depression(pA, d),
  shared_term          = heterosis_partition(pA, pA, d)[["shared"]],
  deviation            = abs(inbreeding_depression(pA, d) -
                             heterosis_partition(pA, pA, d)[["shared"]]))
#> inbreeding_depression      shared_term      deviation
#>          288.2869          288.2869          0.0000
```

The closed form is checked against simulated genotypes rather than against its own algebra, which is what `ref_inbreeding_depression()` exists for: it draws each genotype as identical by descent with probability F and as two independent alleles otherwise, so the agreement is to Monte Carlo error:

```
k <- 1:300
c(closed_form = inbreeding_depression(pA[k], d[k], 0.6),
  simulated   = ref_inbreeding_depression(pA[k], d[k], 0.6, n = 20000, seed = 2))
#> closed_form      simulated
#>    27.58563    27.57923
```

Now the part that matters, because the obvious next step is wrong.

It is tempting to argue: α budgets a loss of within-pool gene diversity; gene diversity is $\sum_k p_k q_k$; inbreeding depression is $\sum_k d_k p_k q_k$; therefore spending α costs yield in the current cycle, and *that* is the price of the budget. The algebra does not support it.

```
c(max_dominance_contribution_of_an_inbred_line = max(abs(heterosis_value(U, d))))
#> max_dominance_contribution_of_an_inbred_line
#> 0
```

Zero, exactly. A fully inbred line has no heterozygous loci, so its dominance contribution is zero **whatever the pool's allele frequencies are**. In a programme whose parents are inbred lines the depression was paid in full when the lines were made; it is a constant, and it does not move when the pool subsequently loses diversity. The formula above describes the gap between a hypothetical random-mating pool and its inbred derivatives, and that gap *shrinks* as $p_k q_k$ shrinks – but nothing in the programme stands at the random-mating end of it.

! Inbreeding depression is not the current-cycle price of the diversity budget

This is the same fact that makes GD_{WI} exactly zero in Chapter 11.'s partition, and it has the same cause: with inbred parents there is no within-individual variation left to lose. Two consequences.

The cost of α in this system remains what Chapter 19. measures – an option value, realised cycles later through the variance available to select on (Chapter 3.), not a yield loss booked this season.

And inbreeding depression is the live quantity wherever the population is only partially inbred: F2, F3 and F4 in Chapter 5., where F is still climbing and the depression is still accumulating. That is the chapter it belongs to, not this one.

Note also what this does *not* say. Expected heterosis in the product is a separate question. Section 10.7 answers it for the neutral case: under drift the shared term lost and the divergence term gained cancel exactly, so hybrid heterosis is invariant. Under selection the frequencies move directionally, and what happens to expected heterosis is then the alignment question of Section 10.8, which can go either way. In neither case is the diversity loss itself an immediate dominance cost. The bill arrives later, and it arrives in the additive variance.

10.6 When genetic distance predicts heterosis exactly

Before the corrections start, one special case is worth having, because it is the only case in which marker distance and heterosis are the same number.

Set $d_k \equiv d$, constant across loci. Then H stops being a weighted count and becomes a plain one: d times the number of loci at which the two parents differ. For inbred lines that count is exactly what the coancestry kernel of Chapter 1. measures, so for a cross between lines a and b

$$H_{ab} = d m (1 - f_{ab}^L) = d m \text{MRD}_{ab}^2 \quad (10.9)$$

```
h1 <- heterosis_value(ctx$X, rep(1, ctx$m))
fL <- molecular_coancestry(L / 2)
ab <- cbind(ctx$ped$a, ctx$ped$b)
c(vs_coancestry = max(abs(h1 - ctx$m * (1 - fL[ab]))),
  vs_mrd        = max(abs(h1 - ctx$m * mrd_matrix(L / 2)[ab]^2)))
#> vs_coancestry      vs_mrd
#> 1.136868e-13      2.273737e-13
```

Both residuals are at machine precision. Under flat dominance, heterosis and molecular coancestry are one statistic seen twice – Trap 1 of Chapter 13. read from the other side – and the diversity kernel this whole book runs on turns out to be the right object for the heterosis question as well.

Flat d is also precisely the condition under which H is a function of marker distance alone. It is a strong condition, and it fails as soon as d varies: two crosses at the same distance then differ by *which* loci they differ at. With the gamma-distributed deviations used everywhere else in this chapter the relationship survives but stops being an identity:

```
c(cor_uniform_d = cor(h1, 1 - fL[ab]),
  cor_real_d    = cor(heterosis_value(ctx$X, d), 1 - fL[ab]))
#> cor_uniform_d      cor_real_d
#> 1.0000000          0.7809147
```

A correlation of 0.78 is a useful predictor and a poor identity. Two things move it, and both are properties of the germplasm rather than of the statistic: how d is spread over the loci at which the parents differ, which is what Section 10.8 is about, and – once the third premise is relaxed – how well the panel tags the causal loci. Under flat d on a panel that is the causal set the correlation is 1 by the identity above. Departing from either condition generally lowers it – not always, since a d that varies only across loci at which no pair of parents differs changes no H_i at all – and by how much is something this chapter does not attempt to predict.

10.7 Divergence is not a bonus

Read *divergence* as free money and you will reach for F_{ST} . Do not.

Draw the pools further apart and watch all three terms. Nothing here is selected. Each point re-draws both pools from the *same* ancestral frequencies – the ancestral vector is fixed by `seed = 5` and shared across the sweep, while the two pool-frequency vectors are fresh Balding-Nichols draws at each `fst` (Appendix A.) – so what varies along the sweep is how far the draw scatters and nothing else:

```
sweep_het <- function(fst) {
  cf <- modifyList(book_cfg, list(fst = fst, seed = 5))
  Lf <- simulate_lines(cf)
  pl <- rep(c("A", "B"), c(cf$n_pool_A, cf$n_pool_B))
  Pf <- pool_freqs(Lf / 2, pl)
```

```

hp <- heterosis_partition(Pf["A", ], Pf["B", ], rep(1, cf$m)) / cf$m
cm <- pool_complementarity(Lf / 2, pl)
data.frame(fst = fst, shared = hp[["shared"]], divergence = hp[["divergence"]],
           total = hp[["total"]],
           frac_opposite_fixed = cm[["frac_opposite_fixed"]],
           mean_abs_delta_p = cm[["mean_abs_delta_p"]],
           mean_sq_delta_p = cm[["mean_sq_delta_p"]])
}
sw <- do.call(rbind, lapply(c(0.02, 0.05, 0.10, 0.20, 0.30, 0.45), sweep_het))
fmt(sw, 4)

```

fst	shared	diver- gence	total	frac_opposite_fixed	mean_abs_delta_p	mean_sq_delta_p
0.02	0.3549	0.0107	0.3656	0.000	0.1138	0.0214
0.05	0.3489	0.0171	0.3659	0.000	0.1414	0.0341
0.10	0.3374	0.0246	0.3620	0.000	0.1715	0.0493
0.20	0.3231	0.0433	0.3664	0.002	0.2252	0.0867
0.30	0.3042	0.0595	0.3638	0.012	0.2589	0.1191
0.45	0.2792	0.0846	0.3639	0.041	0.3012	0.1693

Across the sweep the divergence term rises by a factor of 7.9. The comparison that matters is in absolute units, not in ratios: divergence gains 0.0739 per locus and the shared term loses 0.0756. The total does not move.

10.7.1 The two terms, term by term

The algebra is one line. Pool A and pool B are drawn *independently* around a common ancestral frequency p , so $\mathbb{E}[p_A] = \mathbb{E}[p_B] = p$ and $\mathbb{E}[p_A p_B] = p^2$. Therefore

$$[\mathbb{E}_A + p_B - 2p_A p_B] = 2p(1 - p), \quad (10.10)$$

which contains no F_{ST} at all. With ancestral frequencies drawn uniformly on $[0.05, 0.95]$ the predicted level is $\mathbb{E}[2p(1 - p)] = 0.365$, against a simulated mean of 0.3646 across the whole sweep.

That the *sum* is flat is the headline, but the Balding-Nichols model predicts each half separately, which is the sharper statement. Under it $\text{Var}(p_A) = \text{Var}(p_B) = F_{ST} p(1 - p)$, so $\text{Var}(\bar{p}) = F_{ST} p(1 - p)/2$ and $\text{Var}(y) = 2F_{ST} p(1 - p)$. Substituting:

$$[\mathbb{E}\bar{p}(1 - \bar{p})] = 2p(1 - p) \left(1 - \frac{F_{ST}}{2}\right), \quad [\mathbb{E}^2/2] = F_{ST} p(1 - p) \quad (10.11)$$

The shared term loses $F_{ST} p(1 - p)$ and the divergence term gains $F_{ST} p(1 - p)$: the same expression with opposite signs, not approximately and not on average over some range. Both predictions are parameter-free once the ancestral distribution is fixed, so they can be laid straight on top of the sweep:

```

Ep <- 2 * (0.5 - (0.9^2 / 12 + 0.25)) # E[2p(1-p)] for p ~ U(0.05, 0.95)
sw$pred_shared <- Ep * (1 - sw$fst / 2)
sw$pred_divergence <- Ep * sw$fst / 2
fmt(sw[, c("fst", "shared", "pred_shared", "divergence", "pred_divergence")], 4)

```

fst	shared	pred_shared	divergence	pred_divergence
0.02	0.3549	0.3614	0.0107	0.0036
0.05	0.3489	0.3559	0.0171	0.0091
0.10	0.3374	0.3468	0.0246	0.0182
0.20	0.3231	0.3285	0.0433	0.0365

fst	shared	pred_shared	divergence	pred_divergence
0.30	0.3042	0.3102	0.0595	0.0548
0.45	0.2792	0.2829	0.0846	0.0821

```

plot(sw$fst, sw$total, type = "b", pch = 19, col = "#1b7837", ylim = c(0, 0.40),
     xlab = "target fst between pools", ylab = "expected heterosis per locus (d = 1)")
lines(sw$fst, sw$shared, type = "b", pch = 17, col = "#2166ac")
lines(sw$fst, sw$divergence, type = "b", pch = 15, col = "#b2182b")
lines(sw$fst, sw$pred_shared, lty = 2, col = "#2166ac")
lines(sw$fst, sw$pred_divergence, lty = 2, col = "#b2182b")
abline(h = mean(sw$total), lty = 2, col = "grey45")
legend("right", c("total", "shared", "divergence", "predicted"), bty = "n",
     cex = 0.8, pch = c(19, 17, 15, NA), lty = c(NA, NA, NA, 2),
     col = c("#1b7837", "#2166ac", "#b2182b", "grey45"))
grid()

```




Figure 10.1: Expected heterosis and its two components as the pools are driven apart by drift alone. Dashed lines are the Balding-Nichols predictions for each component separately, with no fitted parameter. Each measured component is offset from its prediction – divergence above, shared below – by sampling error in the pool frequencies; that offset narrows along the axis, and the text quantifies both it and the measured gap. The divergence term grows several-fold; the shared term pays for all of it, unit for unit; the total is flat. Nothing is selected here – this is what neutral divergence does.

The measured divergence sits above its prediction at every `fst`, and the measured shared term about the same distance below it. That is not drift, it is estimation. The pool frequencies are estimated from 25 and 25 sampled lines, and a frequency estimated from n lines carries sampling variance $V = pq/n$. Conditional on the true pool frequencies, $\mathbb{E}[\hat{y}^2/2] = y^2/2 + (V_A + V_B)/2$ and $\mathbb{E}[2\hat{p}(1 - \hat{p})] = 2\bar{p}(1 - \bar{p}) - (V_A + V_B)/2$: equal and opposite, so the total estimated from sample frequencies is unbiased while each half is not.

The offset is not constant along the sweep. It is proportional to $\mathbb{E}[p_A q_A + p_B q_B]$, which the pools themselves erode as they diverge, so it shrinks with $(1 - F_{ST})$: about 0.0072 per locus at `fst` = 0.02, down to 0.004 at `fst` = 0.45, against measured gaps of 0.007 and 0.0025.

Those two pairs do not agree equally well. At $f_{ST} = 0.02$ prediction and measurement are within 2% of each other; at $f_{ST} = 0.45$ the measured gap is 38% *below* the predicted offset. The formula gives an expectation; each row of the sweep is one draw of 2000 loci at one seed, and the scatter of that draw is not constant either. $\text{Var}(y)$ is $2F_{ST}p(1-p)$, so the per-locus divergence $y^2/2$ has a variance that grows with F_{ST} while the number of loci averaging it stays fixed – the high- f_{ST} rows are the noisy ones, and the agreement at the bottom of the sweep is the more informative of the two. The offset is the right order and the right sign at both ends; it is not a per-point prediction, and nothing below rests on it being one.

What changes by an order of magnitude along the sweep is not the offset but the signal it sits on: at $f_{ST} = 0.02$ the offset exceeds the predicted divergence of 0.0036, and by $f_{ST} = 0.45$ it is a small fraction of it. A pool pair's *apparent* divergence overstates its true divergence, most severely when the true divergence is small and the pools are few.

! Every unit of divergence is paid for out of the shared term

Under neutral divergence the trade is exactly one for one. Drifting the pools apart does not create heterosis; it relocates it from the shared term to the divergence term. A programme that maximises F_{ST} has bought nothing and spent its within-pool diversity to do it.

This is the mechanism behind Chapter 13. Trap 3. That section states the symptom – “ F_{ST} keeps rising as pools drift apart at loci that are *already* divergent, which buys no additional heterosis” – and points here for the reason. The two displays above are that reason.

10.8 What actually buys heterosis

If divergence per se is worth nothing, why do heterotic groups work?

Because real pools are not diverged at random with respect to d . The divergence term is $\sum_k d_k y_k^2/2$, a *weighted* sum. It grows when the loci that diverge are the loci that carry dominance; under drift the two are independent, which is exactly why the total was flat.

10.8.1 Alignment is a covariance

That sentence can be a number rather than a word. A sum of products is m times a mean times a mean, plus m times a covariance – the population covariance with divisor m , which is why the chunk rescales `cov()`:

$$\sum_k d_k y_k^2/2 = \frac{m}{2} [\overline{dy^2} + \text{cov}(d, y^2)] \quad (10.12)$$

```
y2 <- (pA - pB)^2
c(divergence = sum(d * y2 / 2),
  from_parts = ctx$m / 2 * (mean(d) * mean(y2) + cov(d, y2) * (ctx$m - 1) / ctx$m))
#> divergence from_parts
#> 32.94521 32.94521
```

The first bracketed term is what F_{ST} can see: average divergence times average dominance. The second is what it cannot. **Alignment is $\text{cov}(d, y^2)$** , and no statistic computed from allele frequencies alone can contain it, because d does not appear in one.

10.8.2 The scenarios

Hold the pools fixed, hold F_{ST} fixed, hold the multiset of dominance deviations fixed, and change only which locus gets which d_k . `aligned` gives the largest d_k to the largest y_k^2 , `anti-aligned` does the reverse, and `independent` is the draw as made:

```
mk <- function(dec) { v <- numeric(ctx$m); v[order(y2)] <- sort(d, decreasing = dec); v }
scen <- list(`anti-aligned` = mk(TRUE), `independent` = d, `aligned` = mk(FALSE))
al <- do.call(rbind, lapply(names(scen), function(n) {
  h <- heterosis_partition(pA, pB, scen[[n]])
  data.frame(scenario = n, cor_d_y2 = cor(scen[[n]], y2),
    shared = h[["shared"]], divergence = h[["divergence"]], total = h[["total"]])
}))
```

```

}))
al$relative <- al$total / al$total[al$scenario == "independent"]
fmt(al, 3)

```

scenario	cor_d_y2	shared	divergence	total	relative
anti-aligned	-0.608	270.538	11.283	281.821	0.792
independent	0.006	322.723	32.945	355.669	1.000
aligned	0.953	370.139	66.313	436.452	1.227

```

bp <- barplot(rbind(al$shared, al$divergence), beside = FALSE,
              names.arg = al$scenario, col = c("#2166ac", "#b2182b"),
              ylab = "expected heterosis", border = NA)
legend("topleft", c("shared", "divergence"), fill = c("#2166ac", "#b2182b"),
      bty = "n", cex = 0.8, border = NA)
text(bp, al$total, sprintf("%.2fx", al$relative), pos = 3, cex = 0.8, xpd = NA)
grid(nx = NA, ny = NULL)

```




Figure 10.2: The same two pools, the same F_{ST} , and the same set of dominance deviations in all three bars – only their placement across loci differs. Expected heterosis moves by half again between the extremes. F_{ST} cannot see any of this, because d does not enter it.

F_{ST} is identical in all three columns – it is a function of the frequencies only, and d never enters it. Expected heterosis ranges over a factor of 1.55.

10.8.3 Isolating the divergence channel

One honest qualification. Sorting d against y^2 moves the *shared* term as well, because y^2 is not independent of $2\bar{p}(1 - \bar{p})$: a locus with a large frequency difference tends to sit at an intermediate mean frequency. Part of the effect above therefore arrives through the shared half, not the divergence half.

The channel can be separated. Permute d only *within* groups of loci that share a value of the shared weight $2\bar{p}(1 - \bar{p})$. The shared term is then invariant by construction – exactly, not approximately – and whatever moves is the divergence channel alone:

```

hz <- 2 * ((pA + pB) / 2) * (1 - (pA + pB) / 2)
grp <- split(seq_len(ctx$m), round(hz, 9))
strat <- function(dec) {
  v <- numeric(ctx$m)
  for (i in grp) v[i[order(y2[i])]] <- sort(d[i], decreasing = dec)
  v
}
st <- do.call(rbind, lapply(c(aligned = FALSE, `anti-aligned` = TRUE), function(dec)
  as.data.frame(as.list(heterosis_partition(pA, pB, strat(dec))))))
fmt(st, 4)

```

	shared	divergence	total
aligned	322.7234	57.7211	380.4445
anti-aligned	322.7234	14.4716	337.1949

The shared column is identical to the digit; the divergence term still moves by a factor of 4 and the total by 13%. The alignment result survives its own control.

10.9 Reading the S1 panel again

Chapter 13. logs `frac_opposite_fixed` and `mean_abs_delta_p` every cycle as stand-ins for heterosis potential. The sweep above says how good they are.

`mean_sq_delta_p` is not a proxy at all. It is y^2 , so half of it is the unweighted divergence term GD_{BS} , exactly:

```

c(half_mean_sq = mean(sw$mean_sq_delta_p) / 2,
  divergence = mean(sw$divergence))
#> half_mean_sq  divergence
#>      0.039983      0.039983

```

What it cannot carry is the weighting by d , which is the previous section's point; the frequency half of the divergence term it reports without loss.

`frac_opposite_fixed` is a different matter:

```

fmt(sw[, c("fst", "divergence", "frac_opposite_fixed", "mean_abs_delta_p",
  "mean_sq_delta_p")], 4)

```

fst	divergence	frac_opposite_fixed	mean_abs_delta_p	mean_sq_delta_p
0.02	0.0107	0.000	0.1138	0.0214
0.05	0.0171	0.000	0.1414	0.0341
0.10	0.0246	0.000	0.1715	0.0493
0.20	0.0433	0.002	0.2252	0.0867
0.30	0.0595	0.012	0.2589	0.1191
0.45	0.0846	0.041	0.3012	0.1693

Up to `fst = 0.10` it is exactly zero – not one marker in 2000 qualifies – and it first exceeds 1% of markers at `fst = 0.30`, near the top of the sweep. It counts *near-fixed* opposite loci, and the divergence lives at intermediate frequency differences, which the threshold discards.

Log `mean_sq_delta_p`. It is exact, it is monotone across the whole sweep, and it costs one squaring more than the statistic already being logged. `pool_complementarity()` and `s1_monitor()` return it.

10.10 General and specific combining ability

Everything so far has been about $\mathbb{E}[H]$. Now the spread.

The same expansion that produced the shared/divergence split produces the combining-ability decomposition, with no new machinery and no model fitting. Let $u_{ik} \in \{0, 1\}$ be the genotype of pool-A line i at locus k and v_{jk} that of pool-B line j , and write each as its pool frequency plus a deviation, $u = p_A + \delta u$ and $v = p_B + \delta v$. The per-locus heterozygosity of the cross, $u + v - 2uv$, expands to

$$u + v - 2uv = \underbrace{p_A + p_B - 2p_A p_B}_{\text{mean}} + \underbrace{(1 - 2p_B) \delta u}_{\text{parent A}} + \underbrace{(1 - 2p_A) \delta v}_{\text{parent B}} - \underbrace{2 \delta u \delta v}_{\text{both}} \quad (10.13)$$

Four terms: one constant, one linear in each parent, one bilinear. Weight by d_k , sum over loci, and the classical decomposition falls out,

$$H_{ij} = \mu + g_i^A + g_j^B + s_{ij} \quad (10.14)$$

with

$$g_i^A = \sum_k d_k (1 - 2p_{Bk}) \delta u_{ik}, \quad s_{ij} = -2 \sum_k d_k \delta u_{ik} \delta v_{jk} \quad (10.15)$$

and g_j^B symmetrically. What is linear in one parent is general combining ability: it is carried into every cross that parent enters. What is bilinear depends on *which two* lines met, and is specific combining ability. Because the deviations are taken from each pool's own line frequencies they sum to zero over that pool's lines, which is what makes the three components mean zero and mutually orthogonal. `heterosis_gca_sca()` is that expansion and nothing more:

```
heterosis_gca_sca
#> function (U, V, d)
#> {
#>   stopifnot(ncol(U) == ncol(V), length(d) == ncol(U))
#>   pA <- colMeans(U)
#>   pB <- colMeans(V)
#>   cU <- sweep(U, 2, pA)
#>   cV <- sweep(V, 2, pB)
#>   list(mu = heterosis_expected(pA, pB, d), gca_A = drop(cU %*%
#>     (d * (1 - 2 * pB))), gca_B = drop(cV %*% (d * (1 - 2 *
#>     pA))), sca = -2 * (cU * rep(d, each = nrow(U))) %*% t(cV))
#> }
#> <bytecode: 0xc0b224078>
#> <environment: namespace:hybdiv>
```

It reconstructs every hybrid in the factorial exactly, and all three components are mean zero:

```
cs <- heterosis_gca_sca(U, V, d)
H <- heterosis_value(ctx$X, d)
ia <- ctx$ped$a
ib <- ctx$ped$b - nA
c(max_residual = max(abs(H - (cs$mu + cs$gca_A[ia] + cs$gca_B[ib] +
  cs$sca[cbind(ia, ib)]))),
  mu_vs_mean = abs(cs$mu - mean(H)),
  mean_gca_A = mean(cs$gca_A), mean_sca = mean(cs$sca))
#> max_residual mu_vs_mean mean_gca_A mean_sca
#> 7.958079e-13 1.136868e-13 -4.796163e-16 -1.264766e-15
```

This is not a decomposition that resembles the two-way analysis of variance. It *is* the two-way analysis of variance: μ , g^A and g^B are the grand mean and the row and column effects of the H_{ij} table, and s_{ij} is the interaction residual of the additive fit, to machine precision:

```
max(abs(resid(lm(H ~ factor(ia) + factor(ib))) - cs$sca[cbind(ia, ib)]))
#> [1] 6.966872e-12
```

R. E. Comstock, H. F. Robinson, and P. H. Harvey [8] designed reciprocal recurrent selection around the distinction between these two components. The algebra above says what each one is made of – and the next section says which of the two their scheme actually selects on.

10.10.1 GCA is an average effect in the other pool's background

Look at the weight on δu : it is $d_k(1 - 2p_{Bk})$, and $1 - 2p_B$ is $q_B - p_B$. That is the frequency factor of the classical average effect of a gene substitution, $a + d(q - p)$ [16, Ch. 7], with two differences: the additive part a is absent, because it belongs to the parental mean and not to H , and the frequencies are **pool B's**, not pool A's.

This is the whole reason combining ability is not a property of a line. A pool-A line's alleles are always substituted into a pool-B background, so its general combining ability for heterosis is a dominance effect *evaluated at the tester pool's frequencies*. Change the tester pool and every number changes. Two limits follow immediately:

- A locus at which the tester pool sits at $p_B = 0.5$ contributes **exactly zero** to g^A , whatever δu_{ik} is. It still contributes to g^B , through $(1 - 2p_A)\delta v_{jk}$, and to s_{ij} ; what it cannot do is discriminate between pool-A parents.
- A locus at which the tester pool is fixed, $p_B \in \{0, 1\}$, has $\delta v \equiv 0$, so it contributes $\pm d_k \delta u_{ik}$ to g^A and nothing to either g^B or s_{ij} . The cross is then predictable from the pool-A parent by itself.

Between those two limits sits every real programme, and Section 10.11 is about which limit divergence pushes it towards.

This is also the exact sense in which R. E. Comstock, H. F. Robinson, and P. H. Harvey [8]'s scheme is built on the decomposition. Reciprocal recurrent selection evaluates each pool-A candidate on its cross performance against a *sample of the opposite pool*, and symmetrically for pool B. Averaging over testers removes s_{ij} , so of the heterosis in the decomposition above only g_i^A survives as a contrast between pool-A candidates. Note what that does *not* say: cross performance is a genotypic value, not a mid-parent difference, so it also carries the candidate's own additive contribution – the part that cancelled out of H in the mid-parent derivation above and is not in this decomposition at all. Selection on testcross means therefore acts on the additive part and g_i^A together. What the algebra here delimits is the dominance share of that response, and it says the share is general combining ability against the tester pool.

10.10.2 SCA is minus twice a weighted parental inner product

The bilinear term is $s_{ij} = -2 \sum_k d_k \delta u_{ik} \delta v_{jk}$: minus twice the d -weighted inner product of the two parents' *centred* genotypes, a covariance up to the factor $1/m$. For $d_k > 0$, two parents that deviate from their pools in the same direction have a positive inner product and therefore a **negative** SCA.

Read in words: two parents that resemble each other – after removing what is typical of each pool – cross badly, and the resemblance that counts is weighted by d , not by marker count. Similarity enters negatively throughout this chapter – $1 - \theta_{AB}$, $d m (1 - f_{ab}^L)$ – but everywhere else it is similarity measured against a fixed base. Here it is measured against each pool's own mean, which is what makes the term specific to a *pair* rather than a property either parent carries into all its crosses.

It is also Section 10.6 again. Since μ , g^A and g^B are the grand, row and column means of the H_{ij} table, s_{ij} is the double-centred H matrix. Under flat d , $H_{ij} = d m (1 - f_{ij}^L)$, and double-centring annihilates the constant and flips the sign of what remains:

$$s_{ij} = -d m \cdot \text{DC} (f^L)_{ij}, \quad (10.16)$$

with $\text{DC}(M)_{ij} = M_{ij} - \bar{M}_i - \bar{M}_j + \bar{M}_.$ over the $A \times B$ block. The factor $-dm$ is not cosmetic: s_{ij} is on the trait scale and f^L is a probability, and the minus is the same minus as in $-2 \sum_k d_k \delta u \delta v$ – more coancestry than the two margins predict means *less* SCA. So the coancestry kernel determines SCA under exactly the condition it determines heterosis, up to that constant, and stops doing so under exactly the same condition: a d that varies across loci.

10.10.3 The variance components

Because the three components are orthogonal, the variance of heterosis over the factorial splits with them. Under the sampling model of the simulation – each line's genotype an independent Bernoulli(p_{Ak}) draw at every locus, so unrelated lines and linkage equilibrium within each pool – each piece has a closed form in the pool frequencies:

$$\sigma_{GCA_A}^2 = \sum_k d_k^2 (1 - 2p_{Bk})^2 p_{Ak} q_{Ak}, \quad \sigma_{SCA}^2 = 4 \sum_k d_k^2 p_{Ak} q_{Ak} p_{Bk} q_{Bk} \quad (10.17)$$

```

qa <- pA * (1 - pA)
qb <- pB * (1 - pB)
fmt(data.frame(
  component = c("GCA_A", "GCA_B", "SCA"),
  observed   = c(var(cs$gca_A), var(cs$gca_B), var(as.vector(cs$sca))),
  closed_form = c(sum(d^2 * (1 - 2 * pB)^2 * qa),
                  sum(d^2 * (1 - 2 * pA)^2 * qb),
                  4 * sum(d^2 * qa * qb)), 2)

```

component	observed	closed_form
GCA_A	32.00	35.03
GCA_B	35.63	38.07
SCA	70.38	71.48

The closed forms take the sample frequencies as the pool frequencies; the observed column is one realisation of them. The GCA rows are variances over 25 lines, whose relative sampling error is about $\sqrt{2/n} \approx 28\%$, and the agreement is within that.

Note what each one requires. σ_{SCA}^2 carries $p_A q_A p_B q_B$: it needs **both** pools to be polymorphic. Fix either pool at a locus and that locus contributes no specific combining ability at all, whatever its d_k . Within-pool diversity is not merely something the diversity constraint defends; it is the raw material of SCA.

10.11 Divergence destroys SCA faster than it builds GCA

Put the two closed forms next to the drift argument of Section 10.7 and the variance tells a different story from the mean. The mean was flat; the variance is not, and it does not merely change shape.

Under Balding-Nichols, $\mathbb{E}[p_A q_A] = p(1-p)(1-F_{ST})$ and $\mathbb{E}[(1-2p_B)^2] = (1-2p)^2 + 4F_{ST}p(1-p)$, and p_A and p_B are independent given the ancestral p . Per locus, in expectation over the draw,

$$\mathbb{E}[\sigma_{SCA,k}^2] = 4d_k^2 p^2 (1-p)^2 (1-F_{ST})^2, \quad \mathbb{E}[\sigma_{GCA_A,k}^2] = d_k^2 p(1-p)(1-F_{ST})[(1-2p)^2 + 4F_{ST}p(1-p)] \quad (10.18)$$

Read the two expressions before reading the sweep, because they already say what it will show.

Both carry the erosion factor $(1-F_{ST})$, and both are **zero at $F_{ST} = 1$** : a pool fixed at every locus has no two lines that differ, so it has no combining-ability variance of either kind. Divergence therefore cannot transfer variance from one component to the other indefinitely; it destroys both in the limit. SCA carries the factor twice, because it needs both pools polymorphic at the same locus. GCA carries it once, multiplied by the tester factor $(1-2p_B)^2$, whose expectation grows with F_{ST} .

The two expressions therefore behave differently in F_{ST} . The SCA expression is convex and decreasing on $[0, 1]$: it is a positive constant times $(1-F_{ST})^2$, so it only ever falls. The GCA expression is a *concave* parabola, its F_{ST} coefficient being $-4d_k^2 p^2 (1-p_k)^2$, and differentiating it gives the only turnover in this section. GCA variance rises from $F_{ST} = 0$ at a locus if and only if $8p(1-p) > 1$, that is $\frac{1}{2}(1-1/\sqrt{2}) < p < \frac{1}{2}(1+1/\sqrt{2})$, and turns over at

$$F_k^* = \frac{8p_k(1-p_k)-1}{8p_k(1-p_k)}, \quad (10.19)$$

which is negative outside that interval – at such loci GCA variance falls from the start. A sum of concave parabolas is a concave parabola, and its maximum works out to the $d_k^2 (p_k q_k)^2$ -weighted mean of the F_k^* . Whether that mean lands inside $[0, 1]$ is a property of the ancestral frequency distribution, not a theorem: pools founded mostly at extreme frequencies would put it below zero and GCA variance would decline monotonically. Under the $U(0.05, 0.95)$ ancestral draw used here enough loci sit in the interval to place it inside the range, which is why the sweep below straddles a peak. The sweep's grid can show that the peak is interior; it cannot locate it.

What is monotone is the ratio,

$$\frac{\mathbb{E}[\sigma_{SCA,k}^2]}{\mathbb{E}[\sigma_{GCA_A,k}^2]} = \frac{4p(1-p)(1-F_{ST})}{(1-2p)^2 + 4F_{ST}p(1-p)}, \quad (10.20)$$

which falls in F_{ST} at every locus. So the defensible statement is about the *share*, not about a transfer. Expected heterosis, we already know, does not move at all. Sweeping the pools apart – same ancestral frequencies, fresh Balding-Nichols draws at each *fst*, as in Section 10.7 – and evaluating the closed forms on the realised frequencies:

```
reloc <- function(fst) {
  cf <- modifyList(book_cfg, list(fst = fst, seed = 5))
  Pf <- pool_freqs(simulate_lines(cf) / 2,
    rep(c("A", "B"), c(cf$n_pool_A, cf$n_pool_B)))
  a <- Pf["A", ]; b <- Pf["B", ]
  va <- sum(d^2 * (1 - 2 * b)^2 * a * (1 - a))
  vb <- sum(d^2 * (1 - 2 * a)^2 * b * (1 - b))
  vs <- 4 * sum(d^2 * a * (1 - a) * b * (1 - b))
  data.frame(fst = fst, var_GCA = va + vb, var_SCA = vs,
    frac_SCA = vs / (va + vb + vs))
}
rl <- do.call(rbind, lapply(c(0.02, 0.05, 0.10, 0.15, 0.20, 0.30, 0.45), reloc))
fmt(rl, 3)
```

fst	var_GCA	var_SCA	frac_SCA
0.02	60.047	91.807	0.605
0.05	62.025	87.941	0.586
0.10	68.630	75.503	0.524
0.15	71.230	67.599	0.487
0.20	77.956	63.166	0.448
0.30	80.630	45.875	0.363
0.45	77.894	27.767	0.263

```
plot(rl$fst, rl$var_GCA, type = "b", pch = 19, col = "#1b7837",
  ylim = range(0, rl$var_GCA, rl$var_SCA),
  xlab = "target fst between pools",
  ylab = "variance of heterosis over the factorial")
lines(rl$fst, rl$var_SCA, type = "b", pch = 15, col = "#b2182b")
legend("right", c("GCA (both pools)", "SCA"), bty = "n", cex = 0.8,
  pch = c(19, 15), col = c("#1b7837", "#b2182b"))
grid()
```




Figure 10.3: The same neutral divergence sweep, read through the combining-ability decomposition rather than through the mean. Expected heterosis is flat across this entire range (previous figure); its decomposition is not. Specific combining ability falls steeply throughout. General combining ability rises, by much less than SCA loses, and itself turns over inside the range – both components vanish at $F_{ST} = 1$. The two curves do not sum to a constant: most of the SCA variance lost is destroyed rather than relocated, and the exact percentages are in the text.

Take the three columns in order, because the headline is in the arithmetic between them.

SCA variance falls by 70% across the sweep, a drop of 64 units. GCA variance ends 30% above where it started, a gain of 17.8 units, and peaks at an f_{ST} of 0.3 – interior, as the ancestral distribution used here leads one to expect, though a grid whose neighbouring points are 0.2 and 0.45 locates the peak only to that interval and cannot resolve it further. So the GCA gain recovers about 28% of what SCA lost. The rest is gone: total variance in heterosis over the factorial falls from 151.9 to 105.7, a 30% decline, while the mean stays fixed throughout.

Read those four percentages as one realisation, not as constants. They come from a single pair of seeds – `seed = 5` for the pool frequencies, `set.seed(4)` for d – with no replication and no dispersion reported, and they will move if either changes. What does not depend on the seed is the direction, which is already in the two closed forms above: σ_{SCA}^2 is a positive multiple of $(1 - F_{ST})^2$ and can only fall, and both components vanish at $F_{ST} = 1$, so no amount of neutral divergence can hold the total up. The ordering of the columns is algebra; their magnitudes are this draw.

The share of that variance which is specific falls from 0.605 to 0.263. That statement – about the share – is the one the sweep supports; “the dominance expressed as SCA reappears as GCA” is not, and the two columns above are why.

The share result itself is not new. J. C. Reif, F.-M. Gumpert, S. Fischer, and A. E. Melchinger [17] prove that the ratio of dominance to additive variance in a hybrid population falls as its two parent populations diverge, and conclude that superior hybrids then become predictable from general combining ability effects. That half of this section is theirs. What the sweep adds is the denominator, which a ratio hides: a ratio can fall because its numerator falls, because its denominator rises, or both, and the columns above say which. Here the numerator collapses and the denominator recovers only part of it, so the share falls and the total falls with it.

Neutral divergence – drift, the only kind on this sweep – does not buy expected heterosis, and it does not move heterosis from one combining-ability component into the other. It **shrinks the part of hybrid variation that only a cross can reveal**, destroying SCA variance faster than it builds GCA

variance. An additive model looks better against diverged pools mainly because there is less left for it to miss. Divergence that is *aligned* with dominance is a different matter and Section 10.8 is where it is measured; the sweep here holds alignment at zero by construction.

This sharpens a question Chapter 14. had to leave open. That chapter fits one additive effect per marker and notes that such a model cannot represent SCA at all. The sweep says how large the unrepresentable share is, and how it moves: against poorly separated pools most of the variance in heterosis is specific; against well separated ones a minority is. These are shares of the variance in *heterosis*, not of total hybrid value, whose parental-mean part is additive at every F_{ST} . F. Technow, T. A. Schrag, W. Schipprack, E. Bauer, H. Simianer, and A. E. Melchinger [18] report high accuracy for additive hybrid prediction in established, clearly separated dent and flint pools, which is the regime where this model puts the specific share at its lowest – consistent with the mechanism, though a single observational study in one crop cannot test it.

⚠ A smaller SCA share is not a better programme, and this is not a licence to maximise F_{ST}

Three things stay true at every point on that sweep. Expected heterosis does not change. The within-pool diversity spent to buy the divergence is spent for good. And $\sigma_{SCA}^2 \rightarrow 0$ is the limit in which no cross is better than its parents predict – there is nothing left to *find* by testing combinations, which is the activity the whole factorial exists to support. Past the peak of the GCA curve even the predictable half shrinks: at $F_{ST} = 1$ there is no combining-ability variance of either kind, because no two lines in a pool differ.

The falling SCA share is therefore a fact about what a predictor can express, not a gain. A programme reading it as one would be congratulating itself on having less to predict.

So this is a second, independent reason for Chapter 13. Trap 3's prescription to hold F_{ST} in a band rather than at a maximum. Section 10.7 gave the diversity-side reason. This is the prediction-side one, and it points the same way.

10.11.1 What this says about ranking crosses

The F1 of two fixed lines is genetically uniform, so the progeny-variance term that drives the usefulness criterion at Chapter 13. S2 is zero for the hybrid itself, and cross ranking reduces to ranking means. That is true, and it is why this book ranks hybrids on an index rather than on $\mu + ih\sigma$.

It does not follow that the ranking is easy. The *mean* of an F1 contains an SCA term that an additive **G** omits entirely, and Section 10.11 sized its share of the variance in heterosis: a majority of it against weakly separated pools, a minority against well separated ones, and zero only at the degenerate limit where one pool is fixed at every locus. The progeny variance collapses; the dominance deviation does not.

10.12 What this does not change

⚠ The diversity constraint stays additive

Nothing in this chapter reopens the optimisation. θ , GD and α are computed on molecular coancestry and are functions of the frequencies alone; no dominance term enters them, and Chapter 15., Chapter 16. and Chapter 17. are untouched. The combining-ability split of Section 10.10 sits outside the constraint for the same reason: it is a statement about the objective, not about the feasible set.

A dominance model changes the **mean-performance ranking**, not the diversity term. Chapter 21. states the same separation from the other side. Do not conflate them, and do not read this chapter as licence to put a heterosis term inside the constraint.

What it does change is the interpretation of the constraint's output. Two plans with the same α and the same GD_{BS} can differ in expected heterosis by half again, if one of them concentrated its divergence on the dominance loci – that is the scenario table of Section 10.8, whose three columns share every frequency-based statistic. The partition cannot distinguish them. Only an estimate of d can.

10.12.1 What it would take to use d

Chapter 11. leaves a weight λ undetermined between within- and between-pool diversity and observes that the ratio of the two heterosis terms is what λ is trying to express. This chapter can now say what would have to be estimated to set it, and it is not much:

- **λ needs two covariances, not d itself.** Each heterosis term is m times a mean times a mean plus a covariance, so their ratio needs $\text{cov}(d, y^2)$ and $\text{cov}(d, 2\bar{p}(1 - \bar{p}))$, each relative to d , on top of what the frequencies already give. A programme does not need per-locus dominance deviations to weight the two diversity terms against each other; it needs to know whether its divergence is landing on its dominance loci.
- **Ranking crosses needs more.** s_{ij} is a d -weighted inner product between two specific parents, so it needs the deviations locus by locus. This is the extension Chapter 21. names and the book does not attempt.
- **Fitting it needs a factorial, and a testcross file is not one.** The decomposition above is the two-way layout with *identified* rows and columns: each pool-A line crossed to several named pool-B lines and each pool-B line to several named pool-A lines, which is what makes g^A, g^B and s_{ij} separately estimable and what the $\text{lm}()$ above fits. Testcrosses against a single common tester give only a column of $\mu + g_i^A +$ (one tester's contribution) – no SCA is identifiable, because there is no second tester to contrast against. Programmes that already cross two pools in a partial factorial for hybrid evaluation have the layout; programmes that testcross against one inbred tester per pool do not.

10.13 Where this connects

Chapter 2. supplied the partition and Chapter 11. will use it; this chapter supplies the reason the GD_{BS} component is singled out for protection, and the reason protecting it is not the same as maximising it. Chapter 13. Trap 3 already recommended holding F_{ST} in a band: Section 10.7 is why a band, and Section 10.8 is what to watch instead of the band. Section 10.6 closes the loop with the coancestry kernel of Chapter 1.: under flat dominance, the diversity matrix and the heterosis prediction are the same object.

Chapter 14. is the chapter this one supplies a missing term to. It predicts breeding values with an additive model on a trait that is additive by construction, and lists specific combining ability first among what it does not model. Section 10.10 is where that term comes from and Section 10.11 is what share of the variance in heterosis it holds; both are derived here rather than fitted, because on a simulated d they can be.

It also gives a reason why a heterotic *pattern* – which pool crosses well onto which – is not recoverable from how far apart the pools have drifted: Section 10.8's three scenarios share every frequency-based statistic and differ by half again in expected heterosis. [M. R. White, M. A. Mikel, N. de Leon, and S. M. Kaeppler \[19\]](#) describe heterotic patterns and diversity separately in North American proprietary dent germplasm; this chapter says why the two need describing separately.

Forward, the alignment result is the argument for Chapter 19.. Drift moves frequencies independently of d , and selection on a line's own performance cannot respond to d either, because an inbred line's dominance contribution is exactly zero (Section 10.5) – which leaves selection on *cross* performance as the only process in this model that can move frequencies in a way that depends on where the dominance is. The alignment between divergence and dominance is therefore something each cycle either builds or erodes, not a standing asset. A one-cycle constraint cannot see that, and neither can F_{ST} . Chapter 7. runs that loop on a dominance trait and is the pipeline this chapter's algebra describes.

11. A catalogue of metrics on $\mathbf{f}$

Every metric in this chapter has the same signature, `f(idx, ctx) -> scalar`, and every one operates on the molecular coancestry matrix from Chapter 1.. That uniformity is not tidiness: it is what makes metrics interchangeable inside an optimiser, and it is what lets Chapter 12. benchmark them against each other on equal terms.

The chapter also settles the question of what “no loss” means, which turns out to be less obvious than it sounds.

11.1 Group coancestry, and the off-diagonal problem

The group coancestry [20] of a set S of n individuals is the mean of **all** n^2 elements, diagonal included:

$$\theta_S = \frac{1}{n^2} \sum_{i \in S} \sum_{j \in S} f_{ij} = \frac{1}{n^2} \mathbf{1}' \mathbf{f}_S \mathbf{1}. \quad (11.1)$$

It is the probability that two alleles sampled at random *from the whole set* are identical – including the possibility that both come from the same individual. That is why the diagonal enters: it is the chance of sampling the same individual twice, which is entirely real when the set is used as a crossing pool.

Splitting the diagonal mean $\bar{f}_d$ from the off-diagonal mean $\bar{f}_o$,

$$\theta_S = \frac{\bar{f}_d}{n} + \frac{n-1}{n} \bar{f}_o. \quad (11.2)$$

So the off-diagonal mean – the metric many programmes already report – is one of two terms. When $\bar{f}_d$ varies little across individuals, which is the case in an all-F1 population, the two rank almost identically:

```
sels <- list(random      = sel_random(ctx, n_sel),
             trunc_cap   = sel_truncation_cap(ctx, n_sel),
             greedy_w1    = sel_greedy(ctx, n_sel, w = 1),
             greedy_div   = sel_greedy(ctx, n_sel, w = 0),
             truncation   = sel_truncation(ctx, n_sel))

cmp <- t(sapply(sels, function(s)
  c(theta = theta_group(s, ctx), offdiag = mean_offdiag(s, ctx))))
fmt(as.data.frame(cmp), 5)
```

	theta	offdiag
random	0.67426	0.67180
trunc_cap	0.67522	0.67280
greedy_w1	0.68759	0.68539
greedy_div	0.66976	0.66730
truncation	0.68759	0.68539

```
c(rank_correlation = cor(cmp[, "theta"], cmp[, "offdiag"], method = "spearman"))
#> rank_correlation
#> 1
```

The off-diagonal metric ranks correctly. What it does not give you is an interpretable scale or a defensible baseline – which is what the rest of this chapter supplies.

11.2 Gene diversity: θ and H_e are the same quantity

This is the result that simplifies everything downstream. Starting from the definition and swapping the order of summation, the double sum factors:

$$\frac{1}{n^2} \sum_i \sum_j x_{ik} x_{jk} = \left(\frac{1}{n} \sum_i x_{ik} \right)^2 = \bar{p}_k^2, \quad (11.3)$$

with $\bar{p}_k$ the **group's** allele frequency. Hence

$$\theta_S = \frac{1}{m} \sum_k (\bar{p}_k^2 + \bar{q}_k^2), \quad GD_S = 1 - \theta_S = \frac{1}{m} \sum_k 2\bar{p}_k \bar{q}_k \equiv H_e, \quad (11.4)$$

which is exactly Nei's gene diversity [21]. **Group coancestry and expected heterozygosity are not two independent checks. They are one quantity written two ways.**

```
c(via_submatrix = theta_group(idx, ctx),
  via_frequency = theta_from_freq(idx, ctx),
  difference     = theta_group(idx, ctx) - theta_from_freq(idx, ctx))
#> via_submatrix via_frequency difference
#>      0.6734562      0.6734562      0.0000000

c(GD = gene_diversity(idx, ctx), He_Nei = he_nei(idx, ctx))
#>      GD      He_Nei
#> 0.3265438 0.3265438
```

Equality to machine precision is the implementation's strongest sanity check. If it fails, **f** is centered, rescaled, or built on the wrong marker coding.

Three practical consequences:

1. One fewer metric to justify to a technical committee.
2. "Losing 5% of diversity" acquires an exact meaning: 5% of the set's expected heterozygosity.
3. Two computation routes with the same answer and very different costs. The matrix route is $O(n^2)$; the frequency route is $O(nm)$. For $n = 200$ and $m = 25000$ that is forty thousand operations against five million.

```
reps <- 40
t_mat <- system.time(for (i in seq_len(reps)) theta_group(idx, ctx))["elapsed"]
t_frq <- system.time(for (i in seq_len(reps)) theta_from_freq(idx, ctx))["elapsed"]
c(us_matrix = 1e6 * t_mat / reps, us_frequency = 1e6 * t_frq / reps,
  ratio = t_frq / t_mat)
#>      us_matrix us_frequency      ratio
#>          25          175           7
```

The matrix route is the one that belongs in an optimiser's inner loop. Section 12.1 makes it cheaper still.

11.3 Status number N_s

D. Lindgren, L. D. Gea, and P. A. Jefferson [22] express group coancestry on the scale of individual counts:

$$N_s = \frac{1}{2\theta_S}. \quad (11.5)$$

N_s is the number of **unrelated, non-inbred** individuals that would have the same gene diversity as the observed set. The edge cases fix the intuition:

```
fake <- ctx
fake$f <- matrix(0.5, 5, 5) # unrelated, p = 0.5
c(theta = theta_group(1:5, fake), Ns = status_number(1:5, fake))
#> theta      Ns
#>    0.5    1.0
```

```
fake$f <- matrix(1, 5, 5)          # all identical and homozygous
c(theta = theta_group(1:5, fake), Ns = status_number(1:5, fake))
#> theta      Ns
#>  1.0      0.5
```

⚠ N_s is a descriptor, not a drift projection

M. A. Toro, B. Villanueva, and J. Fernández [23] show that the **rate-based** effective size $N_e = 1/(2\Delta F)$ loses meaning under molecular diversity management: optimisation increases diversity in early generations, producing a negative ΔF and therefore a negative N_e , with no long-term asymptote.

This does not invalidate N_s . N_s is a monotonic reparametrisation of θ – a *static descriptor of a set* – while N_e is a *rate between generations*. Use N_s to compare sets; do not extrapolate it to project future drift. For the same reason, do not substitute the linkage-disequilibrium N_e here: that estimates population history, not the composition of a selected sample.

And, from Chapter 9.: θ tracks homozygosity, not drift. GD and N_s answer “how much heterozygosity does this selection preserve”, which is the right question for next season’s hybrid vigour. They do not certify that allele frequencies are moving as unmanaged drift would have moved them, which matters more across multiple recurrent cycles than for a single round of hybrid selection.

11.4 Effective number of parents

A metric independent of f , computed purely from the pedigree. With p_i the frequency with which line i appears as a parent among the selected hybrids,

$$N_{e(\text{par})} = \frac{1}{\sum_i p_i^2}, \quad (11.6)$$

the inverse Simpson index. It captures precisely what coancestry blurs: a set of 200 hybrids can have an acceptable θ while drawing on only 12 lines, if those 12 happen to be genetically distant from each other.

```
c(sapply(sels, ne_parents, ctx = ctx), theoretical_max = ctx$n_lines)
#>      random      trunc_cap      greedy_w1      greedy_div      truncation
#> 39.34426      31.03448      13.01989      37.30570      13.01989
#> theoretical_max
#> 50.00000
```

Truncation collapses the parent base far more than it moves θ . Its cost is a `tabulate` call, which makes it a natural candidate to enter the objective function alongside θ – and in Chapter 12. it does.

11.5 Allelic richness

θ and H_e are averages across markers and are dominated by intermediate-frequency alleles. A lost rare allele barely moves them, yet the loss is irreversible. So count markers that were polymorphic in the population and became rare or fixed in the selection:

$$A_{\text{lost}} = \#\left\{k : \text{MAF}_k^{(S)} < 0.05 \wedge \text{MAF}_k^{(\text{pop})} \geq 0.05\right\}. \quad (11.7)$$

```
sapply(sels, alleles_lost, ctx = ctx)
#>      random      trunc_cap      greedy_w1      greedy_div      truncation
#>      24          43          110          22          110
```

It is the only metric in this catalogue **genuinely non-redundant** with θ – a claim Chapter 12. tests rather than asserts.

E. López-Cortegano, R. Pouso, A. Labrador, A. Pérez-Figueroa, J. Fernández, and A. Caballero [24] compared, in a subdivided simulated population, strategies maximising *expected heterozygosity* (H_T , the same

family as GD) against strategies maximising *allelic diversity* (A_T). Maximising A_T retained more alleles **and** controlled molecular inbreeding better than maximising H_T . The mechanism connects directly to Chapter 9.: heterozygosity-maximising management pushes allele frequencies toward 0.5, which favours rare alleles in the short run but does not stop them drifting to loss once they are no longer the rarest; allele-count management defends every segregating allele directly.

Their recommendation for structured conservation programmes was to make A_T , not H_T , the primary target. In this pipeline’s terms: `alleles_lost` deserves a place next to θ as a *tracked outcome*, not only as a post-hoc sanity check, whenever the selection feeds several future cycles rather than a single hybrid release.

11.6 Distances: E-NE and A-NE

With distance $d_{ij} = 1 - f_{ij}$, Core Hunter 3 [25] defines two measures capturing opposite aspects that θ alone cannot separate:

$$\text{E-NE} = \frac{1}{n} \sum_{i \in S} \min_{j \in S, j \neq i} d_{ij}, \quad \text{A-NE} = \frac{1}{N} \sum_{i=1}^N \min_{j \in S} d_{ij}. \quad (11.8)$$

- **E-NE**, entry-to-nearest-entry, measures **non-redundancy within the selection**. It distinguishes “200 individuals spread out” from “100 near-identical pairs spread out” – configurations that can share the same θ .
- **A-NE**, accession-to-nearest-entry, measures **representativeness of the population**. It detects a selection that optimises θ by concentrating in one region of genetic space and ignoring the rest.

```
fmt(t(sapply(sels, function(s) c(ENE = ene(s, ctx), ANE = ane(s, ctx)))), 4)
```

	ENE	ANE
random	0.2533	0.2471
trunc_cap	0.2536	0.2561
greedy_w1	0.2527	0.2523
greedy_div	0.2607	0.2504
truncation	0.2527	0.2523

Note the direction: A-NE is a distance to the *nearest selected entry*, so **lower is better** for representativeness.

11.7 Effective dimensionality

Over the eigenvalues of the submatrix, the participation ratio

$$D_{\text{eff}} = \frac{(\sum_i \lambda_i)^2}{\sum_i \lambda_i^2} \quad (11.9)$$

counts how many independent directions of genetic variation remain: 1 if all variation lies on a single axis, n if spread evenly.

```
fmt(sapply(sels, eff_dim, ctx = ctx), 2)
#>   random trunc_cap greedy_w1 greedy_div truncation
#>   28.59    23.99    12.72    27.92    12.72
```

The submatrix is **double-centered** before the eigendecomposition. On a raw all-positive $\mathbf{f}$ the leading eigenvalue is just $n\theta$, so an uncentered version would only re-measure θ under a more expensive name – a mistake worth naming because it is easy to make and produces a plausible-looking number.

11.8 Decomposition by heterotic group

In hybrid maize the diversity **between** pools is the engine of heterosis and cannot be spent; the diversity **within** each pool is the fuel for future progress. Chapter 10. derives the first half of that sentence – **GD_BS** weighted by dominance deviations is the divergence half of expected heterosis – and qualifies it: divergence is worth having where it lands on the dominance loci, and worth nothing where it does not. Chapter 19. demonstrates the second half. Following [A. Caballero and M. A. Toro \[1\]](#), compute molecular coancestry between lines and weight it by each line's usage among the selected hybrids:

$$\theta_A = \mathbf{w}'_A \mathbf{f}^L \mathbf{w}_A, \quad \theta_B = \mathbf{w}'_B \mathbf{f}^L \mathbf{w}_B, \quad \theta_{AB} = \mathbf{w}'_A \mathbf{f}^L \mathbf{w}_B. \quad (11.10)$$

```
fmt(t(sapply(sels, theta_pools, ctx = ctx)), 4)
```

	theta_A	theta_B	theta_AB
random	0.7090	0.7095	0.6393
trunc_cap	0.7096	0.7153	0.6380
greedy_w1	0.7581	0.7219	0.6352
greedy_div	0.7053	0.7066	0.6336
truncation	0.7581	0.7219	0.6352

Feeding those into the Caballero & Toro partition of Chapter 2.:

```
fmt(t(sapply(sels, gd_partition, ctx = ctx)), 4)
```

	GD_WI_hyb	GD_T	GD_WI	GD_BI	GD_BS	F_ST
random	0.1801	0.3257	0	0.2907	0.0350	0.1074
trunc_cap	0.1816	0.3248	0	0.2876	0.0372	0.1146
greedy_w1	0.1823	0.3124	0	0.2600	0.0524	0.1678
greedy_div	0.1853	0.3302	0	0.2941	0.0361	0.1095
truncation	0.1823	0.3124	0	0.2600	0.0524	0.1678

Two things to read here. **GD_WI** is **exactly zero**: the parent lines are inbred, so there is no within-individual diversity at line level – the algebra of Chapter 2. predicted this, and it is why **GD_WI_hyb**, the heterozygosity actually carried inside the F1s, is reported separately. And **GD_BS** rises under truncation: selection concentrates on the most complementary line pairs, which pushes the pools further apart.

! A blind optimiser erodes one pool while global θ looks fine

Without this decomposition that goes unnoticed, and the consequence appears only in the following cycle. `ne_lines_pool()` is the cheap version of the same alarm – it needs no `fL`, only the pedigree:

```
t(sapply(sels, ne_lines_pool, ctx = ctx))
#>      Ne_lines_A Ne_lines_B
#> random    18.556701  20.93023
#> trunc_cap   15.517241  15.51724
#> greedy_w1    4.488778  11.84211
#> greedy_div   18.947368  18.36735
#> truncation   4.488778  11.84211
```

11.8.1 Where this could go next

θ_A , θ_B and θ_{AB} as computed here are **diagnostic**: reported, not fed back into the objective. The subdivided-population literature offers a ready-made way to promote them into a tunable target. [A. Caballero](#)

and S. T. Rodríguez-Ramilo [26], used by E. López-Cortegano, R. Pouso, A. Labrador, A. Pérez-Figueroa, J. Fernández, and A. Caballero [24] to manage two heterotic-like pools simultaneously, partition total diversity as

$$H_T = D_G + \lambda H_S, \quad A_T = D_A + \lambda A_S, \quad (11.11)$$

where H_S and A_S are the within-pool components, D_G and D_A the between-pool components, and $\lambda \in [0, \infty)$ is a weight chosen by the breeder. $\lambda = 0$ optimises only between-pool divergence – protects heterosis, ignores each pool’s internal health. $\lambda \rightarrow \infty$ optimises only within-pool diversity – protects each pool’s future options, ignores the cross.

That maps directly onto the two constraints a maize breeder already has in mind but applies only qualitatively: *don’t blend the pools* versus *don’t paint either pool into a corner*. Turning λ into a second explicit knob alongside $\alpha_{\max}$ is the natural next step if pool erosion shows up in practice. The current implementation reports the ingredients and leaves the weighting to the analyst.

λ need not stay a matter of taste. Chapter 10. shows that the value of between-pool divergence is not a free parameter but a d -weighted quantity: the divergence term contributes $\sum_k d_k g_k^2 / 2$ against the shared term’s $\sum_k d_k 2\bar{p}_k(1 - \bar{p}_k)$. Given an estimate of the dominance deviations, the ratio of those two sums is what λ is trying to express, and the reason a purely frequency-based partition cannot pin it down is that d does not appear in the partition at all.

11.9 The correct baseline

11.9.1 Why not compare against the whole population

The natural comparison – coancestry of the selected set against that of the whole population – is **biased by construction**. Applying equation (1) to both:

$$\theta_S - \theta_N = \left(\frac{1}{n} - \frac{1}{N} \right) (\bar{f}_d - \bar{f}_o) > 0, \quad (11.12)$$

since $\bar{f}_d > \bar{f}_o$ whenever individuals are not clones. A **random** sample already looks less diverse than the population it came from, with no selection whatsoever. The bias is pure sampling and grows as the sample shrinks.

```
fd <- mean(diag(ctx$f)); fo <- mean(ctx$f[upper.tri(ctx$f)])
predicted_bias <- (1 / n_sel - 1 / ctx$N) * (fd - fo)

set.seed(11)
observed_bias <- mean(replicate(400, theta_group(sample.int(ctx$N, n_sel), ctx))) -
  theta_group(seq_len(ctx$N), ctx)

c(predicted = predicted_bias, observed = observed_bias)
#> predicted observed
#> 0.002222092 0.002276690
```

The formula is exact; the residual is Monte Carlo error in the resampling.

Note also that the **off-diagonal** mean of a random subset is an unbiased estimator of the population’s off-diagonal mean. So the metric many programmes already use is free of this particular bias – but it is also not comparable across sample sizes on the N_s scale.

11.9.2 Operational definition of loss

The “0% loss” reference is the **distribution of random subsets of the same size**:

$$\alpha = \frac{GD_{\text{random}(n)} - GD_S}{GD_{\text{random}(n)}}. \quad (11.13)$$

$\alpha = 0$ means “as diverse as a random sample of the same size”; $\alpha < 0$ means more diverse than chance. Resampling also delivers the standard deviation, and hence a z-score for any strategy.

```

null <- null_distribution(ctx, n_sel, B = 400, B_full = 60)
c(GD_population = gene_diversity(seq_len(ctx$N), ctx),
  GD_reference   = null$gd_ref,
  sd             = null$gd_sd,
  apparent_loss_if_wrong_baseline =
    (gene_diversity(seq_len(ctx$N), ctx) - null$gd_ref) /
    gene_diversity(seq_len(ctx$N), ctx))
#>                GD_population                GD_reference
#>                0.3280928000                0.3258385719
#>                sd apparent_loss_if_wrong_baseline
#>                0.0007654486                0.0068707028

```

That last number is what a programme would report as a diversity loss if it compared against the full pool. It is entirely a sampling artefact.

```

fmt(sapply(sels, alpha_loss, ctx = ctx, gd_ref = null$gd_ref), 5)
#>    random trunc_cap greedy_w1 greedy_div truncation
#>    0.00032   0.00326   0.04122  -0.01352   0.04122

```

Negative values are selections *more* diverse than chance, which is exactly what the max-diversity greedy is for.

11.9.3 How to set the α limit

The acceptable limit **should not be chosen a priori**. Pure truncation – the greediest possible selection – defines the maximum attainable α . If that ceiling is below your intended limit, the constraint restricts nothing and you have bought a false sense of control.

```

attainable <- alpha_loss(sel_truncation(ctx, n_sel), ctx, null$gd_ref)
c(attainable_alpha_max = attainable,
  intended_limit       = 0.05,
  constraint_is_active = attainable > 0.05)
#> attainable_alpha_max    intended_limit constraint_is_active
#>                0.04121618                0.05000000                0.00000000

```

At the full simulated scale that ceiling comes out near 3.6%, so a 5% limit would be inoperative. The correct procedure is to trace the α -versus-gain frontier and choose its knee, which is Chapter 16..

11.10 Every metric against the null

A metric without a baseline says nothing. The z-score of each metric against the random-subset null is how this book reports all of them:

```

tab <- evaluate(sels, ctx, null$gd_ref)
fmt(as.data.frame(discriminatory_power(tab, null)[
  , c("GD", "Ns", "Ne_parents", "F_hom", "F_drift", "alleles_lost", "ENE", "ANE")]), 1)

```

	GD	Ns	Ne_parents	F_hom	F_drift	alleles_lost	ENE	ANE
random	0.0	0.0	1.4	0.6	-1.2	0.0	-1.0	-0.4
trunc_cap	-1.3	-1.3	-2.7	2.2	3.1	3.3	-0.7	17.9
greedy_w1	-17.9	-17.6	-11.4	9.5	30.1	14.9	-1.5	10.2
greedy_div	6.1	6.2	0.4	-5.2	-0.7	-0.3	5.5	6.3
truncation	-17.9	-17.6	-11.4	9.5	30.1	14.9	-1.5	10.2

Read the rows against each other. Truncation is many standard deviations from random on **F_drift** and on **Ne_parents** while being far less extreme on **GD** – the constraint that would catch it is not the one most programmes apply. That observation is what Chapter 12. turns into a recommendation.

12. Choosing metrics: collapse, cost, redundancy

Chapter 11. built a catalogue. This chapter asks which of it to keep, and answers with two questions the catalogue could not: **what does each metric cost per evaluation**, and **how much of what it reports is already in θ** .

The recommendation is **two** metrics inside the optimiser and **three** at validation.

Role	Metric	Where it runs	Cost per evaluation
Con-straint	Group coancestry theta of the selected set, as proportional loss of gene diversity against a same-size resampling baseline	inside DE, every evaluation	5 us (incremental) / 106 us (full)
Second con-straint	Effective number of parent lines	inside DE, every evaluation	7 us , shares the state vector
Validation	Marker alleles retained (allelic richness)	on the returned plan only	0.9 ms
Validation	Within/between heterotic-pool decomposition of theta	on the returned plan only	62 us
Validation	Status number $N_s = 1/(2 \text{ theta})$	free, algebraic	0

Three findings drive this, and each is derived and verified below rather than asserted.

12.1 Finding 1: the hybrid coancestry matrix is never needed

12.1.1 Derivation

Let lines $a, b, \dots$ be homozygous, with $\mathbf{f}^L$ their molecular coancestry matrix. The hybrid (ab) carries one genome from each parent. Molecular coancestry is the probability that two alleles, one drawn at random from each individual, are identical in state. Drawing at random from hybrid (ab) selects parental genome a or b with probability $\frac{1}{2}$ each, and likewise for (cd) , so

$$f_{(ab),(cd)} = \frac{1}{4} (f_{ac}^L + f_{ad}^L + f_{bc}^L + f_{bd}^L). \quad (12.1)$$

Now let a selection of n hybrids use line a exactly k_a times across all its parental slots, $\sum_a k_a = 2n$, and write $w_a = k_a/(2n)$. Group coancestry of the selected hybrids is the mean of f over all ordered pairs; substituting the identity above and collecting terms by line gives

$$\theta_S = \mathbf{w}^\top \mathbf{f}^L \mathbf{w}, \quad (12.2)$$

with the selection allele frequencies following the same collapse, $\bar{\mathbf{p}} = \mathbf{w}^\top \mathbf{G}^L$.

12.1.2 Four consequences

The $N \times N$ hybrid matrix is never required. Memory drops from $O(N^2)$ to $O(L^2)$. At $L = 300$ lines and $N = 22$, 500 hybrids that matrix would occupy 3.77 GB; in the benchmark below it was deliberately never formed, and every metric that needs it simply has no timing entry at that scale.

The vector w is a sufficient statistic. Group coancestry, gene diversity, status number, effective number of parents, selection allele frequencies and the heterotic-pool decomposition are all functions of w alone. One $O(L)$ tabulation per evaluation serves all of them.

Swapping one hybrid is a rank-2 update. Exchanging hybrid h for h' changes at most four entries of w , so θ updates in $O(L)$ rather than $O(L^2)$.

The discrete hybrid problem is exactly optimal-contribution selection on lines, restricted to the lattice $w_a \in \{0, 1/(2n), 2/(2n), \dots\}$. That connects it directly to the OCS literature of Chapter 15., with the caveat that the lattice restriction makes it combinatorial rather than the continuous quadratic programme OCS usually solves.

12.1.3 Verification

Every identity was checked against literal transcriptions of the reference implementations across $F_{ST} \in \{0.05, 0.15, 0.35\}$, two seeds and $n \in \{20, 100\}$, thirty replicate selections each.

	Identity	Worst deviation
max_pairwise_f	pairwise f: hybrid vs line formula	1.1e-16
max_theta_sub_vs_line	theta: line route vs f submatrix	2.2e-16
max_theta_sub_vs_freq	theta: line route vs allele-frequency route	1.1e-16
max_pbar	selection allele frequencies	0 (exact)
max_He	Nei He	0 (exact)
max_Neparents	effective number of parents	2.1e-14
max_alleles_lost	allele loss at MAF threshold	0 (exact)
max_theta_pools	pool decomposition	1.2e-15
max_swap_update	incremental swap vs full recompute	4.4e-16
bitset_vs_literal	bitset allelic richness vs literal count	0 (exact)

The same identity, live, on the book's own data:

```
ctx <- fast_ctx(simulate_data(book_cfg))
set.seed(51)
idx <- sample.int(ctx$N, 60)
c(via_hybrid_matrix = theta_group(idx, ctx),
  via_line_route     = fast_theta(idx, ctx),
  difference         = theta_group(idx, ctx) - fast_theta(idx, ctx))
#> via_hybrid_matrix via_line_route difference
#>      0.6745897      0.6745897      0.0000000
```

And the incremental swap, which is what makes a local search affordable:

```
st <- theta_state(idx, ctx)
out <- idx[1]; inn <- setdiff(seq_len(ctx$N), idx)[1]
st2 <- theta_swap(st, out, inn, ctx)
c(after_swap_incremental = st2$theta,
  after_swap_recomputed  = theta_group(c(idx[-1], inn), ctx),
  difference              = st2$theta - theta_group(c(idx[-1], inn), ctx))
#> after_swap_incremental after_swap_recomputed difference
#>      6.747225e-01      6.747225e-01      2.220446e-16
```

⚠ Two implementation traps found while building this

R's `$` performs **partial matching** on lists, so `ctx$fl` silently returns `ctx$fL` when the hybrid matrix is absent – a wrong answer with no error. Every accessor in the package uses `ctx[["f"]]`. Any pipeline carrying both matrices in one list is exposed to this.

R integers are signed 32-bit, so `bitwShiftL(1L, 31L)` is `NA` and the usual multiply-based popcount overflows. The `bitset` in `pack_lines()` packs 31 bits per word, and `popcount32()` counts them with a branch-free SWAR routine.

12.2 Finding 2: what each metric costs




Figure 12.1: Cost per fitness evaluation. Panel A: scaling in the number of candidate hybrids. Panel B: hours to complete one million evaluations at production scale.

At $L = 300$ lines, $N = 22,500$ hybrids, $m = 10,000$ markers and $n = 1,000$ selected:

metric	us per evaluation	hours per 1e6 evals
theta: incremental swap	5.0	0.0014
Ne_parents	7.2	0.0020
pool decomposition	62.0	0.0172
theta: LINE route	106.1	0.0295

metric	us per evaluation	hours per 1e6 evals
theta: line route precomp	108.0	0.0300
allelic richness: bitset	891.9	0.2478
allelic richness: float	5091.0	1.4142
allele loss (MAF thr)	5296.9	1.4714
He / gene diversity(freq)	50096.0	13.9156
theta: allele-freq route	51113.1	14.1981

Three thresholds matter:

- **The genotype-scan route is disqualified.** Computing θ from selection allele frequencies costs 14.2 hours per million evaluations. The line route costs 106 us – a **482x reduction** for an algebraically identical number.
- **The incremental swap is effectively free** at 5 us, or 1.4 seconds per million evaluations. $N_{e,par}$ at 7 us shares the same state vector and is free in practice once θ has been computed.
- **Allelic richness cannot go inside the loop, and need not.** The bitset implementation is 5.7x faster than the floating-point version, but still 178x the cost of the line-route θ . It belongs at validation.

The same ordering is visible at the book's scale, where every route is affordable and so the comparison is about ratios, not feasibility:

```
fmt(cost_per_eval(ctx, n_sel = 60, reps = 40), 1)
```

metric	us_per_eval
theta_group	25
theta_from_freq	175
ne_parents	0
alleles_lost	150
ENE	75
ANE	475
eff_dim	100
Gst	1350

12.3 Finding 3: the metric space is nearly one-dimensional

621 selections were generated across four strategy families – random, truncation at varying intensity, greedy construction over a grid of diversity weights, and restricted line pools – and all 17 metrics computed on each.




Figure 12.2: Metric redundancy. A: Spearman correlations. B: PCA of rank-transformed metrics. C: the discriminating case, inside a narrow coancestry band.

PC1 absorbs 88% of the rank variance; PC2 adds 7%. Rank correlations with θ are $|\rho| \geq 0.93$ for status number, H_e , effective dimensionality, $N_{e,par}$ and MAF-threshold allele loss. Reporting these alongside θ adds cost and length, not information.

But the marginal correlation is the wrong test. The question is whether a metric varies **at fixed θ** – because θ is the constraint, and any metric that is a deterministic function of it cannot discriminate among plans the optimiser treats as equivalent. Two analyses agree:

- After regressing each metric's ranks on θ (quadratic), **allelic richness retains 52% of its rank standard deviation** and $N_{e,par}$ retains 36%. Effective dimensionality retains 25% – close to a restatement of θ .
- Inside a band $\theta \in [0.666, 0.670]$ containing 399 selections, **alleles retained still spans 171 markers** (11,680 to 11,851 of 12,000) and $N_{e,par}$ spans **30.1 to 55.8**.

Plans the constraint cannot distinguish differ materially in how much allelic variation they carry forward. That is the empirical justification for the two-metric recommendation.

12.3.1 Reproducing the argument at book scale

```
set.seed(77)
grid <- do.call(rbind, lapply(seq(0, 4, length.out = 30), function(w) {
```

```

s <- sel_greedy(ctx, 60, w = w)
data.frame(w = w, theta = theta_group(s, ctx), Ne_par = ne_parents(s, ctx),
           lost = alleles_lost(s, ctx))
}))
rnd <- do.call(rbind, lapply(1:60, function(i) {
  s <- sample.int(ctx$N, 60)
  data.frame(w = NA, theta = theta_group(s, ctx), Ne_par = ne_parents(s, ctx),
             lost = alleles_lost(s, ctx))
}))
all <- rbind(grid, rnd)

plot(all$theta, all$Ne_par, pch = 19, col = ifelse(is.na(all$w), "grey60", "#b2182b"),
     xlab = expression(theta), ylab = expression(N[e]~"parents"))
legend("topright", c("greedy sweep", "random"), pch = 19,
     col = c("#b2182b", "grey60"), bty = "n")
grid()

```




Figure 12.3: At fixed theta, the effective number of parents still varies by a factor of two. The constraint cannot see the difference; the second metric can.

```

band <- all[abs(all$theta - median(all$theta)) < 0.0015, ]
c(n_in_band = nrow(band),
  theta_span = diff(range(band$theta)),
  Ne_par_min = min(band$Ne_par), Ne_par_max = max(band$Ne_par),
  alleles_lost_span = diff(range(band$lost)))
#>      n_in_band      theta_span      Ne_par_min      Ne_par_max
#>      55.000000000      0.002502847      31.578947368      41.860465116
#> alleles_lost_span
#>      31.000000000

```

i A caveat on transfer

These numbers come from a two-pool simulation at $F_{ST} = 0.15$. The qualitative conclusion – near one-dimensionality, with allelic richness as the exception – is robust in direction, but the exact residual fractions shift with pool divergence and marker density. Re-running the redundancy analysis on real genotypes is a half-hour job and worth doing before treating the recommendation as final.

12.4 Finding 4: the encoding, not the constraint handler, decides whether DE works

The optimiser was run under four simultaneous restrictions: diversity loss $\leq 0.5\%$, $N_{e,par} \geq 50$, at most 6 uses per line, and heterotic-pool balance within 5%. Budget fixed at 60,000 evaluations for every configuration, three seeds each.




Figure 12.4: Differential evolution under four simultaneous restrictions at equal evaluation budget.

encoding	index	loss %	Ne parents	max uses	seconds
hybrid keys / penalty	0.351	0.060	76.634	5.667	6.550
hybrid keys / repair	0.351	0.060	76.634	5.667	9.621
line weights / repair	1.988	0.448	50.084	6.000	12.126

The encoding decides the outcome; the constraint handler does not. One random key per hybrid gives DE a 3,600-dimensional problem and it reaches index 0.351. One weight per *line* – dimension 120, available only because of the collapse – reaches **1.988 at the same budget**, a 5.7x improvement. Penalty

and repair handlers produced numerically identical solutions in the hybrid encoding, because at that dimension DE never approaches the constraint boundary where the handlers differ.

Only the line encoding makes the constraints bind. It finishes with $N_{e,par} = 50.08$ against a floor of 50, max uses exactly 6 against a cap of 6, and loss 0.448% against a 0.5% ceiling. The hybrid encoding leaves every constraint slack – not because it found a better trade-off, but because it never searched far enough from random to reach the boundary.

A slack constraint at the end of a run is a diagnostic of a failed search, not of a conservative plan.

For calibration: unconstrained truncation reaches index 2.337 but violates everything (loss 1.62%, $N_{e,par} = 33.2$, one line used 17 times). The line encoding recovers **85% of the truncation index while satisfying all four restrictions.**

A note on tuning. DEoptim warns that NP should be at least ten times the parameter count; at dimension 3,600 that is 36,000 population members, which is infeasible. At dimension 120 the recommendation is 1,200 and entirely reachable. The reparameterisation does not merely speed the search up – it moves the problem into the regime where DE’s own tuning guidance can be followed.

12.5 Finding 5: the index is not the truth, and its error is not white

Findings 1 to 4 are about the diversity side of the objective. This one is about the other side, and it is the larger uncertainty of the two.

Everywhere in this book `ctx$index` is built from *known* marker effects (Chapter 4.), so every result so far assumes the breeding values are predicted without error. Real ones are not, and the error is not white noise: predictions for close relatives are wrong in the same direction, because they lean on the same training records and share the same haplotypes. The book already has a matrix of exactly that structure – the molecular coancestry **f** – so the error can be simulated without inventing anything. (Chapter 14. fits an actual GBLUP model later and finds the same structure in its real residuals, at a comparable accuracy, without injecting anything.)

```
n_sel <- 60
null <- null_distribution(ctx, n_sel, B = 400, B_full = 60)
gd <- null$gd_ref
u <- ctx$index
Lc <- t(chol(ctx$f + diag(1e-8, ctx$N)))

# A predicted index at accuracy r. `kind` decides whether the error is
# correlated with relatedness or independent between hybrids.
predicted <- function(r, seed, kind = "correlated") {
  set.seed(seed)
  z <- rnorm(ctx$N)
  e <- if (kind == "correlated") as.numeric(Lc %*% z) else z
  e <- (e - mean(e)) / stats::sd(e)
  as.numeric(scale(r * u + sqrt(1 - r^2) * e))
}
```

Select on the prediction, then score the plan on the truth:

```
run <- function(r, seed = 1, kind = "correlated") {
  ch <- ctx; ch$index <- predicted(r, seed, kind)
  sel <- list(truncation = sel_truncation(ch, n_sel),
             greedy = sel_greedy(ch, n_sel, w = 3e-4))
  do.call(rbind, lapply(names(sel), function(nm) data.frame(
    method = nm, r = r,
    reported = mean(ch$index[sel[[nm]]]),
    realised = mean(u[sel[[nm]]]),
    alpha = alpha_loss(sel[[nm]], ctx, gd))))
}
acc <- do.call(rbind, lapply(c(1, 0.9, 0.7, 0.5), run))
fmt(acc, 4)
```

method	r	reported	realised	alpha
truncation	1.0	1.9919	1.9919	0.0412
greedy	1.0	1.5893	1.5893	0.0133
truncation	0.9	1.9247	1.9003	0.0406
greedy	0.9	1.4745	1.3965	0.0092
truncation	0.7	1.7529	1.5950	0.0356
greedy	0.7	1.3851	1.1464	0.0088
truncation	0.5	1.6223	1.1751	0.0301
greedy	0.5	1.2591	0.7965	0.0056

Two things move. The gap between **reported** and **realised** opens up – at $r = 0.5$ truncation reports an index of 1.62 and delivers 1.18. Selection on a noisy value selects partly on the noise, so the plan’s own scorecard is biased upward. Nothing in the pipeline can detect this; it is invisible without the truth.

And α falls as accuracy drops. That is not the constraint working better. Noisy selection is closer to random selection, and random selection is what $\bar{H}_{\text{ref}}$ is defined against – so the plan drifts toward the baseline from which loss is measured.

12.5.1 The correlation structure matters more than the amount

Hold accuracy fixed at $r = 0.6$ and change only whether the error is correlated with relatedness, over twenty draws:

```
cmp <- do.call(rbind, lapply(c("correlated", "independent"), function(k)
  do.call(rbind, lapply(1:20, function(s) {
    ch <- ctx; ch$index <- predicted(0.6, s, k)
    tr <- sel_truncation(ch, n_sel)
    data.frame(kind = k, realised = mean(u[tr]),
               alpha = alpha_loss(tr, ctx, gd), max_line = max_line_use(tr, ctx))
  })))
fmt(aggregate(cbind(realised, alpha, max_line) ~ kind, cmp, mean), 4)
```

kind	realised	alpha	max_line
correlated	1.1508	0.0391	18.15
independent	1.2253	0.0121	12.60

Same accuracy, same method, same number selected. Correlated error costs roughly 3.2 times the diversity of independent error, and leans 1.4 times as hard on its favourite line.

The mechanism is direct. Independent error scatters the mistakes, which accidentally diversifies the plan. Error correlated with relatedness makes whole families look good together, so the plan takes families – and a family is exactly what the diversity constraint exists to stop it taking.

⚠ The worst case is not the noisiest one

Truncation on a perfect index loses 4.12% of gene diversity and earns every bit of the index it reports. Under correlated error at $r = 0.6$ it still loses 3.91% – nearly the same price – while delivering 1.15 instead of 1.99.

You pay the full diversity cost of aggressive selection and collect a fraction of the gain. That is the case to design against, and it is the ordinary case in a real programme.

12.5.2 What this does to the recommendation

Nothing, for the metric set. Findings 1 to 3 are statements about f and hold whatever the index is. But it changes what the numbers in Chapter 21. mean, and it reorders the priorities: the choice between two

diversity metrics is worth a fraction of a percent of index, and prediction accuracy is worth tens of percent. Report α against the plan you actually planted, validate accuracy before tuning the constraint, and treat a reported index as an upper bound.

12.6 The recommendation, stated

12.6.1 Inside the optimiser

Use two metrics, both functions of the line-usage vector $\mathbf{w}$:

1. **Proportional loss of gene diversity**, $(\bar{H}_{\text{ref}} - H_S)/\bar{H}_{\text{ref}}$ with $H_S = 1 - \mathbf{w}^\top \mathbf{f}^L \mathbf{w}$ and $\bar{H}_{\text{ref}}$ the mean over same-size random subsets. Compute it by the line route, never by scanning genotypes.
2. **Effective number of parent lines**, $N_{e,\text{par}} = (\sum k_a)^2 / \sum k_a^2$. It costs 7 us, shares the state vector, and blocks concentration onto a few lines – a failure mode θ tolerates.

Reparameterise the search on **line weights**, and enforce the per-line usage cap in the decoder rather than by penalty.

12.6.2 At validation, on the returned plan only

3. **Marker alleles retained**, via the bitset route. The one metric carrying information θ does not.
4. **Within/between heterotic-pool decomposition**, to confirm the plan has not preserved diversity by degrading pool separation.
5. **Status number** $N_s = 1/(2\theta)$, free from θ , as the reportable descriptor. Do not report a rate-based N_e under active management.

12.6.3 Not recommended

Nei's H_e as a separate metric (it is $1 - \theta$); effective dimensionality ($|\rho| = 0.97$ with θ , and needs the hybrid matrix); MAF-threshold allele loss ($|\rho| = 0.93$, and 1.47 hours per million evaluations); Core Hunter's E-NE and A-NE (both require the $O(N^2)$ matrix and are unavailable at production scale – they are designed for core-collection assembly, a different problem).

12.6.4 Caveats

The choice of $\mathbf{f}^L$ remains open in the literature. Post-2018 work is consistent that molecular or IBD-flavoured coancestry beats a VanRaden-style drift matrix for diversity management, but reports no consensus on a single preferred formulation, and matrix choice demonstrably changes allele-frequency trajectories (Chapter 18.). The collapse derived here is a property of the F1 structure and holds for **any** symmetric line-level $\mathbf{f}^L$ – so if the matrix choice changes, every cost and identity result in this chapter survives unchanged.

Managing on molecular coancestry controls homozygosity, not drift; these diverge, and a plan optimal on one is not optimal on the other (Chapter 9.). Predicted Δf is biased downward, with the bias growing as the candidate pool shrinks – report realised θ on the chosen plan rather than a predicted rate.

12.7 Drop-in code

```
ctx <- fast_ctx(ctx, GL = line_genotypes, pack = TRUE) # no N x N matrix needed

gd_ref <- mean(replicate(500, fast_gene_div(sample.int(ctx$N, n_sel), ctx)))

fit <- make_fitness_fast(ctx, n_sel, gd_ref, alpha_max = 0.005,
                        ne_par_min = 50, max_use = 6, penalty = 5,
                        encoding = "lines")

res <- DEoptim::DEoptim(fit,
                        lower = rep(0, ctx$n_lines), upper = rep(1, ctx$n_lines),
                        control = DEoptim::DEoptim.control(NP = 200, itermax = 300))
```

```
idx <- decode_lines(res$optim$bestmem, ctx, n_sel, max_use = 6)
c(theta      = fast_theta(idx, ctx),
  loss_pct   = 100 * (gd_ref - fast_gene_div(idx, ctx)) / gd_ref,
  Ne_parents = fast_ne_parents(idx, ctx),
  alleles    = alleles_retained_fast(idx, ctx),
  fast_theta_pools(idx, ctx))
```


13. Metrics at each pipeline stage, and five traps

The previous chapters established *which* metrics to use. This one places them at the four stages of a hybrid maize programme, and names five traps – each one a thing that looks correct, returns a number, and is wrong.

Stage	What is selected
S1	conserving the heterotic groups (pool level)
S2	choosing parents and crosses for DH induction
S3	selecting lines per se
S4	hybrid mining

Two rules govern everything below. Every formula was verified numerically before it was written; where a result is an exact algebraic identity, the maximum deviation from simulation is reported next to it. And the metric set is the one established in Chapter 12.: molecular coancestry as the kernel, $GD = 1 - \theta$ as the scale, N_s as the interpretable size, allelic richness as the one non-redundant addition, and proportional loss measured against a same-size resampling baseline.

```
sim <- simulate_pools(n_A = 30, n_B = 30, m = 3000, fst = 0.15, seed = 12)
ctx <- ctx_from_pools(sim)
GL <- sim$GL; pool <- sim$pool; fL <- sim$fL
c(lines = sim$L, markers = sim$m, hybrids = sim$N)
#>   lines markers hybrids
#>    60    3000     900
```

13.1 The kernel: one matrix, all four stages

Molecular coancestry f_{ij} is the probability that an allele drawn at random from individual i is identical in state to one drawn from j . For inbred lines coded 0/1 at m biallelic markers,

$$f_{ij} = 1 - \frac{1}{m} \sum_{k=1}^m (x_{ik} - x_{jk})^2. \quad (13.1)$$

Group coancestry of a set is the mean of the full submatrix **including the diagonal**, $\theta = n^{-2} \sum_i \sum_j f_{ij}$, with $GD = 1 - \theta$ and $N_s = 1/(2\theta)$.

The diagonal is not a technicality. It is why a small set always looks highly coancestral, and it is the reason the naive greedy constraint in S3 fails.

13.1.1 Trap 1: modified Rogers distance is the same statistic

Reporting both MRD and coancestry as separate evidence is reporting one number twice. The relation is exact:

```
D <- mrd_matrix(GL)
c(max_deviation = max(abs(coancestry_from_mrd(D) - fL)))
#> max_deviation
#> 1.110223e-16
```

$f_{ij} = 1 - \text{MRD}_{ij}^2$. If a diversity report shows a distance tree and a coancestry table as mutual corroboration, they are the same analysis in two coordinate systems.

13.2 S1: conserving the heterotic groups

Diversity in a hybrid programme is **two quantities moving in opposite directions**: within-pool diversity erodes under selection, while between-pool divergence inflates. A single pool-wide statistic cannot express both, and a programme monitoring only total diversity can look stable while both pools collapse internally and drift apart.

The partition is exact for any grouping:

$$\theta_T = \sum_p w_p^2 \theta_{pp} + \sum_{p \neq q} w_p w_q \theta_{pq} \quad (13.2)$$

```
d <- theta_decompose(GL, pool)
c(theta_T_from_partition = d$theta_T,
  theta_T_direct         = mean(fL),
  deviation              = d$theta_T - mean(fL))
#> theta_T_from_partition      theta_T_direct      deviation
#>          0.6614378          0.6614378          0.0000000
```

13.2.1 Trap 2: two F_{ST} definitions differ by up to 2x

Both are in routine use, both are computed from the same coancestry matrix, and they are **not interchangeable**:

$$F_{ST}^{(W)} = \frac{\theta_w - \theta_B}{1 - \theta_B} \quad G_{ST} = \frac{\theta_w - \theta_T}{1 - \theta_T} \quad (13.3)$$

The first references *between-pool* coancestry; the second references *total* pooled-frequency coancestry and equals Nei's $G_{ST} = (H_T - H_S)/H_T$:

```
c(Fst_Wright = d$Fst_Wright,
  Gst_Nei     = d$Gst_Nei,
  Gst_classical = unname(gst_nei(GL, pool)[["Gst"]]),
  ratio       = d$Fst_Wright / d$Gst_Nei)
#>   Fst_Wright   Gst_Nei Gst_classical   ratio
#>   0.1824006   0.1003525   0.1003525   1.8175994
```

Because $\theta_T > \theta_B$ whenever the pools differ at all, $G_{ST} < F_{ST}$ **always**. Across simulated divergence they never converge:

fst_target	Fst_Wright	Gst_Nei	ratio	theta_B	theta_T	GD_within	frac_opp	mean_dp
0.02	0.0446	0.0228	1.9554	0.6372	0.6453	0.3466	0.0000	0.0991
0.05	0.0742	0.0385	1.9258	0.6361	0.6496	0.3369	0.0000	0.1280
0.10	0.1200	0.0638	1.8800	0.6367	0.6585	0.3197	0.0000	0.1607
0.15	0.1676	0.0915	1.8324	0.6337	0.6644	0.3049	0.0003	0.1907
0.20	0.2268	0.1279	1.7732	0.6289	0.6710	0.2869	0.0013	0.2211
0.30	0.3170	0.1883	1.6830	0.6413	0.6982	0.2450	0.0063	0.2504
0.40	0.4106	0.2584	1.5894	0.6341	0.7092	0.2157	0.0230	0.2840

State which definition you used, and never compare your value against a published threshold computed the other way.

13.2.2 Trap 3: F_{ST} is the wrong target to maximise

F_{ST} keeps rising as pools drift apart at loci that are *already* divergent, which buys no additional heterosis. Chapter 10. makes the reason exact: under neutral divergence expected heterosis does not depend on F_{ST} at all, because every unit gained in the divergence term is paid for out of the shared term. What does

move it is divergence landing on the dominance loci. The frequency-side statistics decouple from F_{ST} accordingly:

```
sweep_fst <- do.call(rbind, lapply(c(0.05, 0.15, 0.25, 0.40), function(fst) {
  s <- simulate_pools(n_A = 25, n_B = 25, m = 2000, fst = fst, seed = 3)
  dd <- theta_decompose(s$GL, s$pool)
  cm <- pool_complementarity(s$GL, s$pool)
  data.frame(target_fst = fst, Fst = dd$Fst_Wright, GD_within = dd$GD_within,
    frac_opposite_fixed = cm[["frac_opposite_fixed"]],
    mean_abs_delta_p = cm[["mean_abs_delta_p"]],
    mean_sq_delta_p = cm[["mean_sq_delta_p"]])
}))
fmt(sweep_fst, 4)
```

target_fst	Fst	GD_within	frac_opposite_fixed	mean_abs_delta_p	mean_sq_delta_p
0.05	0.0907	0.3321	0.0005	0.1412	0.0331
0.15	0.1857	0.3003	0.0000	0.1995	0.0685
0.25	0.2711	0.2644	0.0045	0.2379	0.0984
0.40	0.4348	0.2106	0.0370	0.2940	0.1620

Buying oppositely-fixed loci costs within-pool gene diversity, and it costs a great deal of it. **Hold F_{ST} in a band, not at a maximum**, and monitor complementarity directly.

Of the three complementarity columns, `mean_sq_delta_p` is the one to log. Chapter 10. shows it is not a proxy at all but exactly twice the divergence term of expected heterosis, while `frac_opposite_fixed` sits at essentially zero across the whole range of divergence a real dent-versus-flint pair occupies, and so cannot discriminate between programmes.

13.2.3 The S1 monitoring panel

```
fmt(s1_monitor(GL, pool, cycle = 12), 4)
```

cycle	GD_within	GD_total	Fst_Wright	Gst_Ne	Ns_total	GD_AGD	frac_opposite_fixed	mean_abs_delta_p	mean_sq_delta_p	
12	0.3046	0.3386	0.1824	0.1004	0.7559	0.3050	0.3042	0	0.2003	0.068

Log this every cycle and plot it as a time series. Defend `GD_within` per pool; watch `mean_sq_delta_p` for heterosis potential; treat F_{ST} as a descriptor, not an objective. Chapter 19. is what the time series looks like when it is actually run out to fifteen cycles.

The empirical anchor is Kadoumi et al. (2025), who tracked diversity erosion and heterotic-group divergence over decades of a commercial European dent maize programme – and found exactly the two opposite trends this panel is built to separate. See Chapter 18..

13.3 S2: choosing crosses for DH induction

This is where the largest share of pipeline diversity is lost, and also where the cheapest exact algebra is available.

13.3.1 Three verified identities

For a DH family from cross (i, j) of two inbreds, each locus is inherited from i or j with probability $1/2$ and then doubled.

- (a) Coancestry between two DH sibs of the same family: $\mathbb{E}[f] = (1 + f_{ij})/2$.
- (b) Coancestry between a DH of family (i, j) and one of (k, l) : $\mathbb{E}[f] = (f_{ik} + f_{il} + f_{jk} + f_{jl})/4$.
- (c) Group coancestry of a DH pool from crosses $c = 1 \dots C$, each contributing n_c lines, $D = \sum n_c$:

$$\theta_{pool} = \mathbf{w}^\top \mathbf{f}^L \mathbf{w} + \frac{1}{D} \left(1 - \overline{\left(\frac{1 + f_c}{2} \right)} \right) \quad (13.4)$$

with $\mathbf{w}$ the **parent-usage weight vector**. The second term is the self-coancestry correction; it is $O(1/D)$ and worth keeping.

```
set.seed(88)
A <- 1:30
crosses <- cbind(sample(A, 20, replace = TRUE), sample(A, 20, replace = TRUE))
crosses <- crosses[crosses[, 1] != crosses[, 2], , drop = FALSE]

c(with_self_term = dh_pool_theta(crosses, n_per = 40, fL, self_term = TRUE),
  without_self_term = dh_pool_theta(crosses, n_per = 40, fL, self_term = FALSE),
  sib_coancestry_ex = dh_sib_coancestry(fL[crosses[1, 1], crosses[1, 2]]))
#>      with_self_term without_self_term sib_coancestry_ex
#>      0.7010910      0.7008837      0.8493333
```

! The key structural result of S2

The diversity of the DH pool is a quadratic form in *which parents you use and how often*. **The DH lines never need to be genotyped – or to exist – for their pool diversity to be optimised.** You can evaluate a crossing plan before making a single cross.

13.3.2 Progeny variance, and the usefulness criterion

Under an additive model only loci segregating in the cross contribute:

$$\sigma_{DH(i,j)}^2 = \frac{1}{4} \sum_{k: x_{ik} \neq x_{jk}} \beta_k^2 \quad (13.5)$$

```
sig <- sigma_dh_crosses(crosses, GL, sim$beta)
mu <- rowMeans(cbind(GL[crosses[, 1], ] %*% sim$beta,
                     GL[crosses[, 2], ] %*% sim$beta))
uc <- usefulness_criterion(mu, sig, i = sel_intensity(0.05), h = 0.6)
c(sel_intensity_5pct = sel_intensity(0.05), textbook = 2.063)
#> sel_intensity_5pct      textbook
#>      2.062713      2.063000
fmt(head(data.frame(p1 = crosses[, 1], p2 = crosses[, 2],
                    mu = mu, sigma_dh = sig, UC = uc), 6), 3)
```

p1	p2	mu	sigma_dh	UC
23	27	0.930	4.609	6.634
13	11	-5.977	5.032	0.251
22	21	-3.481	5.264	3.034
7	19	-7.343	5.063	-1.076
19	1	-11.512	5.285	-4.971
2	12	-2.152	4.934	3.955

`sel_intensity(0.05)` returns 2.0627 against the textbook 2.063.

13.3.3 Trap 4: genetic distance is not a proxy for progeny variance

It is common practice to rank crosses by marker distance, on the argument that more distant parents segregate for more loci. Once the heterotic pattern is fixed, this fails:

stratum	cor	R2
inter-pool (A x B)	0.103	0.011
within-pool (A x A)	0.219	0.048
pooled	0.576	0.332



Figure 13.1: Traps 4 and 2, on the full simulation. Panel A: progeny standard deviation against modified Rogers distance – the strong pooled correlation only re-separates inter- from within-pool crosses and carries almost no information within either stratum. Panel B: the two F_{ST} definitions on identical data, differing only in the reference in the denominator.

A raw count of segregating loci explains the same variance – distance is adding nothing beyond “are these parents in different pools”, which you already know.

Rank crosses on `sigma_dh_crosses()` computed from marker effects, not on distance.

13.4 S3: selecting lines per se

Two points of practice.

Apply the constraint within each pool separately. A pooled constraint can be satisfied by keeping the pools apart while both erode internally – exactly the failure mode S1 exists to prevent.

Measure loss against a same-size resampling baseline.

```
merit <- as.vector(GL %*% sim$beta)
merit <- (merit - mean(merit)) / sd(merit)
top <- order(merit, decreasing = TRUE)[1:20]
fmt(loss_vs_resampling(top, fL, B = 300), 4)
#>      GD_selected GD_random_mean      loss_pct      percentile
#>      0.3245      0.3271      0.8029      0.0700
```

13.4.1 Trap 5: the incremental greedy constraint silently returns nothing

The natural implementation – “add the next-best line while θ stays below the ceiling” – **cannot work**, and fails silently by returning an empty set. Group coancestry includes the diagonal, so a single line has $\theta = 1$ and any pair has $\theta = (1 + f_{ij})/2$. A ceiling calibrated on a large pool is unreachable at small set sizes; the greedy never gets past the first candidate.

```
ceiling20 <- theta_ceiling(fL, n_sel = 20, alpha = 0.005)
c(ceiling      = ceiling20,
   theta_single_line = fL[1, 1],
   theta_best_pair  = min((1 + fL[upper.tri(fL)]) / 2),
   any_pair_admissible = any((1 + fL[upper.tri(fL)]) / 2 <= ceiling20))
#>      ceiling  theta_single_line  theta_best_pair any_pair_admissible
#>      0.6746394      1.0000000      0.8008333      0.0000000
```

No single line and no pair can satisfy the ceiling. θ falls monotonically as the set grows, so **the constraint is only meaningful at the target size**:

```
sizes <- c(2, 3, 5, 8, 12, 20, 30, 45, 60)
th <- vapply(sizes, function(n) mean(replicate(50, {
  s <- sample.int(nrow(fL), n); mean(fL[s, s]))}), numeric(1))
plot(sizes, th, type = "b", pch = 19, xlab = "set size", ylab = expression(theta))
abline(h = ceiling20, col = "#b2182b", lty = 2, lwd = 2)
legend("topright", c("mean theta of a random set", "ceiling for n = 20"),
      col = c("black", "#b2182b"), lty = c(1, 2), pch = c(19, NA), bty = "n")
grid()
```




Figure 13.2: Group coancestry falls monotonically with set size. A ceiling calibrated at the target size is unreachable at every smaller size, which is why an incremental greedy stops at the first candidate.

The working route is **truncate-then-repair**: take the top n_{sel} by merit, then swap the most-related member for the least-related non-member until the ceiling is met, breaking ties on merit.

```
sel <- select_lines_constrained(merit, fL, n_sel = 20, theta_max = ceiling20)
c(n_selected = length(sel),
  theta      = s3_theta(sel, fL),
  ceiling    = ceiling20,
  met        = s3_theta(sel, fL) <= ceiling20,
  mean_merit = mean(merit[sel]),
  merit_of_top = mean(merit[top]))
#>   n_selected      theta      ceiling      met  mean_merit merit_of_top
#>   20.0000000    0.6755400    0.6746394    0.0000000    1.0724564    1.0724564
```

`theta_ceiling()` expresses the ceiling as a permitted **proportional loss** against the same-size baseline, so the constraint means the same thing at every set size.

13.5 S4: hybrid mining

With $\mathbf{u}$ the line-usage count vector over the n selected hybrids,

$$\theta_{hyb} = \frac{\mathbf{u}^\top \mathbf{f}^L \mathbf{u}}{(2n)^2} \quad (13.6)$$

```
set.seed(91)
hyb <- sample.int(ctx$N, 100)
u <- line_usage(hyb, ctx$parents, ctx$n_lines)
c(theta_from_usage = hybrid_theta_from_usage(u, fL),
  theta_from_matrix = ref_theta(hyb, ctx),
  deviation = hybrid_theta_from_usage(u, fL) - ref_theta(hyb, ctx))
#>   theta_from_usage theta_from_matrix      deviation
#>   0.6620848      0.6620848      0.0000000
```

Deviation exactly zero. This is the collapse of Section 12.1: cost drops from $O(n^2)$ over an $N \times N$ hybrid matrix to $O(L^2)$ with no hybrid matrix formed at all.

Expected heterozygosity of the hybrid set needs no separate computation: for hybrids of fixed lines it is exactly $1 - \theta_{hyb}$.

```
c(theta = hybrid_theta_from_usage(u, fL),
  He     = 1 - hybrid_theta_from_usage(u, fL),
  Ne     = ne_parents_from_counts(u),
  pool_balance = pool_balance(hyb, ctx$parents, pool))
#>   theta      He      Ne pool_balance1 pool_balance2
#>   0.6620848    0.3379152  49.6277916    0.5000000    0.5000000
```

13.6 Where the diversity actually goes

The four stages were simulated end to end over **8 independent founder pools** (120 lines, 6000 markers, target $F_{ST} = 0.15$; 20 crosses, 40 DH each, 80 lines selected, 200 hybrids mined), under two cross-selection strategies: merit only, and merit plus a diversity term on parental coancestry.

Loss at each stage is proportional loss of gene diversity against that stage's **own same-size resampling baseline**, so the stages sit on one scale.

stage	loss_pct.UC + div	share_of_total.UC + div	loss_pct.UC only	share_of_total.UC only
S2 DH induction	3.43	35.51	9.40	77.46
S3 line per se	2.98	28.99	0.32	2.55

stage	loss_pct.UC + div	share_of_total.UC + div	loss_pct.UC only	share_of_total.UC only
S4 hybrid min- ing	3.47	35.49	2.48	19.99




Figure 13.3: Where the diversity goes. DH induction dominates, and a diversity term applied there redistributes loss rather than removing it.

DH induction dominates. Under merit-only cross selection it accounts for 77% of all diversity lost in the pipeline – more than three times the loss at hybrid mining, which is where optimisation effort conventionally goes.

The reason is structural: 20 crosses out of 1770 possible within-pool combinations is a far harsher bottleneck than 200 hybrids out of 1600, and the S2 collapse makes it the cheapest stage to constrain.

13.6.1 What the diversity term buys, and what it does not

Outcome	merit only	merit + diversity	paired t ($n = 8$)
hybrid merit	33.4	29.71	$p = 0.0002$
effective no. parent lines	24.2	25.33	$p = 0.24$
marker alleles retained	11274.0	11314.00	$p = 0.52$

The merit cost is **11%, consistent in all 8 pools**. Neither end-of-pipeline diversity outcome improves significantly. The diversity term redistributes loss across stages – it flattens the S2 spike – but with downstream selection still running on merit alone, the pipeline reconverges. Diversity spent at S2 is partly reclaimed at S3 and S4.

A diversity constraint applied at one stage buys stage-level diversity, not end-of-pipeline richness. If allelic richness in the commercial hybrid set is the target, the constraint has to be carried through all four stages, and the loss ledger is the instrument for allocating a budget across them.

13.7 Cost summary

Stage	Function	Cost	Notes
S1	<code>theta_decompose</code>	$O(L^2)$ once per cycle	monitoring, not an inner loop
S2	<code>dh_pool_theta</code>	$O(L^2)$ per plan	no DH genotypes needed
S2	<code>sigma_dh_crosses</code>	$O(C\ m)$	the real cost at this stage
S3	<code>select_lines_constrained</code>	$O(\text{swaps} \times n)$	swaps typically under 20
S4	<code>hybrid_theta_from_usage</code>	$O(L^2)$ per evaluation	no $N \times N$ matrix

The $O(L^2)$ quadratic form is the recurring shape. With L in the hundreds it is microseconds, which is why the whole pipeline can be optimised inside a metaheuristic. The dominant cost anywhere here is `sigma_dh_crosses` at S2, because it touches the marker matrix – cache it per candidate cross rather than recomputing inside the optimiser loop.

13.8 Checklist

1. Compute the coancestry matrix once per cycle; never report modified Rogers distance alongside it as separate evidence.
2. Monitor within-pool $\overline{GD}$ per pool, not pooled $\overline{GD}$; state which F_{ST} definition you use; hold it in a band.
3. Rank crosses on σ_{DH} from marker effects, never on parental distance.
4. Evaluate the crossing plan's θ_{pool} *before* making the crosses – it is a quadratic form in parent usage.
5. Calibrate every diversity ceiling against a same-size resampling baseline, and apply it at the target set size.
6. Constrain within pool, at every stage – a single-stage constraint does not survive to the hybrid set.
7. At S4, use the line-usage collapse; the hybrid coancestry matrix is never needed.
8. Keep the per-stage loss ledger. It is the only way to see that DH induction, not hybrid mining, is where the diversity goes.

14. Predicting breeding values: GBLUP, BayesA, BayesB

Three earlier chapters flagged the same gap and moved on. Chapter 4. built `ctx$traits` from *known* marker effects and called the result “an optimistic ceiling.” Chapter 12.’s Finding 5 showed what happens once the index carries real error, but manufactured that error rather than fitting a model to produce it. Chapter 21. repeats the caveat and still moves on. This chapter is where the gap closes: it fits actual prediction models to a training population with real phenotypic noise, and hands the result to the same `stage1_build()` contract every later chapter already consumes.

`vanraden_G()` has said since Chapter 1. that its matrix “is the right matrix for GBLUP.” Nothing in the package has used it that way until now.

The chapter also has a second job, which only became apparent once the models were fitted: an accuracy number means nothing until you say what was hidden from the model. Section 14.4 shows that the obvious split – hold out a random 40% of the hybrids – answers the easiest of the three questions a hybrid programme actually asks, and that the fit itself says so, before the truth is consulted.

14.1 A training population with real noise

Everywhere else in this book, `ctx$index` is the truth. Here it plays a different role: it is the value a real breeding programme does not get to see, and can only estimate from phenotypes recorded on a subset of individuals.

```
ctx <- simulate_data(book_cfg)
N   <- ctx$N
g   <- ctx$traits[, 1]          # the true genetic value -- known only to us
```

A phenotype is the genetic value plus environmental noise, $y = g + e$, and heritability is the fraction of phenotypic variance the genetic part accounts for, $h^2 = \text{Var}(g) / (\text{Var}(g) + \text{Var}(e))$. `traits` is built by `scale()`, so $\text{Var}(g) = 1$ exactly, and the noise variance that delivers a target h^2 is $(1 - h^2) / h^2$:

```
h2 <- 0.5
set.seed(9)
e <- rnorm(N, 0, sqrt((1 - h2) / h2))
y <- g + e
```

Split the population into a training set, whose phenotypes the model gets to see, and a candidate set – the hybrids actually being selected among – whose phenotypes it does not:

```
set.seed(9)
train <- sample.int(N, floor(0.6 * N))
cand  <- setdiff(seq_len(N), train)
c(train = length(train), candidates = length(cand))
#>      train candidates
#>      375         250
```

14.2 The truth this chapter is predicting

Two properties of g decide everything that follows, and both are consequences of how Chapter 4. built it rather than choices made here.

14.2.1 The generative model is GBLUP's own assumption

`ctx$traits` is `scale(Z %*% B)` with the marker effects $\mathbf{B}$ drawn independently. For a single trait that is $\mathbf{g} = \mathbf{Z}\beta$ with $\text{Var}(\beta) = \mathbf{I}\sigma_\beta^2$, so

$$\text{Var}(\mathbf{g}) = \mathbf{Z}\mathbf{Z}'\sigma_\beta^2 = \mathbf{G} \underbrace{2 \sum_k p_k(1-p_k)}_{\sigma_u^2} \sigma_\beta^2, \quad (14.1)$$

because $\mathbf{G}$ is *defined* as $\mathbf{Z}\mathbf{Z}'$ over that same divisor (Chapter 1.). The second equality is exact, not approximate:

```
vr <- vanraden_G(ctx$X * 2)
s2 <- 2 * sum(vr$p * (1 - vr$p))
max(abs(tcrossprod(vr$Z) / s2 - ctx$G))
#> [1] 0
```

So the covariance GBLUP *assumes* for the breeding values is, up to the scalar σ_u^2 , the covariance the breeding values actually have. This simulation is the infinitesimal model, exactly and not approximately, and that single fact is most of the answer to “which model wins” below. It is also the sense in which every accuracy in this chapter is a best case: no model here is misspecified.

14.2.2 A hybrid's value is the mean of two parental values

The hybrids are F1s of homozygous lines, so a hybrid's marker row is the average of its two parents' rows, and averaging is linear. For *any* additive effect vector whatsoever, the hybrid's genetic value is therefore the mean of the two parental line values – general combining ability, with no specific combining ability left over:

```
L <- simulate_lines(book_cfg)
bb <- rnorm(ctx$m) # any effects at all; the identity is linear
vL <- as.numeric((L / 2) %*% bb) # the parental line values they imply
max(abs(as.numeric(ctx$X %*% bb) - (vL[ctx$ped$a] + vL[ctx$ped$b]) / 2))
#> [1] 1.563194e-13
```

The residual is at machine precision because the statement is an identity. Two consequences run through the rest of the chapter. The mild one is that a “hybrid-specific” prediction has nothing to predict here: the dominance deviations of Chapter 10. are absent from `ctx$traits` by construction, so SCA is not merely unmodelled, it is not in the data. The sharp one is Section 14.3.2.

14.3 GBLUP: prediction from the relationship matrix

GBLUP treats the breeding value as a random effect $\mathbf{u}$ with covariance proportional to the genomic relationship matrix, $\text{Var}(\mathbf{u}) = \mathbf{G}\sigma_u^2$, and solves the mixed model equations.

14.3.1 Where the mixed model equations come from

The model is $\mathbf{y} = \mathbf{1}\mu + \mathbf{Z}\mathbf{u} + \mathbf{e}$ with $\mathbf{u} \sim (\mathbf{0}, \mathbf{G}\sigma_u^2)$ and $\mathbf{e} \sim (\mathbf{0}, \mathbf{I}\sigma_e^2)$ independent. Treating $\mathbf{u}$ as a random variable rather than a parameter to be estimated, C.R. Henderson [27] maximises the *joint* density of what is observed and what is not,

$$-2 \log p(\mathbf{y}, \mathbf{u}) \propto \frac{(\mathbf{y} - \mathbf{1}\mu - \mathbf{Z}\mathbf{u})'(\mathbf{y} - \mathbf{1}\mu - \mathbf{Z}\mathbf{u})}{\sigma_e^2} + \frac{\mathbf{u}'\mathbf{G}^{-1}\mathbf{u}}{\sigma_u^2}, \quad (14.2)$$

and the two stationarity conditions, multiplied through by σ_e^2 , are

$$\begin{bmatrix} \mathbf{1}'\mathbf{1} & \mathbf{1}'\mathbf{Z} \\ \mathbf{Z}'\mathbf{1} & \mathbf{Z}'\mathbf{Z} + \mathbf{G}^{-1}\lambda \end{bmatrix} \begin{bmatrix} \hat{\mu} \\ \hat{\mathbf{u}} \end{bmatrix} = \begin{bmatrix} \mathbf{1}'\mathbf{y} \\ \mathbf{Z}'\mathbf{y} \end{bmatrix}, \quad \lambda = \sigma_e^2/\sigma_u^2, \quad (14.3)$$

with $\mathbf{Z}$ the incidence matrix mapping training records to individuals. Three things fall out of that derivation that the equations alone do not say.

$\mathbf{u}$ is defined over *all* N individuals, training and candidates together – it enters through the prior, not through the data – so one solve produces predictions for both, and a candidate with no record of its own is predicted entirely by its $\mathbf{G}$ entries against individuals that do have one. Everything Section 14.4 finds is a consequence of that sentence.

λ is a variance *ratio*, not a tuning parameter: $\sigma_u^2 = 1$ because `traits` is standardised, and `e` was built with variance $(1 - h^2)/h^2$, so $\lambda = (1 - h^2)/h^2$ is not a coincidence but the same number entered twice. Nothing in this chapter estimates it; a real programme would, by REML.

And under normality the joint maximiser is the conditional mean, $\hat{\mathbf{u}} = \mathbb{E}[\mathbf{u} \mid \mathbf{y}]$, which is why BLUP minimises mean squared error rather than merely being a sensible shrinkage rule. Dropping normality it remains best among *linear* predictors, which is what the acronym claims.

14.3.2 Why the ridge blend is not optional

Section 1.2 proves $\mathbf{1}'\mathbf{G}\mathbf{1} = 0$ and reads it as the reason $\mathbf{G}$ cannot measure diversity. It has a second consequence, for prediction. $\mathbf{G}$ is positive semi-definite, and a quadratic form of a PSD matrix vanishes only on its null space, so $\mathbf{1}'\mathbf{G}\mathbf{1} = 0$ forces

$$\mathbf{G}\mathbf{1} = \mathbf{0}. \quad (14.4)$$

$\mathbf{G}$ is singular for *every* population, whatever N and m are. The blend is not insurance against an awkward case; without it `solve()` has nothing to invert.

Here it is much worse than one lost dimension. By the identity of Section 14.2 every hybrid row is an average of two line rows, so the whole $N \times m$ matrix lives in the span of the L line genotypes. Each hybrid uses exactly one line from each pool, which makes the pool-A indicator columns sum to the pool-B ones and costs a further dimension, and centering costs one more:

```
c(max_abs_G_times_one = max(abs(ctx$G %*% rep(1, N))))
#> max_abs_G_times_one
#> 1.16334e-13

rank_G <- function(n_per_pool) {
  cfg <- utils::modifyList(book_cfg, list(n_pool_A = n_per_pool, n_pool_B = n_per_pool))
  G <- simulate_data(cfg)$G
  c(lines = 2 * n_per_pool, hybrids = nrow(G),
    rank = sum(eigen(G, only.values = TRUE)$values > 1e-8))
}
```

lines	hybrids	rank
20	100	18
40	400	38
50	625	48

$\text{rank}(\mathbf{G}) = L - 2$, independent of both the number of hybrids and the number of markers. At the book's scale that is 48 dimensions carrying 625 hybrids: 577 of the eigenvalues are exactly zero, and $0.95 * \mathbf{G} + 0.05 * \mathbf{I}$ is what makes the equations solvable at all. Fifty lines were genotyped and crossed into 625 combinations, but there were only ever 48 independent dimensions to measure – and no amount of phenotyping the 625 adds a forty-ninth.

```
G <- ctx$G
gblup <- function(y, train, h2) {
  Gstar <- 0.95 * G + 0.05 * diag(N)
  n_t <- length(train)
  Z <- matrix(0, n_t, N); Z[cbind(seq_len(n_t), train)] <- 1
  X1 <- matrix(1, n_t, 1); lambda <- (1 - h2) / h2
  LHS <- rbind(cbind(crossprod(X1), crossprod(X1, Z)),
               cbind(crossprod(Z, X1), crossprod(Z) + solve(Gstar) * lambda))
  RHS <- rbind(crossprod(X1, y[train]), crossprod(Z, y[train]))
  C <- solve(LHS) # kept: it carries the PEVs
  list(u = (C %*% RHS)[-1],
```

```

    rel = 1 - diag(C)[-1] * lambda / diag(Gstar))
  }
  fit <- gblup(y, train, h2)
  u_hat <- fit$u
  cat("GBLUP accuracy (candidates):", cor(u_hat[cand], g[cand]), "\n")
#> GBLUP accuracy (candidates): 0.9596232

```

14.3.3 Reliability: the accuracy you can compute without the truth

Inverting the coefficient matrix rather than just solving against it costs one more decomposition and returns something a breeding programme can act on. The **u** block of that inverse holds the prediction error variances, $PEV_i = \sigma_e^2 [C^{uu}]_{ii}$, and the reliability of an individual's prediction is the fraction of its prior variance the data removed:

$$r_i^2 = 1 - \frac{PEV_i}{\sigma_u^2 G_{ii}^*}. \quad (14.5)$$

This is the same r as the accuracy reported everywhere in this chapter, with one difference that is the entire point: g does not appear in it. A programme can compute it for candidates it has never phenotyped and never will.

```

c(mean_reliability = mean(fit$rel[cand]),
  implied_accuracy = sqrt(mean(fit$rel[cand])),
  realised_accuracy = cor(u_hat[cand], g[cand]))
#> mean_reliability implied_accuracy realised_accuracy
#>      0.8016810      0.8953664      0.9596232

```

The model under-states itself by about 0.06. Some of that is the blend's doing: $0.95G + 0.05I$ tells the equations that every candidate carries 5% of variance no relative can explain, and each r_i^2 is computed against that inflated prior. Dialling the blend down says how much:

```

blend <- function(w) {
  Gstar <- w * G + (1 - w) * diag(N)
  n_t <- length(train)
  Z <- matrix(0, n_t, N); Z[cbind(seq_len(n_t), train)] <- 1
  X1 <- matrix(1, n_t, 1); lambda <- (1 - h2) / h2
  C <- solve(rbind(cbind(crossprod(X1), crossprod(X1, Z)),
    cbind(crossprod(Z, X1), crossprod(Z) + solve(Gstar) * lambda)))
  sqrt(mean((1 - diag(C)[-1] * lambda / diag(Gstar))[cand]))
}
c(implied_at_0.95 = blend(0.95), implied_at_0.999 = blend(0.999),
  realised = cor(u_hat[cand], g[cand]))
#> implied_at_0.95 implied_at_0.999 realised
#>      0.8953664      0.9272495      0.9596232

```

Roughly half the gap, then, and the other half survives however little identity is mixed in. A mean reliability and a correlation taken across a *set* of candidates are not the same quantity: the second is helped by the spread of true values among the candidates, which the first knows nothing about. Both halves point the same way – the reliability is conservative – and Section 14.4 shows that the gap is this small only because this particular split was easy.

14.4 What the random split was measuring

The split at the top of this chapter took a random 60% of the *hybrids*. There are 625 hybrids but only 50 lines, so a sample that size contains every line several times over: each of the 250 candidates has **both of its parents** sitting in the training set, in other hybrids with recorded phenotypes. Given Section 14.2 – a hybrid's value is the mean of its two parental values – that split asks the model to interpolate inside a factorial whose margins it has already seen. It is not a hard question.

J. Burgueño, G. de los Campos, K. Weigel, and J. Crossa [28] named the schemes that distinguish the cases, and F. Technow, T. A. Schrag, W. Schipprack, E. Bauer, H. Simianer, and A. E. Melchinger [18]

applied them to exactly this design, a factorial between two heterotic pools. Holding out *lines* rather than hybrids splits the candidates three ways: both parents seen (the random split, CV2), one parent unseen (CV1), neither parent seen (CV0). Removing 8 lines from each pool leaves 289 training hybrids and gives 272 CV1 and 64 CV0 candidates; the CV2 row below draws a random training set of the same 289, so the only thing that differs across the three rows is *what* was hidden, not how much.

```
cv_split <- function(seed, k = 8) {
  set.seed(seed)
  hoA <- sample(unique(ctx$ped$a), k)
  hoB <- sample(unique(ctx$ped$b), k)
  seen_a <- !(ctx$ped$a %in% hoA)
  seen_b <- !(ctx$ped$b %in% hoB)
  lin <- gblup(y, which(seen_a & seen_b), h2)      # lines held out
  rnd <- sample.int(N, sum(seen_a & seen_b))      # same size, random hybrids
  rd <- gblup(y, rnd, h2)
  i2 <- setdiff(seq_len(N), rnd)
  i1 <- which(xor(seen_a, seen_b))
  i0 <- which(!seen_a & !seen_b)
  c(CV2 = cor(rd$u[i2], g[i2]), r2 = sqrt(mean(rd$rel[i2])),
    CV1 = cor(lin$u[i1], g[i1]), r1 = sqrt(mean(lin$rel[i1])),
    CV0 = cor(lin$u[i0], g[i0]), r0 = sqrt(mean(lin$rel[i0])))
}
cv <- t(sapply(1:20, cv_split))
```

scheme	accuracy	sd	reliability
CV2 – random hybrids held out	0.932	0.013	0.873
CV1 – one parent unseen	0.653	0.073	0.644
CV0 – both parents unseen	0.122	0.217	0.128

The accuracy the chapter opened with survives only in the first row. Ask the model about a hybrid with one unseen parent and it drops by a third; ask about one with two unseen parents and what is left is 0.12, with a spread across replicates nearly twice that – small, and not reliably distinguishable from nothing. [D. Habier, R. L. Fernando, and J. C. M. Dekkers \[29\]](#) separate the two things a genomic prediction can be using – the relationship between candidates and training individuals, and the marker effects estimated in their own right – and this table is that separation, measured. The first row is almost entirely relatedness. The third row is what is left when relatedness is removed, and it is the honest measure of what the 2000 markers bought.

That the third row collapses is not a defect of the fit. It is $\text{rank}(\mathbf{G}) = L - 2$ again: the 289 training records carry at most 34 lines' worth of independent information, and asking them to pin down 2000 marker effects well enough to score a genotype no relative shares is asking for something that was never in the data. [H. D. Daetwyler, B. Villanueva, and J. A. Woolliams \[30\]](#) give the deterministic accuracy for exactly that regime, $r = \sqrt{nh^2/(nh^2 + M_e)}$, and with the units the simulation actually has – n the independent *lines*, and M_e the marker count, because these markers are drawn independently and carry no linkage to bundle them into segments – it agrees:

```
n_lines_train <- 2 * (book_cfg$n_pool_A - 8)
c(deterministic = sqrt(n_lines_train * h2 / (n_lines_train * h2 + book_cfg$m)),
  measured_CV0 = mean(cv[, "CV0"]))
#> deterministic measured_CV0
#> 0.09180609 0.12191453
```

The most useful column is the last one in the table. The reliabilities computed from the mixed model equations – with no access to g – track the realised accuracies down the collapse, 0.64 against 0.65 for CV1 and 0.13 against 0.12 for CV0. The model knows. A programme that reports reliability alongside every GEBV does not need to discover this the way this chapter did; the warning is in the same solve that produced the prediction.

⚠ Report which parents the training set had seen

An accuracy figure without its cross-validation scheme is not interpretable. The same fit, the same data and the same trait give 0.93, 0.65 or 0.12 here depending only on which hybrids were withheld. In a hybrid programme this is not a technicality: CV2 is the accuracy for filling in a partly-observed factorial, CV1 for a cross onto a new line from one pool, CV0 for a cross between two new lines. They are different questions, and only the first one is easy.

14.5 BayesA and BayesB: prediction from marker effects

GBLUP assumes every marker contributes an equal, infinitesimal share of the genetic variance. The Bayesian alphabet [14], [31] relaxes that one assumption at a time. **BayesA** gives every marker its own variance, drawn from a scaled-inverse- χ^2 prior – effects are shrunk individually rather than uniformly, which matters when a few markers carry much more signal than the rest. **BayesB** adds a point mass at zero: each marker has prior probability π of being excluded entirely, so the model performs variable selection as it fits. Both reduce to GBLUP’s ridge-regression solution as a special case; they differ only in how much of the genetic architecture they let the data shape.

14.5.1 The priors, explicitly

All three models write $\mathbf{g} = \mathbf{Z}\boldsymbol{\beta}$ and differ only in what they assume about β_k . Ridge regression, and therefore GBLUP, uses one variance for every marker; BayesA gives each marker its own, drawn from a scaled inverse chi-square; BayesB puts a spike at zero in front of it:

$$\text{ridge, and so GBLUP: } \beta_k \sim N(0, \sigma_\beta^2) \quad \text{BayesA: } \beta_k \sim N(0, \sigma_k^2), \sigma_k^2 \sim \chi^{-2}(\nu, S) \quad (14.6)$$

$$\text{BayesB: } \beta_k \sim \begin{cases} 0 & \text{with probability } \pi \\ N(0, \sigma_k^2) & \text{otherwise.} \end{cases} \quad (14.7)$$

The scaled inverse chi-square is the conjugate prior for a normal variance, which is the whole reason it appears: it keeps the Gibbs sampler in closed form. Its two hyperparameters are the arguments the fits below pass. ν is `df = 5`, the prior weight in pseudo-observations – small, so the prior stays vague, but not zero, because the posterior is improper at $\nu = 0$. S is set from `R2 = 0.5`, the share of phenotypic variance the prior expects the markers to explain, spread over the markers by their observed variation. Neither is fitted; both are choices, and `R2` in particular is a statement about heritability made before seeing the data.

This book fits both with `bWGR::BayesA()/bWGR::BayesB()` [32] rather than by hand – unlike GBLUP, a Gibbs sampler over marker effects is not the few lines the mixed-model solve above was, and `bWGR` is exactly the tool `vanraden_G()`’s docstring pointed toward without naming:

```
library(bWGR)
M <- ctx$X * 2 # dosage 0/1/2, what bWGR expects
fitA <- BayesA(y[train], M[train, ], it = 1500, bi = 500, df = 5, R2 = 0.5)
fitB <- BayesB(y[train], M[train, ], it = 1500, bi = 500, pi = 0.95, df = 5, R2 = 0.5)
predA <- as.numeric(M[cand, ] %*% fitA$b)
predB <- as.numeric(M[cand, ] %*% fitB$b)
c(BayesA = cor(predA, g[cand]), BayesB = cor(predB, g[cand]))
#> BayesA BayesB
#> 0.9606263 0.9599031
```

Both fits produce a marker-effect vector `$b`; a candidate’s GEBV is just its genotypes dotted with those effects. GBLUP produced the same quantity from the relationship matrix instead of the markers directly – two routes to the same thing, checked next.

14.5.2 What the sampler actually selects

Both fits also return the quantities the priors above are about – BayesA’s per-marker variances in `$vb`, BayesB’s posterior inclusion probabilities in `$d` – and on the polygenic trait there is nothing in them to

look at. Rebuild the oligogenic architecture the cached sweep below uses, three markers of 2000 given effects and everything else exactly zero, and there is:

```
set.seed(book_cfg$seed + 42)
qtl <- sample.int(book_cfg$m, 3)
Bq <- rep(0, book_cfg$m); Bq[qtl] <- rnorm(3)
g_ol <- as.numeric(scale(vr$Z %*% Bq))
set.seed(31)
y_ol <- g_ol + rnorm(N, 0, sqrt((1 - h2) / h2))
round(apply(vr$Z[, qtl], 2, var) * Bq[qtl]^2 /
      sum(apply(vr$Z[, qtl], 2, var) * Bq[qtl]^2), 3) # share of variance each
#> [1] 0.962 0.038 0.001
```

Worth knowing before reading anything into the sweep: the three effects were drawn from `rnorm(3)`, and this draw put 96% of the genetic variance on one of them. “Three causal markers” describes the construction; the trait is effectively monogenic, and the true π is nearer $1 - 1/2000$ than the $1 - 3/2000$ the caption below quotes.

```
oA <- BayesA(y_ol[train], M[train, ], it = 1500, bi = 500, df = 5, R2 = 0.5)
oB <- BayesB(y_ol[train], M[train, ], it = 1500, bi = 500, pi = 0.95, df = 5, R2 = 0.5)
big <- qtl[which.max(abs(Bq[qtl]))]
c(BayesB_d_at_QTL = oB$d[big], BayesB_d_median = median(oB$d),
  BayesA_vb_at_QTL = oA$vb[big], BayesA_vb_median = median(oA$vb))
#> BayesB_d_at_QTL BayesB_d_median BayesA_vb_at_QTL BayesA_vb_median
#> 0.999000013 0.179000005 0.011139854 0.002957477
```

```
par(mfrow = c(1, 2), mar = c(4, 4, 2, 1))
for (v in list(list(oB$d, "BayesB: posterior inclusion"),
               list(oA$vb, "BayesA: marker variance"))) {
  plot(v[[1]], pch = 16, cex = 0.3, col = "grey60", xlab = "marker", ylab = "",
       main = v[[2]], cex.main = 0.95)
  points(big, v[[1]][big], pch = 1, cex = 2, lwd = 2)
}
```




Figure 14.1: The two priors on the same oligogenic trait, with the marker carrying 96% of the genetic variance circled. Both find it. Only BayesB can then remove the other 1999 from the model.

14.5.3 What BayesB does that BayesA cannot

Both panels find the marker, which is worth saying plainly because the usual telling of this comparison implies BayesA does not: its variance for the QTL is about four times the median marker's and sits well clear of the background band. What separates the two priors is not detection but what each can do afterwards.

BayesA estimates a private variance σ_k^2 for every marker, and the only data bearing on σ_k^2 is the single effect β_k drawn from it. One observation never overwhelms a prior, so each σ_k^2 stays near the $\chi^{-2}(\nu, S)$ it started from however many individuals are phenotyped – D. Gianola, G. de los Campos, W. G. Hill, E. Manfredi, and R. Fernando [33] make this the central objection to describing the Bayesian alphabet as “learning the genetic architecture”. The cost is not the missed QTL. It is that the other 1999 markers keep a variance they have not earned, and go on absorbing signal that belongs to the one marker that has. BayesB's spike deletes them, and it can afford to because it learns a single shared π from all 2000 markers at once rather than 2000 private variances from one observation each.

So the advantage is real but bounded, and shows up in accuracy rather than in the figure: on this architecture at $h^2 = 0.5$ the table below gives BayesB 0.952 against BayesA's 0.931 and GBLUP's 0.930. BayesA sits with GBLUP, not with BayesB, which is the practical summary of the paragraph above.

14.6 Verification: two routes agree

GBLUP with a flat prior on every marker's variance and the un-selected member of the Bayesian alphabet, Bayesian ridge regression (`iv = FALSE`, `pi = 0` in `bWGR::wgr()`), are the same model solved two different ways – one through the $N \times N$ relationship matrix, the other through the $N \times m$ marker matrix.

They are the same model in a stronger sense than “they should agree”. Substitute $\mathbf{u} = \mathbf{Z}\beta$ and $\mathbf{G} = \mathbf{Z}\mathbf{Z}'/2\sum_k p_k(1-p_k)$ into the equations above and the two solutions are related by an identity: with t the training rows,

$$\hat{\mathbf{u}} = \mathbf{G}_t(\mathbf{G}_{tt} + \mathbf{I}\lambda)^{-1}(\mathbf{y}_t - \hat{\mu}) = \mathbf{Z}\hat{\beta}, \quad \hat{\beta} = \mathbf{Z}_t'(\mathbf{Z}_t\mathbf{Z}_t' + \mathbf{I}\lambda_\beta)^{-1}(\mathbf{y}_t - \hat{\mu}), \quad (14.8)$$

with the two penalties differing by exactly the divisor that defines $\mathbf{G}$: $\lambda_\beta = \lambda \cdot 2\sum_k p_k(1-p_k)$. Getting that scalar wrong is the usual way this equivalence is reported as “approximate”. It is not approximate:

```
lambda <- (1 - h2) / h2
yc <- y[train] - mean(y[train])
u_G <- G[, train] %*% solve(G[train, train] + diag(lambda, length(train)), yc)
u_M <- vr$Z %*% crossprod(vr$Z[train, ],
  solve(tcrossprod(vr$Z[train, ]) + diag(lambda * s2, length(train)), yc))
max(abs(u_G - u_M))
#> [1] 1.354472e-14
```

Machine precision, over all 625 individuals. Note this route uses the raw $\mathbf{G}$: it needs no inverse of it, so Section 14.3.2's blend never enters, which is also why it is the cleaner statement of the identity.

The Gibbs sampler should land in the same place, and does, though only to sampling error – it is estimating λ from the data rather than being told it, and the blend used in `gblup()` has no exact marker-space equivalent:

```
yobs <- rep(NA_real_, N); yobs[train] <- y[train]
fit_brr <- wgr(yobs, M, iv = FALSE, pi = 0, it = 1500, bi = 500)
c(gblup_vs_gibbs = cor(u_hat, fit_brr$hat - mean(fit_brr$hat)))
#> gblup_vs_gibbs
#> 0.9992759
```

That correlation, 0.999, is the weaker of the two checks and is included because it exercises the code path the Bayesian fits above actually use. The identity is what proves the two routes are one model.

14.7 Which model wins, and when

`ctx$traits` is built from thousands of markers with independent, similarly sized effects – the textbook case GBLUP's infinitesimal-model assumption was designed for. A trait controlled by a handful of large-effect loci is a different regime, and it is the one BayesB's point-mass prior exists for. [book/scripts/](#)

`regenerate.R` `genpred` reruns the fits above 15 times each, at two heritabilities, under both `ctx$traits`'s polygenic architecture and a book-local variant with genetic variance concentrated on three markers out of 2000:

arch	h2	GBLUP	BayesA	BayesB
oligogenic	0.2	0.789	0.782	0.788
polygenic	0.2	0.799	0.795	0.798
oligogenic	0.5	0.930	0.931	0.952
polygenic	0.5	0.935	0.935	0.934

```
sub <- reshape(acc, direction = "long", varying = c("GBLUP", "BayesA", "BayesB"),
               v.names = "accuracy", timevar = "method",
               times = c("GBLUP", "BayesA", "BayesB"))
sub$grp <- interaction(sub$arch, sub$h2, sep = "\nh2=")
boxplot(accuracy ~ method + grp, sub, las = 2, cex.axis = 0.7,
        col = rep(c("grey85", "grey65", "grey45"), 4),
        ylab = "accuracy (candidates)")
```




Figure 14.2: Prediction accuracy by architecture and heritability, 15 replicate splits each.

At $h^2 = 0.5$ under the oligogenic architecture, BayesB's mean accuracy is 0.022 higher than GBLUP's – the gap the model was built to close. Under the polygenic architecture the three methods are indistinguishable within replicate noise: with genetic variance spread over thousands of markers, BayesB's variable selection has nothing to select, and its point-mass prior buys nothing GBLUP was not already doing.

That result is not evidence that the alphabet rarely helps. It is Section 14.2 restated: on the polygenic arm the trait was *generated* from one shared marker variance, so GBLUP's prior is the true one and no competitor can do better than match it. What the comparison establishes is the weaker and more useful claim – that BayesA and BayesB cost nothing when they are wrong, and recover a real gap when they are right.

14.7.1 π was fixed, not tuned

A real analyst does not know the true sparsity in advance. `genpred`'s second sweep holds the oligogenic architecture and $h^2 = 0.5$ fixed and only moves π :

π	BayesB
0.500	0.948
0.900	0.955
0.950	0.958
0.990	0.964
0.995	0.967

Accuracy improves monotonically as π approaches the truth, but the degradation from being off is gradual, not catastrophic – π is worth cross-validating, not agonising over.

The obvious alternative is to stop choosing it. BayesC π puts a prior on π itself and samples it alongside the effects, which ought to make the sweep above unnecessary. Fit it to both architectures:

```

cpi_poly <- BayesCpi(y[train], M[train, ], it = 1500, bi = 500, df = 5, R2 = 0.5)
cpi_olig <- BayesCpi(y_ol[train], M[train, ], it = 1500, bi = 500, df = 5, R2 = 0.5)
c(pi_hat_polygenic = cpi_poly$pi, # truth: 0
  pi_hat_oligogenic = cpi_olig$pi, # truth: ~0.9995
  acc_oligogenic = cor(as.numeric(M[cand, ] %*% cpi_olig$b), g_ol[cand]),
  acc_BayesB_at_095 = cor(as.numeric(M[cand, ] %*% oB$b), g_ol[cand]))
#> pi_hat_polygenic pi_hat_oligogenic acc_oligogenic acc_BayesB_at_095
#> 0.4369196 0.4390892 0.9299389 0.9524220

```

It returns essentially the same π for a trait with 2000 causal markers as for one that is effectively monogenic, and predicts the oligogenic trait slightly worse than BayesB told the right answer. With 375 records and 2000 markers there is little in the data that speaks to π specifically – Section 14.3.2's 48 dimensions again – so what comes back is close to the prior. Estimating π is not a substitute for choosing it here; the sweep above remains the honest route.

14.8 From GEBVs to `stage1_build()`

Everything above ends where Chapter 4's contract begins. The candidates' predicted values, whichever model produced them, are exactly the `traits` argument `stage1_build()` expects – one column per trait, one row per hybrid:

```

hybrids_traits <- data.frame(hybrid = cand, trait1 = predB)
head(hybrids_traits)

```

hybrid	trait1
4	0.1398207
5	0.3260893
6	-0.1489823
11	1.6891953
13	1.3072349
16	0.2173326

Repeat the fit once per trait in the multi-trait index and this data frame is what feeds `stage1_build(X, ped, traits, fL)`. The second consumer is `ocs_quadprog(f, gebv, lambda, ...)` of Chapter 15., whose argument is literally named `gebv` and has until now been handed the truth.

Carry the reliabilities with them. Section 14.4's table is the argument for storing `fit$rel` next to every predicted value rather than discarding it: a candidate whose GEBV is unreliable is not a bad candidate, it is an unmeasured one, and the difference matters to a selection decision.

14.9 Back to Finding 5: a measured accuracy, not an assumed one

Chapter 12’s Finding 5 assumed a prediction accuracy r and manufactured error correlated with relatedness to match it. The GBLUP fit above at $h^2 = 0.5$ gives a *measured* accuracy of 0.96 on the polygenic architecture, at the top of the range Finding 5 explored ($r \in \{0.5, 0.7, 0.9, 1\}$) – but Section 14.4 is the reason not to read that as the typical case. It is the CV2 number, the accuracy for a candidate whose two parents are already phenotyped in other combinations. The CV1 row of that table, 0.65, sits near the *bottom* of Finding 5’s range, and the CV0 row falls well below anything it explored. Between them the two chapters bracket the same interval from opposite ends, and Finding 5’s pessimistic cases are the realistic ones.

The errors are, in fact, structured by relatedness rather than white:

```
err <- g[cand] - u_hat[cand]
fcand <- ctx$f[cand, cand]
ut <- upper.tri(fcand)
cor(outer(err, err)[ut], fcand[ut])
#> [1] 0.1397751
```

Pairs of candidates that are more related have more correlated prediction errors – the mechanism Finding 5 assumed, confirmed here from an actual fitted model rather than injected. It is a modest correlation, not the extreme case Finding 5 stress-tested, because $h^2 = 0.5$ with $r = \text{length}(\text{train})$ training records is a comparatively favourable regime; a smaller or less related training set would push it higher.

i What this chapter still does not model

Closing the “traits are given” gap opens several more, named here so they are not silently dropped:

- **Specific combining ability.** Section 14.2 showed there is none in `ctx$traits` to find: the trait is additive by construction, so a hybrid’s value is exactly the mean of its parents’. SCA is not merely unmodelled here, it is absent from the data. A real hybrid programme has it – Chapter 10’s dominance deviations are where it comes from – and would fit GCA and SCA jointly rather than one additive effect per marker.
- **Estimated variance components.** h^2 is supplied to `gblup()`, never estimated. A real fit gets λ from REML on the training data, and the reliabilities of Section 14.4 inherit that estimate’s uncertainty.
- **Multivariate prediction.** `ctx$traits` has three correlated columns; this chapter predicted one. A multi-trait Bayesian model borrows information across traits and should out-predict fitting each one separately, especially for the traits with less training data.
- **Forward validation.** Section 14.4 splits by line, which is the right axis for a factorial, but every scheme in it still samples training and candidates from one generation. A programme selecting for the future validates forward in time, on a later cycle than it trained on.
- **Retraining across cycles.** Chapter 19. runs the *selection* side of a multi-cycle programme with a fixed, perfect index. A model trained once and never updated goes stale as the population it predicts drifts from the population it was trained on.
- **Genotyping cost and marker density.** Accuracy was measured here at a fixed 2000 markers; a real programme trades marker density against genotyping cost per individual, not just against training-set size.

IV

Optimisation

15	Optimal contributions and the α constraint	155
15.1	Constraint, not a weighted sum	155
15.2	Relation to the standard OCS formulation	156
15.3	Setting $\alpha_{\max}$	158
15.4	A caveat carried over from the OCS literature	161
15.5	Quadratic programming, when contributions are continuous	162
15.6	What belongs in the inner loop	162
16	Combinatorial selection	165
16.1	Four baselines before any optimiser	165
16.2	Greedy, with incremental updates	166
16.3	The encoding problem	168
16.4	Differential evolution	169
16.5	The full frontier	172
16.6	Which strategy to use	173
17	Advanced selection: structure, bounds and scale	175
17.1	Three facts about this problem	175
17.2	What lazy greedy would not buy	176
17.3	The move that was already paid for	176
17.4	The bound was not a bound	179
17.5	Rounding the relaxation into a plan	181
17.6	Exact, and how far the heuristics really are	181
17.7	Scaling to 200 of 10,000	182
17.8	Multi-objective, and why not NSGA-II	184
17.9	What not to reach for	184
17.10	Which method to use	184

15. Optimal contributions and the α constraint

Everything so far has been measurement. This chapter turns the measurement into a constraint an optimiser can respect, and connects it to the optimal-contribution literature it specialises. `ctx$index` below is still the known-marker-effects version built in Chapter 4.; Chapter 14. shows how a real programme gets a `traits` matrix to hand `stage1_build()` in the first place, and what accuracy it can expect once it does.

15.1 Constraint, not a weighted sum

With $\mathbf{u}$ the multi-trait index and S the selected set of size n :

$$\max_{S, |S|=n} \frac{1}{n} \sum_{i \in S} u_i \quad \text{subject to} \quad \alpha(S) \leq \alpha_{\max}. \quad (15.1)$$

The common alternative, $\max \bar{u} - \lambda \theta_S$, requires calibrating λ , which has no direct interpretation. The constraint form preserves an operational meaning: “*I accept losing at most $\alpha_{\max}$ of diversity.*” That is a sentence a breeding committee can approve or reject. “Diversity is worth $\lambda = 3.7$ index units” is not.

The two are not equivalent in practice either. A weighted sum lets the optimiser buy index by spending diversity anywhere on the curve; a constraint fixes the budget and makes the optimiser find the best plan *inside* it. Since the acceptable loss is a policy decision and the exchange rate is not, the constraint form puts the decision in the right hands.

```
make_fitness
#> function (ctx, n_sel, alpha_max, gd_ref, penalty = 1000, max_use = NULL)
#> {
#>   N <- ctx$N
#>   f <- ctx[["f"]]
#>   ix <- ctx$index
#>   L <- ctx$n_lines
#>   pa <- as.integer(ctx$ped$a)
#>   pb <- as.integer(ctx$ped$b)
#>   function(par) {
#>     idx <- decode(par, N, n_sel)
#>     gd <- 1 - mean(f[idx, idx])
#>     viol <- max(0, (gd_ref - gd)/gd_ref - alpha_max)
#>     if (!is.null(max_use)) {
#>       cnt <- tabulate(c(pa[idx], pb[idx]), nbins = L)
#>       viol <- viol + sum(pmax(0, cnt - max_use))/(2 * n_sel)
#>     }
#>     -(mean(ix[idx]) - penalty * viol)
#>   }
#> }
#> <bytecode: 0xc67013dc8>
#> <environment: namespace:hybdiv>
```

The violation enters as a *proportional* penalty, so a plan that overshoots by 10% of the budget is penalised in units the budget defines, whatever the absolute value of $\alpha_{\max}$.

15.1.1 This argument is older than it looks

The paragraphs above read as a claim about human decision-making: budgets are approvable, exchange rates are not. That is true, and it undersells the case. Kempthorne & Nordskog (1959) posed the identical problem for selection indices – maximise aggregate gain subject to **zero expected change** in a specified

direction – and solved it in closed form. Chapter 3. derives their solution; `index_restricted()` implements it. The distinguishing property is not softness:

```
P <- matrix(c(4.0, 1.2, 1.2, 1.0), 2)
G <- matrix(c(2.0, 0.9, 0.9, 0.3), 2)
a <- c(1, 1); R <- cbind(c(0, 1))
c(unrestricted = drop(crossprod(R, G %*% index_smith_hazel(P, G, a))),
  restricted    = drop(crossprod(R, G %*% index_restricted(P, G, a, R))))
#> unrestricted    restricted
#> 6.679688e-01 2.096025e-17
```

Exactly zero, not approximately. A penalty trades the restricted direction away at some price; a restriction refuses to trade it at any price. That is the distinction this section is drawing, and it has stood since 1959.

i What is actually new here

Not restriction over pricing. What is new is *what* is being restricted. Kempthorne and Nordskog restricted a trait – something with a genetic variance, an economic weight and a row in **G**. The α constraint restricts the loss of molecular gene diversity, which has none of those: it is a property of a matrix, not of a phenotype, and it is bounded relative to a resampling baseline rather than held at zero.

The algebra of restriction carries over. The object being restricted does not, which is why the rest of this chapter cannot simply cite 1959 and stop.

15.2 Relation to the standard OCS formulation

F. Gómez-Romano, B. Villanueva, J. Fernández, J. A. Woolliams, and R. Pong-Wong [4] write the general optimal-contribution problem, for a vector of continuous contributions **c**, as

$$\text{minimise } \mathbf{c}'\mathbf{G}\mathbf{c} \quad \text{subject to} \quad \mathbf{c}'\mathbf{s} = 0.5, \mathbf{c}'\mathbf{d} = 0.5, c_i \geq 0, \quad (15.2)$$

with **s** and **d** indicating sex and **G** the coancestry matrix chosen for managing diversity – here, **f**.

The formulation of this book is that problem **restricted to $c_i \in \{0, 1/n\}$** : a fixed number n of candidates contribute equally, none partially. The minimisation of $\mathbf{c}'\mathbf{f}\mathbf{c}$ is turned into a constraint and $\bar{u}$ is promoted to the objective.

That restriction to equal, all-or-nothing contributions is exactly what makes the problem combinatorial, and what pushes the solution method from convex quadratic programming to a metaheuristic. It is also not arbitrary: a commercial hybrid programme plants a whole plot of a hybrid or does not plant it.

i The sex constraint, in plants

The per-sex sum constraint disappears in hermaphrodite and monoecious species (a single sum to 1) and reappears in dioecious crops and in schemes using cytoplasmic male sterility. In hybrid maize the natural problem is not OCS but **optimal cross selection**: **c** stops being an individual contribution and becomes a weight per cross, with the between-pool covariance entering explicitly. And there is an extra restriction the animal literature does not have – the **maximum number of crosses executable in a field**, which makes the problem mixed-integer.

15.2.1 The continuous relaxation as an upper bound

Dropping $c_i \in \{0, 1/n\}$ back to $0 \leq c_i \leq 1/n$ gives the continuous relaxation of *this* restricted problem. Every discrete solution is feasible under the relaxation, so the relaxation's optimum is a valid **upper bound**.

```
project_capped_simplex
#> function (v, cap)
#> {
#>   lo <- min(v) - 1
```

```
#> hi <- max(v)
#> for (i in 1:60) {
#>   tau <- (lo + hi)/2
#>   if (sum(pmin(pmax(v - tau, 0), cap)) > 1)
#>     lo <- tau
#>   else hi <- tau
#> }
#> pmin(pmax(v - (lo + hi)/2, 0), cap)
#> }
#> <bytecode: 0xc693e5a10>
#> <environment: namespace:hybdiv>
```

The cap is what makes the bound valid. Without $c_i \leq 1/n$ the relaxation would put all weight on the single best candidate and the bound would be useless.

Turning the relaxation into a bound on the *constrained* problem needs one more step, and turning it into an actual selection needs a rounding: both are in Chapter 17..

```
lambdas <- 10^seq(-1, 2.5, length.out = 12)
relax <- ocs_relaxation(ctx, n_sel, lambdas)
fmt(as.data.frame(relax), 4)
```

lambda	index	theta	GD
0.1000	1.9919	0.6876	0.3124
0.2081	1.9919	0.6876	0.3124
0.4329	1.9919	0.6876	0.3124
0.9006	1.9919	0.6876	0.3124
1.8738	1.9911	0.6870	0.3130
3.8986	1.9900	0.6866	0.3134
8.1113	1.9885	0.6863	0.3137
16.8761	1.9831	0.6859	0.3141
35.1119	1.9270	0.6838	0.3162
73.0527	1.6173	0.6784	0.3216
151.9911	1.0023	0.6725	0.3275
316.2278	0.5729	0.6704	0.3296

Sweeping λ traces the merit-versus-diversity frontier. Each row is the best *continuous* plan at that exchange rate, and therefore an upper bound on what any discrete plan of the same diversity can achieve.

```
null <- null_distribution(ctx, n_sel, B = 400, B_full = 60)
a_relax <- (null$gd_ref - relax[, "GD"]) / null$gd_ref

sels <- list(random = sel_random(ctx, n_sel),
             trunc_cap = sel_truncation_cap(ctx, n_sel),
             greedy_w1 = sel_greedy(ctx, n_sel, w = 1),
             greedy_w3 = sel_greedy(ctx, n_sel, w = 3),
             truncation = sel_truncation(ctx, n_sel))
pts <- t(apply(sels, function(s)
  c(alpha = alpha_loss(s, ctx, null$gd_ref), index = mean_index(s, ctx))))

plot(a_relax * 100, relax[, "index"], type = "b", pch = 1, lty = 2, col = "grey45",
     xlab = "gene diversity loss, alpha (%)", ylab = "mean index (s.d.)",
     xlim = range(c(a_relax, pts[, "alpha"]) * 100),
     ylim = range(c(relax[, "index"], pts[, "index"])))
```

```
points(pts[, "alpha"] * 100, pts[, "index"], pch = 19, col = "#b2182b")
text(pts[, "alpha"] * 100, pts[, "index", rownames(pts)], pos = 2, cex = 0.7)
abline(v = 0, lty = 3)
legend("bottomright", c("continuous OCS relaxation (upper bound)", "discrete
strategies"),
      col = c("grey45", "#b2182b"), pch = c(1, 19), lty = c(2, NA), bty = "n", cex
= 0.8)
grid()
```




Figure 15.1: The continuous relaxation traces an upper bound. Discrete strategies must sit below and to the left of it; how far below is the convergence diagnostic.

⚠ Too few iterations gives an “upper bound” below the discrete optimum

The relaxation is solved by projected gradient. Under-converged, it returns a value *smaller* than the discrete optimum – which is useless, and worse, silently so. The default of 4000 iterations is the minimum that converged in testing; re-check it whenever the dataset changes.

```
short <- ocs_relaxation(ctx, n_sel, lambdas = 10, iters = 50)
long <- ocs_relaxation(ctx, n_sel, lambdas = 10, iters = 4000)
c(index_50_iters = short[, "index"], index_4000_iters = long[, "index"],
  greedy_w3 = mean_index(sels$greedy_w3, ctx))
#>   index_50_iters.index index_4000_iters.index   greedy_w3
#>           1.605937           1.987272           1.991937
```

15.3 Setting $\alpha_{\max}$

The limit **should not be chosen a priori**. Pure truncation – the greediest possible selection – defines the maximum attainable α . If that ceiling sits below your intended limit, the constraint restricts nothing.

```

attainable <- alpha_loss(sel_truncation(ctx, n_sel), ctx, null$gd_ref)
c(attainable_ceiling = attainable,
  a_5pct_limit_would_be_active = attainable > 0.05,
  a_1pct_limit_would_be_active = attainable > 0.01)
#>      attainable_ceiling a_5pct_limit_would_be_active
#>      0.04121618      0.00000000
#> a_1pct_limit_would_be_active
#>      1.00000000

```

`stage2_select()` computes this ceiling automatically and warns about any requested α above it, rather than silently returning an unconstrained solution that looks constrained.

The correct procedure is to trace the frontier and choose its knee:

```

ws <- c(0, 0.25, 0.5, 0.75, 1, 1.5, 2, 3, 5, 8, 12)
front <- do.call(rbind, lapply(ws, function(w) {
  s <- sel_greedy(ctx, n_sel, w = w)
  data.frame(w = w, alpha = alpha_loss(s, ctx, null$gd_ref),
    index = mean_index(s, ctx), Ne_par = ne_parents(s, ctx))
}))
plot(front$alpha * 100, front$index, type = "b", pch = 19, col = "#2166ac",
  xlab = "gene diversity loss, alpha (%)", ylab = "mean index (s.d.)")
abline(v = attainable * 100, lty = 2, col = "#b2182b")
text(attainable * 100, min(front$index), "truncation ceiling", pos = 2, cex = 0.7,
  col = "#b2182b")
grid()

```




Figure 15.2: The greedy frontier, swept over the merit weight w . The knee is where further diversity costs index steeply.

```
fmt(front, 4)
```

w	alpha	index	Ne_par
0.00	-0.0135	0.1343	37.3057

w	alpha	index	Ne_par
0.25	0.0412	1.9919	13.0199
0.50	0.0412	1.9919	13.0199
0.75	0.0412	1.9919	13.0199
1.00	0.0412	1.9919	13.0199
1.50	0.0412	1.9919	13.0199
2.00	0.0412	1.9919	13.0199
3.00	0.0412	1.9919	13.0199
5.00	0.0412	1.9919	13.0199
8.00	0.0412	1.9919	13.0199
12.00	0.0412	1.9919	13.0199

15.3.1 How far $\alpha_{\max}$ moves when f moves

The ceiling above is quoted to four decimals, which invites more confidence than the marker panel supports. Chapter 9. showed that the MAF spectrum decides how badly the two inbreeding lenses disagree; the same spectrum feeds f , and therefore $\bar{H}_{\text{ref}}$ and the ceiling. It is worth knowing which choices move the budget and which only look as though they should.

Rebuild everything downstream of the marker panel – f , the null, the ceiling – under a filter, a density and an ascertainment scenario:

```
ceiling_for <- function(keep, label, B = 250) {
  Xs <- ctx$X[, keep, drop = FALSE]
  ck <- build_ctx(X = Xs, f = molecular_coancestry(Xs), G = NULL,
                 traits = ctx$traits, ped = ctx$ped, n_lines = ctx$n_lines,
                 pool = ctx$pool, fL = NULL, cfg = ctx$cfg)
  nl <- null_distribution(ck, n_sel, B = B, B_full = 20)
  data.frame(panel = label, markers = ncol(Xs), gd_ref = nl$gd_ref,
             alpha_max = alpha_loss(sel_truncation(ck, n_sel), ck, nl$gd_ref))
}
maf <- pmin(ctx$p0, 1 - ctx$p0)
polyA <- colMeans(simulate_lines(book_cfg)[seq_len(book_cfg$n_pool_A), ] / 2)
sens <- rbind(
  do.call(rbind, lapply(c(0, 0.01, 0.05, 0.10),
    function(t) ceiling_for(maf >= t, sprintf("MAF >= %.2f", t)))),
  ceiling_for(polyA > 0 & polyA < 1, "ascertained in pool A"))
fmt(sens, 4)
```

panel	markers	gd_ref	alpha_max
MAF >= 0.00	2000	0.3258	0.0411
MAF >= 0.01	1934	0.3369	0.0411
MAF >= 0.05	1780	0.3611	0.0404
MAF >= 0.10	1585	0.3881	0.0402
ascertained in pool A	1762	0.3535	0.0421

$\bar{H}_{\text{ref}}$ moves by 19% across these panels. $\alpha_{\max}$ moves by 4.7%. That is not a coincidence: α is a *ratio* of diversities measured on the same panel, so a filter that shifts the absolute level shifts numerator and denominator together and largely cancels out. **Filtering choices move the diversity level; they do not move the budget.**

Density does not move it either, but it does something worse: it makes the estimate noisy without making it wrong. Thinning the panel and averaging over six independent draws of which markers survive:

```
# Every subset is drawn BEFORE the loop. `null_distribution()` calls set.seed()
# internally, so a `sample.int()` written inside the loop would hand every
# replicate after the first the very same markers -- see the warning below.
set.seed(4)
subs <- lapply(c(250, 500, 1000),
  function(mm) lapply(1:6, function(i) sort(sample.int(ctx$m, mm))))
dens <- do.call(rbind, lapply(subs, function(sl) {
  v <- vapply(sl, function(k) ceiling_for(k, "", B = 200)$alpha_max, numeric(1))
  data.frame(markers = length(sl[[1]]), alpha_max = mean(v), sd = stats::sd(v))
}))
fmt(rbind(dens, data.frame(markers = ctx$m, alpha_max = attainable, sd = NA)), 4)
```

markers	alpha_max	sd
250	0.0445	0.0090
500	0.0461	0.0084
1000	0.0417	0.0067
2000	0.0412	NA

The means sit within one standard deviation of the full-panel ceiling all the way down to 250 markers. The standard deviation does not: at 250 it is about 20% of the estimate itself. A budget set from a thin panel is not biased, it is a coin flip – and a single thin-panel study has no way to tell which side it landed on.

So all three axes leave $\alpha_{\max}$ where it was. That is the useful finding, and it is worth stating positively: **the budget is a more portable quantity than the diversity level it is computed from.** $\bar{H}_{\text{ref}}$ is not comparable across panels; α largely is, provided the panel is dense enough for the estimate to be stable.

 `null_distribution()` reseeds, and will freeze a resampling loop

It calls `set.seed(seed)` on entry, with `seed = 7` by default, so the random stream is reset every time it is invoked. Any `sample()` written *after* it inside a loop therefore draws the same values on every pass but the first, and a replication study silently becomes one replicate reported with a fabricated standard deviation.

This is not hypothetical: the table above returned two different wrong answers – once suggesting thin panels understate the ceiling, once that they overstate it – before the subsets were hoisted out of the loop. Draw randomisation up front, or pass an explicit varying `seed`.

 Quote $\alpha_{\max}$ with the panel that produced it

A budget is not a property of the germplasm. It is a property of the germplasm *and* the marker panel, and only the second of those changes when the genotyping platform does. Re-derive the ceiling whenever the panel changes, and do not compare an α across two studies genotyped differently – Chapter 12. Finding 3's caveat on transfer applies here with more force, because this one is a constraint rather than a diagnostic.

15.4 A caveat carried over from the OCS literature

F. Gómez-Romano, B. Villanueva, J. Fernández, J. A. Woolliams, and R. Pong-Wong [4] show that the *predicted* rate of coancestry in the next generation, $E(\Delta f_{t+1}) = (c'G_t c - f_t)/(1 - f_t)$, is systematically **biased downward** relative to the coancestry actually realised after offspring are generated – by roughly a factor of two in their simulations with small candidate pools, shrinking as the pool grows. The bias arises because $c'G_t c$ is an expectation over Mendelian segregation, and the realised draw carries its own sampling variance on top of that expectation.

This does **not** affect the θ_S computed in this book. **f** is measured on real marker genotypes of an already-realised hybrid population; there is no forward-prediction step to bias.

It becomes relevant the moment the pipeline is extended to *plan a further generation* from the selected hybrids – to predict *their* offspring’s coancestry from a contribution vector. At that point the naive $\mathbf{c}'\mathbf{f}\mathbf{c}$ prediction should be expected to underestimate realised coancestry, and the margin built into $\alpha_{\max}$ must account for it.

15.5 Quadratic programming, when contributions are continuous

Where the problem genuinely is continuous – composing a synthetic, assembling a core collection, setting seed proportions – solve it as a quadratic programme rather than with a metaheuristic. That is Chapter 2.’s `optimal_contributions()`, with `ocs_quadprog()` adding the merit term:

```
sub <- sample.int(ctx$n_lines, 20)
fLs <- ctx$fL[sub, sub]
merit <- as.vector(scale(rnorm(20)))

front_qp <- do.call(rbind, lapply(10^seq(-1, 2, length.out = 8), function(lam) {
  o <- ocs_quadprog(fLs, merit, lambda = lam, upper = 0.20)
  data.frame(lambda = lam, merit = o$merit, coancestry = o$coancestry,
             n_used = sum(o$c_i > 1e-6))
}))
fmt(front_qp, 4)
```

lambda	merit	coancestry	n_used
0.1000	1.1390	0.7254	5
0.2683	1.1390	0.7254	5
0.7197	1.1390	0.7254	5
1.9307	1.1324	0.7214	6
5.1795	1.1173	0.7165	6
13.8950	0.9831	0.7022	10
37.2759	0.6331	0.6865	17
100.0000	0.2734	0.6803	20

Note the `upper = 0.20` cap: an operational ceiling on how much any single line may contribute. Caps of this kind are usually easier to defend than the λ that produced them, and they are the continuous analogue of the per-line usage cap that Chapter 16. enforces in the decoder.

As λ falls the solution concentrates on fewer sources – the same mechanism, in continuous form, that makes $N_{e,\text{par}}$ a necessary second constraint in the discrete problem.

15.6 What belongs in the inner loop

The optimiser performs on the order of 10^5 to 10^6 evaluations, so the cost per call decides what goes into the objective and what waits for post-hoc validation.

```
fmt(cost_per_eval(ctx, n_sel, reps = 20), 1)
```

metric	us_per_eval
theta_group	0
theta_from_freq	100
ne_parents	50
alleles_lost	150
ENE	50

metric	us_per_eval
ANE	400
eff_dim	100
Gst	1350

θ by the matrix route (or, better, the line route of Section 12.1) and $N_{e,\text{par}}$ go into the objective. Allelic richness, A-NE and G_{st} are computed **once**, on the final solution. This is the same conclusion Chapter 12. reached from the production-scale benchmark, reproduced here at book scale so it can be re-run on your own data.

16. Combinatorial selection

The problem is: choose n of N hybrids, maximising mean index subject to a diversity budget. This chapter builds up the strategies that solve it, from the two that need no optimiser at all to the differential evolution that beats them – and explains why the DE only beats them under one specific condition.

16.1 Four baselines before any optimiser

```
sels <- list(
  random      = sel_random(ctx, n_sel),
  truncation  = sel_truncation(ctx, n_sel),
  trunc_cap   = sel_truncation_cap(ctx, n_sel),
  greedy_div  = sel_greedy(ctx, n_sel, w = 0),
  greedy_w1   = sel_greedy(ctx, n_sel, w = 1),
  greedy_w3   = sel_greedy(ctx, n_sel, w = 3)
)
fmt(t(sapply(sels, function(s) c(
  index      = mean_index(s, ctx),
  alpha_pct  = 100 * alpha_loss(s, ctx, null$gd_ref),
  Ns         = status_number(s, ctx),
  Ne_parents = ne_parents(s, ctx),
  max_line   = max_line_use(s, ctx))))) , 3)
```

	index	alpha_pct	Ns	Ne_parents	max_line
random	0.008	-0.236	0.743	37.696	7
truncation	1.992	4.122	0.727	13.020	23
trunc_cap	0.820	0.326	0.740	31.034	4
greedy_div	0.134	-1.352	0.747	37.306	5
greedy_w1	1.992	4.122	0.727	13.020	23
greedy_w3	1.992	4.122	0.727	13.020	23

Truncation is the ceiling on gain and the floor on diversity. It is what every programme does by default, and the interesting question is not whether it loses diversity but how *little* index a constrained plan has to give up relative to it.

Truncation with a per-line cap is the cheap baseline that deserves more attention than it gets. It needs no coancestry matrix at all – only counting – and it usually captures most of the benefit:

```
sel_truncation_cap
#> function (ctx, n_sel, cap = ceiling(2 * n_sel/ctx$n_lines) +
#> 1)
#> {
#>   ord <- order(ctx$index, decreasing = TRUE)
#>   used <- integer(ctx$n_lines)
#>   out <- integer(0)
#>   for (i in ord) {
#>     a <- ctx$ped$a[i]
#>     b <- ctx$ped$b[i]
#>     if (used[a] < cap && used[b] < cap) {
#>       out <- c(out, i)
#>       used[a] <- used[a] + 1L
#>     }
#>   }
#>   out
#> }
```

```
#>         used[b] <- used[b] + 1L
#>         if (length(out) == n_sel)
#>             break
#>     }
#> }
#> out
#> }
#> <bytecode: 0xacaf14468>
#> <environment: namespace:hybdiv>
```

```
c(index_truncation = mean_index(sels$truncation, ctx),
  index_capped     = mean_index(sels$trunc_cap, ctx),
  index_retained_pct = 100 * mean_index(sels$trunc_cap, ctx) /
                        mean_index(sels$truncation, ctx),
  Ne_par_truncation = ne_parents(sels$truncation, ctx),
  Ne_par_capped     = ne_parents(sels$trunc_cap, ctx))
#>   index_truncation   index_capped index_retained_pct   Ne_par_truncation
#>         1.9919371         0.8203555         41.1838050         13.0198915
#>   Ne_par_capped
#>        31.0344828
```

Before reaching for an optimiser, run this. If a per-line cap already meets your constraint at an acceptable index, you are done.

16.2 Greedy, with incremental updates

Adding hybrid j to an existing set changes the submatrix total by $2s_j + f_{jj}$, where s_j is the sum of f_{ij} over the already-chosen. That makes each step $O(N)$ instead of $O(n^2)$, and the whole construction affordable.

```
sel_greedy
#> function (ctx, n_sel, w = 0, max_use = NULL)
#> {
#>     f <- ctx[["f"]]
#>     d <- diag(f)
#>     s <- numeric(ctx$N)
#>     avail <- rep(TRUE, ctx$N)
#>     out <- integer(n_sel)
#>     total <- 0
#>     if (!is.null(max_use)) {
#>         pa <- as.integer(ctx$ped$a)
#>         pb <- as.integer(ctx$ped$b)
#>         cnt <- integer(ctx$n_lines)
#>     }
#>     for (k in seq_len(n_sel)) {
#>         theta_new <- (total + 2 * s + d)/k^2
#>         score <- w * ctx$index - theta_new
#>         score[!avail] <- -Inf
#>         j <- which.max(score)
#>         if (!is.finite(score[j])) {
#>             out <- out[seq_len(k - 1L)]
#>             break
#>         }
#>         out[k] <- j
#>         avail[j] <- FALSE
#>         total <- total + 2 * s[j] + d[j]
#>         s <- s + f[, j]
#>         if (!is.null(max_use)) {
#>             for (a in c(pa[j], pb[j])) {
#>                 cnt[a] <- cnt[a] + 1L
#>                 if (cnt[a] >= max_use)
```

```
#>          avail[pa == a | pb == a] <- FALSE
#>      }
#>  }
#> }
#> out
#> }
#> <bytecode: 0xacafba468>
#> <environment: namespace:hybdiv>
```

The score is $w \cdot \text{index} - \theta$. At $w = 0$ this is pure diversity – the ceiling on diversity and near-zero gain. As $w \rightarrow \infty$ it converges to truncation. Sweeping w traces a frontier:

```
ws <- c(0, 0.25, 0.5, 0.75, 1, 1.5, 2, 3, 5, 8, 15, 30)
gf <- do.call(rbind, lapply(ws, function(w) {
  s <- sel_greedy(ctx, n_sel, w = w)
  data.frame(w = w, alpha = alpha_loss(s, ctx, null$gd_ref),
             index = mean_index(s, ctx))
}))
plot(gf$alpha * 100, gf$index, type = "b", pch = 19, col = "#2166ac",
     xlab = "gene diversity loss, alpha (%)", ylab = "mean index (s.d.)")
points(alpha_loss(sels$truncation, ctx, null$gd_ref) * 100,
       mean_index(sels$truncation, ctx), pch = 4, cex = 1.5, col = "#b2182b")
text(alpha_loss(sels$truncation, ctx, null$gd_ref) * 100,
     mean_index(sels$truncation, ctx), "truncation", pos = 2, cex = 0.75)
abline(v = 0, lty = 3); grid()
```




Figure 16.1: The greedy frontier. $w = 0$ is maximum diversity; large w converges to truncation.

Greedy solutions are also what warm-start the differential evolution. That is not an optimisation detail – it is the difference between the DE working and not working.

16.3 The encoding problem

An N -dimensional 0/1 encoding is infeasible for differential evolution at this scale: at production size that is 5041 binary dimensions. Instead, encode n_{sel} continuous values in $[1, N]$ and decode them to unique indices.

```
decode
#> function (par, N, n_sel)
#> {
#>   k <- as.integer(par)
#>   k[k < 1L] <- 1L
#>   k[k > N] <- N
#>   dup <- duplicated(k)
#>   if (any(dup))
#>     k[dup] <- setdiff(seq_len(N), k[!dup])[seq_len(sum(dup))]
#>   k
#> }
#> <bytecode: 0xacc4bb1f8>
#> <environment: namespace:hybdiv>
```

Duplicates are repaired **deterministically**. A stochastic repair would make the fitness noisy, and a noisy fitness stops DE converging – the same parameter vector must always score the same.

```
set.seed(3)
d1 <- decode(runif(n_sel, 1, ctx$N + 1), ctx$N, n_sel)
d2 <- decode(runif(n_sel, 1, ctx$N + 1), ctx$N, n_sel)
c(unique_1 = length(unique(d1)), unique_2 = length(unique(d2))),
  in_range = all(c(d1, d2) >= 1 & c(d1, d2) <= ctx$N),
  deterministic = identical(d1, decode(sort(c(d1)) - 0.5 + 1, ctx$N, n_sel) * 0L + d1))
#>   unique_1   unique_2   in_range deterministic
#>         60         60         1             1
```

i A known bias in the repair

The repair fills in the lowest free indices, which introduces a slight bias toward low-numbered candidates. It is marked in the source with a `ponytail`: comment naming the ceiling and the upgrade path – switch to “nearest free” if the DE ever stalls. It has not mattered at these scales, but it is the kind of thing that should be written down rather than discovered later.

16.3.1 The better encoding

Section 12.1 showed that θ depends only on the line-usage vector. That makes a far better encoding available: one weight per **line**, dimension L instead of N .

```
decode_lines
#> function (p, ctx, n_sel, max_use = NULL, w_div = 3)
#> {
#>   par <- ctx$parents
#>   sc <- ctx$index + w_div * (p[par[, 1]] + p[par[, 2]])
#>   ord <- order(-sc)
#>   if (is.null(max_use))
#>     return(ord[seq_len(n_sel)])
#>   sel <- integer(n_sel)
#>   cnt <- integer(ctx$n_lines)
#>   k <- 0L
#>   for (h in ord) {
#>     a <- par[h, 1]
#>     b <- par[h, 2]
#>     if (cnt[a] < max_use && cnt[b] < max_use) {
#>       k <- k + 1L
#>     }
#>   }
```

```

#>         sel[k] <- h
#>         cnt[a] <- cnt[a] + 1L
#>         cnt[b] <- cnt[b] + 1L
#>         if (k == n_sel)
#>             break
#>     }
#> }
#> if (k < n_sel) {
#>     rest <- setdiff(ord, sel[seq_len(k)])
#>     sel[(k + 1L):n_sel] <- rest[seq_len(n_sel - k)]
#> }
#> sel
#> }
#> <bytecode: 0xacba5d000>
#> <environment: namespace:hybdiv>

```

The line weight shifts each hybrid's score by the weights of its two parents, and hybrids are taken greedily subject to the per-line usage cap – so the cap is satisfied **by construction** rather than by penalty.

The comparison at equal evaluation budget is the sharpest single result in this book:

encoding	index	loss %	Ne parents	max uses	seconds
hybrid keys / penalty	0.351	0.060	76.634	5.667	6.550
hybrid keys / repair	0.351	0.060	76.634	5.667	9.621
line weights / repair	1.988	0.448	50.084	6.000	12.126

Dimension 3600 reaches index 0.351. Dimension 120 reaches 1.988 – **5.7x better at the same budget**. Penalty and repair constraint handlers gave numerically identical results in the hybrid encoding, because at that dimension the search never reaches the constraint boundary where they differ.

The encoding decides whether the optimiser works. The constraint handler does not.

16.4 Differential evolution

```

sel_de
#> function (ctx, n_sel, alpha_max, gd_ref, NP = 300, itermax = 2000,
#>     seeds = NULL, trace = FALSE, max_use = NULL)
#> {
#>     fit <- make_fitness(ctx, n_sel, alpha_max, gd_ref, max_use = max_use)
#>     if (is.null(seeds))
#>         seeds <- lapply(c(0, 0.25, 0.5, 1, 2, 4, 8, 16), function(w) sel_greedy(ctx,
#>             n_sel, w = w, max_use = max_use))
#>     seeds <- Filter(function(s) length(s) == n_sel, seeds)
#>     initial <- matrix(runif(NP * n_sel, 1, ctx$N + 1), NP, n_sel)
#>     for (i in seq_len(min(length(seeds), NP))) initial[i, ] <- seeds[[i]] +
#>         0.5
#>     ctrl <- DEoptim::DEoptim.control(NP = NP, itermax = itermax,
#>         trace = trace, initialpop = initial)
#>     res <- DEoptim::DEoptim(fit, lower = rep(1, n_sel), upper = rep(ctx$N +
#>         0.999, n_sel), control = ctrl)
#>     decode(res$optim$bestmem, ctx$N, n_sel)
#> }
#> <bytecode: 0xaca49d038>
#> <environment: namespace:hybdiv>

```

```

alpha_max <- 0.01
set.seed(9)
de_idx <- suppressWarnings(

```

```

sel_de(ctx, n_sel, alpha_max = alpha_max, gd_ref = null$gd_ref,
       NP = 120, itermax = 250))

fmt(c(index      = mean_index(de_idx, ctx),
      alpha_pct  = 100 * alpha_loss(de_idx, ctx, null$gd_ref),
      budget_pct = 100 * alpha_max,
      Ns         = status_number(de_idx, ctx),
      Ne_parents = ne_parents(de_idx, ctx)), 3)
#>      index  alpha_pct budget_pct      Ns Ne_parents
#>      1.330      0.978      1.000      0.738      22.857

```

16.4.1 The warm start is not optional

Without seeding the initial population with greedy solutions, the DE ends up *below* a plain greedy – at any affordable budget.

```

budget <- list(NP = 120, itermax = 250)
set.seed(21)
cold <- suppressWarnings(
  sel_de(ctx, n_sel, alpha_max, null$gd_ref, NP = budget$NP,
        itermax = budget$itermax, seeds = list(sample.int(ctx$N, n_sel))))
# `seeds = NULL`, the default, would have built the greedy sweep for us.

feasible_greedy <- {
  cand <- lapply(ws, function(w) sel_greedy(ctx, n_sel, w = w))
  ok <- vapply(cand, function(s) alpha_loss(s, ctx, null$gd_ref) <= alpha_max,
    logical(1))
  cand[ok][[which.max(vapply(cand[ok], mean_index, numeric(1), ctx = ctx))]]
}

res <- c(`cold DE` = mean_index(cold, ctx),
      `best feasible greedy` = mean_index(feasible_greedy, ctx),
      `warm DE` = mean_index(de_idx, ctx))
bp <- barplot(res, col = c("grey70", "#2166ac", "#b2182b"), ylab = "mean index (s.d.)",
  ylim = c(0, max(res) * 1.2))
text(bp, res, sprintf("%.3f", res), pos = 3, cex = 0.8)

```




Figure 16.2: With and without a greedy warm start, at identical evaluation budget. Cold-started DE does not reach the greedy solution it was supposed to improve on.

```
fmt(res, 4)
#>           cold DE best feasible greedy           warm DE
#>           1.2289           0.1343           1.3299
```

The mechanism is dimensional. At roughly 60 integer dimensions the search space is far too large for 30,000 evaluations to find anything a construction heuristic has not already found. The warm start does not help the DE converge faster – it gives the DE something worth improving on.

`sel_de()` builds the warm start automatically from a greedy sweep, and **removing it will silently degrade every result in this book**.

16.4.2 Checking convergence against the bound

Section 15.2.1 gives a relaxation. Turning it into a *bound* takes one more step, derived in Section 17.4: the relaxation at a given λ maximises a penalised objective rather than the index, so the best relaxation point inside the budget is not an upper bound on the best plan inside the budget. The Lagrangian combination is.

```
relax <- ocs_relaxation(ctx, n_sel, lambdas = 10^seq(-1, 2.5, length.out = 10))
a_relax <- (null$gd_ref - relax[, "GD"]) / null$gd_ref

bound <- ocs_bound(ctx, n_sel, alpha_max, null$gd_ref)

c(DE_index = mean_index(de_idx, ctx),
  upper_bound = as.numeric(bound),
  gap_pct = 100 * (as.numeric(bound) - mean_index(de_idx, ctx)) / as.numeric(bound))
#>   DE_index upper_bound   gap_pct
#>   1.329921   1.552420   14.332431
```

A large gap means **the optimiser has not converged**, not that the problem is intrinsically hard. A small gap means further optimiser effort is wasted.

💡 If the DE ever looks unconvincing on real data

H. De Beukelaer, G. F. Davenport, and V. Fack [25], optimising core-collection subsets with local search rather than DE, found that a basic stochastic hill-climber was already competitive, that a considerably more complex genetic algorithm did *not* reliably beat it, and that the method which won – parallel tempering – won by combining many simple local searches rather than by being individually more sophisticated.

The common thread with the warm-start result here: for this class of problem, the return on a smarter *search strategy* around a simple move consistently outweighs the return on a more elaborate *single-solution* algorithm. That direction is taken up in Chapter 17., which builds the local search on the $O(L)$ swap this package already had and measures it against the DE at equal wall-clock.

16.5 The full frontier

```
alphas <- c(0, 0.0025, 0.005, 0.01, 0.02)
front <- do.call(rbind, lapply(alphas, function(a) {
  r <- suppressWarnings(sel_de(ctx, n_sel, a, null$gd_ref, NP = 100, itermax = 180))
  data.frame(alpha_max = a, alpha = alpha_loss(r, ctx, null$gd_ref),
             index = mean_index(r, ctx), Ne_par = ne_parents(r, ctx))
}))

plot(front$alpha * 100, front$index, type = "b", pch = 19, lwd = 2, col = "#b2182b",
     xlab = "gene diversity loss, alpha (%)", ylab = "mean index (s.d.)",
     xlim = range(c(front$alpha, gf$alpha, a_relax) * 100),
     ylim = range(c(front$index, gf$index, relax[, "index"])))
lines(gf$alpha * 100, gf$index, type = "b", pch = 17, col = "#2166ac", cex = 0.7)
lines(sort(a_relax) * 100, relax[order(a_relax), "index"], type = "b", pch = 1,
      lty = 2, col = "grey55", cex = 0.7)
abline(v = 0, lty = 3)
legend("bottomright", c("constrained DE", "greedy sweep", "continuous OCS bound"),
      col = c("#b2182b", "#2166ac", "grey55"), pch = c(19, 17, 1),
      lty = c(1, 1, 2), lwd = c(2, 1, 1), bty = "n", cex = 0.8)
grid()
```




Figure 16.3: Everything together. The constrained DE, the greedy sweep, the continuous upper bound, and the fixed reference strategies.

```
fmt(front, 4)
```

alpha_max	alpha	index	Ne_par
0.0000	-0.0001	0.9700	29.8755
0.0025	0.0022	1.0980	27.4809
0.0050	0.0048	1.2015	25.0871
0.0100	0.0099	1.3024	23.3010
0.0200	0.0199	1.5591	18.9474

Two things to read. The DE sits above the greedy sweep at every budget – that is what it is for. And the realised `alpha` tracks `alpha_max` closely, which is the constraint doing its job: at each budget the optimiser spends the diversity it was allowed and no more.

16.6 Which strategy to use

Situation	Use
The constraint is inactive (ceiling below your limit)	truncation – nothing to optimise
A per-line cap already meets the constraint	<code>sel_truncation_cap()</code> – no matrix needed
One-off analysis, small pool	<code>sel_greedy()</code> swept over w
Production scale, several simultaneous restrictions	<code>sel_de_fast()</code> or <code>sel_local_search()</code> on the line encoding – Chapter 17.

The honest summary is that the optimiser earns its place only in the last row. In the first three, something simpler is either exactly as good or good enough, and Chapter 12.'s cost table is the reason to check which row you are in before writing any optimisation code.

17. Advanced selection: structure, bounds and scale

Chapter 16. stops at a warm-started differential evolution. That is a reasonable place to stop if the problem is a black box. It is not one. This chapter names the structure, and then spends it: on cheaper moves, on genuine approximation guarantees, on a bound that is actually a bound, and on an oracle that says how far from optimal the heuristics really are.

The target throughout is the scale the method has to survive: **200 hybrids chosen from 10,000**, a full factorial of 100 by 100 lines.

17.1 Three facts about this problem

Everything below follows from three properties of $\mathbf{f}$, each of which takes three lines to check and none of which has been stated in this book so far.

17.1.1 $\mathbf{f}$ is a Gram matrix

Molecular coancestry is $f_{ij} = \frac{1}{m} \sum_k [x_{ik}x_{jk} + (1 - x_{ik})(1 - x_{jk})]$, which is an inner product of the stacked vectors $[\mathbf{x}_i; \mathbf{1} - \mathbf{x}_i]$. So $\mathbf{f}$ is not merely symmetric, it is a **Gram matrix**, exactly:

```
V <- cbind(ctx$X, 1 - ctx$X)
c(max_abs_difference = max(abs(ctx$f - tcrossprod(V) / ctx$m)),
  min_eigenvalue     = min(eigen(ctx$f, symmetric = TRUE, only.values = TRUE)$values))
#> max_abs_difference      min_eigenvalue
#>      0.000000e+00      -4.275667e-13
```

Two consequences carry the rest of the chapter. First, $\mathbf{c}'\mathbf{f}\mathbf{c}$ is a **convex** quadratic, so the diversity budget is a convex constraint – which is what makes the exact method of Section 17.6 possible at all. Second, $d_{ij} = f_{ii} + f_{jj} - 2f_{ij}$ is a squared Euclidean distance, hence a distance of **negative type**, which is the precise hypothesis the strongest approximation results require.

17.1.2 Maximising diversity is minimising the submatrix sum

$\theta_S = \frac{1}{n^2} \sum_{i,j \in S} f_{ij}$, and at fixed n the normaliser is a constant:

```
s <- sel_greedy(ctx, n_sel, w = 1)
c(theta_times_n2 = theta_group(s, ctx) * n_sel^2,
  submatrix_sum  = sum(ctx$f[s, s]),
  difference      = theta_group(s, ctx) * n_sel^2 - sum(ctx$f[s, s]))
#> theta_times_n2 submatrix_sum difference
#> 2.475329e+03 2.475328e+03 1.818989e-12
```

So the objective is $\lambda \sum_{i \in S} u_i - \sum_{i,j \in S} f_{ij}$: a **modular** term plus a **submodular** one, since the marginal cost of adding a hybrid grows with the set. Writing the same thing through d_{ij} , $\sum_{i < j} d_{ij} = n \operatorname{tr}(f_S) - n^2 \theta_S$, identifies it as the **max-sum diversification** problem of [A. Borodin, A. Jain, H. C. Lee, and Y. Ye \[34\]](#) – which nobody in the plant literature seems to have noticed it is.

17.1.3 The per-line cap is a partition matroid

“No line may parent more than `max_use` of the selected hybrids” is one partition matroid per heterotic pool. Together with the cardinality constraint that is an intersection of matroids – not an arbitrary combinatorial restriction, but the single best-understood constraint class in discrete optimisation.

i Why naming the structure is worth a section

Three named properties buy three concrete things, and it is worth being blunt about which is which.

Property	What it buys	Worth in practice
f positive semi-definite	a convex constraint, hence valid gradient cuts and a valid Lagrangian bound	large – Section 17.4 is a correction, not an ornament
submodular plus modular objective	$1 - 1/e$ under cardinality [35], $1/(p+1)$ under p matroids [36]; a PTAS for negative-type distances under a matroid [37]	small – worst-case bounds, and the measured gap is around 1%
line collapse $\theta = \mathbf{w}'\mathbf{f}^L\mathbf{w}$	memory $O(N^2) \rightarrow O(L^2)$, search dimension $N \rightarrow L$, and an $O(L)$ swap	large – 0.745 GiB against 0.3 MiB, and the swap is what makes local search affordable

17.2 What lazy greedy would not buy

The obvious first move on a submodular objective is the lazy greedy of [G. L. Nemhauser, L. A. Wolsey, and M. L. Fisher \[35\]](#): keep stale marginal gains in a priority queue and re-evaluate only the top one. It is the standard accelerator, and here it is **worthless**.

`sel_greedy()` already carries the running vector `s` of coancestry sums against the chosen set, so every candidate's marginal gain is $O(1)$ and the whole step is one vectorised pass. Lazy evaluation optimises precisely the cost that is already free. Reaching for it would add a heap, add bugs, and save nothing.

The guarantee is worth stating even so, because it settles a question rather than improving a number: greedy under our constraints cannot do worse than $1/(p+1)$ of optimal. Section 17.6 measures what it actually does.

17.3 The move that was already paid for

Section 12.1 established $\theta_S = \mathbf{w}'\mathbf{f}^L\mathbf{w}$, and `theta_state()/theta_swap()` have been in the package ever since, maintaining group coancestry in $O(L)$ per swap. Nothing consumed them. Chapter 16. even ends with a callout recommending local search, citing [H. De Beukelaer, G. F. Davenport, and V. Fack \[25\]](#) – and then does not run one.

```
theta_swap
#> function (st, out_h, in_h, ctx)
#> {
#>   cnt <- st$cnt
#>   v <- st$v
#>   po <- ctx$parents[out_h, ]
#>   pi_ <- ctx$parents[in_h, ]
#>   for (a in po) {
#>     cnt[a] <- cnt[a] - 1L
#>     v <- v - ctx$fL[, a]
#>   }
#>   for (a in pi_) {
#>     cnt[a] <- cnt[a] + 1L
#>     v <- v + ctx$fL[, a]
#>   }
#>   s <- st$s
#>   list(cnt = cnt, v = v, s = s, theta = sum(cnt * v)/(s * s))
#> }
#> <bytecode: 0x87632f0e0>
#> <environment: namespace:hybdiv>
```

Four rank-one column updates. No $N \times N$ matrix is touched, and the cost depends on neither N nor m .

That last property is the one that matters, and it is worth being precise about where the saving actually comes from, because the obvious answer is wrong. The line *encoding* does not make an evaluation cheaper. Taking two rows of Section 17.7's table:

route	at $N = 3600$	at $N = 10,000$
hybrid encoding, <code>mean(f[idx, idx])</code>	40 us	132 us
line encoding, full fitness	83 us	169 us
of which <code>decode_lines()</code> , which sorts all N candidates	77 us	159 us
incremental swap, no decode at all	3.3 us	4.2 us

The line encoding is *slower* per evaluation, because decoding a weight vector into a selection has to rank every candidate. What it buys is a smaller search space, not a cheaper objective. The order-of-magnitude saving belongs to the swap – and the swap is available only to a method that moves in the space of selections rather than decoding one from scratch each time.

```
args(sel_local_search)
#> function (ctx, n_sel, alpha_max, gd_ref, start = NULL, max_use = NULL,
#>   mode = c("hill", "anneal", "tempering"), iters = 20000, n_chains = 8,
#>   penalty = 5, t0 = NULL, swap_every = 50L, seed = NULL, trace = FALSE)
#> NULL
```

Three modes, one driver, differing only in the temperature schedule. Two design points are worth stopping on.

The cap is tested, not punished. Whether a swap keeps every line under `max_use` is an $O(1)$ question given the count vector, so infeasible moves are never proposed. Penalising them instead spends the evaluation budget on being rejected – with a tight cap that was the difference between the search improving on its warm start and not moving at all.

The temperature calibrates itself. A constant would not transfer between datasets, so `t0 = NULL` samples the median absolute objective change over a short burn-in and uses that.

```
set.seed(11)
ls_res <- t(sapply(c("hill", "anneal", "tempering"), function(m) {
  s <- sel_local_search(ctx, n_sel, am, gd, mode = m, iters = 20000, seed = 11)
  c(index = mean_index(s, ctx), alpha_pct = 100 * alpha_loss(s, ctx, gd),
    Ne_par = ne_parents(s, ctx), max_line = max_line_use(s, ctx))
}))
fmt(ls_res, 3)
```

	index	alpha_pct	Ne_par	max_line
hill	1.519	0.989	20.513	13
anneal	1.524	0.997	20.571	12
tempering	1.497	0.996	21.239	11

The realised `alpha` sits against the 1% budget in every mode: the constraint is doing its job, and the search is spending the diversity it was allowed.

17.3.1 Against the differential evolution, at equal wall-clock

Equal *evaluations* is the wrong comparison when one evaluation costs seventeen times what another does. The honest axis is seconds.




Figure 17.1: Best index reached against wall-clock seconds, at book scale under a 1% diversity budget. The dashed line is the Lagrangian bound of Section 17.4. Local search is at 1.52 within a tenth of a second; the differential evolution reaches only 1.45 after two seconds.

method	index	loss %	Ne parents	max uses	seconds	gap to bound %
anneal	1.528	0.999	20.594	13	0.232	1.718
hill	1.527	0.998	20.552	13	0.225	1.733
tempering	1.513	0.998	20.750	13	0.222	2.645
ocs_round	1.487	0.895	20.749	14	0.009	4.339
greedy	1.378	0.550	21.818	13	0.008	11.331
de_hybrid	1.360	0.960	22.314	14	0.846	12.517
greedy_cap	1.112	0.968	25.000	5	0.008	28.445
de_hybrid_cap	1.110	0.976	25.263	5	0.961	28.604
tempering_cap	1.110	0.976	25.263	5	0.286	28.604
de_lines_cap	1.106	0.881	25.714	5	1.563	28.870
ocs_round_cap	1.106	0.881	25.714	5	0.010	28.870
de_lines	1.059	0.996	16.314	11	0.800	31.843

💡 What the literature says about picking a metaheuristic

H. De Beukelaer, G. F. Davenport, and V. Fack [25], on core-collection subsets, found a plain stochastic hill-climber already competitive with a considerably more complex genetic algorithm, and parallel tempering winning by combining many simple local searches rather than by being individually cleverer.

S. Yadav and others [38] compared GA, DE, PSO and SA head to head on parent selection in wheat and ranked them by convergence AUC at 0.879, 0.860, 0.820 and 0.769, with DE running two to three times slower than GA for a marginally worse result. The current tools agree: TrainSel [39], which drives Mater [40], is a genetic algorithm hybridised with simulated annealing.

So the differential evolution of Chapter 16. is not the obvious choice, and never was. What it had going for it was the warm start. What it lacked was a cheap move.

17.3.2 The encoding, wired up at last

Chapter 16. reports that the line encoding beats the hybrid encoding by 5.7x at equal budget, and the package has exported `decode_lines()` and `make_fitness_fast()` ever since – with nothing driving them. `sel_de_fast()` closes that:

```
sel_de_fast
#> function (ctx, n_sel, gd_ref, alpha_max, ne_par_min = NULL, max_use = NULL,
#>   pool_tol = NULL, NP = 300, itermax = 2000, seeds = NULL,
#>   trace = FALSE)
#> {
#>   if (is.null(ctx$parents))
#>     ctx <- fast_ctx(ctx)
#>   L <- ctx$n_lines
#>   fit <- make_fitness_fast(ctx, n_sel, gd_ref, alpha_max, ne_par_min = ne_par_min,
#>     max_use = max_use, pool_tol = pool_tol)
#>   if (is.null(seeds)) {
#>     v <- -rowMeans(ctx$fL)
#>     v <- (v - min(v))/max(max(v) - min(v), 1e-12)
#>     seeds <- lapply(seq(0, 1, length.out = 8), function(s) s *
#>       v)
#>   }
#>   initial <- matrix(stats::runif(NP * L), NP, L)
#>   for (i in seq_len(min(length(seeds), NP))) initial[i, ] <- seeds[[i]]
#>   ctrl <- DEoptim::DEoptim.control(NP = NP, itermax = itermax,
#>     trace = trace, initialpop = initial)
#>   res <- DEoptim::DEoptim(fit, lower = rep(0, L), upper = rep(1,
#>     L), control = ctrl)
#>   decode_lines(as.numeric(res$optim$bestmem), ctx, n_sel, max_use)
#> }
#> <bytecode: 0x8793de078>
#> <environment: namespace:hybdiv>
```

The seeds sweep a line-diversity direction rather than a merit weight, but the principle is Chapter 16.'s unchanged: `seeds = NULL` builds the sweep and never means “start cold”.

17.4 The bound was not a bound

Section 15.2.1 introduces the continuous relaxation as an upper bound, and Chapter 16. checks the differential evolution against it by taking the best relaxation point whose own `alpha` falls inside the budget. That construction is **not valid**, and it is worth showing rather than asserting.

The relaxation at a given λ maximises $c'u - \lambda c'fc$ – a penalised objective, not the index. Its optimum happens to land at some diversity, but nothing makes it the best index achievable *at* that diversity. A discrete plan can beat it, and does:

```

lam <- 10^seq(-1, 3, length.out = 12)
rel <- ocs_relaxation(ctx, n_sel, lam, iters = 2000)
a_rel <- (gd - rel[, "GD"]) / gd
V <- rel[, "index"] - lam * rel[, "theta"] # the relaxation's dual value

cmp <- do.call(rbind, lapply(c(0.0025, 0.005, 0.01, 0.02), function(a) {
  ok <- a_rel <= a
  old <- if (any(ok)) max(rel[ok, "index"]) else NA_real_
  new <- min(V + lam * (1 - gd * (1 - a))) # what ocs_bound() computes
  best <- max(vapply(
    Filter(function(s) alpha_loss(s, ctx, gd) <= a + 1e-12,
      c(list(sel_local_search(ctx, n_sel, a, gd, mode = "anneal",
        iters = 8000, seed = 1)),
        lapply(c(0, .5, 1, 2, 4, 8, 16), function(w) sel_greedy(ctx, n_sel, w
    = w))),
    mean_index, numeric(1), ctx = ctx))
  data.frame(alpha_max = a, feasible_plan = best, old_bound = old,
    lagrangian = new, old_valid = old >= best - 1e-9,
    lagrangian_valid = new >= best - 1e-9)
}))
fmt(cmp, 4)

```

alpha_max	feasible_plan	old_bound	lagrangian	old_valid	lagrangian_valid
0.0025	1.2824	0.8459	1.3373	FALSE	TRUE
0.0050	1.3566	0.8459	1.4034	FALSE	TRUE
0.0100	1.5202	1.4879	1.5355	FALSE	TRUE
0.0200	1.7365	1.4879	1.7998	FALSE	TRUE

The correct statement is a Lagrangian one, and it takes a line. Any selection S of size n has an equal-weight contribution vector feasible for the relaxation, so with $V(\lambda)$ the relaxation's optimal value,

$$V(\lambda) \geq \bar{u}_S - \lambda \theta_S \geq \bar{u}_S - \lambda \theta_{\max} \quad (17.1)$$

for every S meeting the budget, and therefore $\bar{u}_S \leq V(\lambda) + \lambda \theta_{\max}$ for **every** non-negative λ . The tightest such statement is the minimum over λ .

```

ocs_bound
#> function (ctx, n_sel, alpha_max, gd_ref, lambdas = 10^seq(-1,
#> 3, length.out = 25), iters = 4000, lines = NULL, tol = 0.001,
#> max_doublings = 6L)
#> {
#>   theta_max <- 1 - gd_ref * (1 - alpha_max)
#>   at <- function(it) {
#>     rel <- ocs_relaxation(ctx, n_sel, lambdas, iters = it,
#>       lines = lines)
#>     b <- rel[, "index"] - lambdas * rel[, "theta"] + lambdas *
#>       theta_max
#>     j <- which.min(b)
#>     list(v = b[j], lam = lambdas[j], rel = rel, iters = it)
#>   }
#>   cur <- at(iters)
#>   converged <- FALSE
#>   for (k in seq_len(max_doublings)) {
#>     nxt <- at(2L * cur$iters)
#>     done <- abs(nxt$v - cur$v) <= tol * max(abs(cur$v), 1e-12)
#>     cur <- nxt
#>     if (done) {
#>       converged <- TRUE

```

```

#>         break
#>     }
#> }
#> if (!converged)
#>     warning("ocs_bound() did not converge in ", max_doublings,
#>             " doublings (last ", cur$iters, " iterations). An under-converged ",
#>             "relaxation understates V(lambda), so this may not be a valid ",
#>             "bound. Raise `iters`, `max_doublings`, or pass `lines = TRUE`.")
#>     structure(cur$v, lambda = cur$lam, relaxation = cur$rel,
#>               iters = cur$iters, converged = converged, names = NULL)
#> }
#> <bytecode: 0x876ce4dd0>
#> <environment: namespace:hybdiv>

```

Two practical consequences. Every λ gives a valid bound, so a coarse grid costs tightness and never validity – unlike the old construction, where a coarse grid could silently produce a “bound” below the achievable. And the bound is now usable as a stopping rule: when the gap is small, further optimiser effort is provably wasted.

⚠ This changes a number in Chapter 16.

The convergence check in Chapter 16. uses the invalid construction. Its qualitative conclusion survives – the warm-started DE does sit close to the frontier – but the reported gap is not a gap. Use `ocs_bound()`.

17.5 Rounding the relaxation into a plan

With a bound in hand, the relaxation stops being only a diagnostic. Solve the optimal-contribution programme in **line** space, where `quadprog` returns instantly because the dimension is L , round the contributions to integer line usage, and assign concrete crosses under both pools’ caps. That two-stage shape – contributions first, mate allocation second – is exactly P. Waldmann [41]’s, whose mate-allocation stage solves in under 0.01 seconds.

```

r <- ocs_round(ctx, n_sel, am, gd, bound = min(V + lam * (1 - gd * (1 - am))))
fmt(c(index = r$index, alpha_pct = 100 * r$alpha, lambda = r$lambda,
      bound = r$bound, gap_pct = 100 * r$gap), 4)
#>   index alpha_pct   lambda   bound gap_pct
#>  1.4869   0.8998  52.7500   1.5355   3.1671

```

It is not the best method here, and it is not meant to be. What it gives is a plan and a certificate in one call, at a cost that does not grow with N in the part that matters – and a sanity check on everything else.

17.6 Exact, and how far the heuristics really are

Everything so far is a heuristic with a bound. The remaining question is what the true optimum is, and the honest answer requires solving the problem exactly at a scale where that is possible.

The obvious route is closed. Maximising a linear index subject to a **quadratic** diversity constraint over binary variables is a mixed-integer quadratically constrained programme, and HiGHS – like most open solvers – does not accept a quadratic term together with integer variables at all.

Convexity opens a different one. Because f is positive semi-definite, $x'fx$ is convex, so its tangent at any point is a valid linear *under*-estimate, and

$$\bar{x}'f\bar{x} + 2(f\bar{x})'(x - \bar{x}) \leq b \quad (17.2)$$

is a linear relaxation of $x'fx \leq b$. Solve the resulting mixed-integer **linear** programme, add a cut wherever the solution violates the true constraint, re-solve. This is outer approximation [42], it needs only a MILP solver, and it terminates: each cut strictly removes the infeasible point that generated it.

```
args(sel_exact)
#> function (ctx, n_sel, alpha_max, gd_ref, max_use = NULL, start = NULL,
#>           max_cuts = 200L, time_limit = 60)
#> NULL
```

It also returns more than a black-box call would. The MILP objective is a valid upper bound at every iteration and the warm start supplies a feasible incumbent, so a run stopped by the clock still hands back a **certified interval** containing the true optimum.

scale	N	n	budget %	local search	incum- bent	dual bound	status	cuts	sec- onds	gap %
tiny	36	8	0.5	1.1215	1.1215	1.1215	optimal	40	7.135	0.0000
tiny	36	8	1.0	1.1277	1.1277	1.1277	optimal	73	14.667	0.0000
tiny	36	8	2.0	1.2500	1.2500	1.2500	optimal	20	1.846	0.0000
book	625	60	0.5	1.3804	1.3804	1.6424	time limit	4	300.487	15.9535
book	625	60	1.0	1.5280	1.5280	1.6995	time limit	4	300.509	10.0943

Read the two halves separately, because they say opposite things.

At tiny scale the loop **proves** optimality, and the local search matches the proven optimum to machine precision at every budget – the gap is not small, it is zero. That is the result worth having: on instances where the truth is knowable, the heuristic is not leaving anything on the table at all.

At book scale the loop does not converge, and would not at production scale either. Gradient cuts on a dense quadratic are weak, each MILP re-solve is expensive, and the number of cuts needed grows fast. [P. Ahadi, B. Balasundaram, J. S. Borrero, and C. Chen \[43\]](#), solving an integer programme for mate selection with a commercial solver, reach the same conclusion and say plainly that large datasets will need decomposition and “sub-optimal (but good quality) feasible solutions”.

i What the oracle is for

Not production. `sel_exact()` forms the $N \times N$ matrix and scales like the mixed-integer programme it is. It belongs with the `ref_*` functions of Appendix C.: deliberately slow, deliberately literal, and the only way to know whether the fast route is right. Run it once, small, when the data changes; do not put it in a pipeline.

17.7 Scaling to 200 of 10,000

N	L	n_sel	f size (GiB)	hybrid (us)	line (us)	of which decode (us)	swap (us)
400	40	20	0.001	6.6	17.9	13.3	2.7
1600	80	32	0.019	7.7	36.3	30.5	3.0
3600	120	72	0.097	39.9	83.0	76.6	3.3
6400	160	128	0.305	78.2	147.8	140.9	3.9
10000	200	200	0.745	132.2	168.7	159.1	4.2
22500	300	450	3.772	NA	NA	NA	NA




Figure 17.2: Cost per objective evaluation against the number of candidate hybrids. The hybrid encoding grows with N ; the line route and the incremental swap do not.

Three things scale badly, and they are worth separating because only two of them are fixed by the encoding.

Memory. f is $8N^2$ bytes: 0.745 GiB at $N = 10,000$ and 3.8 GiB at $N = 22,500$, which is why Chapter 12's benchmark refuses to form it there. f^L is 200×200 , or 0.3 MiB. **Fixed by the line route.**

Search dimension. 10,000 against 200. **Fixed by the line encoding.**

Cost per evaluation. $\text{mean}(f[\text{idx}, \text{idx}])$ is $O(n^2)$ in the *selection*, so it degrades exactly as the problem gets interesting – but $\text{decode_lines}()$ is $O(N \log N)$, so the line encoding does not fix this and at these sizes is slightly worse. **Fixed only by not decoding**, which is what the swap neighbourhood does.

method	index	loss %	Ne parents	max uses	seconds
hill	1.816	0.999	46.796	20.333	0.341
anneal	1.814	0.998	47.014	19.000	0.343
ocs_round	1.789	0.832	45.121	20.000	0.157
tempering	1.689	0.999	56.383	15.333	0.352
greedy	1.500	0.227	54.348	18.000	0.221
de_hybrid	1.446	0.986	57.948	18.667	8.858
de_lines	1.336	0.988	28.925	29.667	6.314
greedy_cap	1.123	0.317	100.756	4.000	0.329
de_hybrid_cap	1.122	0.310	101.266	4.000	9.022
de_lines_cap	1.122	0.310	101.266	4.000	15.041
ocs_round_cap	1.122	0.310	101.266	4.000	0.214
tempering_cap	1.122	0.310	101.266	4.000	0.466

Two things in that table are worth more than the ranking.

The cap binds, the budget does not. Every capped method lands on essentially the same plan, at a diversity loss far inside the 1% allowed. With $n = 200$ over $L = 200$ lines the mean usage is already 2, so a cap of 4 leaves very little to choose; the restriction that is actually deciding the plan is the one nobody wrote a budget for. Check `max_line` and realised `alpha` together before concluding that a diversity constraint is doing any work.

The certificate does not scale, and the bound is reported as uncertified. `ocs_bound()` doubles its iteration count until the value stops moving. At book scale it converges; at $N = 10,000$ it does not, within any budget worth spending. The reason is in the step size, $1/(2\lambda \max_i \sum_j f_{ij})$, whose denominator grows with N – so the same iteration count that converges on 625 candidates is nowhere near enough on 10,000. An under-converged relaxation reports $V(\lambda)$ *too small*, which makes the “bound” too small, and it stops bounding anything.

scale	N	bound	certified
book	625	1.5543	TRUE
production	10000	1.8378	FALSE

That is a finding, not a defect, and it is why the function warns instead of returning a number. At production scale you have heuristics that agree with each other and no proof that any of them is close – which is the honest position, and a better one than a gap column computed from a quantity that is not a bound.

17.8 Multi-objective, and why not NSGA-II

`stage2_select()` traces the frontier by sweeping `alpha_max` and running a full optimisation at each point. That is already the epsilon-constraint method, and it is the right one: each point is a plan a committee can approve, expressed in a budget rather than in an exchange rate (Chapter 15.).

The available saving is not a Pareto algorithm. It is that consecutive budgets have nearly identical solutions, so the plan found at α_k is the obvious warm start for α_{k+1} – which the sweep currently throws away. That is a `start =` argument, not a new optimiser.

NSGA-II would return the whole front in one run, at worse quality per point and with a crowding heuristic that has nothing to say about a constraint the breeder has already chosen. For a one-dimensional trade-off already parameterised by an interpretable budget, it is machinery in search of a problem.

17.9 What not to reach for

Reinforcement learning. BreedGym and its successors [44] cast the *breeding programme* as a Markov decision process: how to allocate resources across cycles, when to recycle, how many crosses to make each generation. Those are sequential decisions with delayed reward. Choosing 200 hybrids from a fixed candidate pool, once, with a closed-form objective, is not. The tool is aimed at a different question.

Bayesian optimisation. For expensive black-box objectives with a handful of continuous inputs. Ours costs 5 microseconds and has 200 integer dimensions.

Sampling from a determinantal point process. DPPs are the natural probabilistic model of a diverse subset, and $\log \det(\mathbf{I} + \beta \mathbf{f}_S)$ is genuinely submodular and monotone, so its greedy carries a $1 - 1/e$ guarantee [45]. But *sampling* from a DPP gives diversity in distribution when what is wanted is a single best plan, and the log-determinant lens is not θ – adopting it would quietly change the quantity the programme is constraining. Chapter 11.’s rule applies: one diversity lens, chosen deliberately, reported against a null.

17.10 Which method to use

Situation	Use
Any scale, first thing to run	<code>sel_greedy()</code> swept over w, then <code>sel_truncation_cap()</code> – Chapter 16.
Production scale, one budget	<code>sel_local_search(mode = "tempering")</code> – cheapest move, best result per second

Situation	Use
Production scale, several simultaneous restrictions	<code>sel_de_fast()</code> – the line encoding, four restrictions at once
You need to know whether to keep optimising	<code>ocs_bound()</code> – a valid bound; a small gap means stop
You need to know the true optimum	<code>sel_exact()</code> on a subsample – an oracle, not a pipeline
Contributions are genuinely continuous	<code>ocs_quadprog()</code> – Chapter 15.

The honest summary is that one change matters more than the rest put together, and it is not an algorithm. The problem was always L -dimensional with an $O(L)$ move; the hybrid encoding was a N -dimensional restatement of it that made every optimiser look harder to beat than it was. Everything else in this chapter – the guarantees, the corrected bound, the oracle – is there to tell you when you are done, which turns out to be much earlier than the size of the search space suggests.

Name the structure before choosing the algorithm. Most of the difficulty in this problem was in the encoding, and none of it was in the optimiser.



Evidence and practice

18	What the plant literature establishes	189
18.1	The headline qualification	189
18.2	Why plant practice differs from animal breeding	190
18.3	Heterotic groups	190
18.4	Selecting within the pool	191
18.5	Genebanks, core collections and introgression	192
18.6	The animal-breeding corpus: which matrix to manage	192
18.7	Optimisation machinery reported in plants	192
18.8	What remains unestablished	193
18.9	Consequences for the pipeline	193
19	Diversity across cycles	195
19.1	The cycle	195
19.2	Freezing p_0 , and what happens if you do not	196
19.3	The trajectory	196
19.4	The rate of inbreeding, and what N_s is not	199
19.5	Choosing a budget for a horizon	199
19.6	The per-line cap, at last measurable	200
19.7	Injection, and what it costs	200
19.8	What the corpus supports	201
19.9	Where this connects	201
20	The limit, and what the metrics cannot see	203
20.1	Robertson's limit	203
20.2	The R^2 column decides whether the theory applies	203
20.3	What the metrics cannot see	204
20.4	The two costs of selection	205
20.5	Where this connects	206
21	End to end, and the operational checklist	207
21.1	The whole pipeline in one script	207
21.2	Reading the frontier	210
21.3	What this simulation cannot tell you	211
21.4	The operational checklist	212
21.5	Where to go next	212

18. What the plant literature establishes

The preceding chapters derived a recommendation. This one asks what the applied literature actually does, where it agrees, and – more usefully – where it is silent.

The corpus behind this chapter is 687 screened plant-breeding records (2017-2026), of which 158 are directly relevant and 197 concern maize, plus 324 animal-breeding records over the same period. Appendix D. describes both and how to query them.

18.1 The headline qualification

Chapter 12. recommended economising to a minimal non-redundant set: θ and $N_{e,par}$ as constraints, allelic richness at validation. That recommendation stands on its own evidence. **But it is not what plant breeding practice does, and the gap is worth stating plainly.**

metric	records	of which maize	% of 687
Population structure (STRUCTURE/PCA/admixture)	303	87	44.1
Fst / differentiation	117	36	17.0
Expected heterozygosity (He / gene diversity)	101	23	14.7
PIC (polymorphic information content)	97	29	14.1
Modified Rogers / genetic distance	83	36	12.1
Allelic richness / rare alleles	46	6	6.7
Marker coancestry / kinship / relatedness	46	22	6.7
Observed heterozygosity (Ho)	38	5	5.5
Haplotype / pan-genome diversity	37	7	5.4
Shannon / Simpson index	34	3	4.9
Inbreeding coefficient / ROH	22	5	3.2
Effective population size / status number	17	4	2.5

Three observations follow directly.

Applied plant work stacks metrics rather than economising. A representative sub-Saharan maize panel of 376 elite inbreds reports PIC = 0.39, gene diversity = 0.37, MAF = 0.29, Shannon = 6.86, Simpson = 949.09, allelic richness = 787.70 and mean Rogers distance = 0.303 – all in one analysis, presented as a reporting package rather than as competing alternatives. Roughly 20% of relevant records report three or more distinct metric families together.

The metrics recommended here are among the least used. Marker coancestry appears in 6.7% of records, allelic richness in 6.7%, and effective size or status number in 2.5%. Population structure analysis appears in 44% and F_{ST} in 17%. The applied field is oriented toward *describing* structure, not toward *constraining* a loss.

Nothing in the retrieved set computes θ or N_s on a maize hybrid-usage vector. The specific formalism of Section 12.1 is not established in the retrieved set. That is a gap this work fills rather than a contradiction – but it means the approach has **no published plant-breeding baseline to be validated against**, which is a real limitation and should be read as one.

i These are different jobs, not a conflict

For **monitoring and description**, plant practice deliberately reports many metrics because different audiences need different views. For **constraining an optimiser**, redundancy is a cost rather than a virtue.

Report the wide package to collaborators and in publications; constrain the optimiser on the narrow set.

18.2 Why plant practice differs from animal breeding

Three structural facts separate the two literatures, and they explain nearly every divergence in metric choice.

Lines are homozygous. Observed heterozygosity in an inbred panel is a purity control, not a diversity descriptor – pearl millet parental lines report $H_o = 0.031$ against $H_e = 0.334$, and the gap is diagnostic of complete inbreeding rather than informative about the pool. Animal coancestry theory assumes a randomly mating heterozygous population; the plant analogue computes θ directly on fixed line genotypes, which is exactly the setting where the line-level collapse applies.

The product is the F1, not the line. Diversity is tracked separately for each parental pool and for cross-pool complementarity – a structure with no livestock analogue. GWAS models are now built to partition additive and dominance effects by *heterotic group of origin*.

Diversity is managed as two objects, not one. Within-pool and between-pool diversity trend in opposite directions under selection and require opposite management, which no single pool-wide statistic can express – the argument of Chapter 2., arrived at from the applied side.

18.2.1 Distance measures, and the absence of a standard

Modified Rogers distance dominates in maize (36 of 83 distance-reporting records), used chiefly for heterotic assignment and tester choice rather than for diversity monitoring. A recent methods paper works explicitly to unify similarity and distance measures for inbred comparison – the field still treats the choice as open. Recall from Chapter 13. that MRD and coancestry are the same statistic: $f_{ij} = 1 - \text{MRD}_{ij}^2$.

One result is worth carrying into practice: in 298 African highland maize inbreds, **91% of line pairs differ at 30-36% of scored alleles, yet only 32% of pairs have kinship below 0.05**. The authors read the gap as allelic redundancy from repeated use of narrow source germplasm.

A pool can look diverse by marker counts while being genealogically shallow. That is a direct argument for a coancestry-based constraint over a distance-based one.

Other calibration benchmarks: North American ex-PVP dent inbreds resolve into 8 admixture subgroups aligned with company of origin; China summer maize (490 inbreds, 525,141 SNPs) shows 66.05% of pairwise kinship values near zero; sweet corn shows systematically lower diversity, PIC and MAF than field corn across 514 versus 181 lines.

18.3 Heterotic groups

18.3.1 Assignment is not a solved problem

GCA-based, SCA-based, combined (HSGCA) and SNP-distance classification are compared repeatedly with no consistent winner. Several studies report HSGCA as most consistent with testcross performance; others find SNP-based and testcross-based classification **disagree**. Notably, SSR markers linked to yield QTL give **no advantage** over random SSRs for grouping – against the intuition that linked markers should be more predictive.

This matters for the pipeline: the line-level collapse assumes pool structure is given. Group assignment is itself contested, so the pool-balance constraint inherits whatever uncertainty the assignment method carries. Treat it as an uncertain input, not a fixed given.

18.3.2 The empirical trend that justifies the constraint

The single most load-bearing empirical result in this corpus: in a commercial European dent maize programme, **differentiation between Stiff Stalk and Non-Stiff Stalk groups increased significantly over time while genetic diversity within both groups declined**. This is longitudinal commercial data, not simulation.

It confirms the constraint is doing necessary work rather than solving a hypothetical problem, and it confirms the two-object framing: what erodes is within-pool diversity; what must be protected from erosion *and* from excessive inflation is between-pool divergence. That is exactly the panel `s1_monitor()` logs in Chapter 13..

Why divergence should be maintained rather than maximised now has a mechanistic basis: additive and dominance QTL effects for hybrid performance are **background-specific**, depending on allelic ancestry. Divergence pays through complementarity at specific loci, not through generic distance – so pushing it further does not add hybrid value and risks removing shared favourable alleles. That is Trap 3 of Chapter 13., with a mechanism attached.

The complication worth knowing: deliberate crossing *across* the heterotic boundary for germplasm enhancement generated substantial usable variation, with **10.8% of hybrids from one sub-population exceeding a high-yielding check**. The boundary is maintained for commercial hybrid-parent pools while being selectively porous for pre-breeding.

18.4 Selecting within the pool

18.4.1 The constraint alone is not sufficient

Optimal contribution and optimal cross selection are established across crops. But a load-bearing qualifier appears repeatedly: diversity depletion in closed elite programmes is described as “**ineluctable**” at realistic population sizes even under active management, with bridging populations argued as the necessary remedy.

A θ constraint on the existing pool cannot indefinitely substitute for new germplasm. Plan periodic external injections rather than treating the line pool as closed.

18.4.2 Cycle length, and what goes unreported

Rapid cycling delivers large gain-rate benefits: 12-88% relative gain advantage for oligogenic architectures at 2-3 cycles per year versus one; cycle time cut from 2-5 years to 1 year in intermediate wheatgrass with an explicit diversity-preserving sampling step. A maize two-cycle genomewide selection scheme reports larger gains at equal cost and time but **no diversity metric at all**.

The diversity cost per unit of cycle-time reduction is not established in the retrieved set. Since shortening cycles is the single most common current intervention in commercial programmes, that is a conspicuous hole.

Simulation work on hybrid rice reciprocal recurrent selection reports that larger population sizes give higher gain *and* less depletion of genetic variance – and that diversity is managed at the parental-pool level that feeds hybrids, which is exactly the line-usage level this book constrains.

18.4.3 What is hybrid-specific, and what is not

This distinction decides which literature transfers. **Every usefulness-criterion and progeny-variance study retrieved operates upstream, at the inbred-line development stage** – cross potential selection, UCPC for multi-parental crosses, barley validation against realised populations, soybean family mean and variance prediction, wheat progeny-variance formulas. None operates on F1 hybrids of fixed inbreds.

That confirms rather than challenges the decision to drop the progeny-variance term: the term is real and actively predicted *before* fixation, and vanishes after it. It is why `usefulness_criterion()` belongs at S2 and nowhere else.

One caution that is genuinely hybrid-specific: hybrid **mean** prediction benefits materially from additive-plus-dominance models, since dominance and SCA change predicted performance versus additive-only GEBVs. That is a separate axis from the diversity constraint – it does not reopen the variance term, but it means line-usage optimisation on θ should not be conflated with mean-performance ranking of the same crosses.

18.5 Genebanks, core collections and introgression

Core-subset construction converges on a shared objective – maximise retained diversity at minimum accession count – but diverges on method: GenoCore selecting 310 from 3,300+ sorghum accessions; distance-based optimisation on cassava; rice cores of 247 and 30 from 2,242 accessions; a lettuce core of 268 from 811 capturing 99.4% of total variation.

The quantitatively strongest result: **marker-based core extraction captures 3-78-fold more haplotypic diversity than geography- or environment-based surrogate sampling** across Arabidopsis, Populus and sorghum panels. If cores are being assembled, assemble them on markers.

Maize-specific core-collection algorithm choice is **not established in the retrieved set** – maize genetic-resource work is oriented toward bridging and targeted introgression instead:

- DNA-pool genotyping of 156 landraces (2,340 individuals) against 327 elite lines spanning the 20th century produced two indicators of a landrace's genomic contribution to elite germplasm, used to rank donors rather than introduce material blind.
- Haplotype mining in ~1,000 landrace-derived doubled haploids found landrace haplotypes **outperforming elite alternatives at the same loci**, stable across populations and environments.
- Rapid-cycling genomic selection *within* a landrace population (899 DH lines, 600K array) improves donor material before it contacts the elite pool.
- For capturing landrace variation, the **doubled-haploid route beats gamete capture** – higher prediction accuracy, and no masking of landrace variation behind a capture parent's background.

A θ - or N_s -based accounting of introgression cost against a same-size resampling baseline is not established in the retrieved set.

18.6 The animal-breeding corpus: which matrix to manage

On the animal side there is broad convergence that some form of optimal contribution or optimal cross selection on a **marker-based** coancestry matrix is the right optimisation *framework*. There is **no consensus on which specific genomic matrix** to place inside it, and multiple post-2018 papers report that the choice changes allele-frequency trajectories and diversity outcomes:

- IBD versus homozygosity- versus drift-based matrices give substantially different results unless management is IBD-based (Chapter 9).
- A deviation-based coancestry matrix retained more genetic variability, while VanRaden's kept allele frequencies closer to the base – a genuine trade-off between “more heterozygosity” and “less frequency drift”.
- A linkage-and-pedigree-informed matrix had the highest correlation with true IBD, the most accurate ΔF , and was the only scheme keeping frequency drift down while meeting the inbreeding target.
- ROH-based genomic relationship outperformed SNP-based relationship for limiting inbreeding accumulation at equal gain, in pigs.

The consensus, such as it is: molecular or IBD-flavoured coancestry outperforms simple VanRaden-style relationship for diversity management, but the field has not converged on a single preferred matrix. As Chapter 12. notes, the line-level collapse holds for **any** symmetric $\mathbf{f}^L$, so this open question does not destabilise the computational results.

On allelic richness the post-2018 work sharpens rather than contradicts the recommendation: allelic variation is far more sensitive to bottlenecks than heterozygosity, there is a linear relationship between census-size reduction and reduction in expected number of alleles at drift-mutation equilibrium, and standard N_e corrections do **not** predict allelic loss well. One paper explicitly recommends expanding the $N_e > 500$ conservation guideline to include an allelic-variation criterion.

And on effective size, two cautions beyond $N_s = 1/(2\theta)$: under *active* optimal molecular management, rate-based N_e is numerically pathological – negative in early generations, with no long-term asymptote – and N_e however estimated is informative about heterozygosity loss but not about allelic loss. Both reinforce reporting only a static status-number descriptor under active management.

18.7 Optimisation machinery reported in plants

This is where the corpus is thinnest, and it confirms the computational gap.

- **Optimal contribution methods** are applied across crops, but no abstract states solver type, runtime, or complexity class.
- **Core-subset search** is greedy or distance-based heuristic; only input and output set sizes are reported, never runtime.
- **Evolutionary algorithms** appear mainly *outside* plants. Within plants the closest analogue is a black-box numerical optimiser for progeny allocation, with no named algorithm, problem size, or convergence statistic.
- Problem sizes are reported as candidate-cross counts (330,078 possible barley parent combinations) or collection sizes, never as optimiser performance.

No retrieved plant study applies differential evolution – or any metaheuristic – to an OCS-type hybrid selection problem, and none reports a runtime or complexity comparison against quadratic-programming OCS solvers.

On the animal side, AlphaMate remains the most cited applied tool. The most directly relevant comparison benchmarked genetic algorithms, differential evolution, particle swarm and simulated annealing for parent selection in a 583-line wheat dataset: GA gave the best convergence speed and fitness; **DE was robust but required longer runtime and careful tuning**. That is consistent with Chapter 12.'s finding that the encoding, not the algorithm, is what decides the outcome.

18.8 What remains unestablished

Stated strictly, from the screened abstracts. These are the honest boundaries of what this book can claim support for. The distinction that matters here is between *unestablished in the literature* and *unaddressed in this book*: two of these items now have a simulated answer in Chapter 10. and Chapter 19., which narrows what the book leaves open without changing what the corpus establishes. A simulation under stated assumptions is not evidence from an operating programme, and the items below are still the honest boundary of the second.

- Application of DE or any metaheuristic to a θ -constrained line-usage selection problem in maize.
- Runtime or complexity comparison between OCS quadratic solvers and metaheuristics for any plant breeding problem.
- Empirical proportional diversity loss measured against a same-size resampling baseline in an operating maize hybrid programme. (Chapter 19. simulates the multi-cycle consequence and finds a gain crossover; no published maize programme reports the measurement itself.)
- Longitudinal joint reporting of allelic richness alongside θ or N_s for the same hybrid maize programme.
- Hybrid-genome-level (rather than parental-pool-level) coancestry optimisation for F1 maize.
- Whether a dominance-informed term would change optimal-cross outcomes relative to the mean-only approach. (Chapter 10. derives the term and shows expected heterosis is invariant to F_{ST} under neutral divergence, so a dominance-informed ranking must differ from a frequency-based one. Whether it changes *realised* outcomes in a programme is untested either way.)
- Quantitative diversity cost per unit of cycle-time reduction under rapid-cycling hybrid maize.
- ROH-based diversity metrics applied to maize hybrid pools (22 records use ROH, only 5 in maize, none for hybrid diversity management).
- Per-evaluation computational cost of diversity metrics, as a function of population size or metric type, anywhere in either corpus.
- Diversity metrics defined on hybrids rather than on lines or parents.
- Combinatorial equal-contribution (0/1) selection compared directly against continuous-contribution designs.

18.9 Consequences for the pipeline

1. Keep the narrow metric set inside the optimiser and report the wide package outside it.
2. Plan for periodic germplasm injection. A closed elite pool depletes regardless of how well the within-pool constraint is managed.
3. Prefer bridging over direct introduction for exotic material, and rank donors by their genomic contribution to the elite pool rather than by raw diversity.
4. Treat heterotic-group assignment as an uncertain input – the pool-balance constraint is only as good as the assignment behind it.
5. Constrain between-pool divergence within a band rather than maximising it.

6. Do not conflate the diversity constraint with mean prediction. Hybrid means need additive-plus-dominance models; the diversity constraint needs neither.
7. Report the diversity cost per cycle explicitly when shortening cycles.
8. If assembling a core or donor set, do it on markers rather than passport data.
9. Watch for allelic redundancy: high marker-based distance can coexist with shallow genealogy, so a coancestry constraint is safer than a distance-based one.
10. Validate on real genotypes before treating the redundancy result as settled.

19. Diversity across cycles

Chapter 21. will end with a list of things its numbers cannot tell you, and the last item on it is the one that matters most: *one cycle, not a trajectory. A one-cycle constraint is not a multi-cycle policy.*

Every result in this book so far has been a single round of hybrid selection. But nobody pays an index cost for diversity in order to have more diversity. They pay it because of what next cycle can do with it, and that argument has been taken on faith for eleven chapters. This chapter runs the programme forward and checks.

19.1 The cycle

Selection happens at the hybrid level, on $A \times B$ crosses. Recycling happens *within* pool. That asymmetry is the whole design, and getting it backwards would quietly destroy the thing Chapter 10. says the programme is selling: a doubled haploid derived from an $A \times B$ hybrid is a blend of both pools, and a programme that recycled hybrids would merge its heterotic groups within a few cycles.

What crosses the pool boundary is *information* – which lines proved themselves in hybrid combination – and never germplasm.




Figure 19.1: One cycle. Hybrids are evaluated across the pool boundary; lines are recycled inside it. The usage weights are the only channel between the two, and they are what `line_weights()` returns.

`run_cycles()` is the loop. The assumptions it runs under, stated once:

C1. Loci are unlinked. Every marker segregates independently, so no linkage disequilibrium builds up and there is no Bulmer effect on the genic side. This is the same abstraction Chapter 4. already makes; Chapter 5. is where linkage is modelled properly, and `build_map()/meiosis()` are there if this ever has to change.

C2. Recycling is by doubled haploid. Two parents are drawn within pool with probability proportional to usage, and one DH is derived per pair. Because the lines are inbred, $(GL[i,] + GL[j,]) / 2$ lies in $\{0, 0.5, 1\}$ and a single Bernoulli draw per locus is an *exact* gamete, not an approximation.

C3. Pool sizes are constant. 25 lines per pool, every cycle. A real programme varies this, and the variation is itself a diversity decision.

C4. Marker effects are fixed. The same β throughout, so the index is comparable across cycles and gain accumulates on the founding cycle's scale. No new mutation, no changing selection target, no genotype-by-environment.

C5. The base p_0 is frozen at cycle 0 and carried forward. This is the subject of the next section.

C6. The pool is closed unless a germplasm injection is scheduled.

C7. The index is known without error. Chapter 12. Finding 5 measures what relaxing this does within a cycle; the effect compounds across cycles and is not modelled here. Every gain figure below is therefore an optimistic ceiling, for all budgets equally.

C8. Selection is on the hybrid only. No per-se line selection, no stage structure. Chapter 13. is where the multi-stage version lives.

19.2 Freezing p_0 , and what happens if you do not

`build_ctx()` sets `p0 <- colMeans(X)` on every call. Run it fresh each cycle and the drift lens is re-anchored to the population it is supposed to be measuring drift *away from*, which reports approximately zero loss forever while the pool erodes underneath.

The demonstration needs two generations and nothing else:

```
sim <- simulate_pools(n_A = 25, n_B = 25, m = 2000, seed = 1)
p0 <- colMeans(sim$X) # the founding base, frozen here
GL5 <- sim$GL
set.seed(7)
for (t in 1:5) { # five generations of drift
  GL5 <- rbind(next_pool(GL5[1:25, ], rep(1, 25), 25),
               next_pool(GL5[26:50, ], rep(1, 25), 25))
}
ped5 <- expand.grid(a = 1:25, b = 25 + 1:25)
X5 <- (GL5[ped5$a, ] + GL5[ped5$b, ]) / 2

drift <- function(X, base) {
  p <- colMeans(X); k <- base > 1e-6 & base < 1 - 1e-6
  mean((p[k] - base[k])^2 / (base[k] * (1 - base[k])))
}
c(frozen_p0 = drift(X5, p0),
  reanchored = drift(X5, colMeans(X5)))
#> frozen_p0 reanchored
#> 0.08693168 0.00000000
```

The second number is zero by construction: a population is never drifted from itself. Five generations of real drift have been reported as none.

! This is the most common way to get a multi-cycle report wrong

Chapter 9. named it as the single most frequent error in applying `F_hom` and `F_drift`, and Chapter 21. repeats it in the operational checklist. `run_cycles()` therefore takes `p0` once and overwrites it on every rebuilt context. It is one assignment, and every drift number in this chapter depends on it.

19.3 The trajectory

Three budgets, fifteen cycles, five seeds each. `alpha_max = Inf` is the unconstrained programme: it takes the best hybrids it can find and accepts whatever that costs.




Figure 19.2: Mean index of the selected hybrids, in founding-cycle standard deviations, averaged over five seeds. The unconstrained programme leads for eight cycles and then stops climbing. Neither constrained programme has stopped by cycle 15.

cycle	$\alpha \leq 0.005$	$\alpha \leq 0.02$	unconstrained
0	1.16	1.45	1.73
3	4.05	4.38	5.13
5	5.47	6.05	6.81
8	7.40	8.44	8.42
10	8.66	9.68	9.00
12	9.89	10.72	9.36
15	11.14	11.88	9.55

The unconstrained run wins the first eight cycles and loses every cycle after that. By cycle 15 the $\alpha \leq 0.02$ programme is 2.33 standard deviations ahead – 24% more gain, from a programme that gave up 17% of its gain in year one.

The mechanism is visible in the diversity panel:




Figure 19.3: Left: gene diversity of the selection. Right: F_{drift} against the frozen founding base. The unconstrained programme spends nearly its whole diversity endowment by cycle 8 and has little left to convert afterwards.

The unconstrained programme is not out-selected. It is out of material. Its F_{drift} reaches 0.41 by cycle 8 and only 0.457 by cycle 15 – it has converted almost everything it had into gain, early, and then has nothing to convert.

19.3.1 It is not an artefact of the cheap selector

The trajectories above use a greedy sweep, which is what makes fifteen cycles by five seeds by three budgets affordable. Chapter 16's differential evolution finds slightly better plans per cycle. Re-running one seed with it:

cycle	0.020 de	0.020 greedy	unconstrained de	unconstrained greedy
0	1.39	1.37	1.53	1.53
4	5.53	5.98	5.22	6.60
8	8.44	9.56	7.40	9.08
12	10.47	11.73	7.99	9.84

The crossover reproduces. Under DE the constrained programme is ahead by 2.48 at cycle 12, against 1.89 under greedy. The cheap selector understates the case rather than manufacturing it.

There is a second thing in that table worth not glossing over: DE reaches a *lower* absolute index than greedy by cycle 12 in both columns, despite finding a better plan in every individual cycle. It uses more of the budget each time – $\alpha = 0.0200$ against greedy's 0.0177 at cycle 0 – and spends the difference out of the next cycle's inheritance. **Per-cycle optimality is not multi-cycle optimality.** This is one seed and should be read as a flag rather than a result, but it is the same trade the whole chapter is about, one level up.

19.4 The rate of inbreeding, and what N_s is not

Chapter 21.'s checklist says to regress $\log(1 - F)$ on cycle and check that it is linear, because curvature means a non-constant rate. Doing it:

```
rate <- do.call(rbind, lapply(unique(traj$budget), function(b) {
  d <- aggregate(F_drift ~ cycle, traj[traj$budget == b, ], mean)
  f <- lm(log(1 - F_drift) ~ cycle, d[-1, ])
  data.frame(budget = b, dF_per_cycle = -coef(f)[["cycle"]],
             Ne = 1 / (-2 * coef(f)[["cycle"]]), R2 = summary(f)$r.squared)
}))
fmt(rate, 4)
```

budget	dF_per_cycle	Ne	R2
unconstrained	0.0292	17.1146	0.8534
0.020	0.0358	13.9841	0.9931
0.005	0.0265	18.8898	0.9976

The diagnostic fires, and it fires on the unconstrained run: $R^2 = 0.853$ against 0.998 for the tightest budget. The unconstrained rate is not constant because it is decelerating – there is progressively less left to lose. Read naively its fitted ΔF is the *lowest* of the three, which is exactly backwards; the fit is meaningless for a saturating series, and the R^2 column is what says so.

⚠ N_s does not project this

Section 11.3 insists that the status number is a static descriptor and not a drift projection. Here is the size of the discrepancy: N_s computed on the selected hybrids averages 0.67 across all of these runs, while the realised effective size implied by the drift they actually accumulated is between 14 and 19.

They are not the same quantity and are not off by a constant. Report N_s as a descriptor of a plan, and measure ΔF from the trajectory.

19.5 Choosing a budget for a horizon

Every seed crosses over, and the cycle at which it happens is the useful number:

seed	alpha <= 0.02	alpha <= 0.005
1	8	11
2	12	15
3	10	12
4	7	10
5	5	9

The tighter budget always crosses later and always ends higher. That is the whole trade, and it makes the choice of $\alpha_{\max}$ a function of the planning horizon rather than a matter of taste:

Planning horizon	Budget	Why
3 cycles or fewer	do not constrain	the crossover has not happened yet; the constraint is pure cost
about 5 to 8 cycles	alpha ~ 0.02	the looser budget has crossed and the tighter one has not
10 cycles or more	alpha ~ 0.005 to 0.02	both have crossed; the tighter budget is still climbing

Planning zon	hori-	Budget	Why
indefinite		constrain, and in- ject	no closed budget holds forever – see below

The caveat matters more than the numbers. `n_sel = 60` out of 625 candidates from 25 + 25 lines is a small, intensely selected programme, and the crossover arrives sooner the harder the selection is. What transfers is the shape: an unconstrained programme leads early, saturates, and is passed.

19.6 The per-line cap, at last measurable

Chapter 16. carries `max_use` through every selector, defaulting to `NULL` everywhere, because within a single cycle a usage cap only costs index. Across cycles it is buying something, and this is where it can be seen. Same budget ($\alpha \leq 0.02$), same seeds, three caps:

max_use	cycle	index	GD	F_drift	max_line
3	0	0.491	0.329	0.005	3.0
6	0	1.305	0.323	0.025	6.0
none	0	1.446	0.324	0.030	12.4
3	5	2.423	0.296	0.104	3.0
6	5	5.788	0.270	0.194	6.0
none	5	6.047	0.272	0.210	15.4
3	10	4.058	0.268	0.185	3.0
6	10	9.216	0.233	0.316	6.0
none	10	9.677	0.233	0.348	15.4
3	15	5.532	0.247	0.250	3.0
6	15	11.268	0.210	0.395	6.0
none	15	11.877	0.207	0.431	16.4

`max_use = 6` costs 5% of cycle-15 gain and holds F_{drift} to 0.395 against 0.431, while the uncapped programme leans on one line 16 times in its final cycle. `max_use = 3` is too tight: it halves the gain and is not compensated.

Note *which* lens the cap moves. Gene diversity at cycle 15 is essentially the same with the cap and without it (0.21 against 0.207), because α is already pinning that. The difference shows up in F_{drift} , which is the lens Chapter 9. established the constraint does *not* pin. The cap is a handle on the uncontrolled lens, which is exactly why it is worth having in addition to the budget and not instead of it.

This changes no default. `max_use` remains `NULL` throughout the package.

19.7 Injection, and what it costs

No closed budget holds indefinitely: all three trajectories are still losing diversity at cycle 15, only at different rates. Chapter 18. and Chapter 21. both prescribe periodic germplasm injection, and Chapter 6. is the machinery that performs it – introgression is how new material enters an elite pool without dragging the rest of the genome with it.

Replacing six of the twenty-five lines in each pool with founding-generation material at cycle 8:

regime	cycle	index	GD	F_drift
closed	7	7.697	0.255	0.270
injected	7	7.697	0.255	0.270
closed	8	8.440	0.246	0.301

regime	cycle	index	GD	F_drift
injected	8	6.469	0.283	0.182
closed	10	9.677	0.233	0.348
injected	10	7.865	0.270	0.233
closed	12	10.722	0.222	0.385
injected	12	9.391	0.254	0.291
closed	15	11.877	0.207	0.431
injected	15	11.283	0.234	0.364

The cost is immediate and the return is slow. Gene diversity jumps from 0.246 to 0.283 and F_{drift} falls by 40%, but the index drops 1.97 standard deviations on the spot, and **has not repaid by cycle 15** (11.28 against 11.88).

That is the honest result and it should not be dressed up. Within this horizon injection buys a restored resource and not more gain; it is the same trade as the diversity budget itself, at a longer wavelength and a steeper entry price. The run that would repay it is longer than fifteen cycles, and the founders used here are the programme's own – unadapted material from a genebank would cost considerably more.

19.8 What the corpus supports

Appendix D.'s screened corpora hold 687 plant and 324 animal records, of which a substantial block concerns recurrent selection and long-term gain: two-part and rapid-cycling programme designs, optimal cross selection evaluated over cycles rather than within one, and simulation studies of the long-term effects of genomic selection on both gain and variance.

The qualitative shape found here – early advantage to unconstrained selection, saturation, then a crossover – is the standard finding in that literature. [G. Gorjanc, R. C. Gaynor, and J. M. Hickey \[46\]](#) optimise cross selection explicitly for long-term gain in two-part programmes and report the same tension between per-cycle and multi-cycle optimality that the DE comparison above ran into. [A. Allier, C. Lehermeier, A. Charcosset, L. Moreau, and S. Teyssèdre \[47\]](#) separate the short-term and long-term components directly. [Y. C. J. Wientjes, P. Bijma, M. P. L. Calus, B. J. Zwaan, Z. G. Vitezica, and J. van den Heuvel \[48\]](#) track what genomic selection does to additive variance over many generations, which is the mechanism behind the saturation seen here, and [A.-C. Doublet, P. Croiseau, S. Fritz, A. Michenet, C. Hozé, and C. Danchin-Burge \[49\]](#) measure the diversity cost in three dairy cattle breeds where the cycles have actually been run. [J. E. Rutkoski \[50\]](#) is the counterpart caution: realised rates of gain are harder to estimate from an operating programme than a simulation makes them look. [D. Müller, P. Schopp, and A. E. Melchinger \[51\]](#) is the closest thing in the corpus to the selection criterion used here, chosen for its effect on later cycles rather than the current one. What Chapter 18. listed as unestablished was something narrower and remains so: nobody has published the crossover *for a θ -constrained hybrid maize programme measured against a same-size resampling baseline*. This chapter demonstrates it in simulation, under C1-C8, which is not the same as establishing it in an operating programme.

19.9 Where this connects

Chapter 10. argued that what builds a heterotic pattern is divergence aligned with dominance, and that selection on cross performance is the only process in the programme that searches for it. This chapter is that process running. The GD_{BS} column of `cycles_trajectory.csv` rises in every run – the pools do separate under reciprocal selection, as they should – while GD within them falls, which is the trade Chapter 11.'s partition was built to expose.

Chapter 21. is the single cycle, in operational detail. Read it as cycle 0 of this chapter, and read its final checklist – freeze p_0 , keep the per-stage ledger, check $\log(1 - F)$ for curvature, plan germplasm injection – as the four things this chapter has now shown the consequences of skipping.

20. The limit, and what the metrics cannot see

Chapter 19. ran the programme forward and found the crossover: a constrained programme trails an unconstrained one for about eight cycles and leads thereafter. That is a measurement. This chapter asks the two questions a measurement leaves open – could it have been *predicted*, and is the quantity being managed the quantity that actually matters.

The answers are yes-with-a-condition, and no-not-entirely.

20.1 Robertson's limit

Chapter 3. wrote down the response to one cycle of selection. Robertson (1960) wrote down the sum of all of them. Selection fixes alleles; a finite population fixes some of the favourable ones in the wrong direction; and the total advance before everything is fixed is approximately

$$R_{\text{total}} \approx 2N_e R_1 \quad (20.1)$$

with a half-life of about $1.4N_e$ generations. The result is the theoretical ancestor of every optimal-contribution scheme in this book: it is *why* N_e is the thing to manage, and it says exactly what a unit of N_e is worth in units of cumulative gain.

Both inputs are already in Chapter 19.' trajectory. R_1 is the first cycle's response, and N_e comes from the rate of inbreeding, exactly as Chapter 19. computed it:

```
rob <- do.call(rbind, lapply(unique(traj$budget), function(b) {
  d <- aggregate(cbind(index, F_drift) ~ cycle, traj[traj$budget == b, ], mean)
  f <- lm(log(1 - F_drift) ~ cycle, d[-1, ])
  Ne <- 1 / (-2 * coef(f)[["cycle"]])
  R1 <- d$index[2] - d$index[1]
  data.frame(budget = b, R2 = summary(f)$r.squared, Ne = Ne, R1 = R1,
    limit = 2 * Ne * R1, half_life = 1.4 * Ne,
    observed_15 = d$index[16] - d$index[1],
    predicted_15 = 2 * Ne * R1 * (1 - (1 - 1 / (2 * Ne))^15))
}))
fmt(rob, 3)
```

budget	R2	Ne	R1	limit	half_life	observed_15	predicted_15
unconstrained	0.853	17.115	1.277	43.697	23.960	7.812	15.688
0.020	0.993	13.984	1.222	34.167	19.578	10.431	14.378
0.005	0.998	18.890	0.802	30.281	26.446	9.980	10.032

Read the last two columns against each other, and then read both against the first.

20.2 The R^2 column decides whether the theory applies

For the tightest budget the prediction is 10.03 and the observation is 9.98. Robertson's theory, published in 1960 and derived for a single locus under idealised drift, predicts fifteen cycles of a genomic hybrid programme to better than one percent.

For the unconstrained programme it predicts 15.69 against an observed 7.81 – too high by about half.

The theory did not fail arbitrarily, and Chapter 19. already said why. Robertson assumes a *constant* rate of inbreeding, and the regression of $\log(1 - F)$ on cycle is where that assumption is testable. Its R^2 is 0.998 for the tightest budget and 0.853 for the unconstrained run, whose drift decelerates because there is

progressively less left to lose. That low R^2 is not a nuisance. It is the theory declining to apply, in advance, and in a column you can read before you trust the number.

! Do not quote a Robertson limit without its R^2

N_e here is a *fitted* quantity, and Chapter 19. showed the fit is meaningless for a saturating series – its fitted ΔF comes out lowest for the run that is actually losing diversity fastest, which is backwards. Feeding that N_e into $2N_e R_1$ propagates the error into a number that looks like a long-term plan.

Report the R^2 next to every limit, and treat a limit computed from a fit below about 0.99 as an illustration rather than a projection.

That yields an argument for the constraint that this book did not previously have. Chapter 19. justified α by where the trajectories cross. The table above adds: the constrained programme is the one whose future can be *computed*. An unconstrained programme is not merely worse after cycle 8, it is unforecastable, because the theory that would forecast it needs a constant ΔF that it does not have.

20.3 What the metrics cannot see

Now the second question. Every diversity metric in this book – θ , GD , N_s , allelic richness, F_{hom} , F_{drift} – is a function of allele frequencies **one locus at a time**. Chapter 3. showed that additive variance is not:

$$\text{Var}(\mathbf{X}\beta) = V_{\text{genic}} + 2 \sum_{k < l} \beta_k \beta_l \text{Cov}(x_k, x_l) \quad (20.2)$$

Directional selection drives that covariance term negative (Bulmer (1971)): it couples loci with like effects in repulsion, hiding genic variance from the next generation without moving a single allele frequency. No metric in the catalogue can detect it.

Chapter 19. assumed the effect away, in assumption **C1**: *loci are unlinked, so no linkage disequilibrium builds up and there is no Bulmer effect on the genic side*. That is worth testing rather than assuming, because Bulmer's 1971 treatment covers unlinked loci explicitly – free recombination halves the disequilibrium each generation, but selection regenerates it every generation, and the equilibrium is not zero.

20.3.1 The experiment

The gap $\text{Var}(\mathbf{X}\beta) - V_{\text{genic}}$ is **not** on its own a measure of the Bulmer effect. A finite founder base leaves disequilibrium in the very first cycle, and two divergent pools add more. So the comparison is against a control: an arm identical in every respect – same founders, same seeds, same recycling, same population sizes – in which the hybrids are chosen at random instead of on the index.

```
gapc <- aggregate(gap ~ regime + cycle, vr, mean)
w <- reshape(gapc, idvar = "cycle", timevar = "regime", direction = "wide")
names(w) <- sub("gap.", "", names(w))
w$selection_effect <- w$unconstrained - w$random
fmt(head(w[, c("cycle", "random", "unconstrained", "constrained",
               "selection_effect")], 7), 4)
```

	cycle	random	unconstrained	constrained	selection_effect
1	0	-0.0261	-0.0261	-0.0261	0.0000
4	1	-0.0716	-0.2686	-0.1984	-0.1970
7	2	-0.0328	-0.2167	-0.1930	-0.1839
10	3	-0.0320	-0.2616	-0.0264	-0.2296
13	4	0.0935	-0.0931	-0.0976	-0.1866
16	5	-0.0453	-0.1069	-0.1171	-0.0615
19	6	0.0512	-0.0711	-0.1130	-0.1223

Cycle 0 is identical across all three arms, to the digit. Nothing has been selected yet, so the control is doing its job.

From cycle 1 the arms separate. Paired within seed and cycle:

```
p <- reshape(aggregate(gap ~ regime + seed + cycle, vr, mean),
                idvar = c("seed", "cycle"), timevar = "regime", direction = "wide")
p <- p[p$cycle > 0, ]
dd <- p$gap.unconstrained - p$gap.random
c(mean_effect = mean(dd), se = sd(dd) / sqrt(length(dd)),
  t = mean(dd) / (sd(dd) / sqrt(length(dd))), n_pairs = length(dd))
#> mean_effect      se      t    n_pairs
#> -0.08568088  0.02108980 -4.06266819 96.00000000
```

⚠ Assumption C1 of Chapter 19. is too strong, and is restated here

Selection builds negative disequilibrium among **unlinked** loci, at $t = -4.06$ over 96 paired observations. Unlinked loci do not prevent the Bulmer effect; they make it decay fast.

What C1 should say is that with free recombination the disequilibrium is halved every generation, so it reaches a low equilibrium and does not accumulate across cycles the way frequency change does. That is what makes Chapter 19. *frequency* trajectories trustworthy, and it is the correct reason – not the absence of the effect.

The effect is largest in the early cycles and decays toward zero late, which is the expected shape: it is proportional to the genic variance still available to be hidden, and by cycle 12 there is much less of it.

20.3.2 With real linkage

The same question with real meiosis behind it. Chapter 7. *reciprocal recurrent selection* runs on a coalescent founder history with a genetic map, and AlphaSimR reports `varA` (which carries the disequilibrium) and `genicVarA` (which does not), so this arm needs no variance code of its own.

```
lk$rel <- ((lk$varA - lk$genicVarA) + (lk$varA_B - lk$genicVarA_B)) /
          (lk$genicVarA_A + lk$genicVarA_B)
pl <- reshape(lk[, c("regime", "seed", "cycle", "rel")],
                idvar = c("seed", "cycle"), timevar = "regime", direction = "wide")
pl <- pl[pl$cycle > 0, ]
ddl <- pl$rel.selected - pl$rel.random
c(mean_effect = mean(ddl), se = sd(ddl) / sqrt(length(ddl)),
  t = mean(ddl) / (sd(ddl) / sqrt(length(ddl))), n_pairs = length(ddl))
#> mean_effect      se      t    n_pairs
#> -0.16493928  0.04672916 -3.52968674 32.00000000
```

Same sign, larger in magnitude, and measured with fewer replicates so the evidence is weaker. Linkage does what linkage should: it slows the decay, so the disequilibrium selection creates persists longer than it does under free recombination.

i A reproducibility trap in the linked arm

`alphasimr_founder_pools()` is **not** reproducible from `cfg$seed`. `AlphaSimR::runMacs2()` hands the coalescent to MaCS, which seeds itself outside R's RNG, so two calls with the same configuration return different founder haplotypes. Everything downstream of the founders is seeded normally.

The consequence for an experiment like this one is severe: founding the pools separately for each arm would have compared two different base populations and charged the difference to selection. `book/scripts/regenerate.R` founds once per replicate and branches both arms off that single result, which is why cycle 0 of the table above matches exactly.

20.4 The two costs of selection

The chapter's result stated once.

Selection pays for gain twice. It builds negative covariance between loci, which is a real loss of usable variance, and **recombination repays it** – halved every generation under free recombination, more slowly under linkage. And it changes allele frequencies, moving them toward fixation, and **nothing repays that** short of mutation or migration.

The first cost is invisible to every metric in this book and is largely self-correcting. The second is what every metric in this book measures, and is permanent.

That is the justification for constraining the second and ignoring the first, and it is a stronger statement than “the constraint is what we can compute”. It is worth being explicit that this ordering is a *consequence* of the two decay rates, not an assumption: if the disequilibrium did not decay – tight linkage, few chromosomes, short cycle times – the ranking would change, and a constraint written on allele frequencies would be managing the smaller of the two bills.

20.5 Where this connects

Chapter 3. supplies the variance decomposition and the breeder’s equation this chapter runs forward. Chapter 19. supplies the trajectories, and has one assumption restated as a result of the test above. Chapter 7. supplies the linked arm, which is the only place in the book where a genetic map, rather than an assumption about one, decides how fast disequilibrium decays.

Chapter 21.’s checklist gains one line from this chapter: report the R^2 of the $\log(1 - F)$ regression next to any projection built on N_e , and treat a projection from a poor fit as an illustration.

21. End to end, and the operational checklist

Everything the book has built, run once, on one dataset, in the order a programme would run it. Then the checklist.

21.1 The whole pipeline in one script

```
st1 <- stage1_simulate(book_cfg)
c(hybrids = nrow(st1$X), markers = ncol(st1$X),
  f_min = min(st1$f), f_max = max(st1$f))
#>   hybrids   markers    f_min    f_max
#> 625.00000 2000.00000   0.64475   0.83175
```

```
n_sel <- round(0.10 * nrow(st1$X))
res <- suppressWarnings(stage2_select(
  X = st1$X, f = st1$f, hybrids = st1$hybrids,
  n_sel = n_sel,
  alphas = NULL,      # NULL = automatic grid, 0 to the attainable ceiling
  weights = NULL,     # NULL = equal weight across traits
  fL = st1$fL,        # optional: enables theta_A / theta_B / theta_AB
  B_null = 500, B_full = 80, NP = 100, itermax = 200, verbose = FALSE))
```

i The same pipeline as a script

This chapter runs at the book's reduced scale so it renders in seconds. The identical pipeline at the **production** scale of `sim_config`, writing its tables to `report/`, is `inst/scripts/run_all.R`:

```
Rscript inst/scripts/run_all.R          # ~4 min
Rscript inst/scripts/run_benchmark.R    # ~6 min; the numbers quoted in README.md
```

With real data only the first line changes:

```
st1 <- stage1_build(X = my_hybrid_matrix, ped = my_pedigree,
  traits = my_predicted_traits, fL = my_line_coancestry)
```

21.1.1 The reference values

```
unlist(res$ref)
#>      gd_ref      gd_se      alpha_max      n_sel population_bias
#> 3.259845e-01 3.556567e-05 3.938109e-02 6.200000e+01 6.467346e-03
```

`alpha_max` is the attainable ceiling from Section 11.9: the loss pure truncation would incur. Any requested budget above it constrains nothing, and `stage2_select()` says so rather than returning a solution that merely looks constrained.

21.1.2 Metrics per scenario

```
fmt(res$metrics[, c("scenario", "alpha", "index", "Ns", "Ne_parents",
                    "Ne_lines_A", "Ne_lines_B", "max_line", "alleles_lost")], 4)
```

scenario	alpha	in- dex	Ns	Ne_parents	Ne_lines_A	Ne_lines_B	max_line	alleles_lost
DE_alpha_0.00	0.0000	1.1004	0.7418	27.4571	12.4805	15.2540	10	63
DE_alpha_0.98	0.0098	1.4141	0.7383	22.3488	9.1962	14.2370	14	81
DE_alpha_1.97	0.0195	1.6767	0.7349	18.0894	7.3359	11.7914	16	106
DE_alpha_2.95	0.0295	1.8523	0.7314	15.8843	5.9505	11.9379	20	116
DE_alpha_3.94	0.0357	1.9600	0.7292	14.3701	5.2948	11.1744	21	125
ref_truncation	0.0394	1.9701	0.7280	13.4877	4.6878	12.0125	23	124
ref_trunc_cap	0.0028	0.7769	0.7408	31.6379	15.7541	15.8843	4	47
ref_max_div	-0.0131	0.1100	0.7465	36.9615	18.6602	18.3048	5	36
ref_random	0.0011	0.0682	0.7414	37.3204	17.9626	19.4141	6	31

Read the `index` column against `alpha`. That is the price list: what each percentage point of diversity costs in index units, on this dataset, at this selection intensity.

Now read `Ne_parents` against `Ns` in the same rows. This is the diagnostic Chapter 9. called for, and the reason `Ne_parents` is a second constraint rather than a report line:

```
d <- res$metrics
fmt(data.frame(scenario = d$scenario,
               Ns = d$Ns,
               Ne_parents = d$Ne_parents,
               Ns_relative = d$Ns / max(d$Ns),
               Ne_par_relative = d$Ne_parents / max(d$Ne_parents)), 3)
```

scenario	Ns	Ne_parents	Ns_relative	Ne_par_relative
DE_alpha_0.00	0.742	27.457	0.994	0.736
DE_alpha_0.98	0.738	22.349	0.989	0.599
DE_alpha_1.97	0.735	18.089	0.984	0.485
DE_alpha_2.95	0.731	15.884	0.980	0.426
DE_alpha_3.94	0.729	14.370	0.977	0.385
ref_truncation	0.728	13.488	0.975	0.361
ref_trunc_cap	0.741	31.638	0.992	0.848
ref_max_div	0.747	36.962	1.000	0.990
ref_random	0.741	37.320	0.993	1.000

The status number barely moves across scenarios while the effective number of parents collapses. A programme reporting only N_s would see a healthy pool and a shrinking parent base it never measured.

21.1.3 Both lenses, and the pool partition

```
fmt(res$metrics[, c("scenario", "alpha", "F_hom", "F_drift", "cov_diag",
                    "rare_retained", "GD_WI_hyb", "GD_BS", "F_ST")], 4)
```

scenario	alpha	F_hom	F_drift	cov_diag	rare_retained	GD_WI_hyb	GD_BS	F_ST
DE_alpha_0.00	0.0000	0.0038	0.0161	-0.0123	0.8636	0.1834	0.0393	0.1205
DE_alpha_0.98	0.0098	0.0132	0.0248	-0.0116	0.8312	0.1829	0.0417	0.1293
DE_alpha_1.97	0.0195	0.0224	0.0348	-0.0124	0.7727	0.1828	0.0452	0.1414
DE_alpha_2.95	0.0295	0.0352	0.0422	-0.0069	0.7143	0.1818	0.0471	0.1490
DE_alpha_3.94	0.0357	0.0377	0.0490	-0.0113	0.7143	0.1821	0.0499	0.1586
ref_truncation	0.0394	0.0412	0.0534	-0.0121	0.7338	0.1822	0.0514	0.1640
ref_trunc_cap	0.0028	0.0154	0.0114	0.0039	0.7273	0.1816	0.0370	0.1138
ref_max_div	-0.0131	-0.0124	0.0063	-0.0187	0.9481	0.1853	0.0362	0.1096
ref_random	0.0011	0.0081	0.0066	0.0015	0.9481	0.1805	0.0357	0.1097

`alpha` constrains the first lens only. `F_drift` is the unmonitored bill, and `cov_diag` is exactly the gap between them – verify it:

```
c(max_deviation_from_Eq3 =
  max(abs(d$cov_diag - (d$F_hom - d$F_drift)), na.rm = TRUE))
#> max_deviation_from_Eq3
#> 6.678685e-17
```

`GD_BS` rising across scenarios is selection concentrating on the most complementary line pairs, which pushes the pools further apart. That is not automatically bad – it is heterosis – but Chapter 13. is the reason to watch it rather than celebrate it.

21.1.4 Distance from the random null

```
fmt(res$z_scores[, c("scenario", "z_GD", "z_Ne_parents", "z_F_drift",
  "z_rare_retained", "z_GD_BS", "z_ANE")], 2)
```

scenario	z_GD	z_Ne_parents	z_F_drift	z_rare_retained	z_GD_BS	z_ANE
DE_alpha_0.00	0.00	-5.63	7.71	-4.27	6.57	6.96
DE_alpha_0.98	-5.20	-8.56	14.78	-5.69	11.05	9.97
DE_alpha_1.97	-10.35	-11.00	22.79	-8.24	17.33	19.86
DE_alpha_2.95	-15.63	-12.26	28.76	-10.80	20.87	17.76
DE_alpha_3.94	-18.88	-13.13	34.27	-10.80	25.81	16.12
ref_truncation	-20.85	-13.64	37.81	-9.94	28.51	9.80
ref_trunc_cap	-1.47	-3.23	3.95	-10.23	2.44	17.27
ref_max_div	6.91	-0.18	-0.16	-0.57	0.96	5.91
ref_random	-0.59	0.03	0.08	-0.57	0.12	-0.43

Every metric in standard deviations from a random selection of the same size. A metric near zero is not distinguishing this plan from chance, and cannot be serving as a constraint.

21.1.5 The selection table

```
scn <- setdiff(names(res$selection), c(names(stl$hybrids), "n_scenarios"))
head(res$selection[order(-res$selection$n_scenarios),
  c("hybrid", "line_A", "line_B", scn, "n_scenarios")], 8)
```





Figure 21.1: The price list. Each point is one diversity budget; the curve is what index it buys.

Choose the knee, not the endpoint. Past the knee, each additional unit of index costs disproportionately more diversity, and the diversity is not recoverable next cycle while the index gain is realised once.

21.3 What this simulation cannot tell you

Stated plainly, because the numbers above are precise and precision invites more confidence than the model supports.

The GEBVs carry no prediction error. The traits are built from *known* marker effects (Chapter 4.). Chapter 12. Finding 5 measures what happens when they do not: the reported index is biased upward, and error *correlated with relatedness* costs several times the diversity that independent error of the same size does. The plans below should be read as an optimistic ceiling.

The model is additive. An additive **G** misses heterosis and specific combining ability, which is precisely what a hybrid programme sells. The diversity constraint does not need a dominance model; the *mean-performance ranking* does (Chapter 10.). Do not conflate them, and do not use these *index* values as performance predictions.

The pools are symmetric by construction. Both are simulated from the same ancestral frequencies at the same F_{ST} . Real heterotic pools differ in size, history and internal structure, and a real programme will find one pool eroding faster than the other – which is why `ne_lines_pool()` is reported per pool.

N_s is not a drift projection. It is a static descriptor (Section 11.3). Do not extrapolate it forward.

One cycle, not a trajectory. Every result here is a single round of hybrid selection. Chapter 13. showed that a constraint applied at one stage buys stage-level diversity, not end-of-pipeline richness – and the same logic applies across cycles. A one-cycle constraint is not a multi-cycle policy: Chapter 19. runs fifteen of them and finds the unconstrained programme ahead for eight cycles and behind for the rest.

p_0 must be frozen. In this book it is set once in `build_ctx()`. In a real multi-cycle programme it must be the founder cycle's frequencies, propagated forward – the single most common error in applying `F_hom` and `F_drift` (Chapter 9.). Chapter 19. shows the failure directly: re-anchored, five generations of real drift are reported as exactly none.

21.4 The operational checklist

21.4.1 Setting up

- The base p_0 is frozen at the founding cycle and documented.
- Diversity is computed on molecular coancestry f , never on a centered G . Check `sum(G) == 0` once, to see why.
- The marker panel's MAF spectrum is known – it decides how severely the two lenses diverge.
- Heterotic-group assignment is recorded, with the method that produced it, and treated as uncertain.

21.4.2 Choosing the budget

- The attainable α ceiling is computed *before* choosing $\alpha_{\max}$. A budget above the ceiling constrains nothing.
- Loss is measured against a **same-size resampling baseline**, never against the full pool.
- The constraint is applied at the target set size, not incrementally.
- The frontier is traced and the knee chosen, rather than a round number.

21.4.3 Running the optimiser

- The search is encoded on **line weights**, not hybrid keys.
- The differential evolution is **warm-started** from greedy solutions.
- $N_{e,\text{par}}$ is a second constraint, not a report line.
- The per-line usage cap is enforced in the decoder, not by penalty.
- The result is checked against the continuous relaxation. A large gap means the optimiser has not converged; a slack constraint means the search never reached the boundary.

21.4.4 Validating the plan

- `F_hom` and `F_drift` are computed **separately** and reported together, with `cov_diag` as the diagnostic.
- Allelic richness is computed on the returned plan – it is the one metric carrying information θ does not.
- The within/between pool decomposition confirms diversity was not preserved by degrading pool separation.
- The number of selected parent lines is biologically plausible.
- N_s is reported; a rate-based N_e is **not**, under active management.
- Every metric is reported against the random null, in standard deviations.

21.4.5 Across cycles

- The per-stage loss ledger is kept. DH induction, not hybrid mining, is usually where the diversity goes.
- The regression of $\log(1 - F)$ on cycle is checked for linearity. Curvature means a non-constant rate.
- Periodic germplasm injection is planned. A closed elite pool depletes however well the constraint is managed.
- The diversity cost per cycle is reported when cycle time is shortened.

21.5 Where to go next

Three extensions the material here supports but does not cover, each needing new work rather than new derivation:

Real data ingestion. `load_data()` is the single swap point and is currently a stub. A worked path from a VCF to the `X/f/hybrids` contract is what most readers will need first, and it is the one extension here that is pure engineering rather than new derivation.

Estimating the dominance deviations. Chapter 10. shows that expected heterosis depends on the alignment between divergence and d , and that no frequency-based statistic can see d . Everything downstream of that – ranking crosses on more than the mean, putting a number on the λ of Chapter 11. – waits on an estimate of d the book does not attempt.

Prediction error compounded across cycles. Chapter 12. Finding 5 measures it within one cycle and Chapter 19. assumes it away (C7). The two have not been run together, and the interaction is not

obviously benign: correlated error concentrates the selection, and Chapter 19. shows concentration is what erodes the programme's future.

Appendices

A. Notation and glossary

A.1 Symbols

Symbol	Meaning	In the code
N	candidate hybrids in the pool	<code>ctx\$N</code>
n, n_{sel}	hybrids selected	<code>n_sel</code>
L	parent lines	<code>ctx\$n_lines</code>
m	markers	<code>ctx\$m</code>
x_{ik}	within-individual reference-allele frequency of individual i at marker k , in 0/0.5/1	<code>ctx\$X[i, k]</code>
$\mathbf{X}$	hybrid marker matrix, $N \times m$	<code>ctx\$X</code>
$\mathbf{M}$	hybrid dosage matrix, 0/1/2	-
$\mathbf{G}^L$	line genotype matrix, $L \times m$, coded 0/1	<code>ctx\$GL</code>
$p_k, \bar{p}_k$	allele frequency; the bar denotes a group mean	-
$p_{0,k}$	frozen candidate-pool allele frequency: the base	<code>ctx\$p0</code>

A.1.1 Matrices

Symbol	Meaning	Note
$\mathbf{f}$	molecular coancestry of the hybrids , $N \times N$, in $[0, 1]$	the diversity kernel; never centered
$\mathbf{f}^L$	molecular coancestry between lines , $L \times L$	the sufficient statistic in disguise – see the collapse
$\mathbf{G}$	VanRaden genomic relationship, centered – prediction only	sum of all elements is exactly 0; unusable for diversity
$\mathbf{A}$	pedigree relationship	expectation of IBD, no markers
$\mathbf{G}_{0.5}$	VanRaden with the centering frequency fixed at 0.5	this is $\mathbf{f}$, arrived at from the VanRaden family

A.1.2 Coancestry and diversity

Symbol	Definition	Function
θ_S	$n^{-2} \sum_i \sum_j f_{ij}$, diagonal included	<code>theta_group()</code> , <code>fast_theta()</code>
GD, H_e	$1 - \theta$; these are the same quantity	<code>gene_diversity()</code> , <code>he_nei()</code>
N_s	$1/(2\theta)$	<code>status_number()</code>

Symbol	Definition	Function
$N_{e,\text{par}}$	$1/\sum_i p_i^2$ over parent-line usage	<code>ne_parents()</code>
f_{ii}	coancestry within subpopulation i ; $1 - f_{ii} = H_{e,i}$	<code>ct_partition()</code>
s_i	self-coancestry; $F_i = 2s_i - 1$	<code>ct_partition()</code>
$\tilde{x}$	weighted mean within subpopulations	-
$\bar{x}$	mean over all pairs , between-subpopulation pairs included	-

Always $\bar{f} \leq \tilde{f} \leq \tilde{s}$. The differences between those three numbers *are* the partition of diversity.

A.1.3 The partition

Component	Definition	Meaning	In hybrid breeding
GD_T	$1 - \bar{f}$	total	the pool's diversity
GD_{WI}	$1 - \tilde{s}$	within individuals	0 for inbred lines; use <code>GD_WI_hyb</code> for the F1s
GD_{BI}	$\tilde{s} - \tilde{f}$	between individuals, within subpopulation	segregating variation available to selection
GD_{WS}	$1 - \tilde{f}$	within subpopulations	what erodes – defend it
GD_{BS}	$\tilde{f} - \bar{f}$	between subpopulations	the heterosis engine – maintain , do not maximise (Chapter 10.)

with $GD_T = GD_{WI} + GD_{BI} + GD_{BS}$ and $F_{ST} = GD_{BS}/GD_T$.

A.1.4 Heterosis and cycles

Symbol	Meaning	Function
d_k	dominance deviation at locus k	argument to <code>heterosis_value()</code>
H_i	mid-parent heterosis of hybrid i ; $\sum_k d_k$ over its heterozygous loci	<code>heterosis_value()</code>
$\bar{p}_k$	$(p_{Ak} + p_{Bk})/2$, the merged-pool frequency	inside <code>heterosis_partition()</code>
y_k	$p_{Ak} - p_{Bk}$, the divergence at locus k	<code>pool_complementarity()</code> \$mean_sq_delta_p is y^2
g_i^A, g_j^B	general combining ability for heterosis; weighted by the opposite pool's frequencies	<code>heterosis_gca_sca()</code> \$gca_A
s_{ij}	specific combining ability; $-2\sum_k d_k \delta u_{ik} \delta v_{jk}$	<code>heterosis_gca_sca()</code> \$sca
σ_{SCA}^2	$4\sum_k d_k^2 p_{Ak} q_{Ak} p_{Bk} q_{Bk}$; falls as $(1 - F_{ST})^2$	closed form; see Section 10.11
λ	weight on within- against between-pool diversity	not implemented; see Chapter 11.
t	cycle index	<code>run_cycles()</code>
$p_0^{(0)}$	base frequencies, frozen at the founding cycle	<code>ctx\$p0</code> , set once

Expected heterosis splits as $\sum_k d_k [2\bar{p}_k(1 - \bar{p}_k) + y_k^2/2]$, whose two halves are GD_T and GD_{BS} when $d_k \equiv 1$, and whose sum is $1 - \theta_{AB}$ (Chapter 10). Over a complete factorial the realised values split orthogonally as $H_{ij} = \mu + g_i^A + g_j^B + s_{ij}$.

A.1.5 Inbreeding, the three lenses

Sym- bol	Lens	Definition	Function
F_{hom}	homozygosity	$1 - \overline{H_t/H_0}$; can be negative	<code>freq_metrics()</code>
F_{drift}	drift	$\overline{\delta p^2/[p_0(1-p_0)]}$; never negative	<code>freq_metrics()</code>
F_{IBD}	identity by descent	segments inherited from a defined base	requires pedigree + linkage
F_{ROH}	hybrid	genome fraction in runs of homozygosity	not implemented
cov	-	$2 \text{cov}(\delta p/s, (p_0 - \frac{1}{2})/s) = F_{hom} - F_{drift}$	<code>freq_metrics()\$cov_diag</code>

A.1.6 Selection and optimisation

Symbol	Meaning	Function
u_i	multi-trait index of hybrid i	<code>ctx\$index</code>
α	$(GD_{\text{ref}} - GD_S)/GD_{\text{ref}}$; negative means more diverse than chance	<code>alpha_loss()</code>
$\alpha_{\max}$	the diversity budget – a policy decision	argument to <code>stage2_select()</code>
c	continuous contributions, summing to 1	<code>ocs_quadprog()</code>
w	line-usage weights, $w_a = k_a/(2n)$	<code>line_weights()</code> , <code>line_usage_freq()</code>
k_a	times line a is used as a parent	<code>line_counts()</code>
λ	weight on the coancestry penalty in continuous OCS	<code>ocs_relaxation()</code>
σ_{DH}^2	progeny variance of a DH family	<code>sigma2_dh()</code>
UC	usefulness criterion, $\mu + ih\sigma$	<code>usefulness_criterion()</code>

A.1.7 Genomic prediction

Symbol	Meaning	Note
u , $\hat{\mathbf{u}}$	breeding values; the BLUP of them	defined over all individuals, not just the phenotyped ones
β	marker effects; $\mathbf{u} = \mathbf{Z}\beta$	the same fit, two parametrisations – exactly (Chapter 14.)
σ_u^2, σ_e^2	genetic and residual variance	supplied here, never estimated; REML in practice
λ	σ_e^2/σ_u^2 – not the OCS penalty above	a variance ratio, not a tuning knob
h^2	$\sigma_u^2/(\sigma_u^2 + \sigma_e^2)$	the one number the whole fit is conditioned on
r	accuracy: $\text{cor}(\hat{\mathbf{u}}, \mathbf{g})$, needs the truth	meaningless without its CV scheme (Section 14.4)
r_i^2	reliability: $1 - \text{PEV}_i/(\sigma_u^2 \mathbf{G}_{ii}^*)$, does not	the early warning: it collapses before the accuracy does

Symbol	Meaning	Note
PEV_i	prediction error variance, from the inverse of the MME	$\sigma_e^2[C^{uu}]_{ii}$
π	prior probability a marker has no effect (BayesB)	BayesC π estimates it; at book scale, badly
ν, S	degrees of freedom and scale of the χ^2 prior; df , R2	hyperparameters, not fitted
M_e	effective number of independent chromosome segments	unlinked markers here, so M_e is the marker count

CV0, **CV1**, **CV2** name what the training set had already seen of a candidate hybrid's *parents*: neither, one, or both. In a factorial between two pools these are three different questions, and only CV2 is easy – see Section 14.4.

A.1.8 Population sizing

Symbol	Meaning	Function
H	panel heterozygosity; 0.125 in an F4	<code>panel_het()</code>
D	diversity retained, $\overline{2p(1-p)}/0.5$	<code>diversity_retained()</code>
N_{eq}	$(1-H)/(1-D)$ – a communication device, not a Wright N_e	<code>n_equivalent()</code>
L (Morgans)	total map length	<code>map\$total_cM / 100</code>
A_t	junctions per haplotype at generation t	<code>expected_junctions()</code>
r	recombination fraction, Haldane	<code>build_map()</code>

A.1.9 Response to selection

Sym- bol	Meaning	In the code
R	response to one cycle of selection	-
i	selection intensity for a selected fraction	<code>sel_intensity()</code>
r	accuracy of the selection criterion	-
σ_A	additive standard deviation	-
V_{genic}	variance if the loci were independent; the realised variance minus this is the disequilibrium term	<code>genic_var()</code>
β_k	marker effect at locus k	<code>sim\$beta</code>
d_k	dominance deviation at locus k	<code>d</code>
ID	inbreeding depression, $2F \sum d_k p_k q_k$; a constant with fully inbred parents	<code>inbreeding_depression()</code>
b	selection index weights	<code>index_smith_hazel()</code>
a	economic weights on the aggregate genotype	-
R	restriction matrix: directions held at zero expected gain	<code>index_restricted()</code>
$2N_e R_1$	Robertson's limit; half-life $1.4N_e$	-

$R = ir\sigma_A$ is a bivariate-normal result and runs above the realised response in a structured hybrid population – see Chapter 3., which measures the size of the gap rather than assuming it away.

A.2 Conventions

Coancestry always includes the diagonal. Group coancestry in the sense of Cockerham. That is what makes $1 - \theta$ exactly Nei's H_e and makes the algebra close. Drop the diagonal and none of the identities hold.

Everything is computed per locus and then averaged over loci. Never average frequencies first.

Marker coding. Lines are 0/1 (or 0/2) because they are homozygous. Hybrids are 0/1/2 as dosage, rescaled to 0/0.5/1 as within-individual frequency. `stage1_build()` auto-detects and rescales.

f is never centered, G always is. A negative coancestry is a symptom of the wrong matrix.

Loss is always relative to a same-size random baseline. Never to the full pool.

A.3 Terms

Balding-Nichols model. A Beta distribution with mean p and variance $F_{ST}p(1-p)$, used to draw pool-specific allele frequencies from an ancestral frequency. The parameter p is the expected divergence.

Group coancestry. The probability that two alleles sampled at random from a whole set are identical in state, including the case of sampling the same individual twice.

Heterotic group / pool. A set of lines maintained as a breeding unit, crossed to lines of a complementary group to produce hybrids. Divergence between groups is where heterosis comes from – but only the divergence that coincides with dominance. Divergence produced by drift alone generates none, because it is paid for exactly out of the within-pool term. See Chapter 10..

Mid-parent heterosis. The F1's value minus the mean of its two parents. With inbred parents the additive contributions cancel exactly, leaving only the dominance deviations at the F1's heterozygous loci.

General and specific combining ability (GCA, SCA). The two halves of a hybrid's value that a two-way layout separates: what a line contributes to every cross it enters, and what depends on which two lines met. For heterosis both are exact functions of d and the pool frequencies, with GCA weighted by the *opposite* pool's frequencies – so combining ability is a property of a line against a tester pool, never of the line alone. See Section 10.10.

Reciprocal recurrent selection. Selecting on performance *in cross* and recycling the winners within their own pool. Selection information crosses the pool boundary; germplasm does not. See Chapter 19..

Identity by state (IBS) vs by descent (IBD). IBS is what markers observe: two alleles are the same allele. IBD adds the requirement of common ancestry from a defined base. The gap between them is where Chapter 9.'s whole argument lives.

Line-level collapse. The exact identity $\theta_S = \mathbf{w}'\mathbf{f}^L\mathbf{w}$ for F1 hybrids of homozygous lines, which removes the need for the $N \times N$ hybrid matrix. See Section 12.1.

Status number. $N_s = 1/(2\theta)$, the number of unrelated non-inbred individuals with the same gene diversity. A static descriptor.

Truncation selection. Take the top n by index, ignore everything else. The ceiling on gain and the floor on diversity.

Warm start. Seeding an optimiser's initial population with heuristic solutions. Not optional here: without it the differential evolution finishes below a plain greedy.

B. The **hybdiv** package

Every function used in this book lives in one package, whose source is the `R/` directory of this repository. This appendix is the index.

```
#> exported_objects
#> 193
```

B.1 Installation

```
install.packages(c("DEoptim", "quadprog"))
devtools::install(".") # from the repository root
library(hybdiv)
```

Only `DEoptim` is a hard dependency, used solely by `sel_de()`. `quadprog` is suggested and needed only for `optimal_contributions()` and `ocs_quadprog()`. Everything else, including every plot, is base R.

B.2 File map

File	Role	Chapter
<code>R/00_data.R</code>	simulation, the two matrices, the <code>ctx</code> builder	Chapter 4.
<code>R/01_metrics.R</code>	the diversity metrics, all <code>f(idx, ctx) -> scalar</code>	Chapter 11.
<code>R/02_benchmark.R</code>	null distribution, discriminatory power, cost, plots	Chapter 11.
<code>R/03_de_select.R</code>	selection strategies: truncation, greedy, DE, OCS relaxation	Chapter 16.
<code>R/04_fast_metrics.R</code>	the line-level fast route – no $N \times N$ matrix	Chapter 12.
<code>R/05_pipeline_metrics.R</code>	stage-specific metrics for $S1$ to $S4$	Chapter 13.
<code>R/06_f2size.R</code>	F1 to F4 population sizing simulator	Chapter 5.
<code>R/07_caballero_toro.C</code>	Caballero & Toro subdivided-population analysis	Chapter 2.
<code>R/08_mabc.R</code>	marker-assisted backcrossing: drag, the selection index, sizing	Chapter 6.
<code>R/09_reference.R</code>	slow literal oracles, for verification	Appendix C.
<code>R/10_stage1.R</code>	stage 1: lines to the $X / f /$ hybrids contract	Chapter 4.
<code>R/11_stage2.R</code>	stage 2: constrained selection across an alpha grid	Chapter 21.
<code>R/12_search.R</code>	structure-aware search: local search, the bound, the exact oracle	Chapter 17.
<code>R/13_dominance.R</code>	heterosis under a dominance model, and its link to the partition	Chapter 10.
<code>R/14_cycles.R</code>	recurrent selection across cycles	Chapter 19.
<code>R/15_response.R</code>	genic variance, and the Smith-Hazel and restricted indices	Chapter 3.

Files are numbered because the package has no `Collate` field and sourcing order is alphabetical.

B.3 The two-stage contract

Everything downstream of stage 1 talks only through three objects. This is the seam for swapping in real data.

```
stl <- stage1_simulate()           # or stage1_build(X, ped, traits, fL)
# $X          N x m hybrid markers, 0/0.5/1
# $f           N x N molecular coancestry
# $hybrids     N x (3 + n_traits): id, line_A, line_B, traits
# $fL          L x L line coancestry (optional)

res <- stage2_select(stl$X, stl$f, stl$hybrids, n_sel = 100, fL = stl$fL)
# $selection   0/1 column per scenario, plus n_scenarios
# $metrics     one row per scenario
# $z_scores    each metric in s.d. from the random null
# $ref         gd_ref, its Monte Carlo s.e., alpha_max, n_sel
```

B.4 Function index

B.4.1 Simulation and context

Function	Does
<code>sim_config</code>	simulation scale: lines per pool, markers, traits, divergence
<code>simulate_lines()</code>	Balding-Nichols line genotypes in two divergent pools
<code>simulate_data()</code>	lines, all A x B crosses, both matrices, predicted traits
<code>simulate_pools()</code>	lighter simulator returning raw pieces, with GL in 0/1
<code>build_ctx()</code>	assemble the context; freezes <code>p0</code> and the index
<code>build_ctx_from_stage1()</code>	same, from the stage-1 contract
<code>ctx_from_pools()</code>	context from <code>simulate_pools()</code> , carrying <code>parents</code> and GL
<code>load_data()</code>	the single swap point for real data
<code>stage1_simulate()</code> <code>stage1_build()</code>	/ produce the stage-1 contract
<code>stage2_select()</code>	the whole constrained selection pipeline

B.4.2 The two matrices

Function	Does
<code>molecular_coancestry()</code>	the diversity kernel ; one <code>tcrossprod</code> , stays in $[0, 1]$
<code>vanraden_G()</code>	prediction only ; <code>sum(G) == 0</code> by construction
<code>mrd_matrix()</code> <code>coancestry_from_mrd()</code>	/ modified Rogers distance, the same statistic

B.4.3 Diversity metrics – all `f(idx, ctx) -> scalar`

Function	Does
<code>theta_group()</code>	group coancestry, diagonal included
<code>gene_diversity()</code> / <code>he_nei()</code>	1 - <code>theta</code> ; the same quantity, two routes

Function	Does
<code>theta_from_freq()</code>	theta via allele frequencies – validation only, ~100x slower
<code>status_number()</code>	$N_s = 1/(2 \text{ theta})$; a static descriptor
<code>ne_parents()</code>	effective number of parent lines; independent of f
<code>ne_lines_pool()</code>	the same, per pool; needs only the pedigree
<code>alleles_lost()</code>	markers polymorphic in the pool, rare in the selection
<code>freq_metrics()</code>	one marker sweep : H_e , F_{hom} , F_{drift} , cov_diag , $rare_retained$
<code>line_weights()</code> / <code>theta_pools()</code>	per-pool usage weights and coancestries
<code>gd_partition()</code>	Caballero & Toro Eq. 8 over the heterotic pools
<code>ene()</code> / <code>ane()</code>	Core Hunter entry- and accession-to-nearest-entry
<code>eff_dim()</code>	effective dimensionality, on the double-centered submatrix
<code>gst()</code>	Nei's G_{st} , selected against non-selected
<code>mean_offdiag()</code>	the common production metric, kept for comparison
<code>metrics_cheap()</code> <code>metrics_full()</code>	/ the inner-loop panel and the post-hoc panel

B.4.4 The fast line route

Function	Does
<code>fast_ctx()</code>	add <code>parents</code> , <code>GL</code> and the packed bitset to a context
<code>line_counts()</code> / <code>line_usage_freq()</code>	the sufficient statistic
<code>fast_theta()</code> / <code>fast_gene_div()</code> / <code>fast_status_num()</code>	$\text{theta} = w' fL w$; no $N \times N$ matrix
<code>make_theta_fun()</code>	precompiled closure for ~1e6 evaluations
<code>theta_state()</code> / <code>theta_swap()</code>	$O(L)$ incremental update per swap
<code>fast_ne_parents()</code> / <code>ne_parents_from_counts()</code>	free once the counts exist
<code>fast_pbar()</code> / <code>fast_he_nei()</code> / <code>fast_alleles_lost()</code>	frequencies from the line matrix
<code>pack_lines()</code> / <code>popcount32()</code> / <code>alleles_retained_fast()</code>	bitset allelic richness
<code>fast_theta_pools()</code>	pool decomposition from the same counts
<code>decode_rank()</code> / <code>decode_lines()</code>	the two encodings; use <code>decode_lines()</code>
<code>make_fitness_fast()</code>	DE fitness with several simultaneous restrictions

B.4.5 Selection strategies

Function	Does
<code>sel_random()</code>	the null strategy
<code>sel_truncation()</code>	ceiling on gain, floor on diversity

Function	Does
<code>sel_truncation_cap()</code>	per-line cap; no matrix needed, usually most of the benefit
<code>sel_greedy()</code>	incremental greedy; <code>w = 0</code> is max diversity
<code>decode()</code> / <code>make_fitness()</code> / <code>sel_de()</code>	differential evolution, warm-started by default
<code>project_capped_simplex()</code> / <code>ocs_relaxation()</code>	the continuous upper bound
<code>optimal_contributions()</code> / <code>ocs_quadprog()</code>	quadratic-programming OCS

B.4.6 Benchmarking

Function	Does
<code>null_distribution()</code>	the correct baseline: random subsets of the same size
<code>alpha_loss()</code>	proportional diversity loss against that baseline
<code>evaluate()</code> / <code>discriminatory_power()</code>	metric panel and its z-scores
<code>cost_per_eval()</code>	microseconds per metric; decides what goes inside the loop
<code>plot_null()</code> / <code>plot_correlation()</code> / <code>plot_frontier()</code>	base-R diagnostics; run by <code>inst/scripts/run_benchmark.R</code> , not by the book

B.5 Runnable scripts

The book runs everything at a reduced scale so it renders quickly. Three scripts ship with the package and run the same work at the production scale of `sim_config`, writing their tables and plots to `report/`:

```
Rscript inst/scripts/run_all.R      # the two-stage pipeline, ~4 min
Rscript inst/scripts/run_benchmark.R # the metrics benchmark, ~6 min
Rscript inst/scripts/run_mabc.R     # the MABC scans, ~3 min
```

`run_benchmark.R` is the source of the numeric results quoted in `README.md`, and the only place the three `plot_*` helpers above are exercised. From an installed copy of the package:

```
source(system.file("scripts", "run_all.R", package = "hybdiv"))
```

B.5.1 Pipeline stages S1-S4

Function	Does
<code>theta_decompose()</code> / <code>gst_nei()</code>	both F_{ST} definitions, and Nei's G_{ST}
<code>pool_freqs()</code>	allele frequencies per pool, one row each
<code>pool_complementarity()</code> / <code>s1_monitor()</code>	the S1 cycle panel; <code>mean_sq_delta_p</code> is the exact one
<code>dh_sib_coancestry()</code> / <code>dh_cross_coancestry()</code> / <code>dh_pool_theta()</code>	the three DH identities
<code>sigma2_dh()</code> / <code>sigma_dh_crosses()</code>	progeny variance from marker effects
<code>usefulness_criterion()</code> / <code>sel_intensity()</code>	UC – applies at S2, not S4

Function	Does
<code>s3_theta() / alleles_retained()</code>	line-level metrics
<code>loss_vs_resampling() / theta_ceiling()</code>	calibrate against a same-size baseline
<code>select_lines_constrained()</code>	truncate-then-repair; the incremental greedy fails
<code>hybrid_theta_from_usage() / line_usage() / pool_balance()</code>	S4 on the line route
<code>loss_ledger()</code>	loss per stage, on one scale

B.5.2 Heterosis and cycles

Function	Does
<code>heterosis_value()</code>	mid-parent heterosis of each hybrid, from <code>d</code> alone
<code>heterosis_expected()</code>	the same, expected over a pool pair
<code>heterosis_partition()</code>	shared and divergence terms – <code>GD_T</code> and <code>GD_BS</code>
<code>run_cycles()</code>	recurrent selection, one row per cycle
<code>next_pool()</code>	one generation of a pool: weighted intermating, DH derivation
<code>sel_greedy_budget()</code>	greedy sweep returning the best plan inside a budget

B.5.3 Caballero & Toro

Function	Does
<code>ct_partition()</code>	subpopulation coancestries, F statistics, the Eq. 7 partition
<code>ct_loss_gain()</code>	loss or gain from removing each subpopulation

B.5.4 Population sizing

Function	Does
<code>maize_chr_len / f2_distorters / f2_opt</code>	defaults; replace the distorters with your data
<code>build_map() / marker_index()</code>	Haldane map on a regular grid
<code>meiosis() / meiosis_sel() / cross_indices() / make_F1()</code>	the vectorised engine
<code>make_distortion() / zyg_fitness()</code>	gametophytic and zygotic distortion, by exact rejection
<code>run_scheme() / reps_for_N() / bulk_weights()</code>	SSD, bulk, Syn schemes
<code>panel_freq() / panel_het() / diversity_retained()</code>	panel metrics
<code>min_minor_freq() / fixed_fraction() / skew_fraction()</code>	worst-region metrics
<code>n_equivalent()</code>	equivalent N – a communication device, not a Wright N_e

Function	Does
<code>block_lengths()</code> / <code>junctions_per_hap()</code> / <code>expected_junctions()</code>	map expansion
<code>multilocus_recovery()</code> / <code>f1_sizing()</code>	the objectives that keep paying for larger N
<code>f2_run_experiment()</code>	the full grid; minutes

B.5.5 Reference oracles

Function	Does
<code>ref_theta()</code> , <code>ref_gene_div()</code> , <code>ref_theta_freq()</code> , <code>ref_he_nei()</code>	literal transcriptions
<code>ref_status_num()</code> , <code>ref_ne_parents()</code> , <code>ref_alleles_lost()</code> , <code>ref_alleles_fixed()</code>	-
<code>ref_ene()</code> , <code>ref_ane()</code> , <code>ref_eff_dim()</code> , <code>ref_theta_pools()</code>	-

B.6 Full export list

<code>alleles_lost</code>	<code>fast_alleles_lost</code>	<code>mean_offdiag</code>	<code>reps_for_N</code>
<code>alleles_retained</code>	<code>fast_ctx</code>	<code>meiosis</code>	<code>rrs_config</code>
<code>alleles_retained_fast</code>	<code>fast_gene_div</code>	<code>meiosis_sel</code>	<code>run_bc_program</code>
<code>alpha_loss</code>	<code>fast_he_nei</code>	<code>metrics_cheap</code>	<code>run_cycles</code>
<code>alphasimr_founder_pools</code>	<code>fast_ne_parents</code>	<code>metrics_full</code>	<code>run_scheme</code>
<code>alphasimr_gd</code>	<code>fast_pbar</code>	<code>min_minor_freq</code>	<code>s1_monitor</code>
<code>alphasimr_make_hybrids</code>	<code>fast_status_num</code>	<code>molecular_coancestry</code>	<code>s3_status_num</code>
<code>alphasimr_pipeline</code>	<code>fast_theta</code>	<code>mrd_matrix</code>	<code>s3_theta</code>
<code>alphasimr_rrs_cycle</code>	<code>fast_theta_pools</code>	<code>multilocus_recovery</code>	<code>sel_de</code>
<code>alphasimr_to_stage1</code>	<code>fixed_fraction</code>	<code>n_equivalent</code>	<code>sel_de_fast</code>
<code>ane</code>	<code>foreground_ok</code>	<code>n_for_count</code>	<code>sel_exact</code>
<code>backcross</code>	<code>freq_metrics</code>	<code>ne_lines_pool</code>	<code>sel_greedy</code>
<code>BITS_PER_WORD</code>	<code>gd_partition</code>	<code>ne_parents</code>	<code>sel_greedy_budget</code>
<code>block_lengths</code>	<code>gene_diversity</code>	<code>ne_parents_from_counts</code>	<code>sel_intensity</code>
<code>build_ctx</code>	<code>genic_var</code>	<code>next_pool</code>	<code>sel_local_search</code>
<code>build_ctx_from_stage1</code>	<code>gst</code>	<code>null_distribution</code>	<code>sel_random</code>
<code>build_map</code>	<code>gst_nei</code>	<code>ocs_bound</code>	<code>sel_truncation</code>
<code>bulk_weights</code>	<code>haldane_r</code>	<code>ocs_quadprog</code>	<code>sel_truncation_cap</code>
<code>coancestry_from_mrd</code>	<code>he_nei</code>	<code>ocs_relaxation</code>	<code>select_lines_constrained</code>
<code>col_cumsum</code>	<code>heterosis_expected</code>	<code>ocs_round</code>	<code>select_parent</code>
<code>cost_per_eval</code>	<code>heterosis_gca_sca</code>	<code>optimal_contributions</code>	<code>selfing_step</code>
<code>cp_hi</code>	<code>heterosis_partition</code>	<code>p_het</code>	<code>sigma_dh_crosses</code>

cp_lo	heterosis_value	p_same_segment	sigma2_dh
cross_indices	hybrid_theta_from_usagpack_lines		sim_config
ct_loss_gain	ibd_elite	panel_freq	simulate_data
ct_partition	inbreeding_depression	panel_het	simulate_lines
ctx_from_pools	index_restricted	plot_correlation	simulate_pools
decode	index_smith_hazel	plot_frontier	sizing_table
decode_lines	junctions_per_hap	plot_null	skew_fraction
decode_rank	line_counts	pool_balance	stage1_build
dh_cross_coancestry	line_usage	pool_complementarity	stage1_simulate
dh_pool_theta	line_usage_freq	pool_freqs	stage2_select
dh_sib_coancestry	line_weights	popcount32	status_number
discriminatory_power	load_data	project_capped_simplex	strat_rec_until
diversity_retained	loss_ledger	ref_alleles_fixed	summarise_reps
donor_dosage	loss_vs_resampling	ref_alleles_lost	target_indices
donor_specific	mabc_opt	ref_ane	theta_ceiling
drag_total	maize_chr_len	ref_eff_dim	theta_decompose
eff_dim	maize_targets	ref_ene	theta_from_freq
ene	make_distortion	ref_gene_div	theta_group
evaluate	make_F1	ref_genic_var	theta_pools
expected_donor	make_fitness	ref_he_nei	theta_state
expected_drag	make_fitness_fast	ref_heterosis	theta_swap
expected_ibd	make_shared_ibd	ref_inbreeding_depression	usefulness_criterion
expected_junctions	make_theta_fun	ref_ne_parents	vanraden_G
f1_sizing	marker_index	ref_status_num	zyg_fitness
f2_distorters	marker_weights	ref_theta	
f2_opt	max_line_use	ref_theta_freq	
f2_run_experiment	mean_index	ref_theta_pools	

C. Verification

A derivation the reader cannot re-run against a naive implementation is a claim, not a verification. This appendix says how each result in the book was checked, and how to re-run the checks.

C.1 The method: slow oracles

Every fast route in the package has a **deliberately slow, literal** counterpart in `R/09_reference.R` that transcribes the definition with no algebraic shortcut. The fast route must match it to machine precision.

```
ref_theta
#> function (idx, ctx)
#> mean(ctx[["f"]][idx, idx])
#> <bytecode: 0xa33b3c580>
#> <environment: namespace:hybdiv>
theta_group
#> function (idx, ctx)
#> mean(ctx$f[idx, idx])
#> <bytecode: 0xa33b7b968>
#> <environment: namespace:hybdiv>
fast_theta
#> function (idx, ctx)
#> {
#>   cnt <- line_counts(idx, ctx)
#>   u <- which(cnt > 0L)
#>   cu <- cnt[u]
#>   s <- sum(cu)
#>   as.numeric(crossprod(cu, ctx$fL[u, u, drop = FALSE] %*% cu))/(s *
#>     s)
#> }
#> <bytecode: 0xa33bad3f0>
#> <environment: namespace:hybdiv>
```

Three implementations of one quantity: a literal mean over the submatrix, the packaged version, and the line-route collapse that never forms the submatrix at all. They must agree.

```
sim <- simulate_pools(n_A = 20, n_B = 20, m = 1500, seed = 7)
ctxp <- ctx_from_pools(sim)
set.seed(3); sel <- sample.int(ctxp$N, 50)

c(oracle = ref_theta(sel, ctxp),
  packaged = theta_group(sel, ctxp),
  line_route = fast_theta(sel, ctxp),
  max_dev = max(abs(c(theta_group(sel, ctxp), fast_theta(sel, ctxp)) -
    ref_theta(sel, ctxp))))
#>      oracle      packaged      line_route      max_dev
#> 6.708316e-01 6.708316e-01 6.708316e-01 1.110223e-16
```

The oracles exist so that this comparison is possible for a reader who does not trust the derivation – which is the correct posture.

```
fmt(rbind(
  he_nei = c(oracle = ref_he_nei(sel, ctxp), packaged = he_nei(sel, ctxp)),
  theta_freq = c(ref_theta_freq(sel, ctxp), theta_from_freq(sel, ctxp)),
  status_num = c(ref_status_num(sel, ctxp), status_number(sel, ctxp)),
  ne_parents = c(ref_ne_parents(sel, ctxp), ne_parents(sel, ctxp)),
```

```

alleles_lost = c(ref_alleles_lost(sel, ctxp), alleles_lost(sel, ctxp)),
ENE          = c(ref_ene(sel, ctxp), ene(sel, ctxp)),
ANE          = c(ref_ane(sel, ctxp), ane(sel, ctxp))), 10)

```

	oracle	packaged
he_nei	0.3291684	0.3291684
theta_freq	0.6708316	0.6708316
status_num	0.7453435	0.7453435
ne_parents	26.4550265	26.4550265
alleles_lost	43.0000000	43.0000000
ENE	0.2581333	0.2581333
ANE	0.2503233	0.2503233

C.2 The identity sweep

The line-level collapse was checked across $F_{ST} \in \{0.05, 0.15, 0.35\}$, two seeds and two selection sizes, thirty replicate selections each.

	Identity	Worst deviation over the sweep
max_pairwise_f	pairwise f: hybrid vs line formula	1.1e-16
max_theta_sub_vs_lin	theta: line route vs f submatrix	2.2e-16
max_theta_sub_vs_freq	theta: line route vs allele-frequency route	1.1e-16
max_pbar	selection allele frequencies	0 (exact)
max_He	Nei He	0 (exact)
max_Neparents	effective number of parents	2.1e-14
max_alleles_lost	allele loss at MAF threshold	0 (exact)
max_theta_pools	pool decomposition	1.2e-15
max_swap_update	incremental swap vs full recompute	4.4e-16
bitset_vs_literal	bitset allelic richness vs literal count	0 (exact)

C.3 The test suite

The suite is `tests/testthat/`, run with `devtools::test()`. It takes about a minute. What it asserts, and why each check earns its place:

Check	If it fails
theta via submatrix == theta via allele frequencies	f is wrong: centered, rescaled, or the wrong marker coding
1 - theta == Nei's He	the identity underpinning every 'loss of X%' statement is broken
f is bounded to [0, 1]	a centered matrix is being used as coancestry
sum(VanRaden G) == 0	the demonstration in Section 1.2 no longer holds on this data

Check	If it fails
theta and Ns analytical edge cases	the diagonal is being dropped somewhere
greedy < random < truncation on theta	there is no trade-off to optimise; the whole benchmark is pointless
decode() is unique, in range, deterministic	the DE fitness is noisy and cannot converge
fast line route == hybrid-matrix route	the collapse is broken ; production scale becomes infeasible
incremental swap == full recomputation	any local search built on the state vector is wrong
popcount32 == naive bit count	allelic richness counts are wrong
F_hom and F_drift are 0 when the selection is the base	the base p0 is not frozen, or is being recomputed
cov_diag == F_hom - F_drift	one of the three lens formulas is wrong
max-diversity drives F_hom negative, F_drift positive	the two-lens argument of Chapter 9. does not reproduce
freq_metrics reproduces he_nei and alleles_lost	the consolidated marker sweep changed an existing number
the Caballero & Toro partition closes	the Eq. 8 algebra is broken
GD_WI == 0 for inbred lines, GD_WI_hyb large for their F1s	lines are not actually inbred, or the hybrid diagonal is wrong
eff_dim in (0, n - 1]	the double-centering was dropped and eff_dim re-measures theta
the two-stage output contract	downstream code reading <code>\$selection</code> or <code>\$metrics</code> will break
stage 2 runs without fL	the optional-fL path is broken
stage1_build accepts both marker codings	real data in the other coding would give different answers
genic_var == its literal oracle	the genic-variance route diverged from its definition
realised minus genic variance == the disequilibrium term	the disequilibrium decomposition of Chapter 3. is wrong
inbreeding depression == the shared term of the heterosis partition	Falconer's symmetry is broken, or d is being applied inconsistently
inbreeding depression == simulated genotypes (Monte Carlo)	the closed form does not describe genotypes anyone could actually draw
the restricted index moves the restricted trait exactly zero	the restriction is a penalty , and the Chapter 15. analogy is wrong
no restriction reduces the restricted index to Smith-Hazel	the degenerate case silently returns something other than Smith-Hazel
MRD and coancestry are one statistic	Trap 1 of Chapter 13. no longer holds
Gst < Fst always, and Gst matches the classical route	Trap 2 is wrong, or the two definitions were conflated
the DH identities	the S2 algebra is wrong and crossing plans cannot be pre-evaluated

Check	If it fails
hybrid theta from usage == the hybrid matrix	Trap 5's repair route is broken
constrained line selection meets its ceiling	the constrained selection silently returns an infeasible plan
the oracles agree with the packaged metrics	a fast route diverged from its definition
map length and Haldane spacing	the sizing chapter's map is not the map it claims
heterozygosity halves per selfing	the simulator is not doing what selfing does
$D == 1 - (1 - H)/N$ under neutrality	the identity every sizing number rests on is broken
n_equivalent inverts diversity_retained	N_eq is not the inverse it is described as
a Syn cycle adds half a map length of junctions	the intercrossing budget in Chapter 5. is wrong by 2x
distortion shifts ancestry away from 0.5	the distortion mechanism is inactive
f1_sizing halves the loss probability per plant	the F1 sizing argument is wrong
heterosis_value matches the literal oracle	the fast heterosis route diverged from its definition
additive effects cancel from the mid-parent difference	the derivation of Chapter 10. is wrong, or the parents are not inbred
shared and divergence == GD_T and GD_BS	the central identity of Chapter 10. is broken
$2 * GD_WI_hyb == GD_T + GD_BS$	the hybrid-level and line-level partitions have come apart
neutral divergence leaves expected heterosis unchanged	Chapter 10.'s correction is wrong and F_ST would predict heterosis
alignment between d and divergence is what moves heterosis	the alignment result, and the argument for Chapter 19. with it, is wrong
GCA + SCA reconstruct heterosis exactly and are mean-zero	the combining-ability decomposition of Chapter 10. is not a decomposition
SCA is the interaction residual of the two-way additive fit	SCA is a proxy for the interaction, not the interaction – the section overclaims
the closed-form GCA and SCA variance components match	the variance-component result of Section 10.11 rests on a wrong closed form
uniform-d heterosis is the coancestry kernel, exactly	genetic distance no longer predicts heterosis under flat dominance
$GD_T + GD_BS == 1 - \theta_{AB}$	the between-pool coancestry reading of the identity is wrong
pool_freqs reproduces what its callers computed inline	a pool frequency route diverged after the refactor
the DH draw is an unbiased gamete; selfing an inbred returns it	recycling fabricates or destroys diversity; every trajectory is wrong
p0 is frozen, not re-anchored, across cycles	both inbreeding lenses report no loss forever – the trap of Chapter 19.

Check	If it fails
a closed programme loses diversity and gains index	the cycle loop is not selecting, or not recycling what it selected
the budget is respected; alpha_max = Inf is truncation	the budget is not binding across cycles
injection raises diversity relative to a closed run	injection is not actually introducing founder material

```
Rscript -e 'devtools::test()'
```

C.4 Provenance of every figure and table

The book runs light examples live and reads heavy results from cache. This is the complete accounting.

C.4.1 Computed live at render

Chap-Runs live
ter

Chap- both matrices, the sum(G) demonstration, the f diagonal
ter 1.

Chap- the whole worked example, reproduced against the published tables
ter 2.

Chap- the breeder's equation against the realised response and its decomposition, the genic-variance
ter 3. oracle check, the GCA/SCA variance identity, both selection indices

Chap- the simulation, the fst sweep, the contract, the sampling-bias histogram
ter 4.

Chap- the theoretical identities, heterozygosity halving, map expansion, f1_sizing
ter 5.

Chap- the founder-engine reproducibility and LD comparison, the varA-against-genicVarA trajectory,
ter 7. the zero F1 variance of inbred parents, the selfing-versus-DH homozygosity trajectory, the
breeder's equation over a selected-fraction sweep, sel_intensity against AlphaSimR selInt,
Smith-Hazel against a numerical optimum, the RRS trajectory and the stage-2 metrics

Chap- the staged programme in both arms, the paired phenotypic-versus-genomic comparison over 24
ter 8. seeds, the gene-diversity ledger

Chap- the Eq. 3 identity, the four-strategy lens panel
ter 9.

Chap- the heterosis oracle check, the additive cancellation, the shared/divergence identity, the coances-
ter 10. try-kernel identity, the fst sweep and its Balding-Nichols prediction, the alignment scenarios
and their stratified control, the GCA/SCA decomposition and its variance components, the SCA/
GCA relocation sweep

Chap- every metric, the sampling-bias formula, the null distribution
ter 11.

Chap- the collapse, the swap update, cost at book scale, the fixed-theta band, the prediction-error sweep
ter 12.

Chap- all five traps, the S1 panel, the DH identities, the S4 collapse
ter 13.

Chap- Runs live
ter

Chap- the OCS relaxation, the convergence check, the greedy frontier, the marker-panel sensitivity
ter 15.

Chap- all baselines, the encodings, the DE, the warm-start comparison
ter 16.

Chap- the Gram identity and PSD check, the submatrix-sum identity, the three local-search modes, the
ter 17. invalid-versus-Lagrangian bound comparison, ocs_round

Chap- the frozen-p0 demonstration, the delta-F regression
ter 19.

Chap- Robertson's limit against the observed trajectory, the paired selected-versus-random variance
ter 20. test in both the unlinked and the linked arm

Chap- the full two-stage pipeline
ter 21.

Ap- the oracle comparisons
pen-
dix C.

C.4.2 Read from cache

File	Why
data/theta_identity_checks_full.csv	the full identity sweep, 12 configurations
data/cost_per_eval.csv	production scale: L = 300, N = 22500, m = 10000
data/metric_redundancy.csv	621 selections x 17 metrics
data/de_constraints.csv	9 DE runs at 60,000 evaluations each
data/fst_definitions.csv	the divergence sweep
data/dist_vs_variance.csv	600 candidate crosses with simulated DH
data/pipeline_loss_ledger.csv	9 founder pools, four stages, two strategies
data/plant_metric_usage.csv	687 screened plant records
data/bibliography_*.csv	the screened corpora themselves
data/f2_results.csv	5 schemes x 20 sizes, adaptive replication
data/f2_multilocus.csv	multilocus recovery across the grid
data/mabc_*.csv	the MABC grid: 4 policies x 4 gene numbers x 11 sizes x 3 s, 300 reps
figs/mabc_sizing.png	analytical sizing by stage
figs/mabc_policies.png	the four recombinant-selection policies
figs/cost_scaling.png	cost scaling at production scale

File	Why
<code>figs/de_convergence.png</code>	DE convergence under four restrictions
<code>figs/metric_redundancy.png</code>	the redundancy panels
<code>figs/pipeline_loss.png</code>	the loss ledger
<code>figs/metric_pitfalls.png</code>	Traps 4 and 2 on the full simulation
<code>data/search_benchmark.csv</code>	8 methods x 2 scales x 3 seeds, production scale N = 10000
<code>data/search_scaling.csv</code>	per-evaluation cost from N = 400 to N = 22500
<code>data/search_exact_gap.csv</code>	the outer-approximation oracle, proven small and time-limited large (target <code>search_exact</code>)
<code>figs/search_anytime.png</code>	anytime curves at equal wall-clock
<code>figs/search_scaling.png</code>	the encoding scaling wall
<code>data/cycles_trajectory.csv</code>	3 budgets x 16 cycles x 5 seeds (target <code>cycles</code>)
<code>data/cycles_crossover.csv</code>	the crossover cycle, per budget per seed
<code>data/cycles_maxuse.csv</code>	3 caps x 16 cycles x 5 seeds
<code>data/cycles_injection.csv</code>	closed against one injection at cycle 8, 5 seeds
<code>data/cycles_selector.csv</code>	greedy against DE, 2 budgets x 13 cycles, seed 1
<code>figs/cycles_gain.png</code>	cumulative gain by budget
<code>figs/cycles_diversity.png</code>	gene diversity and F_drift by budget

Nothing is read from `report/` or `data/` at the repository root: both are gitignored, so a clean clone would fail to render. Everything cached lives in `book/data` and `book/figs`, which are tracked.

C.5 Regenerating the cached results

```
Rscript book/scripts/regenerate.R      # everything; hours
Rscript book/scripts/regenerate.R f2size # one target
Rscript book/scripts/regenerate.R search # the advanced-selection benchmark
Rscript book/scripts/regenerate.R search_exact # just the exact oracle
Rscript book/scripts/regenerate.R cycles # the recurrent-selection trajectories
```

The script is the single place where the expensive computations live. It is never invoked by `quarto render`.

C.6 Reproducing the book

```
Rscript -e 'devtools::document(); devtools::load_all()'
Rscript -e 'devtools::test()'
Rscript -e 'devtools::check()'
quarto render book
```

Quarto's `freeze: auto` means a chapter re-runs only when its source changes. To force a full recomputation, delete `book/_freeze/`.

C.7 Known limitations of the verification

Stated so they are not mistaken for oversights.

The oracles share the author's understanding. They check that the fast route computes what the slow route computes. They cannot check that the slow route implements the right definition. That is what the reproduction of the published Caballero & Toro tables in Chapter 2. is for – an external check against numbers this project did not produce.

Monte Carlo checks have Monte Carlo error. Where a test asserts agreement with a simulated expectation rather than an algebraic identity, the tolerance is loose by necessity, and a genuinely small bias would pass.

Everything is verified on simulated data. The identities are algebraic and will hold on real genotypes. The *empirical* results – which metrics are redundant, how much index a constraint costs – were derived at $F_{ST} = 0.15$ with simulated markers, and Chapter 12. says explicitly that they should be re-derived on real genotypes before being treated as settled.

D. The bibliographic corpora

Chapter 18. draws on two screened corpora. This appendix describes them so the claims there can be checked, extended, or disagreed with.

D.1 What they are

Corpus	Records	Years	Focus
<code>bibliography_animal</code>	324	2018-2026	optimal contribution, coancestry matrix choice, effective size, allelic richness
<code>bibliography_plant</code>	687	2017-2026	diversity quantification, heterotic structure, genetic resources, optimisation

Both are annotated screens, not exhaustive searches. The plant corpus came from 2004 records across 29 plant-breeding queries on Crossref and Europe PMC, of which 687 were scored relevant.

D.2 Structure

```
names(pl)
#> [1] "doi"           "year"          "authors"       "title"
#> [5] "journal"       "citations"     "relevance"     "crop"
#> [9] "metric_used"   "selection_method" "contribution"  "url"
```

Column	Meaning
<code>relevance</code>	3 = directly relevant, 2 = related background
<code>crop</code>	plant corpus only
<code>metric_used / metric_advocated</code>	the diversity metrics the record reports
<code>selection_method</code>	plant corpus only
<code>theme</code>	animal corpus only; multi-valued, semicolon-separated
<code>contribution / relevance_note</code>	a one-line note on what the record adds

D.3 Relevance and coverage

```
table(plant = pl$relevance)
#> plant
#>  2  3
#> 529 158
table(animal = ani$relevance)
#> animal
#>  2  3
#> 235  89
```

crop	records
maize	54

crop	records
wheat	11
multi-crop	10
rice	9
none	7
sorghum	5
barley	4
none (general crop)	4
potato	3
soybean	3

Maize dominates, which is the intended bias – but it also means every claim in Chapter 18. about crops other than maize rests on a handful of records.




Figure D.1: Publication year of the screened records. Both corpora are deliberately post-2017.




Figure D.2: Which diversity metrics the plant literature actually reports.

D.4 Themes in the animal corpus

The `theme` column is multi-valued. Splitting it:

```
th <- unlist(strsplit(ani$theme[ani$relevance == 3], ";\\s*"))
th <- trimws(th); th <- th[nzchar(th)]
sort(table(th), decreasing = TRUE)
#> th
#>      OCS / contribution optimisation Coancestry / relationship matrix choice
#>                                68                                61
#>      Optimisation algorithms / cost      Allelic richness / rare alleles
#>                                27                                20
#>      Effective size      Heterozygosity / gene diversity
#>                                18                                15
#>      Core collection / subset selection
#>                                9
```

D.5 Querying them

The corpora are ordinary CSVs, so the claims in Chapter 18. are checkable in one line each.

```
# records that report a coancestry or kinship metric
sum(grepl("coancestry|kinship|relatedness", pl$metric_used, ignore.case = TRUE))
#> [1] 12

# records using ROH, and how many of those are maize
roh <- grepl("ROH", pl$metric_used, ignore.case = TRUE)
c(roh_total = sum(roh), roh_maize = sum(roh & grepl("maize", pl$crop, ignore.case = TRUE)))
#> roh_total roh_maize
#>      1      0

# the most-cited directly relevant maize records
mz <- pl[pl$relevance == 3 & grepl("maize", pl$crop, ignore.case = TRUE), ]
head(mz[order(-mz$citations), c("year", "title", "citations")], 6)
```

year	title	citations
2020	Diversity and heterotic patterns in North American proprietary dent maize germplasm	58
2020	Discovery of beneficial haplotypes for complex traits in maize landraces.	46

year	title	citations
2021	Genetic diversity and selection signatures in maize landraces compared across 50 years of in situ and ex situ conservation	44
2021	Application of Genomic Selection at the Early Stage of Breeding Pipeline in Tropical Maize.	38
2020	Efficiencies of Heterotic Grouping Methods for Classifying Early Maturing Maize Inbred Lines	26
2021	Genetic variation and population structure in China summer maize germplasm.	25

```
# the animal-side records on matrix choice
mc <- ani[grepl("Coancestry / relationship matrix choice", ani$theme), ]
head(mc[order(-mc$citations), c("year", "title", "citations")], 6)
```

year	title	ci- tations
2018	Optimal cross selection for long-term genetic gain in two-part programs with rapid recurrent genomic selection	156
2019	The impact of genomic selection on genetic diversity and genetic gain in three French dairy cattle breeds	100
2020	Management of Genetic Diversity in the Era of Genomics.	72
2019	Genomic Tools for Effective Conservation of Livestock Breed Diversity	71
2020	Molecular genetic analysis of spring wheat core collection using genetic diversity, population structure, and linkage disequilibrium	65
2019	Improving Short- and Long-Term Genetic Gain by Accounting for Within-Family Variance in Optimal Cross-Selection.	59

D.6 Limitations

These are abstract screens. Relevance and the metric annotations were scored from titles and abstracts, not from full texts. A paper that computes a metric without naming it in the abstract is invisible to this table – which almost certainly understates how often coancestry is computed as an intermediate step.

The queries shaped the answer. 29 plant-breeding queries were run; a different query set would return a different landscape. The strong maize skew is by design, and the counts for other crops are too small to support a claim about those crops.

“Not established in the retrieved set” means exactly that. Every such statement in Chapter 18. is a claim about *this corpus*, not about the literature as a whole. It is a prompt to look further before assuming novelty, not a proof of it.

Citation counts are a snapshot at screening time and are already stale.

D.7 Classical references

The corpora are deliberately post-2017. The classical results the book rests on are cited from the source literature, and collected in `book/references.bib`:

Result	Source
group coancestry	Cockerham (1967)
molecular / IBS coancestry	Nejati-Javaremi et al. (1997)
gene diversity, <i>Gst</i>	Nei (1973)
the subdivided-population partition	Caballero & Toro (2000, 2002)

Result	Source
status number	Lindgren & Mullin (1997, 1998)
optimal contribution selection	Meuwissen (1997)
genomic relationship matrix	VanRaden (2008)
modified Rogers distance	Wright (1978)
usefulness criterion	Schnell & Utz (1975)
differential evolution	Storn & Price (1997)

References

Bibliography

- [1] A. Caballero and M. A. Toro, "Analysis of genetic diversity for the management of conserved subdivided populations," *Conservation Genetics*, vol. 3, no. 3, pp. 289–299, 2002, doi: [10.1023/A:1019956205473](https://doi.org/10.1023/A:1019956205473).
- [2] A. Caballero and M. A. Toro, "Interrelations between effective population size and other pedigree tools for the management of conserved populations," *Genetics Research*, vol. 75, no. 3, pp. 331–343, 2000, doi: [10.1017/S0016672399004449](https://doi.org/10.1017/S0016672399004449).
- [3] A. Nejati-Javaremi, C. Smith, and J. P. Gibson, "Effect of total allelic relationship on accuracy of evaluation and response to selection," *Journal of Animal Science*, vol. 75, no. 7, pp. 1738–1745, 1997, doi: [10.2527/1997.7571738x](https://doi.org/10.2527/1997.7571738x).
- [4] F. Gómez-Romano, B. Villanueva, J. Fernández, J. A. Woolliams, and R. Pong-Wong, "The use of genomic coancestry matrices in the optimisation of contributions to maintain genetic diversity at specific regions of the genome," *Genetics Selection Evolution*, vol. 48, p. 2, 2016, doi: [10.1186/s12711-015-0172-y](https://doi.org/10.1186/s12711-015-0172-y).
- [5] M. A. Toro, B. Villanueva, and J. Fernández, "Genomics applied to management strategies for conservation of farm animal genetic resources," *Animal Genetics*, vol. 45, no. 3, pp. 1–10, 2014, doi: [10.1111/age.12140](https://doi.org/10.1111/age.12140).
- [6] T. H. E. Meuwissen, A. K. Sonesson, G. Gebregiorgis, and J. A. Woolliams, "Management of genetic diversity in the era of genomics," *Frontiers in Genetics*, vol. 11, p. 880, 2020, doi: [10.3389/fgene.2020.00880](https://doi.org/10.3389/fgene.2020.00880).
- [7] R. C. Gaynor, G. Gorjanc, and J. M. Hickey, "AlphaSimR: an R package for breeding program simulations," *G3: Genes, Genomes, Genetics*, vol. 11, no. 2, 2021, doi: [10.1093/g3journal/jkaa017](https://doi.org/10.1093/g3journal/jkaa017).
- [8] R. E. Comstock, H. F. Robinson, and P. H. Harvey, "A breeding procedure designed to make maximum use of both general and specific combining ability," *Agronomy Journal*, vol. 41, no. 8, pp. 360–367, 1949, doi: [10.2134/agronj1949.00021962004100080006x](https://doi.org/10.2134/agronj1949.00021962004100080006x).
- [9] G. K. Chen, P. Marjoram, and J. D. Wall, "Fast and flexible simulation of DNA sequence data," *Genome Research*, vol. 19, no. 1, pp. 136–142, 2009, doi: [10.1101/gr.083634.108](https://doi.org/10.1101/gr.083634.108).
- [10] I. J. Mackay, J. Cockram, and P. Howell, "Understanding the classics: the unifying concepts of transgressive segregation, inbreeding depression and heterosis and their central relevance for crop breeding," *Plant Biotechnology Journal*, vol. 19, no. 1, pp. 26–34, 2021, doi: [10.1111/pbi.13481](https://doi.org/10.1111/pbi.13481).
- [11] O. Kempthorne and A. W. Nordskog, "Restricted selection indices," *Biometrics*, vol. 15, no. 1, pp. 10–19, 1959, doi: [10.2307/2527598](https://doi.org/10.2307/2527598).
- [12] S. Michel, F. Löschnerberger, C. Ametz, B. Pachler, E. Sperry, and H. Bürstmayr, "Combining grain yield, protein content and protein quality by multi-trait genomic selection in bread wheat," *Theoretical and Applied Genetics*, vol. 132, no. 10, pp. 2767–2780, 2019, doi: [10.1007/s00122-019-03386-1](https://doi.org/10.1007/s00122-019-03386-1).
- [13] R. C. Gaynor *et al.*, "A two-part strategy for using genomic selection to develop inbred lines," *Crop Science*, vol. 57, no. 5, pp. 2372–2386, 2017, doi: [10.2135/cropsci2016.09.0742](https://doi.org/10.2135/cropsci2016.09.0742).
- [14] T. H. E. Meuwissen, B. J. Hayes, and M. E. Goddard, "Prediction of total genetic value using genome-wide dense marker maps," *Genetics*, vol. 157, no. 4, pp. 1819–1829, 2001, doi: [10.1093/genetics/157.4.1819](https://doi.org/10.1093/genetics/157.4.1819).
- [15] M. R. Labroo, A. J. Studer, and J. E. Rutkoski, "Heterosis and hybrid crop breeding: a multidisciplinary review," *Frontiers in Genetics*, vol. 12, 2021, doi: [10.3389/fgene.2021.643761](https://doi.org/10.3389/fgene.2021.643761).
- [16] D. S. Falconer and T. F. C. Mackay, *Introduction to Quantitative Genetics*, 4th ed. Harlow: Longman, 1996.
- [17] J. C. Reif, F.-M. Gumpert, S. Fischer, and A. E. Melchinger, "Impact of interpopulation divergence on additive and dominance variance in hybrid populations," *Genetics*, vol. 176, no. 3, pp. 1931–1934, 2007, doi: [10.1534/genetics.107.074146](https://doi.org/10.1534/genetics.107.074146).
- [18] F. Technow, T. A. Schrag, W. Schipprack, E. Bauer, H. Simianer, and A. E. Melchinger, "Genome properties and prospects of genomic prediction of hybrid performance in a breeding program of maize," *Genetics*, vol. 197, no. 4, pp. 1343–1355, 2014, doi: [10.1534/genetics.114.165860](https://doi.org/10.1534/genetics.114.165860).
- [19] M. R. White, M. A. Mikel, N. de Leon, and S. M. Kaeppler, "Diversity and heterotic patterns in North American proprietary dent maize germplasm," *Crop Science*, vol. 60, 2020, doi: [10.1002/csc2.20050](https://doi.org/10.1002/csc2.20050).
- [20] C. C. Cockerham, "Group inbreeding and coancestry," vol. 56, no. 1, 1967, pp. 89–104.

- [21] M. Nei, "Analysis of gene diversity in subdivided populations," *Proceedings of the National Academy of Sciences*, vol. 70, no. 12, pp. 3321–3323, 1973, doi: [10.1073/pnas.70.12.3321](https://doi.org/10.1073/pnas.70.12.3321).
- [22] D. Lindgren, L. D. Gea, and P. A. Jefferson, "Status number for measuring genetic diversity," *Forest Genetics*, vol. 4, no. 2, pp. 69–76, 1997.
- [23] M. A. Toro, B. Villanueva, and J. Fernández, "The concept of effective population size loses its meaning in the context of optimal management of diversity using molecular markers," *Journal of Animal Breeding and Genetics*, vol. 137, no. 4, pp. 345–355, 2020, doi: [10.1111/jbg.12457](https://doi.org/10.1111/jbg.12457).
- [24] E. López-Cortegano, R. Pouso, A. Labrador, A. Pérez-Figueroa, J. Fernández, and A. Caballero, "Optimal management of genetic diversity in subdivided populations," *Frontiers in Genetics*, vol. 10, p. 843, 2019, doi: [10.3389/fgene.2019.00843](https://doi.org/10.3389/fgene.2019.00843).
- [25] H. De Beukelaer, G. F. Davenport, and V. Fack, "Core Hunter 3: flexible core subset selection," *BMC Bioinformatics*, vol. 19, no. 1, p. 203, 2018, doi: [10.1186/s12859-018-2209-z](https://doi.org/10.1186/s12859-018-2209-z).
- [26] A. Caballero and S. T. Rodríguez-Ramilo, "A new method for the partition of allelic diversity within and between subpopulations," *Conservation Genetics*, vol. 11, no. 6, pp. 2219–2229, 2010, doi: [10.1007/s10592-010-0107-7](https://doi.org/10.1007/s10592-010-0107-7).
- [27] C. R. Henderson, "Best linear unbiased estimation and prediction under a selection model," *Biometrics*, vol. 31, no. 2, pp. 423–447, 1975, doi: [10.2307/2529430](https://doi.org/10.2307/2529430).
- [28] J. Burgueño, G. de los Campos, K. Weigel, and J. Crossa, "Genomic prediction of breeding values when modeling genotype $\times$ environment interaction using pedigree and dense molecular markers," *Crop Science*, vol. 52, no. 2, pp. 707–719, 2012, doi: [10.2135/cropsci2011.06.0299](https://doi.org/10.2135/cropsci2011.06.0299).
- [29] D. Habier, R. L. Fernando, and J. C. M. Dekkers, "The impact of genetic relationship information on genome-assisted breeding values," *Genetics*, vol. 177, no. 4, pp. 2389–2397, 2007, doi: [10.1534/genetics.107.081190](https://doi.org/10.1534/genetics.107.081190).
- [30] H. D. Daetwyler, B. Villanueva, and J. A. Woolliams, "Accuracy of predicting the genetic risk of disease using a genome-wide approach," *PLoS ONE*, vol. 3, no. 10, p. e3395, 2008, doi: [10.1371/journal.pone.0003395](https://doi.org/10.1371/journal.pone.0003395).
- [31] G. de los Campos, J. M. Hickey, R. Pong-Wong, H. D. Daetwyler, and M. P. L. Calus, "Whole-genome regression and prediction methods applied to plant and animal breeding," *Genetics*, vol. 193, no. 2, pp. 327–345, 2013, doi: [10.1534/genetics.112.143313](https://doi.org/10.1534/genetics.112.143313).
- [32] A. Xavier, W. M. Muir, B. Craig, and K. M. Rainey, "bWGR: Bayesian whole-genome regression," *Bioinformatics*, vol. 36, no. 6, pp. 1957–1959, 2019, doi: [10.1093/bioinformatics/btz794](https://doi.org/10.1093/bioinformatics/btz794).
- [33] D. Gianola, G. de los Campos, W. G. Hill, E. Manfredi, and R. Fernando, "Additive genetic variability and the Bayesian alphabet," *Genetics*, vol. 183, no. 1, pp. 347–363, 2009, doi: [10.1534/genetics.109.103952](https://doi.org/10.1534/genetics.109.103952).
- [34] A. Borodin, A. Jain, H. C. Lee, and Y. Ye, "Max-sum diversification, monotone submodular functions, and dynamic updates," *ACM Transactions on Algorithms*, vol. 13, no. 3, p. 41, 2017, doi: [10.1145/3086464](https://doi.org/10.1145/3086464).
- [35] G. L. Nemhauser, L. A. Wolsey, and M. L. Fisher, "An analysis of approximations for maximizing submodular set functions – I," *Mathematical Programming*, vol. 14, no. 1, pp. 265–294, 1978, doi: [10.1007/BF01588971](https://doi.org/10.1007/BF01588971).
- [36] M. L. Fisher, G. L. Nemhauser, and L. A. Wolsey, "An analysis of approximations for maximizing submodular set functions – II," *Mathematical Programming Studies*, vol. 8, pp. 73–87, 1978, doi: [10.1007/BFb0121195](https://doi.org/10.1007/BFb0121195).
- [37] A. Cevallos, F. Eisenbrand, and R. Zenklusen, "Max-sum diversity via convex programming," in *32nd International Symposium on Computational Geometry (SoCG)*, 2016, p. 26:1–26:14. doi: [10.4230/LIPIcs.SoCG.2016.26](https://doi.org/10.4230/LIPIcs.SoCG.2016.26).
- [38] S. Yadav and others, "Optimising parent selection in plant breeding: comparing metaheuristic algorithms for genotype building," *Theoretical and Applied Genetics*, vol. 138, no. 9, p. 242, 2025, doi: [10.1007/s00122-025-05028-1](https://doi.org/10.1007/s00122-025-05028-1).
- [39] D. Akdemir, S. Rio, and J. Isidro y Sánchez, "TrainSel: an R package for selection of training populations," *Frontiers in Genetics*, vol. 12, p. 655287, 2021, doi: [10.3389/fgene.2021.655287](https://doi.org/10.3389/fgene.2021.655287).
- [40] J. Fernández-González, S. M. Metwally, and J. Isidro y Sánchez, "MateR: a novel genomic mating framework," *Genetics*, vol. 232, no. 4, p. iyag013, 2026, doi: [10.1093/genetics/iyag013](https://doi.org/10.1093/genetics/iyag013).
- [41] P. Waldmann, "Genomic optimum contribution selection and mate allocation using JuMP," *Bioinformatics Advances*, vol. 5, no. 1, p. vbaf259, 2025, doi: [10.1093/bioadv/vbaf259](https://doi.org/10.1093/bioadv/vbaf259).
- [42] M. A. Duran and I. E. Grossmann, "An outer-approximation algorithm for a class of mixed-integer nonlinear programs," *Mathematical Programming*, vol. 36, no. 3, pp. 307–339, 1986, doi: [10.1007/BF02592064](https://doi.org/10.1007/BF02592064).

- [43] P. Ahadi, B. Balasundaram, J. S. Borrero, and C. Chen, “Development and optimization of expected cross value for mate selection problems,” *Heredity*, vol. 133, pp. 80–90, 2024, doi: [10.1038/s41437-024-00697-y](https://doi.org/10.1038/s41437-024-00697-y).
- [44] O. G. Younis, L. Corinzia, I. N. Athanasiadis, A. Krause, J. M. Buhmann, and M. Turchetta, “Breeding programs optimization with reinforcement learning,” *arXiv preprint arXiv:2406.03932*, 2024, doi: [10.48550/arXiv.2406.03932](https://doi.org/10.48550/arXiv.2406.03932).
- [45] L. Chen, G. Zhang, and H. Zhou, “Fast greedy MAP inference for determinantal point process to improve recommendation diversity,” in *Advances in Neural Information Processing Systems 31*, 2018, pp. 5622–5633.
- [46] G. Gorjanc, R. C. Gaynor, and J. M. Hickey, “Optimal cross selection for long-term genetic gain in two-part programs with rapid recurrent genomic selection,” *Theoretical and Applied Genetics*, vol. 131, 2018, doi: [10.1007/s00122-018-3125-3](https://doi.org/10.1007/s00122-018-3125-3).
- [47] A. Allier, C. Lehermeier, A. Charcosset, L. Moreau, and S. Teyssèdre, “Improving short- and long-term genetic gain by accounting for within-family variance in optimal cross-selection,” *Frontiers in Genetics*, vol. 10, 2019, doi: [10.3389/fgene.2019.01006](https://doi.org/10.3389/fgene.2019.01006).
- [48] Y. C. J. Wientjes, P. Bijma, M. P. L. Calus, B. J. Zwaan, Z. G. Vitezica, and J. van den Heuvel, “The long-term effects of genomic selection: 1. Response to selection, additive genetic variance, and genetic architecture,” *Genetics Selection Evolution*, vol. 54, 2022, doi: [10.1186/s12711-022-00709-7](https://doi.org/10.1186/s12711-022-00709-7).
- [49] A.-C. Doublet, P. Croiseau, S. Fritz, A. Michenet, C. Hozé, and C. Danchin-Burge, “The impact of genomic selection on genetic diversity and genetic gain in three French dairy cattle breeds,” *Genetics Selection Evolution*, vol. 51, 2019, doi: [10.1186/s12711-019-0495-1](https://doi.org/10.1186/s12711-019-0495-1).
- [50] J. E. Rutkoski, “Estimation of realized rates of genetic gain and indicators for breeding program assessment,” *Crop Science*, vol. 59, 2019, doi: [10.2135/cropsci2018.09.0537](https://doi.org/10.2135/cropsci2018.09.0537).
- [51] D. Müller, P. Schopp, and A. E. Melchinger, “Selection on expected maximum haploid breeding values can increase genetic gain in recurrent genomic selection,” *G3: Genes, Genomes, Genetics*, vol. 8, 2018, doi: [10.1534/g3.118.200091](https://doi.org/10.1534/g3.118.200091).
- [52] A. Robertson, “A theory of limits in artificial selection,” *Proceedings of the Royal Society of London B*, vol. 153, no. 951, pp. 234–249, 1960, doi: [10.1098/rspb.1960.0099](https://doi.org/10.1098/rspb.1960.0099).
- [53] M. G. Bulmer, “The effect of selection on genetic variability,” *The American Naturalist*, vol. 105, no. 943, pp. 201–211, 1971, doi: [10.1086/282718](https://doi.org/10.1086/282718).